

JORDAN BOYD-GRABER

QUESTIONING ARTIFICIAL INTELLIGENCE

UNIVERSITY OF MARYLAND

Copyright © 2026 Jordan Boyd-Graber

[TUFTE-LATEX.GITHUB.IO/TUFTE-LATEX/](https://tufte-latex.github.io/tufte-latex/)

2022

Contents

1	<i>Introduction: The Unanswered Questions</i>	5
	<i>I Question Answering Before Computers</i>	9
2	<i>Question Answering: The Foundation and Savior of Civilization</i>	11
3	<i>The Turing Test: A Game Show Pitch that Defined Artificial Intelligence</i>	17
4	<i>The Cranfield Paradigm: How a University with an Airstrip made Google Possible</i>	21
5	<i>The Manchester Paradigm: The Art of Asking the Perfect Question</i>	31
	<i>II Question Answering Present</i>	45
6	<i>Watson on Jeopardy!: Unquestioned Answers from IBM's tour de force</i>	47
7	<i>Formats Beyond Stand-Alone QA</i>	67
8	<i>Datasets: Build or Adopt?</i>	91

9	<i>The Tools to Answer Questions</i>	113
10	<i>Siri takes the SAT</i>	117
11	<i>Finding and Dealing with Bad Questions</i>	141
	<i>III Question Answering Future</i>	143
12	<i>Fulfilling Turing's Gameshow</i>	145
13	<i>What AI Dystopias You Should be Afraid of</i>	147

Thanks

1

Introduction: The Unanswered Questions

Since this is a book about question answering, I should try to answer some questions to start things out. I'm going to try to answer: "what's this book about", "who am I", and "who can use this book".

AI has been defined by asking questions even before it was "a thing": Alan Turing designed a question answering procedure to determine if machine is intelligent. In the years since, AI has continued to be defined by questions: answering questions on *Jeopardy!* and on our phones.

This book looks at not just the *how* of computers answers questions but also *why* answering questions is so important for AI. Answering that "why" questions requires us to go back to mythological, civil service exams, and game shows and connect them to AI question answering.

We call someone smart if they can answer questions: the pass a test, win on *Jeopardy!*, or offer a witty retort in a debate. This is not just a recent phenomenon. Question answering is deeply embedded in our culture: in myth Oedipus answered the riddle of the sphynx; getting into a school depends on passing a standardized test; good governance depends on civil service exams; a trivia subculture has grown around answering all manner of questions.

AI has been defined by asking questions even before it was "a thing": Alan Turing designed a question answering procedure to determine if machine is intelligent. In the years since, AI has continued to be defined by questions: answering questions on *Jeopardy!* and on our phones. From IBM Watson on *Jeopardy!* to Siri and Alexa in every home, the ability of computers to answer is a common yardstick of how smart computers are.

As we'll see this is important in both contemporary civilization and is fundamental to artificial intelligence. The goal of this book is to help explain how computers answer questions on your phone, on game shows, and to fight fake news. Importantly, from my perspective, this won't just be about computers answering questions: humans and computers have a lot to learn from each other, and I hope we'll be able to discover new things about those interactions.

This book looks at not just the *how* of computers answers questions but

also *why* answering questions is so important for AI. Answering that “why” questions requires us to go back to mythological, civil service exams, and game shows and connect them to AI question answering.

To be transparent, this book has an agenda: we can better understand artificial intelligence by looking at its ability to answer questions through this historical lens. If we want to call a computer smart—indeed, smarter than a human—we should make sure the competition is fair. Moreover, writing and asking questions is an art that has been refined over decades.

Part of the story I want to tell is how we got here. And why little things like the decisions made at a tiny British library in Bedfordshire fifty years ago have locked the artificial intelligence community into a particular way of evaluating whether a computer is smart. These historical decisions have shaped the systems that we have today; as a consequence, we are stuck with the crappy answers from our smartphone.

But it didn’t have to be this way! The very definition of computing and artificial intelligence was based a thought experiment that is closer to how computers and humans should interact with each other. However, that definition—from a boffin in Manchester called Alan Turing—was too fanciful to really be implemented in the lab in the twentieth century. I’ll argue that we could do something closer to Alan Turing’s vision, and you, the reader, can judge whether it’s (still) too fanciful.

This is not just a narrow question for those building question answering systems. I think that the way we interact with our smartphones and computers will define the future of human-computer interaction and thus the shape of the modern, AI-infused economy.

1.1 Who am I?

I’m excited to be writing this because it’s something that I’m passionate about: answering questions is central not just to my research program but also to my day job as a professor at the University of Maryland: teaching and guiding students. And it’s also important to me as a hobby: even before this became part of my research program, I spent weekends writing and answering silly trivia questions. I’m hoping that this will help more people appreciate the joy and beauty of posing and answering questions.

In my day job, I have built question answering systems that used the structure of Wikipedia to better answer questions, was one of the first to apply deep learning to the QA task, and played a lot of exhibition games against trivia whizzes like Yogesh Raut, Ken Jennings, and Roger Craig. I’m also a professor of computer science, and one of the courses that I teach is a course on question answering (and this book is the textbook for it).

As a hobby, I’m also a trivia enthusiast. I played QB for Caltech (where I got BS degrees in computer science and history), started Princeton’s pub quiz (where I went to grad school), and appeared on *Jeopardy!*. To be clear,

I'm not very good at trivia: our QB teams at best cracked top ten nationally, and I came in second place in the one game of *Jeopardy!* I taped (my students are particularly fond of mocking my poor performance on the video game category). But this has brought me a familiarity with the trivia community and how it works, one that I think gives me a unique (or deranged) look at how computers answer questions.

In contrast, I've won more awards as a researcher. Probably the one I'm most proud of is the Karen Spärk Jones award (we'll learn more about her in Chapter 4) from the British Computing Society, but I've also gotten an NSF CAREER award and "best of" awards from Intelligent User Interfaces, Neural Information Processing Systems, the North American Association for Computational Linguistics, Empirical Methods in Natural Language Processing, and the Conference on Natural Language Learning. I've published over a hundred peer-reviewed publications on interactive machine learning, question answering, and exploring document collections; and this isn't my first book, with David Mimno and Yuening Hu, I wrote *Applications of Topic Models*. Although I should say that because I'm a professor, this is mostly the work of the students that I've worked with over the years (I thank them by name at the end of the book).

But the biggest prize that I've gotten is to have a career (and permissive bosses) that allow me to combine my hobby with these intellectual pursuits. And I'm even more lucky that that career allows me to now share it with you.

1.2 How the Book is Structured

While this book is ostensibly about how *computers* answer questions, that is not where the book begins. This history is important because it explains why humans value answering questions. This provides historical context on why humans equate answering questions correctly with intelligence but also because key elements for artificial intelligence—item response theory and preference models—were developed from looking at how humans answered questions.

The next stop is looking at *today's* question answering, begining with Alan Turing, arguably the father of computing and artificial intelligence. Two wartime efforts—one at air and one at sea—began a rivalry in question answering, which we call the Cranfield and Manchester based on the college towns home to Alan Turing's theorizing on intelligent computing and Cyril Cleverdon's practical evaluation of computing systems. Having set the stage of this rivalry, we then see how these ideas built the foundations of modern AI from the methods that let *Watson* win on *Jeopardy!* to datasets like Google's Natural Questions that set the stage for smart chatbots and the models like GPT.

Finally, the book closes with the *future* of question answering, focusing on what it means for an AI to be intelligent: how to turn Alan Turing's thought

experiment into reality and how to make sure that computers don't trick us into thinking they're intelligent (but just cheating on the test). And if computers are becoming intelligent the question is how to *align* them with users, which sees the return of item response theory and preference models: this time built up so we can better understand why AI says the things it does. Finally, we close with a discussion of the "big questions" that computers cannot answer on their own: how can the models developed from answering questions from humans best work alongside humans for work and play.

Part I

Question Answering Before Computers

2

Question Answering: The Foundation and Savior of Civilization

Despite its central importance to AI, asking and answering questions is not new—it stretches back into our collective unconscious. Riddles and trivia have persisted for millenia not just because they’re entertaining. Indeed, being able to ask and answer questions is close to godliness. From the Sphynx to Gestumblindi, this chapter examines how myths from several cultures connect answering questions to forming identity, unlocking the secrets of the universe, and gaining intelligence. And connecting this to nascent “superhuman” AI. Later, beyond proving worthiness in myth, asking questions helped combat corruption and nepotism, the downfall of multiple civilizations. However, the administration of carefully formed questions saved the China and the United States via the reforms of Wu Zetian to the Pendleton Act. Exams created meritocracy and social mobility while also building an understanding of what it means to write a “fair” and “useful” question.

With many tasks in computer science, asking the question “where did this task come from” has a simple answer usually in the form of an acronym. Character recognition was sparked by the MNIST¹ dataset (Wilkinson et al., 1992; LeCun et al., 1998); entailment was sparked by SNLI (MacCartney and Manning, 2008; Bowman et al., 2015); sentiment classification was defined by Pang and Lee (2008).

And while there are particular *flavors* of question answering that have been defined to make life easier or more discriminative for computers, the basic task of “question in”, “answer out” is natural to any inquisitive three year old. Of course, question answering is not the only “natural” task like this: translation likewise is recognizable to any stranger in a strange land trying to cope with an unfamiliar language: dozens of books have been written on that subject, my favorite is *Le Ton Beau de Marot* (Hofstadter, 1997).

Question answering’s roots likewise run deep. In this chapter, I want to argue that question answering is not just a useful task but one that is woven deep into our Jungian subconscious. Our collective civilization is defined by how questions are answered; it forms an identity, reveals fundamental truths

¹ Throughout, I will do my best to cite the origin of ideas, but MNIST—like other early works in AI—is hard to pin down to a single citation. The OCR group at the United States National Institutes of Standards and Technology (NIST) was charged with helping the post office sort letters automatically by “reading” envelopes, so they organized a conference to try to get people to give it their best shot and released data to let people train models. One of the approaches from LeCunn et al. years later defined what people mean when they actually say MNIST, which then became a part of countless tutorials for people learning machine learning.

of the universe, and creates a shared reality. And beyond these ancient components, it has historically created a historical basis for a shared expectation of competence that created the modern bureaucracy.

2.1 *The Sphynx*

In the Greek of *Oedipus* (Sophocles, 1959), the Sphynx asked everyone who entered the city a riddle. Here's PDQ Bach's interpretation of the riddle (Schickele, 1990):²

What starts out on four legs then goes round on two Then finishes on three
before it's through

And the answer is:

Well a baby then a man, that's plain, then finally an old geezer on a cane

Now I won't spoil the whole story, but by answering the question correctly, Oedipus revealed that he was smarter than the average bear and also unlocked secrets from his past. In a commencement address to Whitman college, Dana Burgess Burgess explained the difference between this and Google's Natural Questions (covered in detail in Chapter 8) and how that relates to the mission of professors (like me):

Nobody told Oedipus who he was; he figured that out for himself. ... You can do a Google search to find out the capital of Arkansas (Little Rock?) but you can't do a Google search to find out who you are, what you're good at, what makes you happy, what matters for your life. Information transferal isn't any help with that stuff.

When folks say that a college professor is like a sphinx, they don't usually mean that in a nice way. It suggests that we know the answers students need, but we keep them secret for the sake of our own power. ... That's part of the riddle of [The Sphynx's] hybrid nature: the curse of being forced to solve a riddle is also the gift of the ability to solve riddles.

2.2 *The Socratic Method*

This aspect of question answering also brings us to another Greek connection to asking questions: teaching through the Socratic method (Trepanier, 2017). Through asking the right questions, a teacher guides the answerer to understanding. While Cranfield questioners are seeking information from a more knowledgeable answerer, Manchester questioners often test less knowledgeable answerers. Similarly, perhaps by asking the right questions, the QA community can coax computers to understand more than they do now (Dunietz et al., 2020; Perez et al., 2020).

Because I'm a professor, I am somewhat biased in my appreciation of this use of questions: both in the classrooms and on exams. If you haven't heard of

² I use this fairly unserious version of the riddle because it is fun and Renger (2013) reports hexameter versions, and the origin (Pindar fragments are the first mention) is unclear. So if one has to arbitrarily pick, a version localized into American vernacular is perhaps the most accessible.

this before, you might have experienced it (especially if you're in my classes). Let's hear how John Houseman explain what it is in the 1971 film *The Paper Chase* ([pap, 1973](#)):

We use the Socratic Method here. I call on you, ask you a question...and you answer it. Why don't I just give you a lecture? Because through my questions, you learn to teach yourselves.

Through this method of questioning, answering...questioning, answering...we seek to develop in you the ability to analyze that vast complex of facts that constitute the relationships of members within a given society. Questioning and answering.

At times you may feel that you have found the correct answer. I assure you that this is a total delusion on your part. You will never find the correct, absolute, and final answer. In my classroom, there is always another question...another question to follow your answer.

Yes, you're on a treadmill. My little questions spin the tumblers of your mind. You're on an operating table. My little questions are the fingers probing your brain.

We do brain surgery here. You teach yourselves the law but I train your mind. You come in here with a skull full of mush...and you leave thinking like a lawyer.

I want to particularly hone in on the assertion that it "trains your brain". While not quite mythological, the term "Socratic Dialog" goes back to ancient Greece. Plato featured Socrates in his dialogues *Euthyphro* and *Ion* ([Plato, 1949](#)), giving the perspective of the student the student, the person answering the questions, struggling with multiple interpretations and "showing their work" as they emerge on the other side, enlightened by the exchange.

And indeed, one of the best moments for me as a professor is to see the aha moment in students eyes after they answer a question correctly. And this is still how we train lawyers, who have to use argument and discussion to keep our society ordered! But as we will see later in this chapter, that is not the only way the questions and answers create order in our civilization.

2.3 Gestumblindi

In Norse myth, Gestumblindi got into some trouble with King Heidrek and asked the Gods for help ([Tolkien, 1960](#)).³ In response, Odin came to Gestumblindi—looking exactly like him—and stood in his place in front of Heidrek. What took place was an interrogation that went in two directions: King Heidrek attempted to determine who this Gestumblindi-looking person was in front of him and Odin attempted to see how smart King Heidrek was.

What I like about this story is that it presages many of the themes that we will see with AI: ordinary people thinking that AI will help them, an interrogation contest to see who is human or not, and an impersonation. It also presages many of our fears about AI: Odin's final question to Heidrek is:

³ I had originally wanted to use Auden's translation (since Auden is more highbrow), but Tolkein's translation also evokes (and reminds the reader) of Gollum's battle of riddles with Bilbo that won him the ring, an example of riddles deciding the fate of a (fictional) civilization.

What said Ódin
in the ear of Baldur,
before he was borne to the fire?

We talk in this book frequently about bad questions (Chapter 11), but this a particularly bad question. Not only is it something that only Odin knows, it probably has not happened yet (depending on what you believe about the circularity of time), and when it does happen it will end the world by bringing about Ragnarök.

Once Heidrek hears this question, he knows that the figure before him is not Gestumblindi but rather Odin. And like a frustrated customer who thought they were talking to a human for ten minutes only to figure out that they're talking to a bot, Heidrek takes out his sword and attacks, only for Odin to turn into a bird and fly away.

This verbal jousting is common for the Norse Gods (Mitchell, 2020); being a God requires Omniscience, so the only way to tell if you are in the company of God is through riddles. The only problem is that—at least in Norse Myth—challenging a God is hazardous to your health: those who challenge Odin to a battle of wits often end up dying. My goal here is not to say that we should not ask probing questions; subsequent chapters will argue that we should ask these questions often and advise on how to ask the questions well. Rather, the lesson of Gestumblindi is that the art of making a riddle is prized. Just as some Christians argued that “cleanliness is next to godliness” (Jackson, 1872), in the Norse tradition, questioning is synonymous with being a God, which will return when we talk about the apocalyptic visions of AI (Chapter (?)).

2.4 *The Foundation of Civilization*

Less exciting than the mythic foundations of civilization are the humdrum matters of administration. But even if it's boring, a pillar of civilization is making sure that those who work for society (in other words, civil servants) are competent. As governments became larger and had to rule an entire continent, countries had to develop civil service examinations. Many countries did this, including the US, whose Pendleton act ended an era of political patronage (Johnson et al., 2007).

But I'd rather talk about the granddaddy of civil service exams, kējū, which for 1300 years determined who became part of the intelligentsia of Imperial China (Qian and Totten, 1982). It encouraged social mobility by getting smart people good jobs and it encouraged good governance by making sure that important jobs were done by smart people. And this is important when you have a diverse, massive, empire that is run by a centralized bureaucracy. Again, this shows that civilization is not possible without effective systems of question answering.

This system worked well for so long because it was built on a foundation of fairness. The organizers thought about the potential for bias: the exams were

transcribed so that bad handwriting wouldn't be judged against an applicant, children of current members of the imperial court had to submit their exams in their home province, etc.

Now, it wasn't perfect, Most notably its quotas led to imbalances between regions. But it was important enough that recent scholarship has suggested that its abolition after the Russo-Japanese war and a turn to "Western-style" education on the Prussian model that turns out good soldiers, helped hasten the end of the Qing dynasty (Bai and Jia, 2016).

As Chu (2015) argues, it's not just that you have a test to make sure that the smartest people are doing important jobs in society. Part of the process is also to explain why people got the questions wrong! This ensures that people trust how the questions are being graded. If it's just a black box, it is not an improvement on the capricious systems of patronage that it's supposed to replace. As evidence of this, during the Dà Míng dynasty, there was furor over the punctuation of graded exams.

So why am I, a computer science professor, talking all of this nonsense about history? We use questions to test the intelligence of both humans and machines. I talk about the Turing Test (Chapter 3), the classic test of computer intelligence, and modern leaderboards (Chapter 10), but in testing computer intelligence we shouldn't forget about the lessons we've learned about human intelligence.

One of those lessons that I think that we've forgotten is that questions are not just an evaluation but also for instruction: in the lingo of artificial intelligence, they're training data (Chapter 8). And because these training data build intelligence, they should be as high quality as possible so that the AI that results is as high quality as possible.

But even if you just take question answering as an evaluation for how smart an AI is, society should be able to trust those evaluations. Just as the civil service exams made people trust their interactions with the government, as AI becomes more tightly integrated into our economy and our society, our vetting of AI will become more important (Chapter 13).

Question answering is just one of many ways to secure trust, but it's analogous to how to ensure lawyers, doctors, and pilots are qualified for their jobs. So these tests should be unbiased, reliable, and the feedback from the tests should be transparent and understandable.

And we should do these things not just because we're trying to make society work like a well-oiled machine or because we want to be confident in our estimates of statistics about artificial intelligence. Just like Socrates, the questions we ask are trying to get to the truth, and scientific inquiry requires openness and a willingness to question the outcomes of a test.

3

The Turing Test: A Game Show Pitch that Defined Artificial Intelligence

In the 1950s, Alan Turing proposed a parlor game that would come to define the artificial intelligence: could a wiley interrogator discern whom they were talking to just through posing clever questions. The eponymous Turing Test is the most durable (but contentious) definition of what an intelligent computer is, and this chapter reviews how it shaped the development of artificial intelligence.

Topics in chapter:

1. **Turing's Legacy.** Why Turing is widely considered the father of computer science for multiple reasons.
2. **The Imitation Game.** A gay man in England—where homosexuality is criminalized—posits a game about the ability of someone to pass in different gender roles. Why this is difficult lays the foundation for why computers still struggle to capture human nuance.
3. **No, the Turing Test has not been Solved.** While many have claimed to pass the Turing Test, it hasn't happened yet. We review some notable claims of passing the Turing Test and why they fell short.
4. **A Rigorous Test.** The ingredients that make a Turing Test legitimate
5. **General Artificial Intelligence.** How the holy grail of AI is defined and how the Turing Test can measure it.

The history of artificial intelligence begins in many ways with a question answering game. This idea came out of the research of a researcher at the University of Manchester named Alan Turing.

3.1 Turing's Legacy

He's probably better known for being one of the scientists who decoded the Nazi Enigma device (which is why his memorial bench has ciphertext underneath his name) and helped bring World War II in Europe to an earlier conclusion.

3.2 *The Imitation Game*

But Turing is also a game designer. Quoting from Bishop's description:

Turing called for a human interrogator (C) to hold a conversation with a male and female respondent (A and B) with whom the interrogator could communicate only indirectly by typewritten text. The object of this game was for the interrogator to correctly identify the gender of the players (A and B) purely as a result of such textual interactions

The players can lie, so the key to correctly deciding the genders of the players is more about determining which player lacks key knowledge about the experience of being a man or a woman.¹

But what does this have to do with AI? Turing then thought, let's replace the man and the woman with a computer and a human. If the interrogator cannot determine which is the computer and which is the machine, then the machine has displayed something that looks like intelligence.

This is hard because there are lots of "tells" that can expose someone who is transgressing against their assigned gender roles. A skilled interrogator can ask about how it feels to have a period, the rules for picking which urinal to use in a crowded restroom, or how to respond to a girlfriend asking about her outfit.

And while an impostor can read many accounts to get an idea of what someone has said about it, simply parroting these stories does not equate to understanding.

3.3 *No, the Turing Test has not been Solved*

But I want to emphasize not just the broad strokes of the Turing Test, but why we need to think critically about things that call themselves the Turing test. Also, it's more fun than talking abstractly about the Turing test.

Let's start with PARRY, a system designed by Kenneth Mark Colby to simulate a paranoid schizophrenic. And when you connected psychiatrists to either real patients or PARRY, they couldn't tell the difference. So here the problem is the judges. Not because they aren't experts—they are—but because their backgrounds prevent them from being effective judges.

Psychiatrists are doctors bound by the Hippocratic oath: they cannot ask probing, in-depth questions that might harm a patient. So, thus, this really isn't a Turing test.

Stu Schieber has a great take on this problem; I'd encourage you to read the whole thing. His take is on a competition called the Loebner prize that purports to implement the Turing Test. But here, the interrogators are limited in the topics they can ask about. So again, competitions that don't have skilled interrogators allowed to ask any question they wish.

Another example that some people claim is an example of AI passing the Turing Test is Google Duplex: Google offers to call a restaurant to make a

¹ What makes the Turing Test more poignant to me, at least, is that as a closeted gay man in a country where homosexuality is illegal, every interaction is dangerous: Turing is pretending to be something that he is not, and has to present himself as if he were a heterosexual man. There's no evidence to suggest that this was an inspiration for his first Turing Test, however.

reservation for you. Their text-to-speech system is very good, but also puts in dysfluencies such as adding “um” and pauses to make it sound more human. Here, the judge is fooled into thinking that they’re talking to a human. But this doesn’t count either because the judge doesn’t know they’re a judge! The poor employee on the other end of the phone call is expecting a human until proven wrong.

All of this doesn’t mean that the Turing Test is flawed. It has remained a part of AI for three quarters of a century because it’s a simple, intuitive test of whether we have achieved artificial intelligence. So although we haven’t had a real Turing Test yet, a judge asking questions of either an AI or a human remains many researchers’ goal.

3.4 A Rigorous Test

So let’s be true to the spirit of Turing’s idea of a parlor game. Let’s make it visible to the public, let’s refine the rules and the judges to make it more realistic and more fun. By putting these games in the public view and letting judges learn the best strategies for discerning humans from computers, both sides can become worthier adversaries.

And I think putting this slow advance of ever more capable computers answering trickier questions should be out in front of the public. Not just to keep the interrogators honest but to also keep the companies and interests that sell AI honest. The public has a vested interest in knowing the limits of AI, and this is a fine way to make that public. But on the other side of the coin, it is also worthwhile for the public to know when AI has really advanced... the public has a right to know how computers react to challenging scenarios. Better to see them first played out for fun in a game than in high-stakes transactions, a doctor’s office, or a courtroom.

But above all, this only works if we have good questions, so if we believe that the Turing test really is the holy grail of AI, we as humans need to know how to ask the right questions and computers need to be able to answer any question that’s thrown at them. In Chapter 12, we outline how this can become a gameshow.

3.5 General Artificial Intelligence

4

The Cranfield Paradigm: How a University with an Airstrip made Google Possible

Information retrieval is the foundation of multi-billion dollar web companies from Baidu to Yahoo! But none of these companies would exist without the ideas of reusable test collections and term-based queries, which came about because of a few crazy experiments that happened in a small UK University in Cranfield. This chapter lays out the history of this methodology (which has become known as the Cranfield paradigm) and how it can answer questions.

Often, when I talk about search engines and how to build them, these ubiquitous portals to the Internet—whether you’re using Bing, Yahoo!, Baidu, or Ask Jeeves—seem like a fundamental law of science. With the onslaught of AI, they almost seem antiquated... my parents use Google, but I use Grok. But it wasn’t always like that: there was a time before search engines, and although there’s a certain charm to card catalogs, that’s not an effective way to find information. But that was essentially the world before search engines: you had to find what you wanted from a list or index.

My goal in this video is to tell the story of the Cranfield paradigm and how it made the twentieth century Internet possible. Most of this is drawn from the article “The Evolution of Cranfield” by [Voorhees \(2019\)](#), which I encourage you to read something with more detail and more references, more focused on information retrieval.

The story begins in 1967 with Cyril Cleverdon at Cranfield University in England: which has an airport right on campus.¹ The University was the heart of British aviation research. Cyril Cleverdon was in charge of their library there. And he was big into computers. He wanted to see if he could measure how good an index was.

Now, when I say “index”, if you’re over 40, you might think of this sort of thing in the back of a book ([Ritchie and Kernighan, 1988](#)):

INDEX 269

read system call 170
readdir function 170

¹ Here at the University of Maryland we have to walk a whole fifteen minutes to get to the oldest continuously operated airport in the world.



Figure 4.1: Photo of Cranfield University (and its adjacent airstrip). Photo by Thomas Nugent [Creative Commons Attribution-ShareAlike 2.0 license].

```

readlines function    170
realloc library function    170
recursion 86, 139, 141, 182, 269
  
```

An author carefully wrote down topics mentioned in a text and told you where you could find them (i.e., which page of the book they were on).²

But if you’re younger and a computer scientist, you probably think of an index as a lookup table of every darn word possible. The reason you think that is because of the Cranfield experiments. His claim, which was quite controversial at the time, was that you could look up documents based on the words in the document. You didn’t need laboriously curated indices you could use a machine, an engine if you will, to search through them. The claim was that the Universal Decimal Classification, a facet classification, and alphabetical subject catalog, and a unigram index were roughly comparable (Cleverdon, 1967).

Before we analyze this, just a reminder of how a search engine query is evaluated. Again, this is old hat now but was revolutionary at the time. To evaluate a system’s ability to search a dataset, you take some queries from users and then find the truly relevant documents to that query. Then, given a system’s results, you compute the precision—out of all of the documents they returned how many were right—and recall—out of all possible right documents how many did they find.

Because these concepts are important, let’s pause to review how success at finding a relevant answer was measured by Keen (1966). First let’s consider—given a query—how many answers you retrieved. Some of those were relevant to the query and gave you information you wanted. Let’s say there are A of

² In this case, they also helpfully told you that “recursion” appeared in the index (the page you’re looking at), a computer science joke.

those. But others were not relevant, these are distractions that you still have to weigh through. Let's say there are B of those.

	Relevant	Non-relevant	Total
Retrieved	A_q	B_q	$A_q + B_q$
Not Retrieved	C_q	D_q	$C_q + D_q$
Total	$A_q + C_q$	$B_q + D_q$	$A_q + B_q + C_q + D_q$

Table 4.1: Contingency table for question q used to define precision and recall.

Clearly it's better to have A be higher than B , and this metric is called *precision*:

$$\text{precision}(q) = \frac{A_q}{A_q + B_q}, \quad (4.1)$$

of the answers you provided, how many were relevant. The higher the better. But let's consider how many of the relevant answers you did find: A_q . It would be better if A_q were *all* the relevant documents, of which there are A_q and C_q . This is called the recall:

$$\text{recall}(q) = \frac{A_q}{A_q + C_q}. \quad (4.2)$$

Both of these numbers are between 0 (bad) and 1 (good); ideally, you would want both to be 1.0, but this is hard to achieve. You can get perfect recall by returning all possible answers, but then your precision sinks. If you return only the answers you're absolutely certain of, your recall suffers.

Different applications (or users) would have different tradeoffs. If you're trying to find a "smoking gun" in a court case, you want recall over all else: it doesn't matter if you have to comb through too many irrelevant documents. But if you're trying to answer a simple question, you don't need every document that tells you what the capital of Arkansas is, you just need one, and you don't want anything irrelevant.

So, if you have a set of queries $q_1 \dots q_N$ you can run your system on each question, get a set of results, compute the precision and recall of each one, and then create a ranking of the systems.

Let's pause for a moment to recognize how revolutionary this was. Once you found the relevant documents, you can evaluate *any* system. Before, you had to run an expensive user experiment every time you tweaked your system: did that output look good? How about this one? With a reusable test collection, you could turn the crank on your retrieval system without pesky humans getting in the way. Just check whether the red box finds the relevant documents!

But not everybody was convinced. The Cranfield paradigm is making a lot of assumptions here: first, if a document is relevant, it's equally relevant: everything is equally relevant: showing you my explanation of topic modeling is as good as a video from Siraj Raval's The users' information needs are static; the answer to the question "who is the president of the united states"

never changes (Nixon is still president, right?). Finally, the Cranfield paradigm assumes that somebody, somehow can find the right answer from thousands or millions of documents.

The academic community at the time was skeptical. There's no way these newfangled search engines could be better than a trained librarian armed with a card catalog! And Cuadra (1968) noted reasonable criticisms: different users have different information needs—I don't need anybody to tell me what the capital of Maryland is, but I don't know the capital of the Yevreyskaya Oblast; the same user has different needs—I might do the same search twice to get the citation for Cuadra and to find a good quote. And Cuadra does have some good quotes, he correctly warns us:

For a time, psychologists devoted their attention not to what was important, but to what was measurable. This is a surefire approach to a certain kind of respectability, which information scientists now desire as much as psychologists did then, but it risks losing the baby with the bathwater.

Despite all those flaws, the assumptions—on average—mostly hold. That's why we're still talking about the Cranfield paradigm fifty years later. And there were some big additions to the Cranfield paradigm. Most prominent was the use of better document representations and creating larger test collections, which was spearheaded by Karen Spärk Jones a little bit west of Cranfield at Cambridge and by Ellen Vorhees and others at the National Institute for Standards and Technology in Gaithersburg, Maryland.

Term-Frequency and Inverse Document Frequency

We will talk about neural representations later when we talk about models like BERT (Section 9.1). Those representations are called *dense*—not because they're dumb, they're quite clever—but they pack a lot of information into a small packet of computer memory. In contrast, the representations that Karen Spärk Jones came up with are *sparse*: they could have an entry for every word in your language, but for any given question, you're only going to have a handful of words present.

And because “word” is ambiguous, we'll call this “term frequency” instead. The entry for a question has a higher weight for a word the more often it appears. E.g., if the question is:

Which of the members of Lincoln's cabinet
did Lincoln previously serve in the House of Rep-
resentative with?

then the words that appear most commonly are going to be more important. In this case, that would be “Lincoln”...but another seemingly important is “the” since it also appears just as often.

The great intuition of Spärk Jones is that while you and I know that “the” is unimportant, and you could make a list of such unimportant words (often

called a “stop” list in the lingo), that list is going to be brittle. For example, “resolved” is an unimportant word in legislative contexts. The thing that ties together unimportant words is that they appear in many different questions.

Mathematically, you want to downweight words that appear in many documents. Spärk Jones motivated inverse document frequency as

[placing] greater emphasis on the value of a term as a means of distinguishing one document from another than on its value as an indication of the content of the document itself

In other words, you want to combine the term frequency (this term is important to the document) with the inverse document frequency (this word is important overall). The most common way to do this (Robertson and Zaragoza, 2009, there are countless combinations with funny names like BM25); the details are important in practice but not for understanding the theory. The most common way to do this is to weight a term w in document d as number of times w is in document d relative to the total number of terms in d , and inversely by the number of documents the word w appears in N_d . The term frequency is

$$\text{tf}(w, d) = \frac{n(w, d)}{\sum_{w' \in d} n(w', d)}, \quad (4.3)$$

where the denominator sums over all of the other words in the document, and the inverse document frequency is

$$\text{idf}(w) = \log \frac{N}{N_d}, \quad (4.4)$$

where we use log to not so carefully distinguish very frequent words (“the”) from fairly frequent words (“very”); both are about as unimportant.

These took the theoretical insights of the original Cranfield experiments and made them into something that could actually work. Again, this was revolutionary: as a result the British Computing Society named their annual award for researchers in this area after Karen Spärk Jones.

4.1 Old-Fashioned AI

Given the Cranfield paradigm, we now have a setup for creating QA system. Build a collection of questions and see how well you can answer them. And this defined the evaluation of how many computers answered questions in the mid twentieth century.³ But the approaches were not just using the term weightings above to find relevant documents: could we not just find documents to help a reader answer a question but generate the answer given a question. These approaches went deeper and often used representations that required logic and reasoning to answer questions.

Arguably, the history of computer question answering begins with answering questions about the American sport of baseball. Since one of my big interests is comparing human vs. computer ability, it’s worth comparing

³ This is sometimes called “good ol’ fashioned AI” or GOFAI for short.

it to the state of the art in human question answering in the mid-twentieth century ([Sedman, 2011](#)):

Abbott: Goofy, huh? Now let's see. We have on the bags - we have Who's on first, What's on second, I Don't Know's on third.
 Costello: That's what I wanna find out.
 Abbott: I say Who's on first, What's on second, I Don't Know's on third—
 Costello: You know the fellows' names?
 Abbott: Certainly!
 Costello: Well then who's on first?
 Abbott: Yes!
 Costello: I mean the fellow's name!
 Abbott: Who!
 Costello: The guy on first!
 Abbott: Who!
 Costello: The first baseman!
 Abbott: Who!
 Costello: The guy playing first!
 Abbott: Who is on first!
 Costello: Now whaddya askin' me for?
 Abbott: I'm telling you Who is on first. ([Light, 2016](#), pages 1002–1004)

That's going to be tough to beat. Let's take a look at the `BASEBALL` system from MIT. Despite the odd phrasing, the `BASEBALL` system of Green et al. is essentially a natural language interface to an existing database ([Green et al., 1961](#)).

It does this by taking each word in the question, trying to match it to a field in its database. The challenge then is to figure out what the question is looking for, which is usually found in the question word (this is quite a bit simpler than the lexical answer type analysis of the Watson system (Chapter 6.2), where you need to figure out things like “This Argentinian Author” could be Borges, Cortazar, or Puig). For baseball, the scope is relatively narrow: a team, a location, or a number (number of runs or wins).

Like the search over large document collections the `BASEBALL` system specifies what it is looking for using what it calls a specification list. If the question is

Where did the Red Sox play on July 7?

the specification is

```
Place = ?
Team = Red Sox
Month = July
Day = 7
```

and then it finds all of the corresponding records where the Red Sox played a game on July 7 and uses that to construct an answer.

But if you know a database language like SQL, you recognize that this is essentially a bunch of SELECTs and WHEREs. This isn't even taxing the full

breadth of SQL. To go the next step, we turn to a system called LUNAR (Woods et al., 1971).

Which, as you might guess, was about the moon. It could answer much more complicated questions, such as those that require averages or averages filtered by some criterion. It can do this by converting the questions into a rich, recursive grammar. But first, it needs to analyze the grammatical structure of the input question.

For example, the question

How many samples are there?

is analyzed with the following parts of speech:

```
S Q
NP DET HOWMANY
N SAMPLE
NU PL
AUX TNS PRESENT
VP V EXIST
```

Unlike BASEBALL it can handle previous contexts (like conversational QA, Chapter 7.2).

After that syntactic analysis, it then gets turned into a logical form that can be executed against its database.

In 2022, it's probably not worth going into that. Our databases don't look like this anymore. One thing that I found interesting in rereading this in the perspective of the 21st century, is that it also has something that looks like an IR component, although it's explicitly looking for documents tagged with a particular topic, unlike doing something like tf-idf retrieval.

The last old school QA system I wanted to talk about is probably the most famous of them all: SHRDLU. SHRDLU was described in the thesis of Terry Winograd... remember that name, we'll be seeing it again in Chapter 10.18!

Unlike the LUNAR and BASEBALL which are exactly what they say on the tin, I need to explain the name a little bit more. This isn't going to be easy.

First, since you're reading a book printed⁴ in the twenty-first century, you probably don't know what a linotype system is. It was the "missing link" between moveable type presses (where you had to move the letters in by hand) and sending output from a computer keyboard to a "real" printer.

The most amazing part of this system is that it had a keyboard: when you hit a key, it would jam some metal together to create a beautiful printing piece called a slug.

But one key the machine didn't have was a backspace button: you can't undo melting metal together.

So if you screwed something up, you needed some way to signal don't use this slug. So the operator would button mash. The first two columns of buttons, spell out ETAOIN SHRDLU: a phrase that represents when technology goes

⁴ In all likelihood, you're not actually reading a printed book, but a boy can dream.

wrong...a warning that human intervention is needed to prevent garbage in, garbage out: you don't want your readers to see this, after all. Thus, that became a title of a short story by the author Fredric Brown. That story is about a Linotype machine—the pinnacle of technology at the time—that became self-aware because it understood all of the information being fed into it (Figure 4.2).



(a) Linotype Machine



(b) Linotype Science Fiction

Figure 4.2: The origin of the name "SHRDLU": the most frequent letters on a linotype machine (left, a precursor to computerized word processors) which then became the title of a short story (right) of a machine that ingested the information it printed and became intelligent.

This machine becomes voracious, demanding more and more input until the protagonist of the story selectively feeds it Buddhist philosophy and it achieves Nirvana. Despite its funny name, SHRDLU remains the OG question answering system.

The basic idea is that Instead of just answering questions about a static world, you could move objects around in this 3D world. This is a later rendering of it ... back in the early 70s computer graphics aren't as advanced as they are today.

And if the system couldn't tell what object you're asking about (e.g., it's ambiguous), it asks a follow-up question to clarify which one you meant (which is still a challenge we discuss in Chapter 11). I'm not going to go through the components of SHRDLU, but it is fairly similar to LUNAR: you parse sentences into a logical form and find what matches in the knowledge base. The big difference is that the knowledge base isn't static: it changes as the user interacts with the block world.

The reason that I'm not talking about the details is because you've probably noticed that these systems are talking about narrower and narrower domains. They can just answer questions about baseball games (not even baseball players), rocks from the moon (not rocks generally, that's too much)...SHRDLU goes further to build its own world and only answering questions about that.

And that goes back to the origin of the name SHRDLU: taking text and making it into a reality rather than understanding the world as it is. And the world as it is has ambiguity.

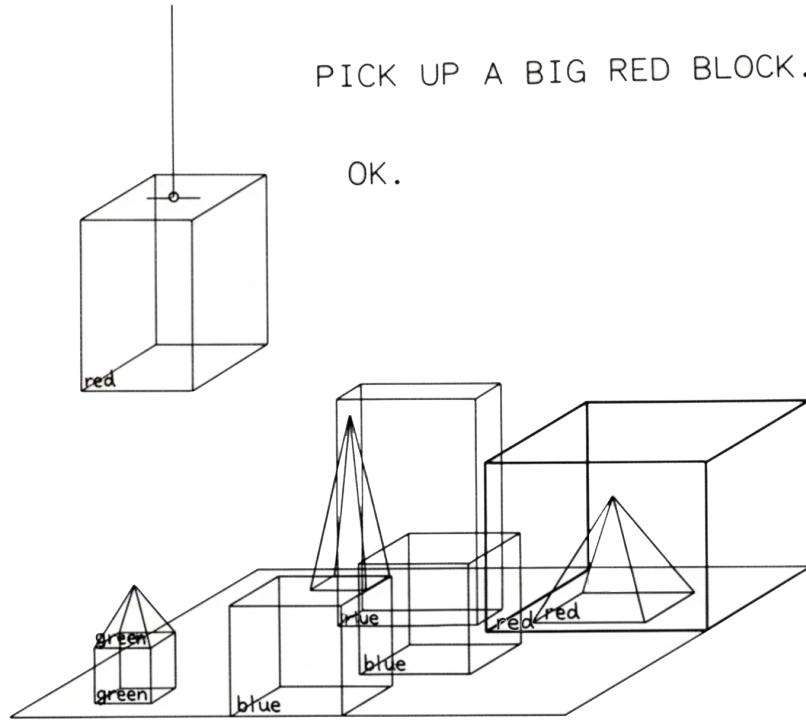


Figure 4.3: A screen shot of the SHRDLU system. Unlike other early QA systems with a static world, you can move the blocks around and ask questions about them. However, it's a very small world.

Let's go back to the "Who's on first" bit that I started this video with. So let me explain⁵ the Abbot and Costello routine: there is a baseball player that has the name "Who".

Costello cannot figure out that "who" is not an interrogative pronoun and never interprets it as a proper noun. Now, usually, that's the right way to do it. That is what all of the systems in this chapter would have done: it's hard coded into their interpretations. But natural language is funny and ambiguous sometimes.

And that's the big difference between modern QA systems and these older systems. They're small and brittle: when they work, they work really well. But they don't work that often.

They're not going to be able to handle random questions from people on the Internet. But that's not to say that logic, semantic parsing, and knowledge bases have no place in 21st century question answering systems: rather than taking a string of text and giving a single answer, there are probabilistic interpretations that provide a distribution over interpretations, and then you train approaches using the Cranfield paradigm to make sure that, yes, on average it can answer the typical user's questions.

⁵ At the risk of explaining the joke to much, there are additional aspects that showcase the difficulties in question answering. Pragmatics can be misinterpreted: "I don't know" can be a statement of fact rather than saying that you *lack* knowledge. And grammar can be critical in deciphering clues: "I want to throw the guy at first base, so I pick up the ball and throw it to Who" makes it clear that "who" is a name, while "throw it to whom" makes it clear that this is an interrogative in the objective case.

4.2 *Leading up to Google*

And that's exactly the long-lasting legacy of the Cranfield paradigm shaped how we get answers from the internet. For that, we need to leave central England and have the mantle get taken up by a new player: NIST. The American National Institute for Standards and Technology.

While the story of the birth of the Cranfield paradigm is an innovative insight with Cyril Cleverdon combined with the vision of Karen Spärk Jones, the story of NIST's development of these test collections is one of slow, methodological iteration. Building up decades of expertise and slowly building up bigger and better test collections, updating them each year so that by the time Google's PageRank ([Brin and Page, 1998](#)) came around, Google used the same evaluations internally to optimize their algorithms.

Just fire up the latest test collection, compute your precision and recall, and call it a day. And this approach toward evaluation is with us in everything we do. Test collections are how we evaluate virtual assistants, new AI systems, and just about every aspect of the modern Internet (Chapter 10). And it's not always for the best. The ubiquity of these reusable test collections leads to overfitting. TREC puts out new datasets every year or so, but that's not fast enough for machine learning in the 21st century! And while the focus on real user data is good, the data you collect is biased by who is using technology (mostly rich people from Western countries) and how they've been trained to ask questions through decades of living with Cranfield-paradigm systems. And unlike people who work on, say, human-computer interaction, those building QA systems often don't actually put their system in front of real users. And I don't think that's what Cyril Cleverdon would have wanted. His goal, as a librarian, was to help users. And yes, reusable test collections help you do that—and they've spawned multibillion dollar companies around the world—but at the end of the day you still need to check whether real users are actually finding the information they need.

5

The Manchester Paradigm: The Art of Asking the Perfect Question

While computer scientists were getting their feet wet answering questions in the wake of World War II, trivia enthusiasts were perfecting how to ask the perfect question. This chapter outlines the conventions and processes of the highest form of question answering—in the biased opinion of the author—and argues for why many of these standards could be a part of creating data for computer question answering. Previously published as Boyd-Graber and Börschinger (2020).

Topics in chapter:

1. **Two British University Towns Alike in Question Answering.**
2. **It Takes a Village to Ask a Question.** The norms of the trivia community
3. **Cranfield vs. Manchester.** The distinction between the two paradigms and why we call it that.

The QA community is obsessed with evaluation. Schools, companies, and newspapers hail new SOTAs and topping leaderboards, giving rise to troubling claims (Lipton and Steinhardt, 2019) that an “AI model tops humans” (Najberg, 2018) because it ‘won’ some leaderboard, putting “millions of jobs at risk” (Cuthbertson, 2018). But what is a leaderboard? A leaderboard is a statistic about QA accuracy that induces a ranking over participants.

Newsflash: this is the same as a trivia tournament. The trivia community has been doing this for decades (Jennings, 2006); Section 8.8 details this overlap between the qualities of a first-class QA dataset (and its requisite leaderboard). The experts running these tournaments are imperfect, but they’ve learned from

their past mistakes (see Appendix ?? for a brief historical perspective) and created a community that reliably identifies those best at question answering. Beyond the format of the *competition*, trivia norms ensure individual questions are clear, unambiguous, and reward knowledge (Section 5.2).

We are not saying that academic QA should surrender to trivia questions or the community—far from it! The trivia community does not understand the

real world information seeking needs of users or what questions challenge computers. However, they have well-tested protocols to declare that someone is better at answering questions than another. This collection of tradecraft and principles can nonetheless help the QA community.

Beyond these general concepts that QA can learn from, Section 5.3 reviews how the “gold standard” of trivia formats, QB can improve traditional QA. We then briefly discuss how research that uses fun, fair, and good trivia questions can benefit from the expertise, pedantry, and passion of the trivia community

(Section ?reference? ??).

Last year, I wrote a position paper about the AI researchers who make question answering systems can learn from trivia nerds. Given the time constraints of the conference and the need to be more “academic”, I couldn’t say everything I wanted to. So I’ll try not to be too repetitive in this video. What I would like to do is talk more about how the trivia community learned hard lessons and how the question answering community is still learning those lessons. I’m hoping by going into a little more detail about how the trivia community made mistakes … might prod computer scientists to see how they might be relevant to how we’re training and evaluating computer question answering systems.

Even if you don’t care about question answering specifically, I think this will help get ideas across about how to do effective evaluation in machine learning. Okay, so what are hard won lessons of the trivia community that the geniuses working on AI have yet to master? 1. Don’t cheat. 2. Reward knowledge 3. Write good questions. While that may seem simple, it took a while for humans to figure these things out for trivia competitions, and the artificial intelligence community is still figuring it out.

GI Bill Footage from CBS This Morning: <https://www.youtube.com/watch?v=fgtvMceoimU>

As veterans returned home from World War II, there was a huge thirst for new forms of entertainment. The Manhattan project and the space race showed the importance of knowledge and expertise, and the GI Bill turned open the taps of knowledge to an information-thirsty populace. These trends were naturally reflected in popular entertainment, as new quiz shows like Twenty One sprang up.

<https://youtu.be/JRKxNpwqBac?t=1154>

What’s the matter with 21? It was rigged, as Herb sandbagged to let his opponent win as documented in the 1994 movie Quiz Show!

<https://youtu.be/bj-m3Ddmn0E?t=83>

The handsome Charles van Doren was able to peek at the answers and did pretty well!

Now, the obvious version of cheating is something that we don’t often see in AI research. But there are subtler ways that systems after a manner cheat. Let’s think about this very abstractly. Let’s say that when you are asked some questions, you get to absorb some information that you can use next time.

Now, in the 21 cheating scandal, the process allegedly went something like this: Let's try practicing these questions! Sure ... Ooh, you got two wrong, here are the correct answers You could then compute the total length of the correct answers and that's how much information "leaked" from the test set. That's a lot!

Now, let's compare that with a more "reasonable" system that resembles what you actually see for a lot of question answering datasets. You can submit a system that answers questions (usually as a docker container to make sure it doesn't phone home), and then you get back a score with how well you did. There's no problem here, right? You don't know what the answer is, right?

Well, don't forget that computers have infinite patience ... think about it like cracking a password. You could try all possible answers for position 1, all possible answers for position 2, etc. But computers can crack more cleverly ... you have a set of answers and have some confidence on each guess. Take your lowest confidence answer, switch the answer. If the score goes up, the guess went from wrong to right. If the score goes down, the guess went from right to wrong. If the score stays the same, the guess stayed wrong ... keep hunting. But in this example, the score went up, so we know that that answer was actually correct. Nobody explicitly told us the answer was Uluru, but the effect is the same. If you do it this way, there are thousands of possibilities to try, not millions. And this is clearly in the realm of possibility for computers: even if you can't probe every answer, you can still crank up your score quite a bit (and don't forget that outcomes are correlated!).

Now, I don't think people would actually do this (or at least admit to it), but there are many black box neural network optimization toolkits that adjust the many parameters of a neural model (often on dev data). And implicitly, even on test data, you get higher numbers for some models than others ... and that's a signal. People try out many different models and just use the highest score as the legitimate representative of a system or a technique.

Jesse Dodge has done some great work on how to more fairly report experimental results without cheating: essentially, you need to have a budget for how many comparisons you make. And then you need to apply the same budget to every system fairly.

And this is why I think the human model for comparing question answering ability (i.e., a trivia tournament) is the model we should follow: you only get to run on the test data.

And, yeah, maybe the reason is that we can't wipe human memories. But I think it's more realistic, even for computers: you only have a single chance to make a first impression on a user. And this is why the model is also what happens for DARPA evaluations.

Now, bad stuff happens in this style of evaluation. We had a particularly embarrassing outing in 2016 when we lost ... very badly against some smart humans. Despite having done much better the year before.

Us losing (badly) to top quiz bowlers: <https://www.youtube.com/watch?v=c2kGD1EdffW>

Ugh, that still hurts. But that's not an argument against these style of evaluations. You should lower your embarrassment threshold and do evaluations more frequently.

And this is why I use surprise questions for my courses to see how well systems react to novel data.

But it's not just the sanctity of the test set. Let's move on to how you ask questions.

5.1 Quiz Bowl

Regular viewers of this channel will know that I am an advocate of the Quizbowl model of asking questions. But let's dig into its evolution a little bit more.

The mythology of Quizbowl says that it was a USO diversion created by Canadian Don Reid as a way of entertaining allied troops in Europe. However, in my reading it seems that a lot of the core ideas were introduced by host Allen Ludden once it became a radio program.

But the key idea is that the questions can be interrupted, meaning that part of the skill is correctly balancing your own uncertainty. Every word is an opportunity for either team to show that it knows more than the other team. I have videos arguing for why this is the right and proper way to evaluate who knows more about a topic, so I will keep this brief.

In that 2020 paper I mentioned before, Ben Boerschinger and I argued that you want question answering datasets to be maximally discriminative.

See the paper for the equations! But the moral of the story is that the Quizbowl format popularized by College Bowl created the right format for reliably determining who knows the most about a topic. Every question is discriminative. In other words, rewarding knowledge rather than buzzer speed or luck (you can see my other videos for why, for instance Jeopardy isn't a good test of human vs. computer intelligence).

But the format is a necessary but not sufficient condition: you also need to have good quality questions.

Let's now turn to the final lesson from the trivia community: writing good questions. During the early nineties, three heroes created the Academic Competition Federation. As a point of Terrapin pride, John and Ramesh were Maryland students. Together, with Carol Guthrie, they created an organization dedicated to creating questions that fully take advantage of the beauty and elegance of the Quizbowl format.

So what innovations did they introduce? First, avoiding ambiguity:

So what does a good question look like? First, you don't want to have ambiguous questions. For example, if you ask the question "What's the capital of Georgia", you answer beautiful Tblisi but the official answer is Atlanta, then you're going to be upset!

Sewon Min recently had a great paper highlighting this problem in modern

question answering datasets in 2020, but this was a big part of the diatribes Carol was making in the early nineties.

Part of ACF's style was to make sure that what they're asking is obvious as a reaction to some of the excesses of ambiguous or misleading college bowl questions. There are many dimensions of this. For example, making sure that you don't just ask "when was the Battle of Hastings" you say things like "in what year" or "day and month required". This is a problem in datasets like SQuAD and Natural Questions where it's not always clear what's required to answer the question: annotators just highlight whatever is convenient.

So why is this a problem? Machine learning algorithms have an objective function that they're trying to optimize. If they give a correct answer, they should be rewarded so that they can repeat the process on the next question. If you get the question, "What is the capital of Georgia" and you're ruled incorrect when you say "Atlanta" (after learning the opposite lesson from the previous training example), you are going to be upset. And this is true whether you're a human or a machine learning algorithm.

Similarly, you need to accept all of the possible answers to a question. Again, we see the same pattern as before: research papers are only now recapitulating the lessons learned by the trivia community 40 years ago. A great undergrad I'm working with showed that this can make machine learning question answering systems more robust. If you only accept "Timothy Donald Cook" as the correct answer for "Who is the CEO of Apple", then that makes "Tim Cook" as wrong as "Tim Apple". Again, this confuses algorithms that are trying to learn how to answer these questions. And even if you just care about climbing the leaderboards, we recently had a paper explaining how to squeeze out a few more points by paying attention to that.

<https://slate.com/culture/2019/04/jeopardy-quiz-bowl-connection-ken-jennings.html>

Another organization that should be mentioned in the promulgation of good Quizbowl is National Academic Quiz Tournaments, founded by R. Hentzel. Not only have they created a system of good, discriminative high-quality tournaments for middle school and high school students in North America, they have also served as the employer of people like Ken Jennings and Larissa Kelly (and full disclosure, me too in the early aughts) and have made Quizbowl more professional. Here's Ken Jennings extolling the virtues of Quizbowl in a NAQT promotional video.

There are many ways that this is relevant: things have different names because of language, history, pen names, nicknames, convenience, mathematical equivalence, or domain-specific rules. The new professional organizations have created systematic evaluations that, because the goal is to reward knowledge, also are the kinds of rules for judging the correctness of machine answers. Even if a user might prefer to get their answers in Fahrenheit, the computer should know that they're the same thing!

But these organizations are not just question writing organizations. Their goal is to create efficient systems for as quickly as possible determining out of

256 teams at a convention center which knows the most. In addition to the Quizbowl format, the selection of appropriately difficult questions asked at the right time to the right teams makes this possible.

Here, I don't even have a "perfect" recent paper to flash on the screen to show that, yeah, computer science knows that it's a problem 40 years to late. We've been working on using item response theory and adversarial interfaces, but there's still quite a ways to go.

Let's take how computer comparisons lag behind those for humans. While there are monthly competitions to see which human can best answer questions, computers are reusing stale data from 2018 to see who's best. And those questions are written by random people on the internet: either crowdworkers or clueless people looking for answers in a search engine. In contrast, domain experts are writing questions lovingly edited by people like Ken Jennings and Larissa Kelly when people go to NAQT or ACF competitions.

And while things like Dynabench are trying to create adversarial datasets (or our own TrickMe interface), these datasets are tiny compared to the tens of thousands of questions written every year for human trivia competitions. I, of course, think that this is a possibility for synergies between the two communities.

Computers can help find patterns and reduce drudgery: make human trivia more exciting and efficient. And the trivia experts, with their depth of knowledge, can help computers understand more of the world. Likewise, the best practices of the hard won victories of the last seventy years of trivia can help machine learning researchers avoid the same mistakes.

And, most importantly, I think it can be fun for everybody, especially us researchers.

5.2 *The Craft of Question Writing*

Trivia enthusiasts agree that questions need to be well written (despite other disagreements). Asking "good questions" requires sophisticated pragmatic reasoning ([Hawkins et al., 2015](#)), and pedagogy explicitly acknowledges the complexity of writing effective questions for assessing student performance ([Haladyna, 2004](#), focusing on multiple choice questions).

QA datasets, however, are often collected from the wild or written by untrained crowdworkers. Crowdworkers lack experience in crafting questions and may introduce idiosyncrasies that shortcut machine learning ([Geva et al., 2019](#)). Similarly, data collected from the wild such as Natural Questions ([Kwiatkowski et al., 2019](#)) or AmazonQA ([Gupta et al., 2019](#)) by design have vast variations in quality. In the previous section, we focused on how datasets as a whole should be structured. Now, we focus on how specific *questions* should be structured to make the dataset as valuable as possible.

Avoiding ambiguity and assumptions

Ambiguity in questions not only frustrates answerers who resolve the ambiguity ‘incorrectly’. Ambiguity also frustrates the goal of using questions to assess knowledge. Thus, the US Department of Transportation explicitly bans ambiguous questions from exams for flight instructors (Flight Standards Service, 2008); and the trivia community has likewise developed rules and norms that prevent ambiguity. While this is true in many contexts, examples are rife in format called QB (Boyd-Graber et al., 2012), whose very long questions¹ showcase trivia writers’ tactics. For example, QB author Zhu Ying (writing for the 2005 PARFAIT tournament) asks participants to identify a fictional character while warning against possible confusion [emphasis added]:

He’s **not Sherlock Holmes**, but his address is 221B. He’s **not the Janitor on Scrubs**, but his father is played by R. Lee Ermey. [...] For ten points, name this misanthropic, crippled, Vicodin-dependent central character of a FOX medical drama.

ANSWER: Gregory House, MD

In contrast, QA datasets often contain ambiguous and under-specified questions. While this sometimes reflects real world complexities such as actual under-specified or ill-formed search queries (Faruqui and Das, 2018; Kwiatkowski et al., 2019), ignoring this ambiguity is problematic. As a concrete example, Natural Questions (Kwiatkowski et al., 2019) answers “what year did the US hockey team win the Olympics” with 1960 and 1980, ignoring the US women’s team, which won in 1998 and 2018, and further assuming the query is about *ice* rather than *field* hockey (also an Olympic event). Natural Questions associates a page about the United States men’s national ice hockey team, arbitrarily removing the ambiguity *post hoc*. However, this does not resolve the ambiguity, which persists in the original question: information retrieval arbitrarily provides one of many interpretations. True to their name, Natural Questions are often under-specified when users ask a question online.

The problem is neither that such questions exist nor that machine reading QA considers questions given an associated context. The problem is that tasks do not explicitly acknowledge the original ambiguity and gloss over the implicit assumptions in the data. This introduces potential noise and bias (i.e., giving a bonus to systems that make the same assumptions as the dataset) in leaderboard rankings. At best, these will become part of the measurement error of datasets (no dataset is perfect). At worst, they will recapitulate the biases that went into the creation of the datasets. Then, the community will implicitly equate the biases with correctness: you get high scores if you adopt this set of assumptions. These enter into real-world systems, further perpetuating the bias. Playtesting can reveal these issues (Section 8.8), as implicit assumptions can rob a player of correctly answered questions. If you wanted to answer 2014 to “when did Michigan last win the championship”—when the Michigan State Spartans won the Women’s Cross Country championship—and you

¹ Like *Jeopardy!*, they are not syntactically questions but still are designed to elicit knowledge-based responses; for consistency, we still call them questions.

cannot because you chose the wrong school, the wrong sport, and the wrong gender, you would complain as a player; researchers instead discover latent assumptions that creep into the data.²

It is worth emphasizing that this is not a purely hypothetical problem. For example, Open Domain Retrieval Question Answering (Lee et al., 2019) deliberately avoids providing a reference context for the question in its framing but, in re-purposing data such as Natural Questions, opaquely relies on it for the gold answers.

Avoiding superficial evaluations

A related issue is that, in the words of Voorhees and Tice (2000), “there is no such thing as a question with an obvious answer”. As a consequence, trivia question authors delineate acceptable and unacceptable answers.

For example, in writing for the trivia tournament Harvard Fall XI, Robert Chu uses a mental model of an answerer to explicitly delineate the range of acceptable correct answers:

In Newtonian gravity, this quantity satisfies Poisson’s equation. [...] For a dipole, this quantity is given by negative the dipole moment dotted with the electric field. [...] For 10 points, name this form of energy contrasted with kinetic.

ANSWER: potential energy (*prompt on energy; accept specific types like electrical potential energy or gravitational potential energy; do not accept or prompt on just “potential”*)

Likewise, the style guides for writing questions stipulate that you must give the answer type clearly and early on. These mentions specify whether you want a book, a collection, a movement, etc. It also signals the level of specificity requested. For example, a question about a date must state “day and month required” (September 11, “month and year required” (April 1968), or “day, month, and year required” (September 1, 1939). This is true for other answers as well: city and team, party and country, or more generally “two answers required”. Despite these conventions, no pre-defined set of answers is perfect, and every worthwhile trivia competition has a process for adjudicating answers.

In high school and college national competitions and game shows, if low-level staff cannot resolve the issue by throwing out a single question or accepting minor variations (America instead of USA), the low-level staff contacts the tournament director. The tournament director, who has a deeper knowledge of rules and questions, often decide the issue. If not, the protest goes through an adjudication process designed to minimize bias:³ write the summary of the dispute, get all parties to agree to the summary, and then hand the decision off to mutually agreed experts from the tournament’s phone tree. The substance of the disagreement is communicated (without identities), and the experts apply the rules and decide.

Consider what happened when a particularly inept *Jeopardy!* contestant⁴

² Where to draw the line is a matter of judgment; computers—which lack common sense—might find questions ambiguous where humans would not.

³ <https://www.naqt.com/rules/#protest>

⁴ http://www.j-archive.com/showgame.php?game_id=6112

did not answer laproscope to “Your surgeon could choose to take a look inside you with this type of fiber-optic instrument”. Since the van Doren scandal (Freedman, 1997), every television trivia contestant has an advocate assigned from an auditing company. In this case, the advocate initiated a process that went to a panel of judges who then ruled that endoscope (a more general term) was also correct.

The need for a similar process seems to have been well-recognized in the earliest days of QA system bake-offs such as TREC-QA, and Voorhees (2008) notes that

[d]ifferent QA runs very seldom return exactly the same [answer], and it is quite difficult to determine automatically whether the difference [...] is significant.

In stark contrast to this, QA datasets typically only provide a single string or, if one is lucky, several strings. A correct answer means *exactly* matching these strings or at least having a high token overlap F_1 , and failure to agree with the pre-recorded admissible answers will put you at an uncontestable disadvantage on the leaderboard (Section 8.8).

To illustrate how current evaluations fall short of meaningful discrimination, we qualitatively analyze two near-SOTA systems on SQuAD V1.1: THE ORIGINAL XLNET (YANG ET AL., 2019) AND A SUBSEQUENT ITERATION CALLED XLNET-123.⁵

Despite XLNET-123’s margin of almost four absolute F_1 (94 vs 98) on development data, a manual inspection of a sample of 100 of XLNET-123’s wins indicate that around two-thirds are ‘spurious’: 56% are likely to be considered not only equally good but essentially identical; 7% are cases where the answer set omits a correct alternative; and 5% of cases are ‘bad’ questions.⁶

Our goal is not to dwell on the exact proportions, to minimize the achievements of these strong systems, or to minimize the usefulness of quantitative evaluations. We merely want to raise the limitation of blind automation for distinguishing between systems on a leaderboard.

Taking our cue from the trivia community, we present an alternative for MRQA. Blind test sets are created for a specific time; all systems are submitted simultaneously. Then, all questions and answers are revealed. System authors can protest correctness rulings on questions, directly addressing the issues above. After agreement is reached, quantitative metrics are computed for comparison purposes—despite their inherent limitations they at least can be trusted. Adopting this for MRQA would require creating a new, smaller test set every year. However, this would gradually refine the annotations and process.

This suggestion is not novel: Voorhees and Tice (2000) accept automatic evaluations “for experiments internal to an organization where the benefits of a reusable test collection are most significant (and the limitations are likely to be understood)” (our emphasis) but

⁵ We could not find a paper describing XLNET-123, the submission is by <http://tia.today>.

⁶ Examples in Appendix ?reference??.

THAT “SATISFACTORY TECHNIQUES FOR [AUTOMATICALLY] EVALUATING NEW RUNS” HAVE NOT BEEN FOUND YET. WE ARE NOT AWARE OF ANY CHANGE ON THIS FRONT—IF ANYTHING, WE SEEM TO HAVE BECOME MORE INSENSITIVE AS A COMMUNITY TO JUST HOW LIMITED OUR CURRENT EVALUATIONS ARE.

Focus on the bubble

WHILE EVERY QUESTION SHOULD BE PERFECT, TIME AND RESOURCES ARE LIMITED. THUS, AUTHORS AND EDITORS OF TOURNAMENTS “FOCUS ON THE BUBBLE”, WHERE THE “BUBBLE” ARE THE QUESTIONS MOST LIKELY TO DISCRIMINATE BETWEEN TOP TEAMS AT THE TOURNAMENT. THESE QUESTIONS ARE THOROUGHLY PLAYTESTED, VETTED, AND EDITED. ONLY AFTER THESE QUESTIONS HAVE BEEN PERFECTED WILL THE OTHER QUESTIONS UNDERGO THE SAME LEVEL OF POLISH.

FOR COMPUTERS, THE SAME LOGIC APPLIES. AUTHORS SHOULD ENSURE THAT THESE DISCRIMINATIVE QUESTIONS ARE CORRECT, FREE OF AMBIGUITY, AND UNIMPEACHABLE. HOWEVER, AS FAR AS WE CAN TELL, THE AUTHORS OF QA DATASETS DO NOT GIVE ANY SPECIAL ATTENTION TO THESE QUESTIONS.

UNLIKE A HUMAN TRIVIA TOURNAMENT, HOWEVER—with finite patience of the participants—this does not mean that you should necessarily remove all of the easy or hard questions from your dataset. This could inadvertently lead to systems unable to answer simple questions like “Who is buried in Grant’s Tomb?” (Dwan, 2000, chapter 7). Instead, focus more resources on the bubble.

5.3 Why QB is the Gold Standard

WE NOW FOCUS OUR THUS FAR WIDE-RANGING QA DISCUSSION TO A SPECIFIC FORMAT: QB, WHICH HAS MANY OF THE DESIRABLE PROPERTIES OUTLINED ABOVE. WE HAVE NO DELUSION THAT MAINSTREAM QA WILL UNIVERSALLY ADOPT THIS FORMAT (INDEED, A MONOCULTURE WOULD BE BAD). HOWEVER, GIVEN THE COMMUNITY’S EMPHASIS ON FAIR EVALUATION, COMPUTER QA CAN BORROW ASPECTS FROM THE GOLD STANDARD OF HUMAN QA.

WE HAVE SHOWN EXAMPLE OF QB QUESTIONS, BUT WE HAVE NOT EXPLAINED HOW THE FORMAT WORKS; SEE RODRIGUEZ ET AL. (2019) FOR MORE. YOU MIGHT BE SCARED OFF BY HOW LONG THE QUESTIONS ARE. HOWEVER, IN REAL QB TRIVIA TOURNAMENTS, THEY ARE NOT FINISHED BECAUSE THE QUESTIONS ARE *INTERRUPTIBLE*.

Interruptible A MODERATOR READS A QUESTION. ONCE SOMEONE KNOWS THE ANSWER, THEY USE A SIGNALING DEVICE TO “BUZZ IN”. IF THE PLAYER WHO BUZZED IS RIGHT, THEY GET POINTS. OTHERWISE, THEY LOSE POINTS AND THE QUESTION CONTINUES FOR THE OTHER TEAM.

NOT ALL TRIVIA GAMES WITH BUZZERS HAVE THIS PROPERTY, HOWEVER. FOR EXAMPLE, TAKE *JEOPARDY!*, THE SUBJECT OF WATSON’S TOUR DE FORCE (FERRUCCI

ET AL., 2010). WHILE *JEOPARDY!* ALSO USES SIGNALING DEVICES, THESE ONLY WORK ONCE THE QUESTION HAS BEEN READ IN ITS ENTIRETY; KEN JENNINGS, ONE OF THE TOP *JEOPARDY!* PLAYERS (AND ALSO A QBER) EXPLAINS IT ON A *PLANET MONEY* INTERVIEW (MALONE, 2019):

Jennings: THE BUZZER IS NOT LIVE UNTIL ALEX FINISHES READING THE QUESTION. AND IF YOU BUZZ IN BEFORE YOUR BUZZER GOES LIVE, YOU ACTUALLY LOCK YOURSELF OUT FOR A FRACTION OF A SECOND. SO THE BIG MISTAKE ON THE SHOW IS PEOPLE WHO ARE ALL ADRENALIZED AND ARE BUZZING TOO QUICKLY, TOO EAGERLY.

Malone: OK. TO SOME DEGREE, *JEOPARDY!* IS KIND OF A VIDEO GAME, AND A CRAPPY VIDEO GAME WHERE IT'S, LIKE, LIGHT GOES ON, PRESS BUTTON—THAT'S IT.

Jennings: (LAUGHTER) YEAH.

JEOPARDY!'S BUZZERS ARE A GIMMICK TO ENSURE GOOD TELEVISION; HOWEVER, QB BUZZERS DISCRIMINATE KNOWLEDGE (SECTION 8.8). SIMILARLY, WHILE TRIVIAQA (JOSHI ET AL., 2017) IS WRITTEN BY KNOWLEDGEABLE WRITERS, THE QUESTIONS ARE NOT PYRAMIDAL.

Pyramidal RECALL THAT EFFECTIVE DATASETS DISCRIMINATE THE BEST FROM THE REST—THE HIGHER THE PROPORTION OF EFFECTIVE QUESTIONS ρ , THE BETTER. QB'S ρ IS NEARLY 1.0 BECAUSE DISCRIMINATION HAPPENS WITHIN A QUESTION: AFTER EVERY WORD, AN ANSWERER MUST DECIDE IF THEY KNOW ENOUGH TO ANSWER. QB QUESTIONS ARE ARRANGED SO THAT QUESTIONS ARE MAXIMALLY PYRAMIDAL: QUESTIONS BEGIN WITH HARD CLUES—ONES THAT REQUIRE DEEP UNDERSTANDING—TO MORE ACCESSIBLE CLUES THAT ARE WELL KNOWN.

Well-Edited QB QUESTIONS ARE CREATED IN PHASES. FIRST, THE AUTHOR SELECTS THE ANSWER AND ASSEMBLES (PYRAMIDAL) CLUES. A SUBJECT EDITOR THEN REMOVES AMBIGUITY, ADJUSTS ACCEPTABLE ANSWERS, AND TWEAKS CLUES TO OPTIMIZE DISCRIMINATION. FINALLY, A PACKETIZER ENSURES THE OVERALL SET IS DIVERSE, HAS UNIFORM DIFFICULTY, AND IS WITHOUT REPEATS.

Unnatural TRIVIA QUESTIONS ARE FAKE: THE ASKER ALREADY KNOWS THE ANSWER. BUT THEY'RE NO MORE FAKE THAN A COURSE'S FINAL EXAM, WHICH—LIKE LEADERBOARDS—ARE DESIGNED TO TEST KNOWLEDGE.

EXPERTS KNOW WHEN QUESTIONS ARE AMBIGUOUS (SECTION 5.2); WHILE “WHAT PLAY HAS A CHARACTER WHOSE FATHER IS DEAD” COULD BE *HAMLET*, *ANTIGONE*, OR *PROOF*, A GOOD WRITER'S KNOWLEDGE AVOIDS THE AMBIGUITY. WHEN AUTHORS OMIT THESE CUES, THE QUESTION IS DERIDED AS A “HOSE” (ELTINGE, 2013), WHICH ROBS THE TOURNAMENT OF FUN (SECTION 8.8).

ONE OF THE BENEFITS OF CONTRIVED FORMATS IS A FOCUS ON SPECIFIC PHENOMENA. DUA ET AL. (2019) EXCLUDE QUESTIONS AN EXISTING MRQA SYSTEM COULD ANSWER TO FOCUS ON CHALLENGING QUANTITATIVE REASONING. ONE OF THE TRIVIA EXPERTS CONSULTED IN WALLACE ET AL. (2019A) CRAFTED A QUESTION THAT TRIPPED UP NEURAL QA BY EMBEDDING THE PHRASE “THIS

AUTHOR OPENS *CRIME AND PUNISHMENT*" INTO A QUESTION; THE TOP SYSTEM CONFIDENTLY ANSWERS FYODOR DOSTOYEVSKI. HOWEVER, THAT PHRASE WAS IN A LONGER QUESTION "THE NARRATOR IN *COGWHEELS* BY THIS AUTHOR OPENS *CRIME AND PUNISHMENT* TO FIND IT HAS BECOME *THE BROTHERS KARAMAZOV*". AGAIN, THIS SHOWS THE INVENTIVENESS AND LINGUISTIC DEXTERITY OF THE TRIVIA COMMUNITY.

A COUNTERARGUMENT IS THAT REAL-LIFE QUESTIONS—E.G., ON YAHOO! QUESTIONS (SZPEKTOR AND DROR, 2013), QUORA (IYER ET AL., 2017) OR WEB SEARCH (KWIATKOWSKI ET AL., 2019)—IGNORE THE CRAFT OF QUESTION WRITING. REAL HUMANS REACT TO UNCLEAR QUESTIONS WITH CONFUSION OR DIVERGENT ANSWERS, EXPLICITLY ANSWERING WITH HOW THEY INTERPRETED THE ORIGINAL QUESTION ("I ASSUME YOU MEANT...").

GIVEN REAL WORLD APPLICATIONS WILL HAVE TO DEAL WITH THE INHERENT NOISE AND AMBIGUITY OF UNCLEAR QUESTIONS, OUR SYSTEMS MUST COPE WITH IT. HOWEVER, ADDRESSING THE REAL WORLD CANNOT HAPPEN BY GLOSSING OVER ITS COMPLEXITY.

Complicated QB IS MORE COMPLEX THAN OTHER DATASETS. UNLIKE OTHER DATASETS WHERE YOU JUST NEED TO DECIDE WHAT TO ANSWER, IN QB YOU ALSO NEED TO CHOOSE WHEN TO ANSWER THE QUESTION.⁷ WHILE THIS IMPROVES THE DATASET'S DISCRIMINATION, IT CAN HURT POPULARITY BECAUSE YOU CANNOT COPY/PASTE CODE FROM OTHER QA TASKS. THE CUMBERSOME PYRAMIDAL STRUCTURE COMPLICATES⁸ SOME QUESTIONS (E.G., WHAT IS LOG BASE FOUR OF SIXTY-FOUR).

⁷This complex methodology can be an advantage. The underlying mechanisms of systems that can play QB (e.g., reinforcement learning) share properties with other tasks, such as simultaneous translation (Grissom II et al., 2014; Ma et al., 2019), human incremental processing (Levy et al., 2008; Levy, 2011), and opponent modeling (He et al., 2016).

⁸But does not necessarily preclude, as the Illinois High School Scholastic Bowl Coaches Association shows:

This is the smallest counting number which is the radius of a sphere whose volume is an integer multiple of π . It is also the number of distinct real solutions to the equation $x^7 - 19x^5 = 0$. This number also gives the ratio between the volumes of a cylinder and a cone with the same heights and radii. Give this number equal to the log base four of sixty-four.

Part II

Question Answering Present

6

WATSON ON JEOPARDY!:

UNQUESTIONED ANSWERS FROM IBM'S TOUR DE FORCE

IT'S BEEN NEARLY A DECADE SINCE IBM WATSON CRUSHED TWO PUNY HUMANS ON JEOPARDY! SOME PEOPLE TOOK THAT TO MEAN THAT COMPUTERS WERE DEFINITIVELY BETTER THAN HUMANS AT TRIVIA. BUT THAT ISN'T THE COMPLETE ANSWER—THIS CHAPTER, INSPIRED BY JEOPARDY!'S GIMMICK OF RESPONDING TO ANSWERS TO QUESTIONS, QUESTIONS SOME OF THE “ANSWERS” THAT EMERGED FROM IBM'S TOUR DE FORCE.

WE BEGIN THE “MODERN” AGE OF COMPUTER QUESTION ANSWERING AND THE RISE OF AI WITH *WATSON*: AN IBM-BUILT AI THAT COULD PLAY THE AMERICAN GAME SHOW *JEOPARDY!*

IT'S BEEN OVER A DECADE SINCE *WATSON* APPEARED ON TV, BUT IT WAS REVOLUTIONARY. THIS CHAPTER WILL TALK ABOUT HOW IT CHANGED MY LIFE, HOW IT CHANGED THE WAY THE PUBLIC THOUGHT ABOUT AI, AND HOW ITS TECHNOLOGY PAVED THE WAY FOR “MODERN” AI AS EXEMPLIFIED BY LLMS LIKE GPT.

BUT THIS ISN'T A LOVE LETTER TO *WATSON*. FOR ALL THE FAME AND GLORY IT WON, THE *WATSON* STORY ISN'T WITHOUT PROBLEMS. INDEED, THIS CHAPTER WILL MAKE THE ARGUMENT THAT THE GAME WAS SUBTLY RIGGED IN A WAY THAT SET BACK RESEARCH ON AI.

6.1 Why IBM Chose Jeopardy! for a Grand Challenge

WATSON IS PART OF A LONG HISTORY OF “GRAND CHALLENGE” PROJECTS MEANT TO SYMBOLIZE PROGRESS ON AI. IBM (INTERNATIONAL BUSINESS MACHINES) IS A NEW YORK-BASED TECHNOLOGY COMPANY THAT FOR MUCH OF THE TWENTIETH CENTURY WAS SYNONYMOUS WITH COMPUTING.

Grand Challenges

IT HAD PREVIOUSLY WOWED THE WORLD BY CREATING “DEEP BLUE”, A CHESS-PLAYING COMPUTER THAT DEFEATED GARY KASPAROV IN 1997, WIDELY CON-

SIDERED THE BEST CHESS PLAYER OF THE TIME ([HSU, 2002](#)). OTHER AI GRAND CHALLENGES INCLUDE THOSE ISSUED BY FUNDING AGENCIES LIKE THE UNITED STATES' DEFENSE ADVANCE PROJECECTS RESEARCH ADMINISTRATION (DARPA) TO BUILD AUTONOMOUS CARS THAT CAN DRIVE A HUNDRED MILES IN THE MOJAVE, ALTHOUGH THE BEST TEAM JUST MANAGED OVER SEVEN MILES ([PATTERSON, 2005](#)). GOOGLE/DEEPMIND HAD A SIMILAR AMBITION WHEN THEY CREATED ALPHAGO, WHICH LIKE DEEPBLUE DEFEATED LEE SEDOL IN 2016, WIDELY CONSIDERED THE BEST GO PLAYER OF THE TIME ([KOCHE, 2016](#)).

THE GOAL OF THESE PROJECTS IS NOT JUST TO ADVANCE TECHNOLOGY BUT TO CREATE A BIG ENOUCH SPLASH THAT PEOPLE CHANGE HOW THEY THINK ABOUT TECHNOLOGY.

IT'S HARD TO IMAGINE AN ALTERNATE WORLD WHERE THESE GRAND CHALLENGES DID NOT HAPPEN, BUT I DON'T THINK IT'S EXAGGERATION TO SAY THAT THESE GRAND CHALLENGES DID INDEED CHANGE THE WORLD. AFTER HIS DEFEAT, GARY KASPAROV TOOK A "IF YOU CAN'T BEAT THEM, JOIN THEM" APPROACH AND THEN BEGAN ARGUING FOR CENTAUR CHESS: HUMANS AND COMPUTERS PLAYING CHESS TOGETHER (THIS WORKS FOR QUESTION ANSWERING TOO, AS WE DISCUSS IN CHAPTER [?REFERENCE?](#) [??](#)). LIKEWISE, THE BEST GO PLAYERS BECAME MORE CREATIVE AND NOVEL MOVES BY INCORPORATING COMPUTER-LIKE PLAY INTO THEIR PLAY STYLE [?](#). AND TWO DECADES AFTER DARPA'S AUTONOMOUS GRAND CHALLENGE, WE NOW HAVE SELF-DRIVING CARS ON THE STREET IN SEVERAL US CITIES.

SO WHAT MADE IBM PICK *JEOPARDY!* FOR ITS GRAND CHALLENGE?

What Jeopardy! is and How it Works

JEOPARDY! IS A GAMESHOW CREATED BY MERV GRIFFIN THAT DEBUTED IN 1967 ([GRIFFIN, 2003](#)). ITS BIG GIMMICK IS THAT THE PLAYER RESPONSES ARE GIVEN IN THE FORM OF A QUESTION IN REACTION TO INFAMOUS HUMAN CHEATING SCANDALS (WE DISCUSS HOW COMPUTERS "CHEAT" IN CHAPTER [10](#)). FOR EXAMPLE, THE CLUE

THE CAPTCHA TEST AGAINST SPAM & ROBOT PROGRAMS IS CALLED THE 'REVERSE' TEST NAMED FOR THIS BRITISH CODE BREAKER

WOULD HAVE RESPONSE WHO IS ALAN TURING. THE CLUES ARE ARRANGED IN A GRID: COLUMNS REPRESENT CATEGORIES AND ROWS REPRESENT DIFFICULTY, WITH THE MORE DIFFICULT QUESTIONS BEING WORTH MORE.

THERE ARE THREE PLAYERS WHO STAND SIDE BY SIDE BEHIND PODIUMS. WHEN A CLUE IS READ, ANY OF THE THREE PLAYERS CAN "BUZZ IN" TO SAY THAT THEY WANT TO GIVE A RESPONSE. IF THEY GIVE THE CORRECT RESPONSE, THEY THEN HAVE "CONTROL OF THE BOARD" AND CAN SELECT THE NEXT CLUE.

ONE ADVANTAGE OF CONTROLLING THE BOARD IS THAT SOME QUESTIONS ARE CALLED "DAILY DOUBLES" WHICH ALLOW THE PLAYER TO POTENTIALLY

DOUBLE THEIR SCORE: A PLAYER CAN WAGER ANY PART OF THEIR SCORE, AND IF THEY GET IT RIGHT THEY GET THAT AMMOUNT ADDED TO THEIR SCORE (BUT A WRONG RESPONSE WILL SUBTRACT IT FROM THEIR SCORE).

How Watson changed my life

I REMEMBER WHEN I FIRST HEARD RUMBLINGS OF WATSON. BECAUSE I HAD A FOOT IN BOTH THE AI AND TRIVIA COMMUNITIES, I HEARD TWO DIFFERENT STORIES. I HEARD RUMORS OF AMAZING WORK IN PARSING AND SEMANTIC ROLE LABELING HAPPENING FROM RESEARCHERS WHO VENTURED NORTH TO WESTCHESTER COUNTY IN NEW YORK (I WAS DOING MY PhD IN NEW JERSEY). FROM THE TRIVIA COMMUNITY, I HEARD OF SOME PEOPLE WHO WERE BEING PAID BY IBM TO PLAY TRIVIA GAMES BUT THAT THEY COULDN'T SAY ANYTHING MORE.

I WAS VERY SCEPTICAL. BY THE TIME THAT WATSON CAME TO FRUITION, I HAD MOVED TO THE UNIVERSITY OF MARYLAND. THEN, MY SCEPTICISM TURNED TO JEALOUSY. I WATCHED, ALONG WITH THE REST OF THE WORLD, ONE OF THE GREATEST ACHIEVEMENTS OF AI UNFOLD IN FRONT OF ME. WHAT WAS I DOING WASTING MY TIME WORKING ON LANGUAGE MODELS IF THIS WAS ALSO LEGITIMATE RESEARCH?

LET ME BE CLEAR THAT THE TECHNICAL TRIUMPHS ARE INDISPUTABLE (AND, IN MY OPINION, UNDER-APPRECIATED). FROM THE WORK ON WAGERING TO SYNTHESIZING MULTIPLE INFORMATION SOURCES, WATSON ([FERRUCCI ET AL., 2010](#)) WAS FROM TOP TO BOTTOM A TOP-NOTCH WELL-OILED MACHINE. AND IT COMPUTED ALL OF THIS IN REAL TIME—SOMETHING THAT WASN'T STRICTLY NECESSARY BUT STILL IMPRESSIVE.

6.2 How Watson Works

WHILE “NEURAL” QUESTION ANSWERING (WHICH WE’LL DISCUSS IN FUTURE CHAPTER) IS THE BIG THING THESE DAYS, WATSON CAME OF AGE IN THE STATISTICAL AGE. TO UNDERSTAND SOME OF HOW WATSON WORKS, ITS HELPFUL TO REVIEW SOME OF THE WORK THAT CAME BEFORE WATSON THAT HELPED INSPIRE IT.

Rule-based Question Answering

PREECEDING THE STATISTICAL AGE OF QA WAS THE RULE-BASED AGE: SYSTEMS THAT WERE IMPRESSIVE ON INDIVIDUAL QUESTIONS ABOUT BASEBALL OR GEOLGY BUT THAT FALTERED AS SOON AS IT SAW A QUESTION THAT WAS UNEXPECTED IN TERMS OF THE DOMAIN, THEY'D FAIL MISERABLY.

ONE OF THE GUIDING PRINCIPLES OF THE STATISTICAL AGE WAS “THE UNREASONABLE EFFECTIVENESS OF BIG DATA”, AND THIS REQUIRED SCALING APPROACHES TO WEB-SCALE DATA. THE FIRST THINGS THAT PEOPLE TRIED WAS SCALING UP THE “GOOD, OLD-FASHIONED” APPROACHES THAT DEFINED THE

RULE-BASED SYSTEMS. PROBABLY THE MOST PROMINENT SYSTEM IN THIS CATEGORY IS START FROM BORIS KATZ AT MIT. THIS SYSTEM FIRST LAUNCHED IN 1993. THIS, LIKE MANY OF THE SYSTEMS IN THIS ERA, TAKE AN APPROACH THAT'S SOMEWHAT SIMILAR TO THE OLD-FASHIONED AI APPROACHES: PARSE THE QUERY WITH A AND LOOK IT UP IN A KNOWLEDGE BASE. FOR EXAMPLE, GIVEN THE QUESTION:¹

WHO DIRECTED GONE WITH THE WIND?

BECOMES THE LOOKUP

```
(get "imdb-movie" "Gone with the Wind (1939)" "DIRECTOR")
=> ("George Cukor" "Victor Fleming" "Sam Wood") (KATZ
ET AL., 2002).
```

SOME OF THESE CAN BE LOOKED UP DIRECTLY IN DATASETS, BUT START FURTHER

BUILDS ON THE SYSTEMS LIKE LUNAR AND BASEBALL (CHAPTER ??). **?reference?** ??.

WHILE THOSE SYSTEMS LUNAR AND BASEBALL ALSO NEED TO TURN THE RAW TEXT OF QUESTIONS INTO A STRUCTURED REPRESENTATION, THE DIFFERENCE OF START IS THAT IS TO DEAL WITH THE MESSINESS OF INTERNET TEXT. AND ITS ABILITY TO SEARCH THE INTERNET TO FIND ANSWERS IS CLEARLY A HUGE INFLUENCE ON THE SUBJECT OF THIS CHAPTER—WATSON—AND LATER DENSE PASSAGE RETRIEVAL LATER (CHAPTER 9.1).

Transforming Questions to find the Answer

IN THE NEXT CHAPTERS ABOUT NEURAL LANGUAGE MODELS, WE WILL DISCUSS HOW NEW MODELS TRANSFORM QUESTIONS INTO NUMBERS (VECTORS) TO FIND OR GENERATE THE ANSWER. BUT BEFORE WE GET THERE, IT IS USEFUL TO TALK ABOUT HOW WE CAN TRANSFORM QUESTIONS INTO TEXT THAT CAN BETTER FIND THE ANSWER. WE BEGIN WITH THE QUESTION ANSWERING USING STATISTICAL MODELS (RADEV ET AL., 2001, QASM).

ONE WAY OF THINKING ABOUT QUESTION ANSWERING IS THAT IT'S A TRANSLATION PROBLEM: QUESTIONS ARE ASKED IN ONE LANGUAGE BUT THE ANSWERS ARE FOUND IN ANOTHER LANGUAGE. QASM'S APPROACH IS TO EXPLICITLY “UNDO” THE PROCESS TO FIND THE CORRECT RESPONSE GIVEN AN INPUT CLUE.

IN CHAPTER ??, WE'LL TALK ABOUT THE CHEATING SCANDAL AROUND THE AMERICAN GAME SHOW 21, WHERE CHARLES VAN DOREN GOT THE ANSWERS IN ADVANCE. PART OF THE MYTHOLOGY OF JEOPARDY! IS THAT MERV GRIFFEN SAID, WHY DON'T WE JUST GIVE THEM THE ANSWERS! SO YOU HAVE CLUES LIKE

DURING WWII THIS COMPUTER SCIENTIST & CODE
BREAKER CONVERTED HIS MONEY INTO SILVER &
BURIED IT; HE NEVER FOUND HIS BURIED TREASURE.

THE CORRECT RESPONSE, WHICH IS PHRASED AS A QUESTION, IS “WHO IS ALAN TURING?” IN OTHER WORDS, THE CENTRAL CONCEIT OF JEOPARDY! IS THAT IT

¹ This is an impressive/complicated question because David O. Selznick removed the first director (George Cukor) to replace him with Victor Fleming (the credited director), but Sam Wood had to fill in when Fleming collapsed on set.

CONVERTS QUESTIONS INTO ANSWERS (AND THEN THE CONTESTANTS NEED TO GIVE A RESPONSE TO THAT ANSWER IN THE FORM OF A QUESTION).

WHAT QASM IS DOING IS DOING THE REVERSE: IT HAS A MODEL THAT TRIES TO TURN QUESTIONS INTO THE SEARCH STRING THAT COULD FIND THE ANSWER TO THE QUESTION (WE'LL SEE MODERN VERSIONS OF THIS WHEN WE TALK ABOUT MULTIHOP QUESTION ANSWERING IN CHAPTER 7). QASM CAN APPLY PLENTY OF TRANSFORMATIONS: SWAPPING WORDS, DELETING WORDS, IGNORING WORDS (E.G., IF YOU HAVE CLEVELAND, OHIO AND IGNORE OHIO, YOU'LL FIND MORE THINGS ABOUT US PRESIDENT GROVER CLEVELAND), ETC.

IN THE END, YOU HAVE A MODEL THAT CREATES SPECIALIZED QUERIES FROM THE ORIGINAL QUESTION. LATER ON, WHEN WE TALK ABOUT DENSE PASSAGE RETRIEVAL, THOSE SYSTEMS ARE DOING SOMETHING SIMILAR, EXCEPT THEY'RE GENERATING THE QUERY IN A CONTINUOUS VECTOR SPACE (CHAPTER 9.1). DOING IT WITH REAL WORDS LIKE QASM IS MORE INTERPRETABLE: A HUMAN COULD UNDERSTAND WHAT'S GOING ON, WHILE A VECTOR QUERY IS INSCRUTABLE IT SEEMS TO WORK BETTER.

BUT HOW DOES THIS ALLOW WATSON TO ANSWER QUESTIONS? WHEN A QUESTION COMES IN, WATSON SEARCHES OVER PREVIOUS JEOPARDY! QUESTIONS, WIKIPEDIA, ARTICLES, AND A SELECT SUBSET OF THE INTERNET TO FIND EVIDENCE THAT CAN ANSWER THE QUESTION.

Buzzing is important for Humans and Computers

LET'S BREAK THIS INTO TWO PHASES: **guessing** AND **buzzing**. A GUESS IS ITS BEST GUESS OF WHAT THE ANSWER COULD BE: LITERALLY ONE OF THOUSANDS OF POSSIBLE ANSWERS. A BUZZ IS A BINARY DECISION OF WHETHER TO TRUST WHETHER THIS PARTICULAR GUESS IS CORRECT OR NOT.

LET'S BEGIN WITH AN EASY CASE: LET'S SAY YOU GET A CLUE THAT'S NEARLY IDENTICAL TO AN EXISTING CLUE:

DOUBLE D WORDS
WALK LIKE A DUCK

HAS APPEARED FOUR TIMES AS A CLUE ON JEOPARDY!: 1987, 1996, 2001, AND 2010. IN ADDITION TO THE ABOVE CATEGORY, THIS CLUE ALSO APPEARED UNDER CATEGORIES LIKE **SIX-LETTER WORDS** AND "**WA**". THE GUESS PROCESS LOOKS TO SEE IF IT HAS EVER SEEN THIS EXACTLY CLUE BEFORE, IT HAS, SO IT PRODUCES THE GUESS WADDLE.

THE BUZZER NEEDS TO DECIDE IF IT SHOULD TRUST THIS GUESS. IT HAS A NUMBER THAT IT ASSOCIATES WITH SEEING THE EXACT CLUE BEFORE; LET'S ARBITRARILY SAY THAT IT'S 3.0. IT THEN COMPARES THAT SCORE WITH ITS THRESHOLD ON WHEN TO BUZZ. AGAIN, LET'S ARBITRARILY SAY THAT THE THRESHOLD IS 1.0. 3.0 IS LARGER THAN 1.0, SO IT BUZZES IN WITH THAT GUESS (WE'LL TALK MORE ABOUT HOW TO SET THESE NUMBERS NOT-SO-ARBITRARILY IN SECTION 6.2).

BUT *JEOPARDY!* DOESN'T HAVE ONE KIND OF CLUE. THERE ARE MANY DIFFERENT TYPES OF QUESTION, AND SOMETIMES YOU GET A CLUE FROM THE CATEGORY ABOUT WHAT KIND OF QUESTION YOU WILL NEED TO ANSWER. LET'S START WITH "POTENT POTABLE RHYME TIME": FROM THE NAME OF THE CATEGORY, YOU KNOW THAT BOTH WORDS IN THE RESPONSE HAVE TO RHYME WITH EACH OTHER AND THAT IT WILL HAVE SOMETHING TO DO WITH ALCOHOL. SO, TO ANSWER THE CLUE:

RICE WINE FOR THE GUY WHO RIDES A RACEHORSE

, YOU ARE ESSENTIALLY DOING A CONSTRAINED SEARCH: YOU CAN FIND INDIVIDUAL PIECES THAT FIT THE CLUES (E.G., REWRITING SYNONYMS TO DISCOVER THINGS LIKE MAKKOLI, BREM, TAPAI, ETC. FOR "RICE WINE" AND "COWBOY" AND "KNIGHT" FOR PEOPLE WHO RIDE A HORSE). SO YOUR GUESSER WOULD GENERATE LOTS OF DIFFERENT POSSIBILITIES, BUT FROM PAST EXPERIENCE YOUR BUZZER WOULD GIVE A VERY LOW SCORE TO ANYTHING THAT DIDN'T HAVE TWO WORDS. FOR EXAMPLE, IF THE CATEGORY IS "RHYME TIME" BUT THERE'S ONLY ONE WORD IN THE ANSWER, THE SCORE GOES DOWN BY 5.0, WHICH BALANCES OUT EVEN IF THE BUZZER LIKES SAKE BY ITSELF AS A RESPONSE. THIS IS IN ESSENCE A CONSTRAINED SEARCH: YOU SEARCH FOR INDIVIDUAL COMPONENTS THAT FIT WITH THE CLUE UNTIL YOU GET TWO THINGS THAT MATCH THE CLUE AND THEN RHYME WITH EACH OTHER.

Using Features

NOW HOW DOES THE BUZZER KNOW TO USE THIS INFORMATION? THERE ARE MANY LITTLE PIECES THAT COULD GO INTO THINKING THAT A GUESS IS GOOD OR NOT. THESE PIECES OF EVIDENCE ARE CALLED FEATURES. WHILE FEATURES AREN'T ALWAYS ACTIVE, WHEN THEY ARE, THEY PROVIDE EVIDENCE OF WHETHER A GUESS IS GOOD. FOR EXAMPLE, YOU COULD HAVE A FEATURE (AND WATSON PROBABLY DID) FOR WHEN A CATEGORY HAS A "QUOTE" IN IT: THIS MEANS THAT THE TEXT INSIDE THE QUOTE SHOULD APPEAR IN THE ANSWER. E.G., IF YOU HAVE A CATEGORY "Ten"-letter words, THEN THE CLUE

PATRICK ROY AND HOPE SOLO PLAYED THIS POSITION

WOULD HAVE MANY POSSIBLE ANSWERS THAT MIGHT SCORE HIGHLY IF YOU DIDN'T TAKE THE CATEGORY INTO ACCOUNT: E.G., GOALIE. BUT THE CATEGORY HAS TWO PIECES OF INFORMATION: THAT THE ANSWER SHOULD BE TEN LETTERS LONG: GOALIE FAILS THAT BUT GOALKEEPER DOES. HOWEVER, THE ADDITIONAL QUOTED PART OF THE CATEGORY TELLS THE BUZZER THAT IT SHOULD GIVE A LOW SCORE TO ANY GUESS THAT DOESN'T HAVE "TEN" IN IT.

IN A COMPUTER PROGRAMMING LANGUAGE, THIS COULD SIMPLY BE AN IF STATEMENT: IF ASCII CHARACTER 34 IS IN THIS STRING, RETURN 1. BUT JUST KNOWING THAT THERE'S A QUOTE IN THE CATEGORY DOESN'T TELL YOU WHICH CLUES TO FAVOR. YOU WOULD NEED TO MAKE THE THE FEATURE MORE SPECIFIC.

LET'S GO TO A QUESTION THAT I GOT WHILE I WAS ON JEOPARDY! AND THIS SHOWS THAT THESE AREN'T JUST A CATEGORY-SPECIFIC PHENOMENON... THIS ALSO APPLY TO INDIVIDUAL CLUES AS WELL. IN A QUITE NORMAL CATEGORY ABOUT GEOGRAPHY, THIS CLUE CAME UP:

"G.I." HOPE YOU KNOW THAT 0 DEGREES LATITUDE
& 180 DEGREES LONGITUDE IS JUST EAST OF THIS
GROUP, PART OF KIRIBATI

NOTICE THAT "G.I." IS IN QUOTATION MARKS. THIS IS A SIGNAL THAT THE CORRECT RESPONSE WILL START WITH G.I., IN THIS CASE THE CORRECT RESPONSE IS "GILBERT² ISLANDS". SO PERHAPS WE NEED TO HAVE MULTIPLE FEATURES TO CAPTURE WHETHER OUR ANSWER IS CONSISTENT WITH THIS CLUE. TO BE CLEAR, WE DON'T KNOW EXACTLY WHAT FEATURES WATSON USED, BUT IT COULD LOOK SOMETHING LIKE THIS:

- YOU MIGHT HAVE ONE FEATURE TO INDICATE THAT THERE IS A QUOTATION MARK IN THE CLUE.
- YOU COULD HAVE ANOTHER FEATURE TO SHOW THAT THE QUOTE MATCHES THE CANDIDATE GUESS.
- PERHAPS YOU HAVE ANOTHER FEATURE THAT INDICATES WHETHER THE FEATURE MATCHES A MULTIWORD GUESS.

WHAT MADE WATSON SUCH A TOUR DE FORCE IS THAT HUMANS HAD TO COME UP WITH EACH OF THESE FEATURES. THIS IS NOT A ONE AND DONE PROCESS. IF YOU'RE BUILDING A SYSTEM LIKE THIS, YOU SHOULD LOOK AT THINGS THE SYSTEM IS STILL GETTING WRONG AND ADD A NEW FEATURE TO CORRECT THE ISSUE. THIS IS THE SAME PROCESS THAT WATSON USED TO BUILD THEIR STATISTICAL SYSTEM. MAYBE THE SYSTEM GOT CONFUSED WHEN SOMEBODY HAS A QUOTE FROM LITERATURE AND YOU NEED TO MAKE SURE THAT DOESN'T GET COUNTED OR ADD A NEW FEATURE TO HANDLE THAT CASE. IT'S VERY RARE TO GET THE FEATURES RIGHT THE FIRST TIME AROUND.

THE STATISTICAL APPROACH HAS SOME ADVANTAGES OVER THE NEURAL APPROACH WE'LL SEE IN THE NEXT CHAPTER: IT IS EASY TO UNDERSTAND WHY A SYSTEM IS DOING WHAT IT IS. AND IT'S EASY TO FIX PROBLEMS AS THEY POP UP.

IN OTHER CASES, THE CONSTRAINT IS THAT THE CORRECT ANSWER STARTS WITH A PARTICULAR LETTER.

THE CATEGORY THAT I FEARED THE MOST WAS ANAGRAMS: WHERE YOU NEED TO REARRANGE THE LETTERS IN THE CLUE TO GET TO THE ANSWER. THIS IS ACTUALLY EASIER FOR A COMPUTER THAN A HUMAN!

EACH OF THESE CAN HAVE DIFFERENT APPROACHES THAT CAN GENERATE GUESSES THAT MIGHT BE THE CORRECT ANSWER. AND YOU MIGHT THINK THAT THIS IS ALL ENCODED IN THE CATEGORY. NOT SO!

AND FOR "NORMAL" CLUES, THIS LOOKS A LOT LIKE THE REFORMULATIONS THAT WE SAW IN QASM: 400TH ANNIVERSARY OF 1898 GETS REWRITTEN TO 27TH MAY 1498, INDIA GETS TRANSFORMED INTO KAPPAD, AND THEN YOU CAN FIND SOMETHING WITH THE CORRECT RESPONSE: VASCO DE GAMA.

² Named of course for Johnny Gilbert, "the voice of *Jeopardy!*". Actually not, it's far more confusing than that: it was named by a German admiral—Adam Johann von Krusenstern—who led the first Russian circumnavigation of the globe. Krusenstern recognized that the British Captain Thomas Gilbert had described each of the islands individually and applied the name to the group of islands.

SO THE FIRST PHASE OF THE WATSON PIPELINE IS TO—IN PARALLEL—GENERATE ALL SORTS OF GUESSES BASED ON DIFFERENT INTERPRETATIONS OF THE QUESTION. SOME TRY TO SOLVE ANAGRAMS, OTHERS LOOK FOR RHYMES, OTHERS RUN AN IR QUERY LIKE THE START SYSTEM WE TALKED ABOUT BEFORE, OTHERS ARE DOING TRANSFORMATIONS LIKE THE QASM APPROACH WE TALKED ABOUT BEFORE.

AND SOMETIMES THE APPROACHES NEED TO TAKE A RECURSIVE APPROACH, COMBINING DIFFERENT SUBSYSTEMS TO GET TO THE RIGHT ANSWER.

Logistic Regression and Gradient Descent

FROM THIS YOU HAVE DOZENS OF POSSIBLE GUESSES. HOW DO YOU KNOW WHICH—if any—to select? THIS ALL GETS FED INTO A LOGISTIC REGRESSION PROBLEM. THIS CAN TAKE INTO ACCOUNT HOW CONSISTENT THE EVIDENCE MATCHES THE CLUE, HOW POPULAR THE RESPONSE IS, AND IT CAN EVEN TAKE INTO ACCOUNT THINGS LIKE JEOPARDY!’S LOVE OF PUNS.

IN THE EXAMPLES ABOVE, WE ASSUMED THAT WE KNEW WHAT THE WEIGHTS WERE FOR EACH OF OUR FEATURES. FOR THE MOMENT, LET’S CONTINUE THAT ASSUMPTION... BUT WE’LL TRY TO IMPROVE THEM SO THAT WE GET MORE BUZZES RIGHT AND FEWER WRONG. THIS WILL REQUIRE A BIT OF MATH, SO IF YOU JUST WANT TO JUMP AHEAD TO THE BIG PICTURE (E.G., IF YOU ALREADY KNOW WHAT SGD FOR LOGISTIC REGRESSION LOOKS LIKE), GO AHEAD TO SECTION 6.2.

SO WHAT DOES IT MEAN TO GET MORE BUZZES RIGHT? FIRST, WE NEED TO DEFINE A LITTLE TERMINOLOGY: WE NEED TO CONSIDER THE CORRECT EXAMPLES \mathcal{X}_c AND ALL THE WRONG EXAMPLES \mathcal{X}_f . WE WANT ALL OF THE CORRECT EXAMPLES TO HAVE SCORES MORE THAN THE THRESHOLD AND ALL OF THE INCORRECT EXAMPLES TO HAVE SCORES LESS THAN THE THRESHOLD. LET’S CALL THE THRESHOLD b , AND THEN WE NEED TO SUM UP ALL OF THE NON-ZERO FEATURES OF AN EXAMPLE TO COMPUTE THE SCORE OF AN EXAMPLE i :

$$\text{SCORE}(\mathbf{x}_i) = \sum_j w_j x_{i,j}. \quad (6.1)$$

IN OTHER WORDS, WE TAKE EVERY FEATURE, MULTIPLY IT BY THE WEIGHT ASSOCIATED WITH IT (E.G., $x_{i,j}$ MIGHT BE DOES THIS CLUE HAVE A CATEGORY WITH A QUOTE IN IT). WE WRITE THIS AS A MULTIPLICATION AND NOT AS JUST ADDING THE FEATURE WEIGHTS DIRECTLY BECAUSE SOMETIMES OUR FEATURES WON’T BE BINARY. AND THIS SHOULD LOOK FAMILIAR: IT LOOKS A LOT LIKE THE DOT PRODUCT THAT WE FIRST SAW IN THE CHAPTER ON INFORMATION RETRIEVAL (CHAPTER 4), WHERE WE COMPUTE THE SIMILARITY OF TWO EXAMPLES BY MULTIPLYING THE TF-IDF SCORES OF TWO DOCUMENTS’ WORDS.

AND THIS IS A GOOD SEGUE TO WHY WE WRITE THIS AS A DOT PRODUCT RATHER THAN AS A SUM OF ALL OF THE RELEVANT FEATURE WEIGHTS. IF ALL OF THE FEATURES OF AN EXAMPLE $\mathbf{x}_{i,j}$ ARE BINARY, THEN YOU TAKE THE NON-ZERO ELEMENTS AND ADD UP THE WEIGHTS (SINCE ONE TIMES THE WEIGHT JUST GIVES YOU THE FEATURE WEIGHTS). BUT ANOTHER VERY GOOD FEATURE MIGHT

BE HOW WELL THE GUESS MATCHES THE CLUE (E.G., DOT PRODUCT BETWEEN TF-IDF SCORES). SO A WAY TO GENERALIZE BOTH THE BINARY FEATURES AND CONTINUOUS FEATURES IS TO TAKE THE DOT PRODUCT OF A VECTOR REPRESENTING ALL OF THE FEATURES AND ALL OF THE WEIGHTS TO COMPUTE A SCORE FOR THE GUESS. THE HIGHER THE SCORE, THE BETTER THE GUESS AND THE MORE LIKELY THE SYSTEM WILL ANSWER.

BUT THIS DOESN'T ANSWER THE QUESTION OF HOW WE CAN LEARN WHAT THE WEIGHTS SHOULD BE. WE NEED TO LOOK AT THE CASES WHERE THE SYSTEM HAS GENERATED A LOT OF GUESSES AND SEEN WHICH OF THEM ARE CORRECT OR NOT. FROM THIS, YOU CAN LEARN THAT WHEN A LETTER IS IN QUOTES AND IT MATCHES, THAT'S A GOOD SIGN FOR GETTING A GUESS RIGHT. MORE CONCRETELY WE WANT THE SCORE (EQUATION 6.1) TO BE HIGH FOR RIGHT GUESSES AND LOW FOR WRONG ONES, AND IF IT'S NOT, WE NEED TO ADJUST THE SCORES BY CHANGING THE WEIGHTS w .

ONE MISSING PIECE OF THE EXPLANATION IS HOW TO TURN OUR WEIGHTS INTO PROBABILITIES. FOR THAT, WE ARE GOING TO USE THE LOGISTIC FUNCTION. THIS IS A FUNCTION THAT CAN TAKE ANY REAL NUMBER AS AN INPUT AND TURN IT INTO A PROBABILITY BETWEEN ZERO AND ONE. WHEN THE INPUT IS ZERO, IT PROVIDES A PROBABILITY IS 0.5, AND WHEN THE INPUT GOES TO POSITIVE INFINITY, THE PROBABILITY APPROACHES ONE; LIKEWISE, WHEN THE INPUT GOES TO NEGATIVE INFINITY, THE PROBABILITY APPROACHES ZERO. FOR CONVENIENCE, WE WILL CALL THIS FUNCTION $\sigma(x)$, BECAUSE SOME PEOPLE CALL THIS THE SIGMOID FUNCTION.




Figure 6.1: This function is called the sigmoid function because the way it transforms all inputs into values makes it look like an “S” shape, with 0 outputs to the left, increasing in the middle (an input of 0 right at the threshold gives an output of 0.5) and 1 outputs to the right.

FOR THIS, WE ARE GOING TO NEED AN *OBJECTIVE FUNCTION*: A MATHEMATICAL EXPRESSION THAT TELLS US HOW FAR WE ARE FROM OUR GOAL. IN THIS CASE OUR OBJECTIVE IS TO MAKE THE PROBABILITY SAYING WE SHOULD BUZZ—COMPUTED FROM OUR WEIGHTS—FOR INCORRECT BUZZES BE AS LOW AS POSSIBLE. SO THIS MEANS THAT WE'RE GOING TO TRANSFORM OUR WEIGHTS INTO A PROBABILITY π_i BETWEEN ZERO AND ONE THAT GETS BIG WHEN THE WEIGHTS PASS THE THRESHOLD:

$$\pi_i = \sigma \left(\sum_j w_j x_{i,j} - b \right) \equiv p(\text{BUZZ} | x) \equiv \frac{1}{1 + \exp - \left(\sum_j w_j x_{i,j} + b \right)}. \quad (6.2)$$

SINCE WE ONLY WANT TO BUZZ WHEN WE'RE RIGHT, THEN WE WANT π_i TO

BE SMALL (CLOSE TO ZERO) WHENEVER THE BUZZ IS WRONG AND BIG WHEN THE BUZZ IS RIGHT (CLOSE TO ONE). ALTERNATIVELY, WE COULD ALSO SAY WE WANT $1 - \pi_i$ TO BE BIG WHEN THE BUZZ IS WRONG, SINCE THIS MEANS THAT π_i IS AS CLOSE TO ZERO. THIS ALTERNATE FORMULATION IS ACTUALLY WHAT WE WANT SINCE WE WILL WANT TO OPTIMIZE A SINGLE FUNCTION GOING IN THE SAME DIRECTION. BECAUSE WE CARE ABOUT ALL EXAMPLES, YOU WOULD NORMALLY MULTIPLY THE PROBABILITY OF JOINT EVENTS TOGETHER, BUT THIS LEADS TO A VERY SMALL NUMBER (SINCE PROBABILITIES ARE LESS THAN ONE), SO FOR MATHEMATICAL CONVENIENCE WE TAKE THE LOG:

$$\mathcal{L}' \equiv= \underbrace{\sum_{i \in \mathcal{X}_c} \log \pi_i}_{\text{HOW RIGHT ARE WE}} + \underbrace{\sum_{i \in \mathcal{X}_f} \log 1 - \pi_i}_{\text{HOW WRONG ARE WE}} . \quad (6.3)$$

SO THIS OVERALL EXPRESSION (THE LOG PROBABILITY) TELLS US HOW GOOD OF A JOB OUR FEATURES ARE DOING IN TELLING US WHEN TO BUZZ.

LATER IN THE BOOK, OBJECTIVE FUNCTIONS LIKE THIS WILL BE CALLED A *LOSS FUNCTION*, SO TO BE CONSISTENT WITH THAT, WE WILL THINK ABOUT HOW TO *MINIMIZE OUR MISTAKES*, SO WHAT WE WILL DO NOW (SOMEWHAT CONFUSINGLY) IS TO MINIMIZE THE NEGATIVE OF \mathcal{L}' , WHICH WE WILL CALL \mathcal{L} :

$$\mathcal{L} \equiv= - \sum_{i \in \mathcal{X}_c} \log \pi_i - \sum_{i \in \mathcal{X}_f} \log 1 - \pi_i . \quad (6.4)$$

JUST TO RECAP: EQUATION 6.2 IS THE PROBABILITY OF THE MODEL SAYING WE SHOULD BUZZ ON EXAMPLE i , EQUATION 6.3 COMBINES ALL OF THOSE TOGETHER TO SEE HOW CLOSE WE ARE TO BEING RIGHT, AND THEN EQUATION 6.4 FLIPS IT AROUND SO THAT THE LARGE IT IS, THE WRONGER WE ARE... WE WANT TO MINIMIZE THAT EXPRESSION.

IF YOU REMEMBER HIGH SCHOOL CALCULUS, THIS IS AN OPTIMIZATION PROBLEM: YOU CAN COMPUTE THE DERIVATIVE (GRADIENT, ACTUALLY) OF EACH OF EACH OF THE VARIABLES YOU HAVE CONTROL OVER—THE FEATURE WEIGHTS w —AND THEN ADJUST THOSE WEIGHTS TO DECREASE THE LOSS FUNCTION REPRESENTED BY \mathcal{L} (EQUATION 6.4). WHEN YOU DID THIS IN HIGH SCHOOL CALCULUS, YOU COULD PROBABLY SOLVE THE EQUATION FOR WHEN THE DERIVATIVE WAS ZERO, SET THE VARIABLE TO MAKE THE DERIVATIVE ZERO, AND THEN YOU'RE DONE. BUT THAT'S NOT POSSIBLE HERE, SO WE NEED TO TAKE LITTLE ADJUSTMENTS TO w TO PUSH \mathcal{L} DOWN AGAIN AND AGAIN.

AN INTUITIVE WAY OF THINKING ABOUT THIS IS THAT THE SHAPE OF \mathcal{L} REPRESENTS A HILL. WHERE YOU STAND ON THE HILL IS REPRESENTED BY HOW YOU'VE SET w : EACH DIMENSION REPRESENTS ONE DIMENSION OF THIS VECTOR.³ IN TWO DIMENSIONS, ONE DIMENSION IS HOW FAR NORTH–SOUTH YOU ARE AND ONE DIMENSION IS HOW FAR EAST–WEST YOU ARE: AND YOU'RE TRYING TO GET TO THE LOWEST VALLEY YOU CAN. THE CATCH IS THAT YOU CAN'T REALLY SEE WHAT THE WHOLE LANDSCAPE LOOKS LIKE; IT'S SO FOGGY YOU CAN ONLY SEE RIGHT AROUND WHERE YOU ARE. THE INTUITION IN THIS CASE IS TO WALK “DOWNHILL”, AND THIS IS EXACTLY WHAT THE GRADIENT OF \mathcal{L} GIVES YOU.

³ The threshold b is also an additional dimension you need to optimize, but we'll gloss over this for now... you can always fudge it by imagining it an additional dimension of w that's always active for every example.

NOW THE PROBLEM IS THAT COMPUTING \mathcal{L} OVER YOUR ENTIRE DATASET IS HARD: YOU HAVE TO GO OVER EVERY SINGLE QUESTION IN THE *JEOPARDY!* TRAINING SET TO COMPUTE ITS CONTRIBUTION TO THE GRADIENT. THE WAY THAT I LIKE TO THINK ABOUT IT IS THAT EACH EXAMPLE IS A PERSON YOU CAN ASK “HOW DO I GO DOWNHILL”. WHILE YOU COULD ASK LOTS OF DIFFERENT PEOPLE, IF YOU ASK LITERALLY EVERYONE, THAT’S PROBABLY GOING TO BE OVERKILL. YOU CAN PROBABLY GET A GOOD SENSE OF THE DIRECTION BY JUST ASKING A HANDFUL OF EXAMPLES OF WHICH WAY YOU SHOULD GO.

IN THE LINGO, THE SAMPLE OF EXAMPLES THAT YOU ASK FOR THE DIRECTION IS THE *MINIBATCH*, AND THIS OVERALL PROCESS OF ASKING A FEW (OR EVEN ONE) EXAMPLE FOR A DIRECTION IS CALLED *STOCHASTIC GRADIENT DESCENT* (BOTTOU, 2004).

GIVEN THAT THE MATH FOR THE DERIVATION IS EVERYWHERE, I WILL SKIP IT HERE (YOU CAN TAKE A LOOK AT CHAPTER 5 OF SPEECH AND LANGUAGE PROCESSING BY JURAFSKY AND MARTIN, WHICH I OFTEN USE IN MY CLASSES AND THUS HAVE ADOPTED THEIR NOTATION), BUT IT IS SO INTUITIVE THAT IF YOU LOOK AT IT, IT JUST “FEELS RIGHT”. FOR A MINIBATCH OF SIZE ONE, IF YOU ASK IT WHICH WAY YOU SHOULD GO, IT TELLS YOU TO UPDATE⁴ YOUR WEIGHT w_i TO BE:

$$w_j^{(new)} = w_j^{(old)} + \lambda \frac{\delta \mathcal{L}}{w_j} = w_j^{(old)} - \lambda \underbrace{(\pi_i - y_i)}_{\text{ERROR}} x_{i,j}. \quad (6.5)$$

THE MOST IMPORTANT PART OF THE EQUATION IS THE ERROR: HOW FAR OFF YOUR PREDICTION FOR EXAMPLE j IS. REMEMBER THAT THE TRUE OUTCOME y_i IS BINARY (EITHER YOU SHOULD HAVE BUZZED— $y_i = 1.0$ —OR YOU SHOULDN’T HAVE— $y_i = 0.0$), BUT THAT THE PREDICTION π_i IS A PROBABILITY AND THUS RANGES BETWEEN 0.0 AND 1.0. THE MORE RIGHT YOU ARE—THE CLOSER YOU ARE TO y_i —THE LESS THE WEIGHTS CHANGE. IN OTHER WORDS, IF YOU GOT THIS EXAMPLE EXACTLY RIGHT, NOTHING CHANGES AT ALL. BUT THE BIGGER YOUR ERROR IS, THE MORE YOU CHANGE. EXACTLY HOW MUCH IS CONTROLLED BY THE STEP SIZE PARAMETER λ , AND THE DIRECTION OF THE CHANGE DEPENDS ON THE SIGN OF THE ERROR. IF THE ERROR WAS POSITIVE, YOUR PREDICTION WAS TOO LOW, SO YOU SHOULD INCREASE ALL OF THE FEATURE WEIGHTS THAT CAUSED YOU TO SAY YOU SHOULD HAVE BUZZED WHEN YOU DIDN’T. IF THE ERROR WAS NEGATIVE, YOUR PREDICTION WAS TOO HIGH, AND SO YOU SLIGHTLY DECREASE ALL OF THE FEATURE WEIGHTS THAT CAUSED YOU TO GET IT WRONG.

THE OTHER INPUT TO THE UPDATE IS HOW MUCH OF THIS FEATURE THE EXAMPLE HAS: $x_{i,j}$. TO MAKE SURE YOUR GRADIENTS DON’T GET TOO BIG OR TOO SMALL, IN PRACTICE WE OFTEN STANDARDIZE THE FEATURE WEIGHTS SO THAT THEY HAVE MEAN ZERO AND UNIT VARIANCE. HOWEVER, BECAUSE THIS CHAPTER IS TRYING EXPLAIN HOW THESE THINGS WORK (AND WE’LL WORK THROUGH AN EXAMPLE NEXT), WE WON’T DO THAT TO MAKE THE MATH MORE STRAIGHTFORWARD.

SO LET’S SEE HOW THIS WOULD WORK IN A MORE REAL-WORLD EXAMPLE.

⁴Because we’re doing stochastic *descent*, we subtract the gradient from the current weight...if we were trying to make the objective function as large as possible, we’d add it to the current weights.

LET'S SAY THAT YOU HAVE A CLUE

US CITIES
ITS LARGEST AIRPORT IS NAMED FOR A WWII HERO.
ITS SECOND LARGEST, FOR A WWII BATTLE.

THAT WATSON FAMOUSLY GOT WRONG AS A FINAL *JEOPARDY!* WHEN IT ANSWERED THE FAMOUS AMERICAN CITY TORONTO. SO IN THIS CASE OUR BUZZER INPUT⁵ IS WHETHER WE SHOULD BUZZ ON THIS CLUE WITH THE RESPONSE TORONTO, FOR WHICH WE GET A π_i . FOR THE SAKE OF CONCRETENESS, LET'S SAY π_i IS 0.4... IT'S NOT HIGH, BUT IT SHOULD BE ZERO BECAUSE MIDWAY AND O'HARE AIRPORTS ARE ACTUALLY IN THE GREAT CITY OF CHICAGO. NOW THIS IS TOO HIGH, SO SOMETHING WENT WRONG WITH OUR FEATURES. LET'S SEE WHAT FEATURES WERE ON AND WHAT THEIR CURRENT WEIGHTS IN TABLE 6.1.

Feature	Weight w_i	x_i
IR Score	2	0.1
<u>Toronto</u> compatibility with Category	1	0.2
Knowledge Base (# Airports in Toronto)	0.1	2
Previous responses (<u>Toronto</u>)	0.005	40
Bias		-0.8

THESE ARE REASONABLE FEATURES: THE IR SCORE ENCODES HOW SIMILAR THIS IS TO QUESTIONS THAT YOU'VE SEEN BEFORE; THE CATEGORY FEATURE ENCODES WHETHER THE ANSWER IS COMPATABLE WITH THE CATEGORY "US CITIES" (AND WHILE YOU IMAGINE THAT IT ISN'T GREAT, YOU COULD IMAGINE SUBSTITUTING "AMERICAN" FOR "US", WHICH WOULD MAKE TORONTO A LOT MORE PLAUSIBLE); THE KNOWLEDGE BASE FEATURE CHECKS TO SEE HOW MANY AIRPORTS ARE IN TORONTO; AND THE FINAL FEATURE CHECKS TO SEE IF TORONTO HAS BEEN AN ANSWER BEFORE (YOU MIGHT IMAGINE THAT MORE FREQUENT ANSWERS IN THE PAST WILL APPEAR AGAIN).

THERE ARE LOTS OF OTHER FEATURES THAT WE'RE NOT SEEING BECAUSE THEY'RE ZERO, E.G., THE FEATURE THAT CHECKS TO SEE IF THERE'S A QUOTATION MARK IN THE CATEGORY DOESN'T SEE ONE, SO IT'S ZERO AND WE CAN SAFELY IGNORE IT ALONG WITH EVERY OTHER FEATURE THAT DIDN'T FIRE.⁶ AND THIS MAKES SENSE: THE FEATURES THAT AREN'T ON FOR THIS FEATURE HAD NOTHING TO DO WITH THE MISTAKE, SO THEY SHOULDN'T BE PUNISHED OR REWARDED.

FIRST, LETS SEE HOW THESE FEATURES TRANSLATE INTO A PROBABILITY:

$$\pi_i = \sigma \left(\underbrace{2 \cdot 0.1}_{\text{IR}} + \underbrace{1 \cdot 0.2}_{\text{CATEGORY}} + \underbrace{0.1 \cdot 2}_{\text{KB}} + \underbrace{0.005 \cdot 40}_{\text{PREV}} - \underbrace{0.8}_{\text{BIAS}} \right) \quad (6.6)$$

$$= \sigma (0.2 + 0.2 + 0.2 + 0.2 - 0.8) = \sigma (0.0) = 0.5 \quad (6.7)$$

FROM EQUATION 6.5, THE ERROR IS $0.5 - \pi_i$ SHOULD HAVE BEEN ZERO, BUT IT WAS 0.5—AND IF WE ASSUME THAT OUR STEP SIZE⁷ IS 1.0, WE CAN NOW WRITE

⁵ In reality, the calculus for Final *Jeopardy!* is a little different because it makes sense to provide some answer no matter what, you don't really buzz in. However, you still want the best answer to have the highest probability, so we can still work through this example. And indeed, given the debugging output *Watson* wasn't very confident in its wrong output, so it probably wouldn't have "buzzed in" during normal play, suggesting that everything was working as intended. We talk about wagers at a cursory level in Section 6.4.

⁶ Lest there be any doubt, I am making up these weights and features for a simple example... *Watson* had many more and better features. Most unbelievably, the first four features are non-zero for this example and the rest are zero. Finally, you normally don't have feature weights with such round numbers, but it makes the example cleaner.

⁷ In practice, this would be a really bad step size—it's way too large, you're likely to jump over the valley in the objective function you're looking for—but we'll go with this because it makes the math easier.

THE OVERALL UPDATE

$$\mathbf{w}^{(new)} = \mathbf{w}^{(old)} - \lambda \nabla_{\mathbf{w}, b} \mathcal{L} \quad (6.8)$$

$$= \begin{bmatrix} 2 \\ 1 \\ 0.1 \\ 0.005 \\ -0.8 \end{bmatrix} - \lambda \begin{bmatrix} \frac{\partial \mathcal{L}}{\partial w_1} \\ \frac{\partial \mathcal{L}}{\partial w_2} \\ \frac{\partial \mathcal{L}}{\partial w_3} \\ \frac{\partial \mathcal{L}}{\partial w_4} \\ \frac{\partial \mathcal{L}}{\partial b} \end{bmatrix} = \begin{bmatrix} 2 \\ 1 \\ 0.1 \\ 0.005 \\ -0.8 \end{bmatrix} - 1.0 \begin{bmatrix} 0.5 \cdot 0.1 \\ 0.5 \cdot 0.2 \\ 0.5 \cdot 2 \\ 0.5 \cdot 40 \\ 0.5 \cdot 1 \end{bmatrix} = \begin{bmatrix} 1.95 \\ 0.75 \\ -0.9 \\ -19.995 \\ -1.3 \end{bmatrix}.$$

THIS SHOWS THE DANGERS OF HAVING UNBALANCED FEATURE WEIGHTS AND TOO LARGE STEP SIZES. THIS CHANGE IS OVERALL TOO BIG; YOU WANT TO BE TAKING LITTLE STEPS ON YOUR OBJECTIVE FUNCTION, NOT TAKING HUGE LEAPS—YOU MIGHT JUMP OVER THE “RIGHT ANSWER” IN FRONT OF YOU AND NOT FIX THE REAL PROBLEM IN THIS EXAMPLE: THAT TORONTO IS INCONSISTENT WITH THE CATEGORY “US CITIES”.

Feature Engineering

HOWEVER, THESE UPDATES CAN ONLY UPDATE THE FEATURES THAT YOU HAVE. THEY CANNOT ADD NEW FEATURES NOR CAN THEY FIX PROBLEMS INTRINSIC IN THE FEATURES. WHEN YOU'RE STARTING OUT BUILDING A SYSTEM LIKE THIS, YOU USUALLY DON'T HAVE VERY MANY FEATURES. BUT AS YOU GO THROUGH THE PROCESS OF TESTING YOUR INITIAL SYSTEM, YOU'LL SEE LOTS THINGS THAT YOUR SYSTEM GETS WRONG BUT *SHOULDN'T*: YOUR SYSTEM HAS ALL OF THE EVIDENCE IT NEEDS TO DO WHAT IT NEEDS, BUT CANNOT ACTUALLY GET THERE.

LET'S FOCUS ON THE TORONTO QUESTION A LITTLE BIT MORE. WE ALREADY TALKED ABOUT HOW THE FEATURE ITSELF MIGHT NEED SOME DEBUGGING (E.G., MAKING SURE THAT “US” DOESN'T EXPAND TO COVER ALL OF THE AMERICAN CONTINENT), BUT HOW COULD WE ADD ADDITIONAL FEATURES THAT COULD HANDLE THIS KIND OF PHENOMENA BETTER?

IN A 2010 TOURNAMENT OF CHAMPIONS GAME, THERE WAS A CATEGORY WITH THE FIRST CLUE:



WHEN THE CATEGORY WAS REVEALED, THE HOST ALEX TREBEK SAID “YOU HAVE TO NAME THE MONTH”, BUT *WATSON* DIDN'T GET THAT HINT. IN A PRESENTATION FROM IBM, THEY SHOWED THAT *WATSON* GOT THAT CLUE WRONG (FIGURE 6.2).

SOMETHING THAT IS PARTICULARLY UNIQUE TO *JEOPARDY!* AND NOT TO THE OTHER QA SETTINGS WE'VE LOOKED AT IS THAT THE CATEGORY (THEY TEXT THAT APPEARS ATOP THE COLUMN A CLUE APPEARS IN) IS A HUGE CONSTRAINT ON WHAT THE CORRECT RESPONSES ARE. A LARGE PART OF THE *JEOPARDY!* BUZZER IS FIGURING OUT THE “LEXICAL ANSWER TYPE”: WHAT KIND OF THING CAN THE ANSWER BE AND ONLY BUZZING IN ON THE THINGS THAT ARE CONSISTENT WITH THAT TYPE. PART OF THIS IS LOOKING AT CLUES IN THE CLUE. THIS CHAPTER

Clue	Type	Watson's Answer	Correct Answer
D-DAY ANNIVERSARY & MAGNA CARTA DAY	day	Runnymede	June
NATIONAL PHILANTHROPY DAY & ALL SOULS' DAY	day	Day of the Dead	November
NATIONAL TEACHER DAY & KENTUCKY DERBY DAY	Day/month(.2)	Churchill Downs	May
ADMINISTRATIVE PROFESSIONALS DAY & NATIONAL CPAs GOOF-OFF DAY	day / month(.6)	April	April
NATIONAL MAGIC DAY & NEVADA ADMISSION DAY	day / month(.8)	October	October

Figure 6.2: Slide from IBM Watson on how a feature to predict the answer type as it works its way through a category's column can improve its ability to buzz in correctly. (Figure Credit: IBM)

HAS USED A LOT OF EXAMPLE CLUES LIKE “THIS WOMAN” (SUE GRAFTON), “THIS CODE BREAKER” (TURING), “THIS POSITION” (GOALTENDER), OR “THIS GROUP” (GILBERT ISLANDS). THESE ARE HINTS ABOUT WHAT KIND OF RESPONSE IS BEING SOUGHT. AND LEST YOU THINK THIS IS AN ARTIFICE OF THE SKILLED WRITERS OF JEOPARDY!, THIS ALSO APPEARS IN MORE “NATURAL” QUESTIONS LIKE THOSE

PEOPLE ASK GOOGLE (CHAPTER **?reference?** ??).

AND PART OF WHAT MADE WATSON A GOOD JEOPARDY! PLAYER IS THAT IT WOULD LEARN AS IT EXPLORED THE CATEGORY. AFTER GETTING A COUPLE OF MONTHS WRONG, WATSON CAN LEARN THAT ALL OF THE ANSWERS SHOULD BE A MONTH. WHILE WE DON’T KNOW FOR SURE HOW THIS WAS IMPLEMENTED IN DETAIL, WE CAN IMAGINE THAT THERE IS A FEATURE THAT SUGGESTS WHETHER THE CLUE IS CONSISTENT WITH AN ANSWER TYPE OF “DAY” OR “MONTH” (E.G., THIS CLUE IS CONSISTENT WITH A “MONTH” AS AN ANSWER, AND JUNE IS A MONTH). AND INFORMATION FROM THE CURRENT COLUMN COULD ALSO BE INCLUDED IN THIS, EITHER DIRECTLY IN THE COMPUTATION OF THE FEATURE OR DIRECTLY ENCODING SOMETHING LIKE “GIVEN THE GUESS APRIL, HOW MANY OF THE PREVIOUS RESPONSES IN THE CATEGORY ARE CONSISTENT WITH THAT TYPE”.

ONE THING THAT MADE WATSON PARTICULARLY GROUNDBREAKING WAS THAT IT NOT JUST COMPUTED RAW ACCURACY BUT ALSO COMPARED AGAINST HUMAN PERFORMANCE. THIS CHART SHOWED HOW WATSON PROGRESSED IN DIFFERENT ITERATIONS OF THE SYSTEM, INCHING UP TO THE CLOUD OF JEOPARDY! CHAMPIONS. KEN JENNINGS IS IN RED THERE, STILL CLEARLY DOMINATING EVEN THE FINAL VERSION OF WATSON. THIS COMPARISON IS EVEN MORE IMPORTANT TODAY (CHAPTER 10) AS PEOPLE CLAIM THAT AIs HAVE SUPER-HUMAN, BUT THE LESSONS OF WATSON CAN HELP INFORM HOW WE CAN JUDGE WHETHER THESE JUDGEMENTS ARE FAIR AND REASONABLE. SO WAS THE WATSON MATCH “THE REAL DEAL”?

6.3 This Game is Rigged, I Tell Ya!

THE NOMINAL SUCCESS OF WATSON HAS BEEN WELL DOCUMENTED (NOT LEAST BY IBM ITSELF, WHO RIGHTLY CELEBRATED THE GREAT TECHNICAL ACHIEVEMENTS AND THE GREAT SHOW THEY PUT ON); HOWEVER, THINGS WERE NOT

PERFECT... THE GAME WAS RIGGED. IT'S USEFUL TO GO OVER THE LIFECYCLE OF AN ENTIRE QUESTION: HOW IT WAS WRITTEN, HOW IT'S COMMUNICATED, HOW PLAYERS ANSWER, AND HOW THE GAME UNFOLDS AFTERWARD. AT EVERY STAGE, THERE'S A SLIGHT BENEFIT TO THE COMPUTER, WHICH TAKEN TOGETHER MAKES THIS AN UNFAIR COMPETITION.

THIS IS A PROBLEM! FIRST, IT'S A PROBLEM SCIENTIFICALLY BECAUSE WE WANT TO HAVE FAIR COMPARISONS OF HUMAN VS. COMPUTER INTELLIGENCE. MORE IMPORTANTLY, I WANT TO HAVE MY TURN HAVING MY QUESTION ANSWERING ROBOTS SIT OPPOSITE AGAINST TRIVIA WHIZZES (CHAPTER 12), AND I CAN'T DO THAT IF EVERYBODY THINKS THAT WATSON'S SPIN ON *JEOPARDY!* SETTLED THE ISSUE (AND IT HASN'T).

BUT FIRST, IN CASE YOU DON'T KNOW HOW *JEOPARDY!* WORKS, WE'LL REVIEW THAT. HOWEVER, IF YOU'VE CALCULATED A CORYAT SCORE BEFORE, YOU CAN GO AHEAD AND SKIP AHEAD TO SECTION 6.3.

The Pool of Questions

PART OF THE AGREEMENT BETWEEN *JEOPARDY!* AND IBM WAS THAT THE COMPETITION WOULD TAKE PLACE ON NORMAL, WRITTEN QUESTIONS. IN THE MEDIA COVERAGE OF THE COMPETITION, THIS FOCUSED ON AVOIDING VIDEO AND PICTURE DAILY DOUBLES (FAIRLY REASONABLE, BUT WE'LL DISCUSS HOW MULTIMEDIA QUESTIONS MIGHT BE MORE FUN IN CHAPTER ?reference? ??). HOWEVER, THIS CAUSES TWO PROBLEMS: THE QUESTIONS ARE TOO EASY AND DO NOT NECESSARILY CHALLENGE COMPUTERS.

SO WHAT MAKES UP "NORMAL" QUESTIONS? EVERY GAME OF *JEOPARDY!* HAS QUESTIONS THAT RANGE IN DIFFICULTY. BECAUSE IT'S A TELEVISION SHOW, MANY QUESTIONS ARE EASY ENOUGH THAT THE AVERAGE VIEWER AT HOME CAN GET THEM. MOREOVER, THE HUMANS ON THE STAGE WITH WATSON ARE NOT NORMAL CONTESTANTS. KEN JENNINGS IS CERTIFIABLY THE GREATEST OF ALL TIME (LOW, 2020, GOAT) *JEOPARDY!* PLAYER, AND BRAD RUTTER ISN'T BAD HIMSELF.

THE AVERAGE "NORMAL" *JEOPARDY!* CONTESTANT, INCLUDING NOT SO GREAT PLAYERS LIKE YOURS TRULY, KNOW A LARGE MAJORITY OF THE CLUES. FOR TOP PLAYERS LIKE BRAD AND KEN, THEY KNOW—WITH A HANDFUL OF EXCEPTIONS—ALL THE CLUES. IN A ONE-ON-ONE FIGHT WITH NORMAL CLUES, KEN AND BRAD WOULD BE FIGHTING OVER EVERY CLUE: IT WOULD COME DOWN TO WHO COULD BUZZ FIRST.

THIS ISN'T FUN TO PLAY. NOR IS IT FUN TO WATCH. THIS IS WHY *JEOPARDY!*'S TOURNAMENT OF CHAMPIONS IS PLAYED ON MUCH MORE DIFFICULT⁸ CLUES (HARRIS, 2006). NONETHELESS, THIS IS THE BATTLEFIELD WHERE WATSON WON: "NORMAL" QUESTIONS THAT DIDN'T CHALLENGE THE HUMAN PLAYERS. INSTEAD, IT ALL CAME DOWN TO THE BUZZER.

⁸ And it the difficulty doesn't just ratchet in one direction, all "special" matches use designated questions: "Celebrity *Jeopardy!*" are easier (as mocked by categories like "States that end with Hampshire" or clues like "You wear these on your face to help you see better" on *Saturday Night Live*), and "College *Jeopardy!*" isn't necessarily easier but does have more college football and popular music. All of these feature questions written with care to be tuned to abilities of the expected players. IBM did not want the computer to be targeted, lest the questions be adversarial against the computer. Chap-

John Henry vs. the Buzzing Machine

UNLIKE QB (CHAPTER 5.1), WHILE *JEOPARDY!* ALSO USES SIGNALING DEVICES, THESE ONLY WORK ONCE THE QUESTION HAS BEEN READ IN ITS ENTIRETY; KEN JENNINGS (ALSO A FORMER QB PLAYER WHILE HE WAS A STUDENT AT BYU) HIMSELF EXPLAINS IT ON A *PLANET MONEY* INTERVIEW (MALONE, 2019):

Jennings: THE BUZZER IS NOT LIVE UNTIL ALEX FINISHES READING THE QUESTION. AND IF YOU BUZZ IN BEFORE YOUR BUZZER GOES LIVE, YOU ACTUALLY LOCK YOURSELF OUT FOR A FRACTION OF A SECOND. SO THE BIG MISTAKE ON THE SHOW IS PEOPLE WHO ARE ALL ADRENALINIZED AND ARE BUZZING TOO QUICKLY, TOO EAGERLY.

Malone: OK. TO SOME DEGREE, *JEOPARDY!* IS KIND OF A VIDEO GAME, AND A CRAPPY VIDEO GAME WHERE IT'S, LIKE, LIGHT GOES ON, PRESS BUTTON—THAT'S IT.

Jennings: (LAUGHTER) YEAH.

JEOPARDY!'S BUZZERS ARE A GIMMICK TO ENSURE GOOD TELEVISION; HOWEVER, QB BUZZERS DISCRIMINATE KNOWLEDGE.

SO HOW DOES THIS INTERACT WITH WATSON? WATSON RECEIVES ALL OF THE CLUES ELECTRONICALLY AND LIKEWISE GETS AN ELECTRONIC SIGNAL TO KNOW WHEN IT IS SAFE TO BUZZ IN, THEN COMPUTES A PROBABILITY OF BEING RIGHT AND BUZZES IN IF IT'S ABOVE THAT THRESHOLD (SECTION 6.2). IN CONTRAST, HUMANS HAVE TO EITHER GUESS WHEN IT IS SAFE TO BUZZ OR WAIT FOR A LIGHT TO TURN ON.

JEOPARDY! GURUS EXPLICITLY ADVISE NEW PLAYERS NOT TO WAIT FOR THE LIGHT—YOUR PUNY HUMAN REFLEXES ARE TOO SLOW. INDEED, ONE OF KEN JENNINGS'S STRENGTHS WAS HIS UNCANNY ABILITY TO INTERNALIZE THE CADENCE OF ALEX'S VOICE AND WHEN A TECHNICIAN WOULD ACTIVATE THE BUZZER (JENNINGS, 2006). IN CONTRAST, WATSON IS LITERALLY AN ELECTROMECHANICAL BUZZING MACHINE THAT COULD GET FIRST CRACK AT EVERY QUESTION IT WANTS.⁹

UNFORTUNATELY, DESPITE SUBJECTING EVERYONE WITHIN EARSHOT TO THESE RANTS, THE COMPUTER SCIENCE COMMUNITY THINKS THAT THE QUESTION IS SETTLED: COMPUTERS ARE BETTER THAN HUMANS AT ANSWERING QUESTIONS. THIS IS DESPITE KEN BASICALLY SAYING THAT IT DID, INDEED, COME DOWN TO THE BUZZER.

MOREOVER, WHAT MAKES FOR A “NORMAL” DIFFICULTY QUESTION FOR A HUMAN DOES NOT ALWAYS APPLY TO A COMPUTER. LET'S FIRST TALK ABOUT WHAT MAKES A CLUE EASIER FOR A COMPUTER AND THEN WE'LL TALK ABOUT WHAT MAKES A CLUE HARDER FOR A COMPUTER.

Easy for a Computer. WHEN I APPEARED ON *JEOPARDY!*, MY FINAL *JEOPARDY!* WAS:

⁹ In practice, this is not always true. Because Watson computed its responses in real time, it could not come up with a response in time for particularly short questions.

AFTER THIS WOMAN'S DEATH, HER DAUGHTER
WROTE, "AS FAR AS WE IN THE FAMILY ARE CON-
CERNED, THE ALPHABET NOW ENDS AT Y"

ALL OF US GOT THE QUESTION RIGHT; IT JUST SO HAPPENED THAT *JEOPARDY!*
USED A VERY SIMILAR CLUE THAT AIRED AS WE WERE RECORDING:

"G" IS FOR GRAND MASTER AS WELL AS THIS
WOMAN WHO RECEIVED THE 2009 GRAND MASTER
AWARD.

THE CORRECT RESPONSE IS OF COURSE SUE GRAFTON. FOR POORLY READ CONTESTANT LIKE MYSELF, ONLY STUDYING PREVIOUS CLUES ALLOWED US TO GET THE ANSWER RIGHT. I'VE NEVER READ A GRAFTON BOOK, BUT I KNOW SHE WRITES MYSTERY BOOKS AND HAS TITLES OF THE FORM "A" IS FOR *ALIBI* (AND THAT'S THE ONLY TITLE I CAN THINK WITHOUT LOOKING AT WIKIPEDIA).

FOR WATSON, THIS IS "MEMORIZATION" IS TRIVIAL: A SINGLE LETTER IMPLIES THAT THE ANSWER IS GRAFTON. BUT JUST LIKE YOU SHOULD NOT THINK THAT I'M SMART FOR GETTING LUCKY TO HAVE SEEN A REUSED CLUE, YOU SHOULD NOT PRAISE WATSON FOR FINDING NEAR-REPEAT QUESTIONS. AND IT'S NOT JUST TRIVIA GAMES: GOOGLE'S DATASET OF QUESTIONS (WHICH WE'LL TALK

ABOUT IN MORE CHAPTER **?reference?** ??) HAVE MANY IDENTICAL QUESTIONS (LEWIS ET AL., 2021), WHICH MAKES IT AN IMPERFECT YARDSTICK. MOREOVER, WATSON CAN ALSO STORE THE ENTIRETY OF WIKIPEDIA, EASILY LOOKING UP CAPITALS, AUTHORS, ETC.

INDEED, WHEN A COMPUTER CAN FIND AN EXACT QUOTE (AS WAS IN MY FINAL *JEOPARDY!* CLUE), THE QUESTION BECOMES EVEN EASIER. THEN THE COMPUTER JUST NEEDS TO FIND THE APPROPRIATE ARTICLE THAT CONTAINS THE QUOTE AND THEN JUST FIND WHATEVER ENTITY IS MOSTLY LIKELY TO BE A *JEOPARDY!* ANSWER.

WHERE SYSTEMS LIKE WATSON STRUGGLE ARE ON COMPUTATION, MATCHING NOVELS AND MOVIES TO PLOTS, COMBINING MULTIPLE CLUES, LATERAL THINKING, AND WORDPLAY (KAUSHIK AND LIPTON, 2018). AND THIS IS NOT JUST A MATTER OF DEGREE: COMPUTERS STRUGGLE WITH ALL SUCH QUESTIONS, EVEN IF THEY'RE IN THE TOP ROW OF *JEOPARDY!* WHILE A COMPUTER IS THEORETICALLY GOOD AT MATH, THE KINDS OF PROGRAMS THAT ANSWER TRIVIA QUESTIONS STRUGGLE ANSWERING MATCH QUESTIONS WITH NUMBERS IN THE DOUBLE DIGITS (WALLACE ET AL., 2019B).

THIS IS WHY THE GOAL OF AI IS GENERAL ARTIFICIAL INTELLIGENCE (CHAPTER 3): WHILE WE CAN BUILD SPECIALIZED SYSTEMS FOR *JEOPARDY!* CLUES OR MATH PROBLEMS, UNLIKE A REASONABLY SMART HUMAN, A SINGLE PROGRAM CAN'T "DO IT ALL". UNLIKE FOR HUMAN CONTESTANTS, THE "DIFFICULTY" OF *JEOPARDY!* QUESTIONS FOR A COMPUTER HAS NO RELATIONSHIP TO THE NOMINAL VALUE. WE TALK ABOUT HOW FACEBOOK/META DEALT WITH THIS PROBLEM IN CHAPTER **?reference?** ??: TELL THE AUTHORS WHAT'S HARD

FOR A COMPUTER.

6.4 *Two Nice Guys, One Computer with no Shame*

GIVEN THE “TOO EASY” QUESTIONS AND THE BUZZER, WHAT DOES THIS ACTUALLY MEAN FOR GAMEPLAY? A QUESTION COMES IN, AND WATSON HAS THE CHOICE OF ANSWERING IT OR NOT: IT CAN WIN EVERY RACE TO THE BUZZER IF IT WANTS. THEN, OF THE THINGS IT CANNOT ANSWER, KEN AND BRAD FIGHT OVER THE SCRAPS. THUS, FOR A COMPUTER TO WIN THIS COMPETITION, IT NEEDS ONLY TO BE ABLE TO ANSWER A THIRD OF THE QUESTIONS CORRECTLY.

NOW, THE *JEOPARDY!* NERDS READING THE BOOK (*I LOVE YOU ALL*), WILL POINT OUT THAT THIS ISN’T TRUE, BECAUSE THE CLUES ARE NOT WEIGHTED EQUALLY: SOME ARE WORTH MORE THAN OTHERS. HOWEVER, AS WE DISCUSSED ABOVE, WHAT’S DIFFICULT FOR A COMPUTER ISN’T ALWAYS DIFFICULT FOR A HUMAN (AND VICE VERSA), SO IT REALLY IS A RANDOM THIRD OF THE QUESTIONS. WHILE A GOOD HUMAN PLAYER MIGHT BE WEAK ON THE BUZZER AND BE CONFIDENT THAT IF THEY KNOW MORE THEY’LL WIN THE HARDER CLUES, THIS ISN’T TRUE FOR A HUMAN FACING OFF AGAINST WATSON.

MOREOVER, THE COMPUTER HAS NO SHAME: IT USES A STRATEGY CALLED THE “FORREST BOUNCE” (ROGAK, 2020, MORE INFAMOUSLY ASSOCIATED WITH JAMES HOLZHAUER AND ARTHUR CHU). RATHER THAN GOING THROUGH THE CATEGORIES TOP TO BOTTOM (EASY TO HARD), WATSON GOES THROUGH THE CLUES SOMEWHAT RANDOMLY, SEARCHING FOR DAILY DOUBLES AND TRYING TO OPTIMIZE ITS SCORE. AGAIN, THERE’S NOTHING WRONG WITH THIS—IT’S THE OPTIMAL STRATEGY! BUT IF IT’S THE “RIGHT” WAY TO PLAY THE GAME, WHY DOESN’T EVERY HUMAN DO IT?

THAT’S BECAUSE HUMANS WANT TO FOLLOW SOCIAL NORMS. THE PRODUCERS OF THE SHOW TELL YOU UP AND DOWN THAT YOU SHOULDN’T PLAY THE GAME LIKE THAT, AND YOU DON’T WANT TO MAKE THEM UNHAPPY WITH YOU... THEY CAN MAKE YOUR LIFE MISERABLE. I REMEMBER WATCHING *JEOPARDY!* WITH MY GRANDMOTHER AND WHEN SOMEONE HUNTED FOR THE DAILY DOUBLE, SHE WOULD ALWAYS SAY “WHO DOES HE THINK HE IS” (AND IT WAS ALWAYS A HE). ALEX TREBEK ALSO WASN’T A FAN (MARCHESE, 2018):¹⁰

WHEN THE SHOW’S WRITERS CONSTRUCT CATEGORIES THEY DO IT SO THAT THERE’S A FLOW IN TERMS OF DIFFICULTY, AND IF YOU JUMP TO THE BOTTOM OF THE CATEGORY YOU MAY GET A CLUE THAT WOULD BE EASIER TO UNDERSTAND IF YOU’D BEGUN AT THE TOP OF THE CATEGORY AND SAW HOW THE CLUES WORKED.
I LIKE THERE TO BE ORDER ON THE SHOW, BUT AS THE IMPARTIAL HOST I ACCEPT DISORDER.

AND NOBODY, NOBODY, WANTS TO MAKE ALEX UNHAPPY.

KEN WOULD SOMETIMES DO A LITTLE HUNDING FOR THE DAILY DOUBLE AGAINST A PARTICULARLY FORMIDABLE OPPONENT, BUT HE WOULD NORMALLY BE WELL-BEHAVED SO AS NOT TO UPSET THE POWERS THAT BE. WATSON, HOWEVER, WAS A SOULLESS MACHINE; AND HAVING A MACHINE ON THE STAGE

¹⁰ Rogak (2020) quotes a saltier take Trebek offered to Howard Stern: “It only works, dickweed, if you know the correct response to everything that’s up there.”

WAS NO EXCITING THAT NOBODY FAULTED IT FOR ITS STRATEGY. IF ANYONE IS TO BLAME, IT'S PROBABLY GARY TESAURO; ADDING HIS STRATEGY FOR PLAYING THE BOARD INCREASED WATSON'S WIN PERCENTAGE CONSIDERABLY ([TESAURO ET AL., 2013](#)). BUT IF YOU PUT HIM IN A ROOM WITH THE WITHERING STARES OF *JEOPARDY!* PRODUCERS (OR WORSE, ALEX TREBEK) AND THAT CODE WOULD BE DELETED IN NO TIME.

6.5 *The Legacy of Watson*

LET'S REVIEW WATSON'S APPEARANCE ON *JEOPARDY!*:

1. ALL QUESTIONS ARE OF "NORMAL" DIFFICULTY;
2. THUS THE TWO HUMAN CONTESTANTS KNOW NEARLY ALL OF THE CLUES; BUT
3. WATSON CAN WIN THE BUZZER RACE WHENEVER IT WANTS.

IF WATSON WINS SUCH A MATCH, DOES THAT MEAN THAT IT IS SUPERIOR TO THESE SUPURB HUMANS?

I HOPE THAT YOU ARE RELUCTANT TO ANSWER "YES" (NOT JUST BECAUSE *JEOPARDY!* HAS TRAINED YOU TO RESPOND TO ANSWERS WITH QUESTIONS). PERHAPS I'VE PLANTED A SEED OF DOUBT: *NO, WE CANNOT YET CONCLUDE COMPUTER SUPERIORITY FROM THIS EXPERIMENT.* AND THIS IS NOT THE END OF THE STORY FOR JUDGING WHETHER COMPUTERS ARE SUPERHUMAN IN THEIR INTELLIGENCE. SINCE THEN, WE'VE SEEN CLAIMS THAT COMPUTERS ARE SMOKING LAWYERS IN TAKING THE LSAT OR IS BETTER THAN RADIOLOGISTS IN READING X-RAYS. WATSON WAS JUST THE BEGINNING OF THE STORY, AS SINCE THEN WE'VE SEEN BOTH A TREMENDOUS IMPROVEMENT IN THE TECHNOLOGIES THAT DRIVE AI (WE REVIEW THESE ADVANCES IN CHAPTER [?reference?](#) ??) AND A GREATER INTEREST IN MEASURING HOW SMART THESE MODELS ARE. WATSON WAS JUST THE BEGINNING OF THE STORY, IN BOTH RESPECTS.

IN THE NEXT CHAPTERS, WE GO BEYOND THE METHODS THAT WATSON USED TO THE MODERN QA HAS ACTUALLY REACHED POSSIBLE PARITY WITH HUMANS AND HOW TO ACTUALLY MEASURE HOW FAR AI HAS COME SINCE KEN JENNINGS LOST TO WATSON (CHAPTER [10](#)).

7

FORMATS BEYOND STAND-ALONE QA

HOW DO YOU ASK A QUESTION: IS IT TEXT, A PICTURE, OR A CONVERSATION? THIS CHAPTER REVIEWS THE DIFFERENT FORMS QUESTION ANSWERING CAN TAKE AND WHAT COMPLEXITIES THAT CAN INTRODUCE.

TOPICS IN CHAPTER:

1. **TASK VS. FORMAT.**
2. **WHAT IS A QUESTION?.**
3. **WHERE DO YOU GET AN ANSWER?.**
4. **WHAT DID YOU MEAN?.**
5. **A PICTURE IS WORTH A THOUSAND QUESTIONS.**
6. **QUESTIONS BEYOND ENGLISH.**
7. **THE SKILLS NEEDED TO ANSWER QUESTIONS.**
8. **A TAXONOMY OF QUESTION ANSWERING.**

7.1 Multihop QA

THUS FAR, WE'VE TALKED ABOUT MODELS THAT CAN TAKE A QUESTION, FIND SOME EVIDENCE, AND EXTRACT THE ANSWER. THAT WORKS FOR THINGS LIKE SQuAD, TRIVIAQA, AND (MANY) NATURAL QUESTIONS.

BUT WHAT ABOUT THE MULTIHOP DATASETS THAT WE TALKED ABOUT BEFORE? HOW WOULD THESE SORTS OF MODELS DO A HOP? JUST AS A REFRESHER (WATCH THE PREVIOUS VIDEO IF YOU WANT MORE), IN DATASETS LIKE HOTPOTQA, YOU NEED TO ANSWER QUESTIONS LIKE "WHERE WAS FACEBOOK FOUNDED" BY HOPPING THROUGH EVIDENCE DOCUMENTS. WHILE THE APPROACHES WE TALKED ABOUT CAN HANDLE ONE HOP JUST FINE, HOW WOULD YOU GO ABOUT SOLVING THESE PROBLEMS WHEN YOU NEED TO FIND *MULTIPLE* PIECES OF EVIDENCE?

WE WON'T GO AS DEEP INTO THE WEEDS ON THESE TECHNIQUES BECAUSE THIS IS RELATIVELY EARLY DAYS; AS I RECORD THIS IN 2022, THE DETAILS ARE PROBABLY GOING TO CHANGE DRAMATICALLY. BUT BARRING SURPRISES, HOPEFULLY THE BROAD STROKES ABOUT THE GENERAL IDEAS OF HOW TO ANSWER THESE MULTIHOP QUESTIONS WILL STAY RELATIVELY SIMILAR. NONETHELESS,

I'LL GIVE THE TITLES AND AUTHORS AS WE GO SO THAT IF YOU'RE INTERESTED IN THE PAPERS YOU CAN DIG IN A LITTLE BIT MORE.

WE TALKED ABOUT THAT BEFORE WITH SO-CALLED "MULTIHOP" QUESTIONS. HOW DO SYSTEMS ACTUALLY ANSWER THOSE QUESTIONS? AS YOU CAN IMAGINE, THIS REQUIRES FINDING MULTIPLE PIECES OF EVIDENCE. LET'S FIRST TALK ABOUT DIFFERENT STRATEGIES FOR GATHERING THE EVIDENCE: YOU CAN FOLLOW STRUCTURED LINKS, YOU CAN GENERATE REAL TEXT QUERIES, OR YOU CAN GENERATE VECTOR-BASED QUERIES.

LET'S FIRST TALK ABOUT FOLLOWING STRUCTURED LINKS. THIS IS PROBABLY THE MOST OBVIOUS AND IS PROBABLY THE MOST SIMILAR TO WHAT YOU'D DO IF YOU AS A HUMAN TYPED IN A QUERY, FOUND A PAGE, AND THEN CLICKED ON THE LINK THAT GOT YOU A LITTLE BIT CLOSER TO WHAT YOU THOUGHT THE ANSWER WOULD BE. OF COURSE, KNOWING WHICH LINK TO EXPLORE TAKES A LITTLE TRIAL AND ERROR. WE DID SOMETHING LIKE THIS IN 2020 WITH A MODEL CALLED DELFT, AND THIS PAPER FROM CHEN ET AL. EXPLORES DIFFERENT WAYS OF FOLLOWING LINKS: FOLLOWING A LINK ON THE WIKIPEDIA PAGE OR RESOLVING A COREFERENCE CHAIN ... WHICH IS A LOT LIKE THE EXAMPLES THAT WE SHOWED IN OUR LAST VIDEO INTRODUCING MULTIHOOP (FOR EXAMPLE, GETTING THE ANSWER TO SUSANAH AND THE ELDERS, REQUIRES FIGURING OUT WHO "HER" IS). THEN MY FORMER STUDENT CHEN ZHAO FOLLOWED UP OUR WORK WITH A COLLABORATION WITH MICROSOFT CALLED TRANSFORMERXH. THIS SEEMS LIKE THE MOST POPULAR APPROACH.

BUT SOMETIMES YOU CAN'T FIND A LINK THAT EASILY. THERE COULD BE A DIFFERENT STRATEGY THAT YOU COULD USE TO SOLVE A MULTIHOOP QUESTION. IF YOU TYPED INTO THE WIKIPEDIA SEARCH BAR TO GET THE FIRST QUERY, WHY NOT JUST TYPE A NEW QUERY? THIS IS THE APPROACH TAKEN BY QI ET AL WITH THE WONDERFULLY NAMED MODEL "GOLDED RETRIEVER". IT LEARNS HOW TO FORMULATE QUERIES *CONDITIONED* ON THE ORIGINAL QUERY, EVIDENCE YOU GOT FROM THE FIRST QUERY, TO GET TO THE NEXT PIECE OF EVIDENCE. FOR EXAMPLE HERE, TO SOLVE THIS QUESTION IT EXPLICITLY CREATES THE QUERY "ARMADA AUTHOR" TO FIGURE OUT THAT IT WAS WRITTEN BY ERNEST CLINE.

ONE NICE THING THIS ABOUT THIS IS THAT THESE APPROACHES ARE PRETTY INTERPRETABLE: KNOWING WHAT SEARCHES AND EVIDENCE LED TO THE ANSWER LETS A USER INSPECT THE CHAIN TO KNOW IF IT'S GOOD OR NOT. THIS PAPER FROM INOUE ET AL. GOES EVEN FURTHER BY DOING ABSTRACTIVE SUMMARIZATION TO CREATE SENTENCES THAT FIND THE MINIMAL INFORMATION THAT'S SUFFICIENT FOR ANSWERING THE QUESTION. SO THEN FOR THE QUESTION "CHARLIE ROWE PLAYS BILLY COSTA IN A FILM BASED ON WHAT NOVEL?" YOU GET THIS NICE ABSTRACTED EXPLAINER THAT SAYS "CHARLIE ROWE PLAYS BILLY COSTA IN THE GOLDEN COMPASS. THE GOLDEN COMPASS IS A FILM BASED ON THE NOVEL NORTHERN LIGHTS." PRETTY GOOD!

BUT DO YOU ACTUALLY NEED TO PUT YOUR FINGERS ON THE KEYBOARD TO GET TO THE EVIDENCE? WE SAW WITH DPR THAT YOU CAN ACTUALLY DO BETTER WITH VECTOR-BASED QUERIES. SO WHY CREATE AN EXPLICIT STRING WHEN IT'S

JUST GOING TO GET TURNED INTO A VECTOR ANYWAY? THAT'S THE APPROACH OF XIONG ET AL WHO JUST DIRECTLY OPTIMIZE THESE VECTOR QUERIES.

THE CHALLENGE IS THAT ALL OF THESE APPROACHES NEED TO BE TRAINED SOMEHOW. LIKE FOR ANY MACHINE LEARNING PROBLEM, YOU NEED TRAINING DATA. EXISTING TECHNIQUES FOR SOLVING THIS USE THE EXPLICIT CORRECT HOPS FROM DATASETS LIKE MultiHop. BUT COULD WE DO THIS IN DATASETS LIKE QUIZ BOWL WHERE THE EVIDENCE IS NOT ANNOTATED?

LET'S THINK THIS THROUGH AND SEE WHERE IT BREAKS DOWN. YOU TAKE YOUR ORIGINAL QUESTION, TURN IT INTO A VECTOR, GET SOME SEARCH RESULTS, AND THEN FROM THAT YOU GENERATE A NEW VECTOR QUERY AND YOU GET SOME RESULTS.

THE FIRST CHALLENGE IS HOW DO YOU KNOW IF A PIECE OF EVIDENCE IS GOOD OR NOT? THIS IS SOMETHING THAT MY FORMER STUDENT CHAN ZHAO LOOKED AT IN AN EMNLP PAPER FROM 2021.

ONE CLUE IS THAT IF THE ANSWER ISN'T IN A PASSAGE, THEN IT CANNOT BE BE RIGHT. AND WE ASSUME THAT WE DO KNOW THE ANSWER TO THE QUESTION ... WE JUST DON'T KNOW HOW TO GET THERE. THIS HOLDS TRUE FOR JUST ABOUT ANY DATASET, SO IT'S NOT TOO CRAZY AN ASSUMPTION. BUT YOU DO NEED TO BE ABLE TO COPE WITH DIFFERENT FORMS OF THE ANSWER. IF THE ANSWER IS "TIM COOK", YOU SHOULD BE ABLE TO ACCEPT "TIMOTHY COOK" OR "TIMOTHY DONALD COOK".

BUT THEN ANY PASSAGE THAT CONTAINS THE ANSWER AS A SUBSTRING IS A GOOD ANSWER. "NEW JERSEY" IS THE ANSWER TO BOTH THE QUESTION "IN WHAT STATE IS RUTGERS LOCATED" AND "WHICH STATE IS THE HOME OF FOUNTAINS OF WAYNE", BUT WE ONLY WANT THE STRING WHEN IT ANSWERS THE APPROPRIATE QUESTION. BUT THANKS TO DATASETS LIKE SQuAD, THERE ARE EXISTING MODELS FOR ANSWERING QUESTIONS *GIVEN* A PASSAGE IN THE STANDARD MACHINE READING CONTEXT, SO WE CAN GET A SCORE FOR EACH POTENTIAL PIECE OF EVIDENCE: HOW WELL DOES IT ANSWER THE QUESTION?

THIS IS IMPORTANT SO LET'S GO OVER THIS WITH AN EXAMPLE. LET'S SAY WE WANT TO ANSWER THE QUESTION "WHAT STATE DOES SANG-WOOK CHEONG WORK AS A MATERIALS SCIENTIST?". WE GET OUR SECOND HOP QUERY AFTER THE FIRST HOP SAYS HE WORKS AT RUTGERS. THE SECOND HOP NEEDS TO TELL US THE RUTGERS IS IN NEW JERSEY. WE CAN ALWAYS SAY THAT EVERYTHING WITHOUT "NEW JERSEY" ISN'T A GOOD PIECE OF EVIDENCE. IN ADDITION, ANYTHING THAT HAS THE "NEW JERSEY" BUT DOESN'T ACTUALLY ANSWER THE QUESTION ISN'T A GOOD PIECE OF EVIDENCE: LIKE THIS THING ABOUT FOUNTAINS OF WAYNE, WHICH IS ALSO IN NEW JERSEY. AND THINGS THAT ANSWER THE QUESTION AND LOOK LIKE AN ANSWER TO THE QUESTION WE KEEP.

THIS NOW FORMS A CYCLE: START WITH *SOME* IR-BASED QA MODEL, GATHER SOME EVIDENCE, LABEL IT AS WE JUST TALKED ABOUT, AND RETRAIN A READER, AND MORE IMPORTANTLY A RETRIEVER. THEN START THE CYCLE AGAIN. AND YOU CAN ALSO TRAIN A RERANKER TO SQUEEZE OUT A FEW MORE POINTS.

SO LET'S SEE IF THIS ACTUALLY WORKS. COMPARED TO TRADITIONAL IR,

TRAINED MODELS THAT USE THE EXPLICIT EVIDENCE FROM HotPotQA FOR FULLY SUPERVISED TRAINING LIKE TRANSFORMER XH FROM MSR CAN DO MUCH BETTER, GETTING THE ANSWER IN THE TOP 10 RESULTS 90

NONETHLESS, IT'S STILL WELL SHORT OF THE CREATIVITY AND THE DEPTH THAT MOTIVATED, SMART HUMANS CAN. IN OTHER WORDS, IT STILL WELL SHORT OF OUR GOAL OF AUTOMATING DAN RUSSELL. AND EVEN WITHIN THE LIMITED WORLD OF HotPotQA, THERE'S SOME RECENT WORK SHOWING THAT MODELS ARE TAKING SHORTCUTS TO GET TO THE ANSWER AND NOT CREATING COMPLETE REASONING CHAINS.

JIANG AND BANSAL SHOW THAT THERE ARE REASONING SHORTCUTS THAT MODELS CAN TAKE. FOR EXAMPLE, IN THIS DOCUMENT, IF YOU HAVE THE QUESTION "WHAT WAS THE FATHER OF KASPER SCHMEICHEL VOTED TO BE BY THE IFFHS IN 1992". FIRST, LET'S TAKE A MOMENT TO STARE AT THE ODD PHRASING: YOU WOULDN'T GET THIS IN A MANCHESTER PARADIGM QUESTION. BUT YOU CAN GRAB THE ANSWER JUST IN A SINGLE DOCUMENT. JIANG AND BANSAL CREATE NEW DOCUMENTS THAT FORCE TRUE MULTIOP REASONING: I'D RATHER CREATE BETTER QUESTIONS RATHER THAN MANUFACTURE EVIDENCE TO MAKE THE PROBLEM HARDER. BUT NONTHELESS, IT SHOWS THE LIMITATIONS OF HotPotQA.

SIMILARLY, TRIVEDI ET AL. CREATE A METRIC CALLED "DiRe" FOR "DISCONNECTED REASONING" TO MEASURE WHETHER A SYSTEM GETS THE WRONG ANSWER WHEN YOU REMOVE RELEVANT EVIDENCE. A GOOD MultiHop DATASET SHOULD BE DIFFICULT ENOUGH TO NOT ALLOW THIS SORT OF CHEATING. IN OTHER WORDS, THIS IS A METRIC TO SHOW HOW OFTEN THE ISSUES THAT JIANG AND BANSAL EXPOSE SHOW UP.

ALL OF THESE OPEN QUESTIONS MAKE MULTIOP QUESTION ANSWERING ONE OF THE MOST EXCITING AREAS FOR QA RESEARCH AT THE MOMENT. BECAUSE THE ANSWERS ARE STILL SHORT (MOSTLY FACTOID), IT DOESN'T HAVE THE SAME ISSUES THAT MAKE LONG FORM QA EVALUATION SO FIENDISHLY DIFFICULT, BUT THE MODELING CHALLENGES ARE STILL REAL. AND NOW THAT THERE ARE APPROACHES THAT GET BEYOND THE REQUIREMENTS OF GOLD EVIDENCE, WE CAN MOVE BEYOND DATASETS WITH GOLD EVIDENCE LABELS.

AND I THINK WE HAVE A LOT TO LEARN FROM HUMANS HERE—EXPLICITLY INCORPORATING HUMAN STRATEGIES FOR SOLVING MULTIOP CHALLENGES TO SYNTHESIZE VECTOR QUERIES, FOLLOWING LINKS, OR SYNTHESIZING RAW TEXT QUERIES.

AND THIS CAN BE PAIRED WITH SOME OF THE FANCY MACHINE LEARNING TOOLS THAT WE HAVE LIKE REINFORCEMENT LEARNING FOR UNCOVERING THE BEST ACTIONS GIVEN SPARSE FEEDBACK ... COULD THIS GET US BETTER MODELS FROM MORE COMPLICATED BUT MORE LIGHTLY ANNOTATED DATASETS?

AND THEN PERHAPS WE BETTER ADDRESS THE "JOY OF SEARCH" PROBLEMS IN A HYBRID COOPERATION BETWEEN HUMANS AND COMPUTERS: WHERE COMPUTERS AND HUMANS COULD JOIN HANDS AND HOP TOGETHER.

MULTIHOP AS A DATASET

[HTTP://SEARCHRESEARCH1.BLOGSPOT.COM/2021/11/ANSWER-WHY-IS-CARQUINEZ-STRAIT-SO.HTML](http://SEARCHRESEARCH1.BLOGSPOT.COM/2021/11/ANSWER-WHY-IS-CARQUINEZ-STRAIT-SO.HTML)

IF YOU'VE WATCHED OUR OTHER VIDEOS ABOUT MACHINE READING QUESTION ANSWERING, WHETHER IT'S MANCHESTER OR CRANFIELD, ALL OF THE EXAMPLES THUS FAR HAVE BEEN A ONE AND DONE AFFAIR: DO A SEARCH, GET SOME TEXT, EXTRACT THE ANSWER FROM THAT. BUT THAT ISN'T ALWAYS REALISTIC. AS ANYONE WHO HAS GONE DOWN A WIKIPEDIA HOLE KNOWS, SOMETIMES YOU NEED SEVERAL PIECES OF INFORMATION TO FULLY ANSWER A QUESTION.

FIRST, WHAT IS THE HUMAN STATE OF THE ART HERE? RESEARCH LIBRARIANS ARE PROBABLY THE BEST SEARCHERS WHEN YOU HAVE A REALLY TRICKY QUESTION YOU HAVE TO ANSWER. FRIEND OF UMD'S LIBRARY SCHOOL GOOGLE'S DAN RUSSEL HAS WRITTEN A BOOK CALLED "THE JOY OF SEARCH" THAT SHOWCASES THE PROCESS FOR FIGURING OUT TRICKY QUESTIONS.

HE GOES THROUGH INDIVIDUAL QUESTIONS ON HIS BLOG, WHICH I ENCOURAGE YOU TO READ: THERE'S A LINK TO THE PROCESS HERE. I DON'T WANT TO GO THROUGH THE WHOLE SOLUTION ... THAT'S DAN'S THING, AND YOU SHOULD CHECK OUT HIS BLOG POST, BUT I DO WANT TO FOCUS ON *WHY* ANSWERING "WHY IS THE CARQUINEZ STRAIT SO UNDERDEVELOPED" IS SO DIFFICULT. FIRST, THE QUESTION IS A COUNTERFACTUAL ABOUT SOMETHING (DEVELOPMENT) THAT *DIDN'T* HAPPEN. SO TO ANSWER THE QUESTION, YOU FIRST NEED TO FIND OUT THE OLD NAME FOR THIS PROPERTY OUTLINED IN RED, WHO OWNED IT IN THE 19TH CENTURY, SEE WHOM THEY SOLD IT TWO IN THE 20TH CENTURY, AND CROSS-REFERENCE THAT TO ZONING DECISIONS THAT GOT MADE AT THE APPROPRIATE COUNTY LEVEL.

THIS IS NOT SOMETHING THAT COMPUTERS CAN CURRENTLY DO. AND I DON'T THINK THAT WE'RE GOING TO GET ANYWHERE CLOSE IN THE NEAR FUTURE. BUT I DO WANT TO TALK TODAY ABOUT HOW WE CAN INCH IN THAT DIRECTION.

THE CURRENT TERM FOR THIS SORT OF QA FOR COMPUTERS IS CALLED "MULTIHOP" QA. THESE QUESTIONS ARE DESIGNED TO TAKE MULTIPLE STEPS—OR HOPS—to GET TO THE RIGHT ANSWER. THE MOST PROMINENT DATASET FOR MULTIHOP QUESTIONS IS HotPotQA. WHICH WHILE NOT AS CHALLENGING AS THE QUESTION ABOUT BAY AREA LAND USE, DO REQUIRE MULTIPLE STEPS TO FIND THE ANSWER. IN PARTICULAR, IN THEORY YOU NEED TO "HOP" FROM ONE WIKIPEDIA PAGE TO ANOTHER TO FIND THE ANSWER. IN THIS EXAMPLE, TO FIND THE CITY FACEBOOK LAUNCHED IN, YOU NEED TO "HOP" TO THE PAGE ON HARVARD.

THE WAY THE HotPotQA DATASET WAS AUTHORED WAS THAT THE AUTHORS SHOWED TWO LINKED WIKIPEDIA PAGES AND ASKED CROWDWORKERS TO WRITE A QUESTION THAT REQUIRED INFORMATION FROM BOTH PAGES TO ANSWER.

THERE'S BOTH AN IR PROBLEM AS WELL AS A MACHINE READING PROBLEM. HopPotQA HAS TWO EVALUATION SETTINGS: A "FULL WIKI" SETTING WHERE YOU HAVE TO FIND THE EVIDENCE YOURSELF TO ANSWER THE QUESTION AND A "DISTRACTOR" SETTING WHERE YOU ARE GIVEN SOME EVIDENCE AND YOU HAVE TO FIGURE OUT HOW TO EXTRACT THE CORRECT HOP OUT OF THAT LIMITED SET.

Now, multihop isn't unique to this dataset. There are multihop questions in, for example, in my favorite question answering format, Quiz Bowl. But the great thing about Quiz Bowl is that if you can't solve the early multihop clues, then you can get the question at the end with clues that don't require multihop reasoning. This makes it, despite its inherent complexity, a more accessible format for long term QA research.

Let's work through trying to solve them, using Wikipedia paragraphs as our source documents. And hopefully this will show you how hard the problem is.

Okay, the first question is "In a play from this country, the audience is given the option to blow up four policemen by the maniac." Okay, so we can find this search result, which is from the page about the play "The Accidental Death of an Anarchist". But we don't know who wrote it, so we can't answer the question. So let's find out more about this play.

Aha, it's by Dario Fo, and what do you know it tells us right here that he's Italian. Oh, but what we're looking for is an exact match answer and our answer key is "Italy"? If you're using an evaluation metric like exact match, that means we have to keep going. So let's find out more about Dario Fo. You have to go three paragraphs deep in his Wikipedia page to find the string "Italy". We'll talk more about different ways to solve this problem when we talk about generation.

But let's take a look at how Quiz Bowl self-annotates these multi-hop links. The same links that we followed in looking up information about Accidental Death of An Anarchist to figure out it was written by Dario Fo to figure out he's Italian are right here at the end of the question. And if you can't do that, you can just look up the nationality of this guy named "Luigi".

Let's move on to another question. This time about this painting. The question is: "A Lorenzo Lotto painting of this woman includes a cartouche showing her response to an accusation." Okay, so let's see what turns up with a search of Wikipedia. The first step is to figure out which one of these is correct.

In the first evidence, you need to know that a signature and a date isn't a response to an accusation, so that's out. Then you need to know that a cartouche on "his" beret is inconsistent with the question's use of "this woman", so that just leaves the last piece of evidence, where you have a woman and dialog in a cartouche. This is why the distractor setting in HotPotQA can be so tricky! A search result can turn up a bunch of seemingly relevant material, but you need to pick the right answer out of it somehow. What makes the distractor setting of HotPotQA more tractable than Quiz Bowl is that the answer is in there *somewhere* which is not guaranteed in

QUIZ BOWL. MOREOVER, IN QUIZ BOWL THE CORRECT PIECES OF EVIDENCE ARE NOT ANNOTATED.

OKAY, NOW THE CHALLENGE IS TO FIGURE OUT WHO THIS “HER” IS HERE. SO IF YOU GO TO THE START OF THIS ARTICLE, YOU CAN SEE THAT THIS IS SUSANNA FROM THE BOOK OF DANIEL. SO OFTEN, THIS MULTIHOP REASONING IS ESSENTIALLY COREFERENCE RESOLUTION EITHER WITHIN A DOCUMENT OR ACROSS DOCUMENTS.

SO HOW CAN THIS BE ACROSS DOCUMENTS? THAT’S NOT HOW PRONOUNS WORK! TO SHOW YOU WHAT I MEAN, LET’S TAKE A LOOK AT A QUIZ BOWL EXAMPLE THAT DOES REQUIRE KNOWING WHO ENTITIES ARE ACROSS DOCUMENTS.

A MECHANIC WAS ARRESTED IN ONE OF THESE TWO COUNTRIES FOR USING WOMEN’S UNDERWEAR SOAKED IN INVISIBLE INK TO PASS SECRETS TO THE OTHER.

WE NEED TO FIGURE OUT WHAT THESE TWO COUNTRIES ARE. SO LET’S DO A QUERY. SO WE GET THIS PRETTY UNAMBIGUOUS RESULT. BUT WE DON’T KNOW THE ANSWER, WHICH NEEDS TO BE TWO COUNTRIES! WHAT’S THIS MOSSAD THING? AND IS CAIRO A COUNTRY? IF IT ISN’T WE NEED TO FIGURE OUT WHAT COUNTRY IT’S IN.

SO PERHAPS YOU AS A HUMAN KNOW THE ANSWER, BUT A COMPUTER NEEDS TO KEEP HOPPING. SO LET’S GET A COUPLE MORE HOPS TO FIGURE OUT WHAT COUNTRIES ARE BEING REFERENCED HERE. THEN WE FIND OUT THAT CAIRO IS THE CAPITAL OF EGYPT AND THAT MOSSAD IS THE CIA, MI6, BND, WHAT HAVE YOU EQUIVALENT FOR ISRAEL.

LET’S AGAIN TAKE A LOOK AT THE GIVEAWAY FOR THIS QUIZ BOWL QUESTION. IT ENDS WITH A REFERENCE TO THE “CAMP DAVID” ACCORDS, WHICH IF YOU LOOK IT UP ON WIKIPEDIA GIVES YOU PLENTY OF REFERENCES TO THE ANSWERS, WHICH ARE EGYPT AND ISRAEL.

I WANT TO NOW RETURN TO THE DISCUSSION OF PARADIGMS OF QUESTION ANSWERING. IF YOU HAVEN’T SEEN THAT VIDEO THAT DEFINES MANCHESTER AS KNOWLEDGE PROBING QUESTIONS AND CRANFIELD AS KNOWLEDGE SEEKING QUESTIONS, CHECK IT OUT, OTHERWISE THIS WON’T MAKE SENSE. BUT LET’S DIG IN TO THIS QUESTION A LITTLE BIT MORE.

MULTIHOP QUESTION ANSWERING IS AN EXAMPLE OF WHERE THE CRANFIELD QA PARADIGM HAS FAILED US! EXISTING SYSTEMS’ INABILITY TO ANSWER THESE QUESTIONS EFFECTIVELY HAVE TRAINED US NOT TO ASK THEM.

I RECENTLY HAD THIS EXPERIENCE. I WAS WATCHING THE MANDALORIAN, AND IN THE EPISODE CALL “THE SHERIFF” THIS GUY SHOWED UP. I REMEMBERED HIM FROM SOME NETFLIX SHOW, BUT I COULDN’T THINK OF THE NAME. BUT DID I ASK GOOGLE TO TELL ME:

“NETFLIX SHOW WITH GUY WHO PLAYED SHERIFF ON THE MANDALORIAN ACTOR”

NO, OF COURSE NOT! DECADES OF USING GOOGLE HAVE TRAINED ME TO DECOMPOSE MY INFORMATION NEED INTO ATOMIC QUERIES. SO I FIRST FIGURE OUT THE NAME OF THE SHERIFF. AHA, COBB VAN TH. THEN I SEE WHO PLAYED

HIM: TIMOTHY OLIPHANT. THEN I ASK FOR THAT NETFLIX SHOW STARRING TIMOTHY OLIPHANT. AHA, THE SANTA CLARITA DIET. THAT'S WHAT IT WAS. HEY, DON'T JUDGE ME.

AND THAT'S THE PROBLEM WITH THE CRANFIELD PARADIGM: BECAUSE A SAVVY USER KNOWS THAT THEY CANNOT ASK THESE MULTIHOP QUESTIONS, YOU WON'T FIND AS MANY OF THESE TYPES OF QUESTIONS IN, SAY, NATURAL QUESTIONS.

SO WHY DON'T PEOPLE JUST USE MANCHESTER QUESTIONS LIKE QUIZ BOWL THAT HAVE PLENTY OF THESE MULTIHOP QUESTIONS YOU ASK? AND I ASK THAT TOO!

THE PROBLEM IS THAT FOR QUIZ BOWL QUESTIONS YOU DON'T HAVE THE GOLD ANNOTATIONS OF WHAT THE GOLD EVIDENCE IS TO ANSWER THE QUESTIONS. AND THAT'S THE BIG BENEFIT OF HotPotQA: IT EXPLICITLY GIVES GOLD ANNOTATIONS FOR WHERE TO FIND THE HOPS AND THE ANSWERS.

BUT IT SITS IN THE LIMINAL SPACE BETWEEN MANCHESTER AND CRANFIELD. THE QUESTIONS ARE SYNTHETIC, NOT REPRESENTING REAL NEEDS, BUT NEITHER DO THEY FOLLOW THE BEST PRACTICES OF THE MANCHESTER PARADIGM TO TEST THE ABILITY OF HUMANS TO DO MULTIHOP QA.

FOR HUMAN QUESTION ANSWERING COMPETITIONS, THERE ARE CLEARER RULES ABOUT HOW TO WRITE "GOOD" MULTIHOP QUESTIONS: YOU TYPICALLY DON'T GO DOWN THE INFORMATION GRADIENT. IN OTHER WORDS, YOU DON'T HOP FROM MORE WELL KNOWN ENTITIES TO LESSER KNOWN ENTITIES. FROM WHAT I'VE SEEN IN, SAY, NATURAL QUESTIONS, THIS OFTEN SEEMS TO BE THE CASE IN CRANFIELD QUESTIONS, SINCE QUESTIONS ARE LIKELY GOING TO BE ABOUT THE THINGS THAT ARE LESS WELL KNOWN: THIS WAS THE CASE IN MY MANDALORIAN EXAMPLE, GOING FROM COBB VANTH TO TIMOTHY OLIPHANT TO SANTA CLARITA DIET. THIS IS OFTEN VIOLATED IN HotPotQA, WHICH OFFENDS MY SENSE OF QUESTION AESTHETICS ... BUT DOES IT MAKE ANY DIFFERENCE IN TRAINING SYSTEMS? WHO KNOWS!

LIKewise, WHEN YOU USE MULTIHOP STRATEGIES IN PYRAMIDAL QUESTIONS, YOU SHOULD SOLVE THE MULTIHOP THROUGH THE JUDICIOUS USE OF COREFERENCE. WE OUTLINED THIS PROCESS IN A 2015 PAPER THAT EXPLICITLY LINKED THE MENTIONS TO THE UNDERLYING ENTITIES BUT NOT, ALAS, TO THE EVIDENCE THAT SOLVED THE KNOWLEDGE REQUIRED FOR THE REFERENCES. AND THAT, IT SEEMS, IS WHAT THE COMMUNITY REALLY WANTED. THAT AND LESS COMPLICATED QUESTIONS!

THERE ARE OTHER QUESTION ANSWERING DATASETS OUT THERE, AND WE'LL TALK ABOUT THEM EVENTUALLY. BUT I WANTED TO START WITH THESE "MACHINE READING" DATASETS BECAUSE IT SHOWS BOTH THE PROMISE OF THESE APPROACHES: THEY CAN ANSWER TONS OF EXCITING QUESTIONS JUST FROM LOOKING AT RAW TEXT. AND THE CHASM BETWEEN HUMAN AND MACHINE ABILITY. THERE IS NO WAY A COMPUTER CAN COME ANYWHERE CLOSE TO A SEMI-SKILLED HUMAN'S ABILITY TO ANSWER THE "JOY OF SEARCH" QUESTIONS WE STARTED THIS VIDEO ABOUT.

SO EVEN AS WE LEARN ABOUT THE METHODS THAT CAN PROVIDE THE ANSWERS TO THESE QUESTIONS, THINK ABOUT WHAT'S PREVENTING COMPUTERS FROM IDENTIFYING WHEN AN ANSWER SATISFIES THE QUESTION, THINKING UP REALLY NOVEL APPROACHES TO REFORMULATE A QUESTION (RATHER THAN COMPUTER BRUTE FORCE), AND RULING OUT PLAUSIBLE BUT WRONG ANSWERS.

SO HERE'S HOPING THAT HOPPING GETS MORE HUMAN.

7.2 *Conversational QA*

THE GOAL OF HAVING INTELLIGENT COMPUTERS IS A PART OF BOTH THE FOUNDATION OF ARTIFICIAL INTELLIGENCE IN THE FORM OF THE TURING TEST AND SCIENCE FICTION.

BUT YOU DON'T INTERACT WITH THESE COMPUTERS THROUGH A COMMAND LINE OR A SEARCH PROMPT. YOU TALK TO THE COMPUTER! AND THE COMPUTER REMEMBERS WHAT YOU SAY, IT HAS SITUATIONAL CONTEXT TO UNDERSTAND, RESOLVE AMBIGUITIES, AND PLACE YOUR QUESTIONS IN CONTEXT. THIS DOESN'T LOOK A LOT LIKE THE EXISTING QUESTION ANSWERING SETTINGS THAT WE'VE DISCUSSED BEFORE. TO ANSWER QUESTIONS IN A CONVERSATION, WE'LL NEED NEW DATASETS, NEW EVALUATIONS, AND NEW MODELS.

LET'S START WITH THE DATASETS. I MENTIONED SOME OF THESE VERY BRIEFLY IN THE DATASETS VIDEO, BUT NOW WE CAN LINGER A LITTLE BIT LONGER ON THESE CONVERSATIONAL QA DATASETS. THERE ARE OTHERS, BUT I WANTED TO TALK ABOUT: QBLINK, CoCa, QUACK, AND CANARD.

NOW YOU MIGHT THINK THAT CONVERSATIONAL QA MIGHT MEAN THAT YOU MIGHT GET A RESPITE FROM MY QUIZBOWL FIXATION, BUT NOT SO FAST! THE TRIVIA COMMUNITY IS ALSO A SOURCE OF FOUND DATA FOR CONVERSATION-LIKE QUESTIONS. WITH MY FORMER STUDENTS, WE CREATED A DATASET OF QUESTIONS THESE QUESTIONS START WITH A SENTENCE TO SET THE CONTEXT AND THEN YOU HAVE TO ANSWER THREE (USUALLY) QUESTIONS WITHIN THAT CONTEXT.

OTHER THAN THE MANCHESTER-STYLE QUESTIONS AND ANSWERS, THERE'S NOTHING SPECIAL LIKE PYRAMIDALITY THAT MAKES THEM "OPTIMAL". I JUST HINK THEY'RE NEAT.

FOR EXAMPLE, YOU HAVE A LEADIN THAT GIVES THE CONTEXT. "ONLY TWENTY ONE MILLION UNITS IN THIS SYSTEM WILL EVER BE CREATED. NAME THIS DIGITAL PAYMENT SYSTEM WHOSE TRANSACTIONS ARE RECORDED ON A "BLOCK CHAIN". NOW, NOTICE THAT YOU CANNOT ANSWER THIS QUESTION *JUST* FROM THE LITERAL QUESTION ... IT DEPENDS ON THE CONTEXT. THERE ARE PLENTY OF CRYPTOCURRENCIES (MORE EVERY DAY), BUT THE 21 MILLION NUMBER IS UNIQUE TO BITCOIN.

THEN WE MOVE ON TO THE NEXT QUESTION: IT WAS INVENTED BY THIS PERSON, WHO, ACCORDING TO A DUBIOUS NEWSWEEK COVER STORY, IS A 64 YEAR OLD JAPANESEAMERICAN MAN WHO LIVES IN CALIFORNIA. NOW, YOU COULD PROBABLY ANSWER THIS JUST FROM THE CLUE HERE, BUT IT HELPS TO

KNOW THAT “IT” HERE IS BITCOIN. OTHERWISE YOU HAVE TO DO A MULTIHOP TO FIGURE OUT THAT THIS IS BITCOIN AND IT WAS INVENTED BY SOMEONE WHO PURPORTED TO BE SATOSHI NAKAMOTO.

FINALLY, THE FINAL QUESTION IS MUCH EASIER IF YOU KNOW THAT BITCOIN IS INVOLVED.

LET’S MOVE ON TO THE NEXT DATASET, QUAC (QUESTION ANSWERING IN CONTEXT) WHICH APPEARED AT THE SAME TIME AS OUR DATASET AND IS NOTICEABLE FOR INTRODUCING EMOJIS IN PAPER TITLES. TWO OF MY FORMER STUDENTS—HE HE AND MOHIT—WORKED ON THIS CRANFIELD PARADIGM DATASET, WHERE TWO CROWDWORKERS PLAY THE ROLE OF A TEACHER AND A STUDENT. THE TEACHER GETS TO SEE A WIKIPEDIA PAGE BUT THE STUDENT DOES NOT. THEY HAVE A CONVERSATION AND THE STUDENT TRIES TO FIGURE OUT AS MUCH AS POSSIBLE ABOUT THAT TOPIC.

AND JUST LIKE THE QB LINK DATASET, YOU NEED TO DO THINGS LIKE RESOLVE COREFERENCE, KNOW THE TOPIC OF THE QUESTIONS, ETC.

AND AGAIN, ONE OF THE BIG PLUSSES IS THAT THE ANSWERS ARE ALL ANCHORED IN THE WIKIPEDIA PAGE (IN THIS CASE DAFFY DUCK), SO THERE’S GOLD EVIDENCE THAT USEFUL FOR MACHINE READING TECHNIQUES. BUT THE CHALLENGE IS THAT YOU ALSO NEED TO EMBED THE CONTEXT SOMEHOW!

ONE APPROACH WOULD BE TO JUST TRANSFORM THE CONVERSATIONAL, CONTEXT-SPECIFIC QUESTIONS INTO “NORMAL” THAT LOOK MORE LIKE NQ OR SQuAD, SO WE HAD CROWDWORKERS REWRITE THE QUESTIONS INTO FORMS THAT DON’T REQUIRE THE CONTEXT. WE CALLED THIS DATASET CANARD, SHORT FOR CONTEXT ABSTRACTION: NECESSARY ADDITIONAL REWRITTEN DISCOURSE, WHERE CANARD IS OF COURSE THE FRENCH WORD FOR “DUCK”.

AND BY CONDITIONING ON THE ENTIRE CONTEXT, YOU CAN INDEED CREATE QUESTIONS THAT DON’T NEED ADDITIONAL CONTEXT, BUT IT’S NOT AS GOOD AS THE HUMAN REWRITES.

THE FINAL DATASET I’D LIKE TO TALK ABOUT IS CoQA, WHICH LOOKS A LOT LIKE QUAC—IT CAME OUT A YEAR AFTER QUAC. BUT RATHER THAN ANCHORING ON A WIKIPEDIA PAGE, IT IS ANCHORED ON A SHORT VIGNETTE, HERE IT’S AN 80TH BIRTHDAY PARTY. LIKE IN QUAC THERE’S A STUDENT AND A TEACHER AND THE STUDENT ASKS QUESTIONS OF THE TEACHER. BUT UNLIKE QUAC THERE ARE ALSO RATIONALES THAT EXPLAIN WHY A QUESTION HAS A PARTICULAR ANSWER.

ONE THING THAT YOU’LL NOTICE ABOUT THESE DATASETS IS THAT THEY’RE NOT REALLY A TRUE CONVERSATION. IT’S KIND OF ONE SIDED. AND PART OF THAT IS THE TRICKY PROBLEM OF EVALUATION. YOU NEED DATASETS TO TRAIN MODELS … BUT TO TRAIN A MODEL YOU NEED A LOSS FUNCTION. TO HAVE A LOSS FUNCTION YOU NEED TO KNOW IF AN ANSWER IS CORRECT.

IN SOME CASES, THE EVALUATION CAN LOOK A LOT LIKE OUR EVALUATIONS FOR “NORMAL” QUESTION ANSWERING. YOU JUST NEED TO GIVE AN ENTITY. THEN YOU CAN USE THE SAME METRICS THAT WE SAW FOR, SAY, MACHINE READING QUESTION ANSWERING: ENTITY EQUIVALENCE, EXACT MATCH, SPAN F1, ETC. THAT’S WHAT YOU CAN DO FOR THE DATASETS THAT WE’VE TALKED

ABOUT THUS FAR.

BUT IS A CONVERSATION JUST SPOUTING THE ENTITIES? NO, OF COURSE NOT. YOU CAN'T JUST HAVE A CONVERSATION JUST THROUGH CULTURAL REFERENCES.

AS WE'LL TALK ABOUT IN THE LONG-FORM QUESTION ANSWERING SEGMENT, EVALUATING LONGER ANSWERS IS REALLY HARD.

SO ARE THERE WAYS THAT YOU CAN MAKE THE PROBLEM EASIER WITHIN A CONVERSATIONAL CONTEXT? RESEARCH IN DIALOG SYSTEMS TYPICALLY DIVIDE THE TASK INTO TWO PHASES: FIGURE OUT A DIALOG ACT THAT FITS IN THE CONTEXT OF THE CONVERSATION, AND THEN GIVEN THE CORRECT DIALOG ACT, YOU THEN NEED TO GENERATE THE RAW TEXT THAT ACTUALLY GOES TO THE USER.

SO WHAT'S A DIALOG ACT? THERE ARE MANY WAYS TO DEFINE DIALOG ACTS, HERE'S ONE EXAMPLE OF AN ONTOLOGY OF DIALOG ACTS CREATED FROM REAL-WORLD CONVERSATIONS BY A TEAM FROM COLORADO AND SRI.

YOU NEED TO BE ABLE TO ANSWER A QUESTION WITH YES OR NO, ASK THOSE QUESTIONS, ASK "WH" QUESTIONS LIKE "WHEN WAS YOUR LAST TRAINING SESSION". AND BECAUSE CONVERSATIONS AREN'T JUST INFORMATION EXCHANGE SESSIONS, SOMETIMES YOU NEED TO AGREE, GIVE AN OPINION, OR JUST LET THE PARTNER KNOW YOU'RE LISTENING.

IT'S SOMETIMES LONELY JUST TALKING INTO THE VOID WITH NOBODY RESPONDING TO YOU ... WHAT DIALOG ACT DOES A COMMENT CORRESPOND TO?

THIS WAS ADAPTED BY RESEARCHERS FROM VIRGINIA TECH TO MAKE IT MORE DETAILED. SOMETIMES YOU NEED TO GIVE AN ANSWER WITH AN EXPLANATION, SOMETIMES YOU NEED TO ANSWER WITH A FOLLOWUP CLARIFICATION QUESTION, OR SAY THAT YOU DON'T KNOW.

SO NOW, TO DO AN EVALUATION, YOU CAN MAKE SURE THAT YOUR SYSTEM CAN RECONSTRUCT THE DIALOG ACTS GIVEN THE CONTEXT. THIS TURNS IT INTO A CLASSIFICATION TASK: DOES YOUR PREDICTED DIALOG ACT MATCH THE GROUND TRUTH? THIS IS MUCH EASIER THAN ACTUALLY EVALUATING LONG, RAMBLING ANSWERS, WHICH AS WE'LL TALK ABOUT WITH GENERATION IS ACTUALLY VERY DIFFICULT.

BUT EVEN SO, THIS IS NOT WITHOUT COMPLEXITY. THESE DATASETS HAVE THE INTERACTIONS BETWEEN TWO PEOPLE WHERE PRESUMABLY EVERYTHING WENT ACCORDING TO PLAN. BUT WHAT HAPPENS IF DURING EVALUATION IF YOU SELECT A DIALOG ACT THAT'S DIFFERENT FROM THE REFERENCE INTERACTION?

THE THING WITH CONVERSATIONAL QUESTION ANSWERING IS THAT MISTAKES ARE NOT THE END OF AN INTERACTION. IDEALLY, A CONVERSATIONAL SYSTEM SHOULD RECOVER FROM IT! LET'S TAKE A LOOK AT THIS REFERENCE CONVERSATION WHERE THE USER ASKS THE SYSTEM WHAT THE CAPITAL OF GEORGIA IS. THE SYSTEM SHOULD GIVE AN ANSWER, AND THE USER CAN GIVE A FOLLOWUP QUESTION.

BUT WHAT HAPPENS IF THE SYSTEM GETS THE QUESTION WRONG? FOR EXAMPLE WITH A WRONG INTERPRETATION OF WHICH GEORGIA THE USER IS TALKING ABOUT. WELL, IF THE SYSTEM GIVES THAT ANSWER, THE USER CAN REACT IN

ALL SORTS OF WAYS! PERHAPS THEY REALIZE THEY CARED ABOUT GEORGIA THE COUNTRY ALL ALONG (THIS IS PROBABLY THE MOST LIKELY OUTCOME IF, SAY, YOUR “USERS” CAME FROM MECHANICAL TURK AND THEY’RE PRETENDING TO ASK QUESTIONS), PERHAPS THEY CORRECT THE SYSTEM, OR PERHAPS THEY WANT MORE INFORMATION: WHY DID YOU SAY “TBLISI”?

IN ANY EVENT, WE NEED TO KNOW HOW BAD IT IS THAT THE SYSTEM SAID TBLISI. AND WE CAN’T REALLY DO THAT WITHOUT TESTING THE SYSTEM ON REAL USERS. MAYBE IT ISN’T GREAT THAT WE SAID TBLISI, BUT IT’S BETTER THAN SAYING THAT THE CAPITAL OF GEORGIA IS “PEACHTREE” OR “WASHINGTON”. BUT AGAIN, WE DON’T KNOW THAT THIS IS A “NEAR MISS” BECAUSE, AGAIN, IT’S NOT IN THE REFERENCE.

AND IT CAN GO IN THE OTHER DIRECTION TOO! LET’S SAY THAT THE WAY WE GENERATED THE REFERENCE DATA WAS QUITE A BIT WORSE THAN WHAT AN “OPTIMAL” CONVERSATION PARTNER COULD DO. LET’S SAY THAT YOUR SYSTEM DOES SOMETHING BETTER THAN THE REFERENCE. PERHAPS IT ASKS WHICH GEORGIA YOU MEAN. THAT’S EVEN BETTER!

JUST BECAUSE YOU DO THE CONVERSATION A LITTLE DIFFERENTLY THAN THE REFERENCE, DOESN’T MEAN THAT IT’S WRONG.

SO THE CHALLENGE IS TO DEFINE AN EVALUATION THAT MEASURES HOW USEFUL AN INTERACTION IS, NOT THE ANSWERS OF INDIVIDUAL TURNS. I SUSPECT THAT THIS WILL NEED TO TAKE THE FORM OF HUMAN-CENTERED EVALUATION. YOU NEED TO HAVE PEOPLE INTERACT WITH A REAL SYSTEM AND GET THEM TO EVALUATE HOW WELL IT WORKS.

AND WE’LL SEE SOMETHING LIKE THAT IN OUR NEXT VIDEO, WHERE WE’LL TALK ABOUT ANOTHER WAY OF MOVING BEYOND SINGLE, SHORT ANSWERS: WHAT HAPPENS IF THE ANSWERS THEMSELVES ARE PRETTY DARN LONG?

LONG FORM QA

AS USUAL, LET’S START WITH HOW COMPLEX QUESTIONS ARE ANSWERED IN THE REAL WORLD. LET’S START WITH APU NAHASAPEEMAPETILON’S ANSWER TO “WHAT WAS THE CAUSE OF THE US CIVIL WAR?”

EXAMINER : ALL RIGHT, HERE’S YOUR LAST QUESTION. WHAT WAS THE CAUSE OF THE CIVIL WAR? APU : ACTUALLY, THERE WERE NUMEROUS CAUSES. ASIDE FROM THE OBVIOUS SCHISM BETWEEN THE ABOLITIONISTS AND THE ANTI-ABOLITIONISTS, THERE WERE ECONOMIC FACTORS, BOTH DOMESTIC AND INTER...

EXAMINER : WAIT, WAIT... JUST SAY SLAVERY. APU : SLAVERY IT IS, SIR.

LIKE THE EXAMINER, IT’S OFTEN EASIER TO GO FOR THE SHORTER ANSWER ... IT’S EASIER TO SCORE, ANYWAY. THERE ARE MANY WAYS TO ANSWER “WHAT WAS THE CAUSE OF THE AMERICAN CIVIL WAR”. EVERYTHING THAT APU SAID WAS CORRECT, AND THERE ARE MANY WAYS TO PHRASE WHAT HE SAID. DETECTING ALL THE THINGS THAT ARE SEMANTICALLY EQUIVALENT TO THAT IS ITSELF A NLP-COMPLETE PROBLEM. TO UNDERSTAND HOW, DESPITE THAT DIFFICULTY, WE’RE TRYING TO BUILD SYSTEMS THAT CAN ANSWER LONG COMPLICATED QUESTIONS, WE’LL GO THROUGH TWO DATASETS THAT CAN HELP COMPUTERS ANSWER HARD QUESTIONS AND WE’LL SEE WHAT MAKES THEM BOTH HARD TO ANSWER AND

HARD TO TELL WHEN THE ANSWER IS RIGHT.

LET'S FIRST BEGIN WITH COMPLEX ANSWER RETRIEVAL, OFTEN CALLED "CAR" FOR SHORT. THIS IS A DATASET FROM TREC AND LAURA DIETZ AT THE UNIVERSITY OF NEW HAMPSHIRE. IF YOU HAVEN'T SEEN OUR PREVIOUS VIDEO ABOUT LIKE MANY OF THE OTHER QUESTION ANSWERING DATASETS WE'VE TALKED ABOUT, THE SOURCE INFORMATION COMES FROM WIKIPEDIA.

THEY START WITH QUESTIONS LIKE "HOW DO YOU MAKE CANDY" OR "IS CHOCOLATE HEALTHY" AND RUN A SEARCH OVER ALL OF WIKIPEDIA. LIKE WITH TRADITIONAL RELEVANCE JUDGEMENTS, ANNOTATORS THEN MARK WHETHER THE PASSAGES RETRIEVED ARE A GOOD ANSWER FOR THESE QUESTIONS. NIST ANNOTATORS MARK WHETHER A PASSAGE *MUST* BE INCLUDED IN THE ANSWER-SET, *CAN* BE INCLUDED, IS ON TOPIC (BUT NOT A GOOD ANSWER) OR JUST IRRELEVANT.

NOW, GETTING ANNOTATORS TO AGREE ON THIS ISN'T TRIVIAL: IF YOU LOOK AT THE KAPPA STATISTICS, IT'S BETWEEN 0.5 AND 0.7. YOU CAN THEN RUN THE STANDARD PRECISION / RECALL STATISTICS FOR AN IR ENGINE. SO IS IT THE CASE THAT LONG FORM QA ISN'T ACTUALLY THAT HARD?

NO, IT'S STILL THAT HARD! THIS CAN ONLY WORK IF WIKIPEDIA HAS THE PERFECT ANSWER FOR EVERY QUESTION ALREADY WRITTEN DOWN. BUT ALTHOUGH WIKIPEDIA IS PRETTY BIG, IT'S NOT *THAT* BIG.

ANOTHER DATASET THAT TRIES TO ADDRESS THE "LONG ANSWER" PROBLEM IS THE ELI5 DATASET. THIS IS BUILT FROM A SUBREDDIT WHERE PEOPLE POSE QUESTIONS AND PEOPLE TRY TO ANSWER THEM USING SIMPLE LANGUAGE. WHAT I LIKE ABOUT REDDIT IS THAT THERE ARE CLEAR RULES ABOUT WHAT QUESTIONS ARE ALLOWED, HOW THEY CAN BE ANSWERED, AND THEN PEOPLE VOTE ON WHICH OF THE ANSWERS IS BEST.

DESPITE THESE BEAUTIFUL RULES AND STYLE GUIDELINES, THE QUESTIONS AREN'T AS BEAUTIFUL AS THE BEST EXAMPLES OF THE MANCHESTER PARADIGM, WHICH ARE CAREFULLY EDITED AND POLISHED UNTIL PERFECTION. HOWEVER, THEY'RE MUCH BETTER THAN THE GRABBAG OF THE WORST OF THE CRANFIELD PARADIGM, SINCE THERE ARE CLEAR RULES: YOU CAN'T ASK QUESTIONS THAT ARE ACTUALLY SIMPLE, THE QUESTIONS CAN'T BE SUBJECTIVE, AND THE QUESTIONS MUST BE ABOUT THE "REAL WORLD".

SO NOW WE HAVE EXAMPLES OF QUESTIONS AND GOOD ANSWERS. BUT WE CAN'T JUST TEST WHETHER THE TWO STRINGS ARE EQUAL OR NOT ... SO HOW DO YOU KNOW IF THE OUTPUT OF A SYSTEM IS GOOD OR NOT?

THIS IS ACTUALLY VERY SIMILAR TO THE PROBLEM FACED IN MACHINE TRANSLATION: YOU WANT TRANSLATIONS TO BE AS CLOSE TO A REFERENCE TRANSLATION AS POSSIBLE. SO THEY DEVELOPED A METRIC CALL BLEU, OR BILINGUAL EVALUATION UNDERSTUDY THAT COUNTS UP HOW MANY WORDS OVERLAP BETWEEN A REFERENCE TRANSLATION AND SYSTEM OUTPUTS.

IT'S CALLED AN "UNDERSTUDY" BECAUSE IT'S MEANT TO MIRROR THE ADEQUACY AND FLUENCY JUDGEMENTS OF HUMAN ANNOTATORS.

BUT WE'RE NOT TALKING ABOUT MACHINE TRANSLATION TODAY, BUT YOU

CAN USE A VERY SIMILAR IDEA FOR COMPARING TWO PIECES OF ENGLISH TEXT. SO BLEU IN FRENCH MEANS “BLUE”; THE OTHER METRIC WE’LL BE TALKING ABOUT IS ROUGE, WHICH MEANS “RED” IN FRENCH, AND STANDS FOR “RECALL-ORIENTED UNDERSTUDY FOR GISTING EVALUATION”. UNLIKE BLEU, WHICH TRIES TO GET THE TRANSLATION EXACTLY RIGHT—IN OTHER WORDS GETTING PRECISION AND RECALL HIGH—ROUGE FOCUSES ON RECALL.

LET’S SEE HOW IT WORKS. LIKE BLEU, ROUGE LOOKS FOR A CERTAIN NUMBER OF N-GRAMS. LET’S START WITH BIGRAMS FOR SIMPLICITY. LET’S SAY YOU WANT TO COMPARE THESE TWO SENTENCES. FIRST, YOU GATHER UP ALL OF THE BIGRAMS. ONE OF THE FOUR MATCHES, SO THIS MEANS THAT IT WOULD HAVE A ROUGE-2 OF $\frac{1}{4}$.

BUT ROUGE-N REQUIRES CHOOSING WHAT N SHOULD BE. A MORE POPULAR VARIANT IS TO USE THE “LONGEST COMMON SUBSEQUENCE” OR LCS. FINDING THE LONGEST COMMON SUBSEQUENCE IS A POPULAR EXAMPLE PROBLEM WHEN YOU’RE FIRST LEARNING DYNAMIC PROGRAMMING … WHILE WRITING A PROGRAM TO COMPUTE IT CAN BE A LITTLE TRICKY, FOR SHORT SENTENCES YOU CAN FAIRLY EASY EYEBALL: WHAT IS THE LONGEST SEQUENCE OF WORDS THAT IS IN *BOTH* SEQUENCES. FOR THESE TWO SENTENCES, THAT WOULD BE “PARENTS ANNOY THE CHILDREN”.

GIVEN THAT LCS, WE CAN THEN COMPUTE THE PRECISION AND RECALL AND THEN WE GET THE ROUGE-L MEASURE. ONE NICE THING ABOUT THIS IS THAT IT’S UPPER BOUNDED BY THE UNIGRAM OVERLAP BETWEEN THE TWO STRINGS.

WE SPENT A LOT OF TIME TALKING ABOUT ROUGE-L … DOES THAT MEAN THAT THIS IS THE RIGHT WAY TO COMPUTE WHETHER A LONG ANSWER IS A GOOD ONE OR NOT? COMPARE THE GENERATED ANSWER TO THE REFERENCE ELI5 ANSWER? A RECENT PAPER FROM MY FORMER STUDENT MOHIT SHOWS THAT … UHH MAYBE NOT!

ELI5, LIKE A LOT OF OTHER CRANFIELD QA DATASETS, HAVE A LOT OF OVERLAP BETWEEN THE TRAINING AND TEST SET (THIS IS LESS OF A PROBLEM IN THE MANCHESTER PARADIGM, AS THE AUTHORS EXPLICITLY AVOID REPETITION). THUS, THE SYSTEM DOESN’T ALWAYS LEARN TO GENERALIZE … IT CAN MEMORIZIE INSTEAD.

I’M NOT GOING TO TALK ABOUT THEIR SYSTEM MUCH—IT’S A GOOD ONE, CHECK OUT THE PAPER LINKED IN THE DESCRIPTION—BUT LIKE YOU MIGHT DO FOR THE CAR TASK, IT ALSO HAS A RETRIEVAL COMPONENT. BUT IT DOESN’T REALLY MATTER WHAT YOU RETRIEVE! GETTING RANDOM RETRIEVALS CAN ADD IN SOME SPECIFIC WORDS THAT CAN UP YOUR ROUGE-L SCORE.

AND IT ALSO SEEMS THAT YOU CAN GAME THE SYSTEM WITH REALLY STUPID BASELINES THAT *SHOULDN’T* GET A GOOD SCORE AND END UP … DOING QUITE WELL. FOR EXAMPLE, JUST COPYING THE ANSWER OVER AND OVER AGAIN OR COPY/PASTING A RANDOM TRAINING ANSWER SOMETIMES BEAT THE GOLD ANSWER.

SO THIS MAKES IT HARD TO PROVE THAT YOUR SNAZZY NEW ALGORITHM IS ACTUALLY GOOD. SO WHAT’S THE ANSWER? I THINK THAT QUESTION ANSWERING

ALWAYS WILL DEPEND ON HUMAN EVALUATIONS, BUT WHILE AUTOMATIC FACTOID QUESTION ANSWERING EVALUATIONS ARE OFTEN “GOOD ENOUGH”, I DON’T THINK—AS I RECORD THIS IN EARLY 2022—WE CAN CONFIDENTLY SAY THAT WE HAVE A “GOOD ENOUGH” AUTOMATIC EVALUATION WITH LONG FOR QUESTION ANSWERING.

SO THAT MEANS THAT YOU REALLY NEED TO DO A HUMAN EVALUATION. THERE’S NOTHING WRONG WITH THAT! AGAIN, I WISH THAT THIS WERE MORE COMMON FOR QUESTION ANSWERING. BUT FOR OUR MODERN DATA-HUNGRY ALGORITHMS AND OUR FAST-PACED RESEARCH CULTURE, PROJECTS DON’T ALWAYS HAVE THE LUXURY TO DO EVALUATION RIGHT … WITH THE FINICKY, EXPENSIVE HUMAN EVALUATIONS OF THESE VERY LONG ANSWERS.

SO WHAT’S THE RIGHT WAY TO QUICKLY EVALUATE QA SYSTEMS THAT OUTPUT MORE THAN A COUPLE OF WORDS? WELL, I HAVE THE ANSWER RIGHT HERE … UNFORTUNATELY IT’S THIS VIDEO IS TOO SMALL TO CONTAIN IT. ANSWERING QUESTIONS ABOUT IMAGES

7.3 *Multimodal*

MOSTLY WE’VE BEEN TALKING ABOUT COMPUTERS THAT CAN TALK TO YOU AND ANSWER QUESTIONS ABOUT TEXT. BUT IS THAT THE EXTENT OF OUR AMBITIONS? TAKE A LOOK AT THIS CLIP FROM KUBRICK’s 2001:

HAL 9000: HAVE YOU BE DOING SOME MORE WORK? DAVE BOWMAN: A FEW SKETCHES? HAL: MAY I SEE THEM? DAVE BOWMAN: SURE. HAL 9000: THAT’S A VERY NICE RENDERING, DAVE.

AND EVEN BEYOND THE FANCIFUL DREAMS OF FUTURE AI, THERE ARE PEOPLE TODAY WHO COULD REALLY USE HELP ANSWERING QUESTIONS ABOUT IMAGES. FOR EXAMPLE, THE VIZWIZ PROJECT FROM JEFF BIGHAM CONNECTED BLIND USERS WITH A QA SYSTEM OVER IMAGES. HERE’S TOM DEKKER, WHO IS BLIND, SHOWING OFF VIZWIZ.

TOM DEKKER: SO NOW WE HAVE VIZWIZ, WHICH I’LL OPEN. PHONE: USE THE CAMERA BUTTON TO TAKE A NEW PHOTO. TOM DEKKER: I’M GOING TO GET POSITIONED ABOVE THESE THINGS SO THEY’RE NICELY VIEWABLE. PHONE: START AND STOP RECORDING A QUESTION TOM DEKKER: WHAT COLOR ARE THESE TWO SOCKS? AND WHAT COLOR IS THE ONE THAT’S FOLDED IN HALF?

AS YOU’LL SEE COMPARED TO THE OTHER DATASETS THAT PEOPLE ARE TRAINING QA SYSTEMS ON, WE’RE NOWHERE NEAR THE ABILITY TO ANSWER THESE QUESTIONS. THE VIZWIZ PROJECT IS WHAT’S CALLED A “WIZARD OF OZ SYSTEM”: IT’S A HUMAN PRETENDING TO BE A MACHINE.

OSCAR ZOROASTER PHADRIG ISAAC NORMAN HENKLE: PAY NO ATTENTION TO THAT MAN BEHIND THE CURTAIN. THE GREAT AND POWERFUL OZ HAS SPOKEN!

AND THESE WORKERS ARE HIRED ON AMAZON MECHANICAL TURK, NAMED AFTER ANOTHER HUMAN PRETENDING TO BE A MACHINE.

TOM DEKKER: IT TOOK ABOUT A MINUTE AND A HALF. NOW LET’S SEE … PHONE: FOLDED SOCK IS BLACK. THE OTHER ONE IS HALF WHITE, HALF GRAY,

DIVIDED DOWN THE MIDDLE. —WEBWORKER PHONE: NEW ANSWERS AVAILABLE.

AND THERE ARE OTHER CASES WHERE IT'S NOT A QUESTION OF ABILITY BUT OF VOLUME: YOU MIGHT BE RESPONSIBILITY FOR WATCHING LOTS OF CAMERA FEEDS AND A VISUAL QA SYSTEM COULD QUICKLY NARROW DOWN THINGS THAT ARE INTERESTING FROM HUGE QUANTITIES OF DATA. FOR EXAMPLE, I COULD LOOK AT MY DOORBELL FOOTAGE AND ASK ... WHEN DID MY NEWSPAPER COME TODAY OR WHO TOOK MY AMAZON PACKAGE?

SO LET'S TAKE A STEP BACK FROM THOSE FANCIFUL AMBITIONS AND TAKE A LOOK AT THE “REAL” DATASETS THAT ARE OUT THERE FOR MULTIMODAL QUESTION ANSWERING. THE MOST PROMINENT IS “VQA” OR VISUAL QUESTION ANSWERING, WHICH HAS BECOME NEARLY SYNONYMOUS WITH THIS GIRL WEARING A BANANA MUSTACHE. THIS DATASET SET UP THE PARADIGM OF LOOKING AT A PICTURE AND ANSWERING THE QUESTION.

THE CROWDWORKERS WHO WROTE THE QUESTIONS IN THE MANCHESTER PARADIGM (SINCE THEY KNEW THE ANSWERS) WERE INSTRUCTED TO WRITE QUESTIONS THAT A HUMAN COULD ANSWER BUT THAT A ROBOT COULD NOT. WHILE IT HAS THE FLAVOR OF AN ADVERSARIAL SETTING, THERE WASN'T A SPECIFIC ADVERSARY ... THE WORKERS HAD TO HAVE THEIR OWN MENTAL MODEL.

HOWEVER, JUST LIKE MANY OF THE DATASETS WE’VE LOOKED AT, THERE WERE MANY SHORTCUTS THAT THE MODELS COULD TAKE: IF THE QUESTION ASKED “IS THERE A” THE ANSWER WAS ALMOST CERTAINLY YES, AND IF THERE WAS A SPORT INVOLVED, IT WAS USUALLY ABOUT TENNIS.

SO TO CORRECT THESE ISSUES, VQA 2.0 PAIRED TWO PICTURES TOGETHER SO THAT THE ANSWER TO THE SAME QUESTION IS DIFFERENT IN EACH OF THE PICTURES. SO FOR INSTANCE HERE THE MAN OR THE WOMAN IS WEARING GLASSES. THIS MAKES IT SO THAT THE DATASET IS INTRINSICALLY BALANCED AND THE MODEL CAN’T JUST WIN THROUGH A TEXT ONLY BASELINE.

SO CAN MODELS ANSWER THESE QUESTIONS? THE ACCURACY HAS BEEN CLIMBING BY 14 POINTS OVER THE LAST FOUR YEARS ... GETTING NEAR 80% ACCURACY. SO WHAT MODELS ARE BEHIND THESE IMPRESSIVE NUMBERS?

LIKE WITH TEXT, TRANSFORMERS WITH MUPPET NAMES LIKE OSCAR, BERT, AND ERNIE ARE TO BLAME. BUT HOW DO TRANSFORMERS DEAL WITH *VISUAL* INPUT? THUS FAR, WE’VE ONLY BEEN TALKING ABOUT WORDS AND WORD EMBEDDINGS. JUST LIKE YOU NEED TO SPLIT A SENTENCE INTO WORD PIECES, THE FIRST STEP IS TO BREAK AN IMAGE INTO PIECES TO CORRESPOND TO OBJECTS.

NOW, I DON’T KNOW HOW THIS WORKS AND EVEN IF I DID, IT WOULD BE OUT OF SCOPE FOR THIS COURSE, SO I WON’T TALK ABOUT HOW ALL THESE FANCY BOXES COME ABOUT. BUT ONCE YOU HAVE THESE BOXES, YOU CAN THEN EMBED THESE BOXES THROUGH A LINEAR OR NON-LINEAR TRANSFORMATION OF THE CONSTITUENT PIXELS WHICH—UNLIKE WORDS—ARE ALREADY CONTINUOUS VECTORS. JUST LIKE WORDS HAVE AN EMBEDDING BASED ON THE WORD’S TYPE AND THE TOKEN’S POSITION, THE VISUAL WORD EMBEDDING ALSO IS A FUNCTION OF THE PIXELS AND ITS LOCATION.

SO THEN YOU CAN GET A COMBINED VISUAL AND TEXT REPRESENTATION OF AN IMAGE, CAPTION OR IMAGE, QUESTION PAIR. AND BECAUSE A TRANSFORMER ISN'T AS TIED TO THE STRICT LINEAR ORDER OF A SENTENCE THANKS TO ITS ATTENTION MAP, IT DOESN'T REALLY MATTER IN WHAT ORDER YOU PLOP DOWN THE OBJECTS.

OKAY, SO HOW DO YOU PRETRAIN TRANSFORMERS WITH VISUAL AND TEXT INPUT? FIRST, YOU CAN DO MASKED WORD PREDICTION, JUST LIKE YOU DID WITH TEXT-BASED BERT. BUT IF YOU CAN DO THAT FOR WORDS, WHY NOT ALSO DO IT FOR OBJECTS? FINALLY, JUST LIKE WE DID NEXT SENTENCE PREDICTION FOR TEXT REPRESENTATIONS, WE CAN MATCH IMAGES WITH THEIR CAPTIONS.

SO THE MODELS (AFTER YOU GET PAST OBJECT RECOGNITION) LOOK A LOT LIKE OUR TEXT-ONLY MODELS. I THINK YOU'D AGREE THAT VISUAL QUESTION ANSWERING CAN BE QUITE A BIT HARDER ... SO WHY ARE THE ACCURACIES SO MUCH HIGHER THAN FOR TEXT-BASED QA?

THAT'S BECAUSE ALL OF THE QUESTIONS IN VQA ARE MULTIPLE CHOICE! IT BASICALLY BOILS DOWN TO SELECTING WHICH OF FOUR ANSWERS IS CORRECT. AND IT'S NOT JUST VQA, OTHER VISUAL DATASETS ARE ALL MULTIPLE CHOICE.

AND WHILE OTHER DATASETS ARE IMPLICITLY MULTIPLE CHOICE—IF YOU ASK SQuAD WHEN SOMETHING HAPPENED, IT'S ESSENTIALLY MULTIPLE CHOICE AMONG ALL THE THINGS THAT LOOK LIKE DATES—it's not as clear-cut as here. And datasets like Quizbowl are a multiple choice task over tens of thousands of entities.

SO WHY IS QA OVER IMAGES ALWAYS MULTIPLE CHOICE? BECAUSE WE HAVE A HARD TIME AGREEING ON WHAT THE ANSWER SHOULD BE! WE SAW THIS PROBLEM FOR LONG-FORM QA AND EVEN FOR FACTOID QA YOU NEED TO BE SURE TO HAVE ALL OF THE RELEVANT NAMES FOR AN ITEM. BECAUSE WE DON'T HAVE TEXT TO GROUND THE ANSWER (AGAIN, WHICH MAKES TEXT-BASED QA MORE LIKE MULTIPLE CHOICE), EVEN FOR EVERYDAY ITEMS YOU NEED TO FIGURE OUT WHAT NAME TO GIVE SOMETHING.

SO, FOR EXAMPLE, TAKE THE QUESTION WHAT IS Mr. Burns WEARING IN THE PICTURE? WHAT SHOULD YOU CALL THESE? MULES? HOUSESHOES? I PERSONALLY THINK THAT SOME OF THESE ARE BETTER THAN OTHERS, BUT ARGUABLY ALL OF THESE SHOULD BE ACCEPTED. AND Mr. Burns SEEMS TO PREFER "SLIPPERS". AND THIS ISN'T JUST A PROBLEM IN ENGLISH: HERE'S A MAP FOR WHAT SLIPPERS ARE CALLED IN THE GERMAN-SPEAKING WORLD. NAMING THINGS FROM A PICTURE IS HARD, DESCRIBING WHAT A THING IS DOING IS EVEN HARDER, AND HARDER STILL IS DESCRIBING *INTERACTIONS*. WHAT IS Mr. Burns DOING TO HIS SLIPPERS? DANCING? BOUNCING? FLAUNTING? TAPPING? GIVEN ALL OF THAT FLEXIBILITY, EVALUATING WHETHER AN ANSWER IS CORRECT IS JUST GENERALLY HARD, AS WE'VE SEEN MANY TIMES BEFORE.

AND IMAGES AREN'T STATIC ... OFTEN WE CARE ABOUT ANSWERING QUESTIONS ABOUT MOVING PICTURES. I.E., MOVIES OR TV. THERE HAVE ALSO BEEN DATASETS ABOUT THESE QUESTIONS AS WELL. JUST LIKE VQA, THESE ARE MULTIPLE CHOICE ... BUT DOES THAT MEAN YOU NEED TO WATCH THE MOVIE TO

ANSWER THE QUESTIONS? NOT NECESSARILY ... MOVIES ALSO HAVE CLOSED CAPTIONS/SUBTITLES, AUDIO DESCRIPTIONS, AND SCRIPTS. JUST IN CASE YOU'RE NOT FAMILIAR WITH THESE, LET'S TALK ABOUT THEM AND DISCUSS WHAT THEY CAN AND CANNOT ANSWER. CAPTIONS ARE THE TEXT THAT APPEAR AT THE BOTTOM OF THE SCREEN AND ARE AVAILABLE FOR PEOPLE WHO ARE HARD OF HEARING OR WHO JUST WANT TO BE ABLE TO ALSO READ WHAT THE PERSON IS SAYING. THEY'RE NOT ALWAYS PURFECT, BUT THEY USUALLY GET THE POINT ACROSS. IT'S A LOT OF WORK TO DO IT MANUALLY! [BUT A GOOD PLACE TO HIDE JOKES!]

WHILE CAPTIONS WERE FIRST INTENDED TO MAKE MOVIES MORE ACCESSIBLE FOR THE DEAF, AUDIO DESCRIPTIONS ARE CREATED FOR THE BLIND ... FOR EXAMPLE HERE'S A CLIP FROM FROZEN. THEY EXPLAIN WHAT'S GOING ON IN THE MOVIE BEYOND THE DIALOG. WHILE HUMANS TYPICALLY GET THIS THROUGH AUDIO, THEY ARE WRITTEN DOWN AT SOME POINT AND SO THESE CAN BE PROVIDED TO THE COMPUTER AS TEXT AS WELL.

FINALLY, MOST MOVIES AND TV SHOWS COME FROM A TEXT CALLED A SCRIPT. THESE CONTAIN NOT JUST THE DIALOG, WHICH GET TURNED INTO CAPTION BUT ADDITIONAL INFORMATION LIKE WHERE A SCENE TAKES PLACE AND THE INSTRUCTIONS TO THE ACTORS ABOUT WHAT ACTIONS THEY NEED TO TAKE. WHILE NOT AS VERBOSE AS THE DESCRIPTIONS, IT NONETHELESS PROVIDES A WAY OF INTERPRETING THE RAW AUDIOVISUAL REPRESENTATION OF A MOVIE.

HERE ARE TWO DATASETS—MOVIEQA AND TVQA—THAT HAVE QUESTIONS (MULTIPLE CHOICE, OF COURSE) BASED ON MOVIES. AND JUST LIKE HOW YOU CAN INSERT BOTH VISUAL AND TEXT INFORMATION INTO BERT, YOU CAN FEED IN RECOGNIZED OBJECTS, THE CAPTIONS (CALLED SUBTITLES IN THIS FIGURE) INTO BERT.

SO ALL OF THIS LOOKS A LOT LIKE HOW YOU DO MACHINE READING QUESTIONS GIVEN A CONTEXT ... BUT THE OTHER KIND OF QUESTION ANSWERING WE SAW WAS "OPEN DOMAIN" QUESTION ANSWERING USING SYSTEMS LIKE DPR?

YES! AND JUST LIKE FOR DPR WHERE WE HAD AN ENCODER FOR THE CONTEXT AND AN ENCODER FOR THE EVIDENCE, WE CAN DO THE SAME THING FOR IMAGES AND TEXT: TRY TO LEARN ENCODERS SUCH THAT THE ENCODING OF THE IMAGE DESCRIPTION AND THE ENCODING OF THE IMAGE ARE AS CLOSE TOGETHER AS POSSIBLE. AS I RECORD THIS IN 2022, THERE ARE TWO BIG MODELS THAT ARE DOING THIS: CLIP FROM OPENAI AND ALIGN FROM GOOGLE. THEY BOTH HAVE FAIRLY SIMILAR ARCHITECTURES.

THEN YOU CAN DO THINGS LIKE FIND GOOD CAPTIONS FOR AN IMAGE AUTOMATICALLY OR GIVEN A SEARCH STRING, FIND APPROPRIATE IMAGES FOR IT.

AS I WRAP UP, I HOPE THAT YOU'VE NOTICED THE PATTERN HERE: THE TOOLS THAT WE CREATED TO HELP PEOPLE WITH VISUAL OR HEARING IMPAIRMENTS END UP NOT JUST BEING A GOOD IDEA FOR EVERYONE, BUT THEY'RE ALSO ONE OF THE KEY TOOLS THAT WE'RE USING TO BUILD UP COMPUTERS' ABILITY TO UNDERSTAND THE INTERACTION BETWEEN TEXT AND IMAGES. SO THE MORAL

OF THE STORY IS THAT IF YOU WANT TO SUPPORT AI, SUPPORT ACCESSIBILITY! AND MAYBE WE'LL HAVE AIs THAT CAN SORT EVERYBODY'S SOCKS AND ANSWER QUESTIONS ABOUT IT AFTERWARD. THE FUTURE CAN'T COME FAST ENOUGH!

THIS IS JUST A SINGLE LECTURE FROM A COURSE.

YOUTUBE LIKES TO SHOW YOU THESE VIDEOS OUT OF ORDER, BUT IF YOU GO TO THE COURSE WEB PAGE LINKED BELOW, YOU CAN SEE THE LECTURES IN THE RIGHT ORDER AND YOU CAN GET RESOURCES LIKE HOMEWORKS OR SUGGESTED READING.

YOU CAN ALSO VISIT QANTA DOT ORG IF YOU WANT TO LEARN ABOUT OUR SYSTEMS FOR CREATING COMPUTERS THAT CAN ANSWER QUESTIONS.

WHERE QANTA STANDS FOR "QUESTION ANSWERING IS NOT A TRIVIAL ACTIVITY".

IF YOU WANT TO HELP THE CHANNEL PROVIDE A BIG GRADIENT TO THE ALGORITHM BY LIKING AND SUBSCRIBING.

HOW PUNTERS CHEATING WITH AI HAS MADE PUB QUIZZES' AUDIO ROUNDS BETTER

IF YOU WENT TO A PUB QUIZ AT THE TURN OF THE CENTURY, ONE STAPLE THAT YOU'D *ALWAYS* SEE WAS MUSIC IDENTIFICATION: THE PUB MASTER PLAYS A SONG AND YOU'RE SUPPOSED TO IDENTIFY IT. BUT THAT'S FAR LESS COMMON THESE DAYS. A BIG REASON IS THAT—AS FAR AS AI IS CONCERNED—it's mostly A SOLVED PROBLEM ... APPS LIKE SHAZAM AND YOUTUBE CONTENT ID CAN DO THIS WITH EXTREMELY HIGH ACCURACY EVEN IN A NOISY BAR. AND THUS IT HAS BECOME REALLY EASY TO CHEAT IN A PUB QUIZ. IT'S ALSO WHY, EVEN THOUGH THIS VIDEO IS ABOUT ANSWER QUESTIONS ABOUT MUSIC, I'M NOT GOING TO PUT ANY POPULAR MUSIC IN THIS VIDEO.

BLEEDING GUMS MURPHY WAS NEVER POPULAR.

SO HOW DO TRIVIA COMPETITIONS ASK QUESTIONS ABOUT MUSIC NOW? IN THE SAME WAY THAT YOUTUBE VIDEOS SPEED UP OR DISTORT MUSIC TO GET AROUND CONTENT ID, PUB QUIZZES HAVE INCREASINGLY ADOPTED USING TRANSFORMATIONS OF SONGS TO ASK QUESTIONS. IDENTIFY THE SONG GIVEN ONLY A SINGLE INSTRUMENT FROM THE TRACK OR IDENTIFY THE SONG FROM THE BLUEGRASS COVER.

IN THE SAME WAY THAT THESE QUESTIONS REWARD DEEPER KNOWLEDGE, THERE HAS ALSO BEEN A MOVEMENT TO REWARD REAL MUSIC KNOWLEDGE:

THIS SYMPHONY'S SECOND MOVEMENT ... FEATURES A RISING HORN CALL OF A C MAJOR ARPEGGIO FOLLOWED AN A-FLAT SIXTEENTH PICKUP TO A G.

[ORCHESTRAL MUSIC WITH A FRENCH HORN SOLO]

AND THIS IS JUST ASKED USING TEXT ... YOU NEED TO KNOW THAT AN ARPEGGIO IS A CHORD PLAYED IN SEQUENCE, YOU NEED TO KNOW WHAT A FRENCH HORN SOUNDS LIKE, AND YOU NEED TO KNOW WHAT A C VS. AN A-FLAT SOUNDS LIKE. MAPPING THIS TO AN UNDERLYING MUSICAL SCORE IS HARD ... THERE'S NO BEAUTIFUL MUSIC TRACK THAT I ADDED IN WITH MY L337 EDITING SKILLS. BUT OF COURSE, AS IS DE RIGUEUR IN THE MANCHESTER PARADIGM, AFTERWARD THE PYRAMIDAL SEQUENCE OF CLUES ENDS WITH THE MUCH EASIER

CLUE: FOR 10 POINTS, IDENTIFY THIS SYMPHONY SOMETIMES CALLED THE “SONG OF THE NIGHT”, WHICH NUMERICALLY PRECEDES ITS COMPOSER’S “SYMPHONY OF A THOUSAND.” REMEMBER THAT THE MANCHESTER PARADIGM THAT WE TALKED ABOUT BEFORE STRUCTURES QUESTIONS TO GO FROM DIFFICULT CLUES TO EASY CLUES, WHICH REWARDS KNOWLEDGE AND BETTER DISCRIMINATES BETWEEN TEAMS.

BUT I WANT TO EMPHASIZE THE VIRTUOUS CYCLE HERE: PUB QUIZ HOSTS WERE WRITING LAZY QUESTIONS THAT ESSENTIALLY AMOUNTED TO PLAYING AN iTUNES PLAYLIST IN A CROWDED BAR DEEP CONVOLUTIONAL NEURAL NETWORKS ALLOWED LAZY PUB QUIZ PARTICIPANTS TO TO CHEAT ON THOSE QUESTIONS OUT OF OUTRANGE AND TO RESTORE EQUITY, PUB QUIZ HOSTS STARTED WRITING BETTER QUESTIONS

THE NEXT STEP—WHICH HASN’T HAPPENED YET—is for AI SYSTEMS TO START USING THAT TRAINING DATA TO BUILD SYSTEMS THAT HAVE TO THINK A LITTLE BIT HARDER TO IDENTIFY THE SONGS.

SO THAT’S WHAT HAS HAPPENED FOR AUDIO QUESTIONS … LET’S NOW TURN TO QUESTIONS ABOUT IMAGES IN THE MANCHESTER PARADIGM. SIMILARLY, PYRAMIDAL QUESTIONS ABOUT VISUAL ART DESCRIBE LESS WELL-KNOWN DETAILS:

A SWORD STUCK IN ITS HOLDER AND THE HEAD OF A CORPSE ARE SEEN NEAR BUSHES ON THE LOWER LEFT.

A RABBIT RUNS TO TAKE COVER FROM THE CENTRAL OBJECT IN THIS WORK, WHICH APPROACHES THE VIEWER.

IF YOU KNOW THE PAINTING WELL, THEN YOU’LL BE ABLE TO ANSWER THE QUESTION. IF NOT, SINCE IT’S A PYRAMIDAL QUESTION, IN THE REST OF THE QUESTION YOU’LL GET THE REST OF THE CLUES.

FTP, NAME THIS PAINTING DEPICTING THE SEACLIFFS OF CRETE AND THE PAIR OF LEGS BELONGING TO A DROWNING YOUTH, BY PIETER BRUEGHEL THE ELDER. AND YOU CAN’T JUST RUN YOUR FAVORITE COMPUTER VISION SYSTEM TO FIGURE OUT THE TITLE HERE … SHIPS IN THE HARBOR? PLOUGHMAN ON THE HILL? NO, THE TITLE OF THIS PICTURE IS “LANDSCAPE WITH THE FALL OF ICARUS”. AND THERE’S NOTHING IN THIS PICTURE THAT ACTUALLY LOOKS LIKE ICARUS HERE … AH, THERE HE IS. HIS LEGS ANYWAY.

OR GOING BACK TO OUR ALLEGED RABBIT …

FTP, NAME THIS 1844 LANDSCAPE DEPICTING A TRAIN CROSSING A BRIDGE, A DEPICTION OF “THE GREAT WESTERN RAILWAY” BY J.M.W. TURNER.

IF YOU KNOW THE ANSWER TO THIS QUESTION, PUT IT AS A COMMENT DOWN BELOW.

AND THIS ISN’T JUST ABOUT VISUAL SALIENCE. THE MOST WELL-KNOWN ASPECT OF BRONZINO’S VENUS, CUPID, FOLLY, AND TIME ISN’T OBVIOUS BUT IS BETTER KNOWN THAN THE REST OF THE PAINTING BECAUSE OF TERRY GILLIAM’S USE OF IT IN MONTY PYTHON.

SO THIS IS HOW TEXTUAL CLUES ARE RENDERED. AND IN THE MANCHESTER PARADIGM, YOU PRESENT THE CLUES IN PYRAMIDAL ORDER, GOING FROM DIF-

FICULT TO EASY. BUT THESE ARE TEXTUAL CLUES. AND YEAH, PERHAPS YOU HAVE TO LOOK AT A PAINTING OR LISTEN TO A PIECE OF MUSIC TO ANSWER IT (AND I THINK IT WOULD BE FUN TO USE VISUAL OR AUDITORY INFORMATION TO BETTER ANSWER THESE QUESTIONS ... PROJECT IDEA). BUT THE QUESTION IS STILL TEXT ... CAN YOU ASK PYRAMIDAL QUESTIONS WHEN THE QUESTION ITSELF IS MULTIMEDIA? THE ANSWER, I THINK, IS YES, BUT BEFORE I GO INTO THAT, I FIRST NEED TO EXPLAIN WHY FOR THE SKEPTICS IN THE BACK.

THE MANCHESTER PARADIGM FOR QUESTION ANSWERING TRIES TO MAKE QUESTIONS THAT ARE DIFFICULT TO ANSWER SO THAT IT REWARDS KNOWLEDGE, UNDERSTANDING, AND ROBUSTNESS. JUST LIKE SHAZAAM ISN'T IMPRESSIVE IF IT ONLY WORKS WITH CLEAR, UNDISTORTED INPUT ... THAT IT WORKS IN SUCH AN ADVERSARIAL ENVIRONMENT LIKE A CROWDED BAR IS PRETTY AWESOME! IF WE WANT ROBUST QUESTION ANSWERING AND TRUE ARTIFICIAL INTELLIGENCE, COMPUTERS NEED TO BE ABLE TO "MAKE DO" WITH LESS AND EMPLOY THE SAME REASONING AND RECALL ABILITY THAT HUMANS ARE ABLE TO DO.

SO LET'S SEE HOW THE MANCHESTER PARADIGM DOES THAT! ONE WAY IS TO HAVE THE NORMAL PRESENTATION OF MATERIAL BUT TO SEVERELY LIMIT THE QUANTITY. FOR EXAMPLE, ON THE GAME SHOW "NAME THAT TUNE", CONTESTANTS BID EACH OTHER DOWN TO SEE WHO CAN NAME THE TUNE WITH THE FEWEST NOTES.

LET'S SEE THAT IN ACTION ...

WASH: I CAN NAME THAT TUNE IN FOUR NOTES SANDY: WASH, I CAN NAME THAT TUNE IN THREE NOTES WASH: SANDY, NAME THAT TUNE JIM LANGE: ALRIGHT, SANDY YOU GET A CHANCE TO NAME THAT TUNE. HERE'S YOUR CLUE ONCE AGAIN. INTRODUCED AS A TITLE TUNE OF A 30S BROADWAY MUSICAL, THIS TUNE LATER APPEARED IN THE FILM YOUNG FRANKENSTEIN AND WAS REVIVED IN 1984. LISTEN CAREFULLY, HERE ARE YOUR THREE NOTES ... [THREE NOTES PLAY] JIM LANGE: SANDY? SANDY: FARMER IN THE DELL? [SAD BUZZER]

THE ANSWER OF COURSE, IS ...

GENE WILDER AS FREDERICK FRANKENSTEEN: IF YOU'RE BLUE AND YOU DON'T KNOW WHERE TO GO TO, WHY DON'T YOU GO TO WHERE FASHION SITS ... PETER BOYLE AS THE MONSTER: ICH HAN MIS PORTMONE VERLOORE!

DID YOU GET IT?

OH, THAT WASN'T VERY CLEAR, LET'S TRY THAT AGAIN ...

TACO: PUTIN' ON THE RITZ

THAT'S RIGHT, PUTTIN' ON THE RITZ.

BUT THIS SHOWS ANOTHER WAY THAT THIS CAN BE PYRAMIDAL ... GO THROUGH PERFORMANCES FROM THE LESSER KNOWN, UNINTELLIGIBLE WARBLING OF PETER BOYLE TO THE INTELLIGIBLE BUT NICHE PERFORMANCE OF TACO TO THE CANONICAL PERFORMANCE OF FRED ASTAIRE.

THIS IS SIMILAR TO HOW WE STRUCTURED THE HUMAN VS. COMPUTER COMPETITION AT THE EFFICIENT QA COMPETITION. GIVE SYSTEMS A SHOT AT HARD EXAMPLES WITH LITTLE EVIDENCE, PROVIDE MORE AND MORE EVIDENCE UNTIL YOU SEE WHERE THE FIRST SYSTEM GETS IT RIGHT. CREATING PYRAMIDAL EXAM-

PLES WITH N VIEWS IS MUCH EASIER THAN CREATING N EXAMPLES OF VARYING LEVELS OF DIFFICULTY.

FOR EXAMPLE, WE COULD CREATE A VISUAL VERSION OF THE LANDSCAPE WITH THE FALL OF ICARUS QUESTION BY ONLY REVEALING PART OF THE IMAGE ... GOING IN REVERSE ORDER OF SALIENCE, JUST LIKE THE TEXT OF THE QUESTION DID: THE HILLS IN THE BACKGROUND, THE SWORD AND SKULL, THE FARMER, AND THEN FINALLY THE HARBOR WITH THE TWO TINY LEGS.

ALTERNATIVELY, YOU COULD PIXELATE THE IMAGE, DECREASING THE PIXELATION UNTIL IT CAN BE IDENTIFIED. AGAIN, IF YOU KNOW WHAT THIS PAINTING IS, LEAVE A COMMENT DOWN BELOW.

BUT WAIT, JORDAN, I HEAR YOU SAY. THIS JUST CORRESPONDS TO DIFFERENT LAYERS OR DIFFERENT POOLS IN A CONVOLUTIONAL NET. THIS WON'T MAKE IT THAT MUCH HARDER FOR A COMPUTER TO UNDERSTAND AN IMAGE.

THAT'S TRUE! AND UNTIL THERE IS A SHAZZAM EQUIVALENT THAT'S AS GOOD AND AS POPULAR FOR IMAGES, WE'RE PROBABLY GOING TO HAVE TO WAIT FOR THE KIND OF INNOVATION WE'VE SEEN FOR AUDIO QUESTIONS. BUT SOME OF THIS IS HAPPENING. FOR EXAMPLE, I'VE SEEN THINGS LIKE "IDENTIFY THE MOVIE OR TV SHOW FROM THE MUPPET-IZED VERSION". SO YOU CAN'T JUST IDENTIFY RAW PIXEL PATTERNS, THEY NEED TO UNDERSTAND WHAT PART OF THE IMAGE IS A MUPPET AND WHICH IS A PART OF THE UNDERLYING MEDIA PROPERTY.

AND IF WE WANT A COMPUTER TO BE ABLE TO DO THAT, WE NEED TO REPEAT THIS CYCLE, PROBABLY MULTIPLE TIMES. BUT NOT JUST FOR IMAGES, BUT FOR ALL OF THE KINDS OF QUESTIONS THAT WE WANT A COMPUTER TO ANSWER.

BUT ONE PART OF THIS PROCESS IS AI HELPING HUMANS CHEAT! ARE YOU CONDONING CHEATING? NO, OF COURSE NOT. BECAUSE WHAT THIS IS REALLY DOING IS HELPING THE SMART, CLEVER PEOPLE WHO CREATE THESE QUESTIONS BE MORE CREATIVE. BUT WE AS AI RESEARCHERS CAN MAKE IT EASIER FOR THAT ADVERSARIAL AUTHORIZING PROCESS TO HAPPEN. EASIER BOTH FOR THE AUTHORS, WHO END UP WITH A BETTER OUTPUT WITH LESS EFFORT AND EASIER FOR THE AI SYSTEMS THAT NEED TO TRAIN ON DATA TO GET THE MOST CHALLENGING DATA THAT REQUIRES ACTUALLY *UNDERSTANDING* THE UNDERLYING INFORMATION.

AND THAT IS WHAT MAKES THIS SORT OF RESEARCH SO FUN. JUST AS A PUB QUIZ SHOULD REWARD KNOWLEDGE AND ADD EXCITEMENT OF COMPETITION, WE CAN USE THAT SAME CREATIVE ENERGY TO MAKE OUR ALGORITHMS LESS BRITTLE AND MORE ROBUST.

WHAT MAKES IT HARD TO ANSWER SPOKEN QUESTIONS? WHEN WE TALKED ABOUT WATSON (LINK IN THE DESCRIPTION), ONE OF MY MINOR CRITICISMS WAS THAT IT DIDN'T LISTEN TO QUESTIONS. NOW WE'RE GOING TO TALK ABOUT WHY THAT'S A SCIENTIFIC ISSUE AND WHAT MAKES THIS PROBLEM DIFFICULT.

NOW, IF YOU'VE JUST BEEN WATCHING THE ERROR RATE FOR AUTOMATIC SPEECH RECOGNITION (ASR FOR SHORT), THE NUMBERS ARE REALLY IMPRESSIVE. AND MICROSOFT HAS CLAIMED THAT THESE RESULTS ARE SUPERHUMAN! BUT THESE HEADLINE NUMBERS ARE TYPICALLY GIVEN AS A WORD ERROR RATE, OR

MORE PRECISELY A TOKEN ERROR RATE. AND IT'S TRUE THAT THE WORD ERROR RATE IS AVERAGE IS LOWER THAN HUMAN NUMBERS, THAT'S HIDING SOMETHING. THE MOST COMMON WORD IN ENGLISH IS THE ... SO JUST GETTING THAT RIGHT DOESN'T MEAN THAT MUCH.

INDEED, THE MOST FREQUENT WORDS IN ENGLISH APPEAR *A LOT*. THIS IS CALLED A ZIPFIAN DISTRIBUTION, WHERE THE PROBABILITY OF A WORD IS INVERSELY PROPORTIONAL TO ITS FREQUENCY RANK. SO GETTING THESE HIGHLY FREQUENT WORDS RIGHT IS EASY AND COUNTS FOR A LOT OF THE WORD ERROR RATE.

BUT ARE THESE THE WORDS THAT MATTER FOR QUESTION ANSWERING? NO! NAMED ENTITIES ARE THE THINGS THAT MATTER, AND THEY'RE MUCH LESS COMMON. LET'S TAKE A LOOK AT QUESTION ANSWERING DATA. EVERY DOT HERE IS AN ERROR, AND WHEN THE ORIGINAL WORD IS VERY COMMON ... THERE AREN'T MANY SPEECH RECOGNITION ERRORS.

LET'S LOOK AT WHERE THE ERROR ACTUALLY ARE: WHEN THE ORIGINAL WORD IS "GUPTA" (E.G., THE EMPIRE OF CHANDRAGUPTA), "CLARENCE" (AS IN HENRY II'S COUNCIL THAT RESTRICTED ECCLESIASTICAL POWER IN ENGLAND) OR "KERMIT" (PROUD UMD ALUM), THESE EITHER ARE UNKNOWN TO THE ASR SYSTEM, AS IN THE CASE WITH OUR GREEN FRIEND HERE, THEY GET RECOGNIZED AS MORE COMMON WORDS. "CLARENCE" BECOMES THE ALLERGY MEDICATION "CLARITIN" AND "GUPTA" BECOMES "GROUP".

PLUS, THEY OFTEN COME FROM FOREIGN LANGUAGES AND ARE EASILY CONFUSED OR MISPRONOUNCED. THIS IS ALSO A PROBLEM OF THE CRANFIELD VS. MANCHESTER PARADIGM: WHILE ALEX TREBEK IS UNLIKELY TO MISPRONOUNCE ANYTHING (HE HAS WATCHDOGS TO RERECORD ANYTHING THAT SOUNDS OFF), A CLUELESS USER MIGHT INDEED MISPRONOUNCE "WHERE WAS MAO ZEDONG BORN?"

SO WHAT CAN YOU DO TO BETTER ANSWER SPOKEN QUESTIONS? FIRST, YOU CAN IGNORE WORDS FROM THE RECOGNIZED TEXT THAT AREN'T CONFIDENT: IF THERE'S ENOUGH REDUNDANCY IN THE QUESTION, YOU MIGHT STILL BE ABLE TO GET IT RIGHT. OR YOU CAN LOOK AT THE DIFFERENT OPTIONS FOR RECOGNIZED WORDS; IN OTHER WORDS, LOOK AT THE "LATTICE" OF THE DIFFERENT OPTIONS. IF THE TOP GUESS ISN'T RIGHT, MAYBE YOU'LL FIND IT UNDER THE OLD ONE.

FINALLY, YOU CAN ALSO LOOK FOR WORDS THAT "SOUND" SIMILAR TO WORDS THAT YOU'RE LOOKING FOR. SO PERHAPS THE ORIGINAL QUESTION HAD "DUMAS" AND THE SYSTEM DETECTED "DUMMY", YOU MIGHT BE ABLE TO PULL UP THE CORRECT AUTHOR BECAUSE IT SHARES THE SYLLABLE "DUM". BUT THIS REQUIRES DIFFERENT ALGORITHMS: YOU NEED TO BE ABLE TO SEARCH BY SYLLABLE RATHER THAN BY JUST WORDS.

AND YOU NEED DATASETS TO DO THIS. ONE CHEAP WAY IS TO USE TEXT TO SPEECH (I.E., HAVE A COMPUTER DO THE SPEAKING) AND THEN TRY TO RECOGNIZE THIS. THIS IS STRICTLY EASIER THAN THE REAL WORLD: THERE'S NO BACKGROUND NOISE, THE PRONUNCIATIONS ARE ALWAYS CONSISTENT, AND THERE'S ONLY ONE SPEAKER.

WE DID THIS FIRST FOR THE MANCHESTER PARADIGM, AND THEN A GROUP FROM CMU DID IT FOR THE CRANFIELD PARADIGM. AS FAR AS I KNOW, THERE'S NO DATASET OF *REAL* SPEECH EVEN THOUGH THERE ARE TONS OF PODCASTS AND YOUTUBE VIDEOS OF PEOPLE READING AND ANSWERING QUESTIONS. AND BECAUSE THE QUESTIONS ARE ALL WRITTEN DOWN, YOU COULD LINE THEM UP. BUT TO THE BEST OF MY KNOWLEDGE, NOBODY HAS DONE THAT: PROJECT IDEA!

AND IF YOU HAVE DIVERSE SPEAKERS, YOU COULD CONTEND WITH THE DISPARITIES IN ASR: MUCH OF THE TRAINING DATA COMES FROM WHITE MEN, SO WHEN TESTED ON OTHER GROUPS, ACCURACY DROPS.

AND THEN YOU COULD ALSO GET TO SOME OF THE OTHER CHALLENGES IN PARSING QUESTIONS: FOR INSTANCE, THIS DATASET FROM GOOGLE TACKLES THE DISFLUENCIES OR CORRECTIONS THAT OFTEN COME UP IN CRANFIELD PARADIGM QUESTIONS.

ANVIL YESTERDAY SETS AND QUESTIONS, REIGNITE IMPROV BROTH ARABLE TUESDAY QUESTIONS AND TO WRECK A NICE BEACH. AND MAYBE GET BETTER CAPTIONS ON YOUTUBE, AS LONG AS PEOPLE ARE GOOD ABOUT UPLOADING GOOD DATA.

8

DATASETS: BUILD OR ADOPT?

THE CLICHE IS THAT DATA ARE THE NEW OIL, POWERING AI. FORTUNATELY, BECAUSE HUMANS NATURALLY ASK QUESTIONS, THERE ARE MANY DATASETS THAT WE CAN FIND 'FOR FREE'. HOWEVER, THESE DATASETS STILL COME AT A COST: MANY OF THESE DATASETS HAVE INHERENT PROBLEMS (E.G., AMBIGUITIES AND FALSE PRESUPPOSITIONS) OR ODDITIES (ONLY TALKING ABOUT AMERICAN MEN) THAT MAKE THEM DIFFICULT TO USE FOR QUESTION ANSWERING. THIS CHAPTER DISCUSSES THESE DATASETS THAT HAVE FORMED THE FOUNDATION OF MUCH OF MODERN AI. IF YOU CAN'T FIND THE DATA YOU NEED, BUILD IT YOURSELF. BUT THIS IS NOT ALWAYS A PERFECT SOLUTION. THIS CHAPTER DISCUSSES THE DATASETS THAT PEOPLE HAVE BUILT AND THE NEW PROBLEMS THAT THIS CAN CREATE.

TOPICS IN CHAPTER:

1. **NATURAL QUESTIONS FIVE YEARS LATER.**
2. **SHOULD QA BE IID?.**
3. **HOW QUESTIONS GO WRONG.**
4. **CONSTRUCTED DATASETS FROM CROWDSOURCING.**
5. **GENERATING QUESTIONS.**

8.1 What Question Datasets are There?

MY GOAL WITH THIS VIDEO IS TO GIVE AN OVERVIEW OF THE KINDS OF QUESTION ANSWERING DATASETS THAT ARE OUT THERE, CONTRAST THEM, AND GIVE MY OPINION ABOUT WHICH IS THE BEST ONE.

FIRST, LET ME START OFF BY SAYING WHAT WE'RE NOT TALKING ABOUT! THERE ARE MANY QUESTION ANSWERING DATASETS OUT THERE THAT ARE ESSENTIALLY LOGIC PROBLEMS OR SQL QUERY LOOKUPS AND WE'RE NOT GOING TO TALK ABOUT THEM. THAT'S A VERY VALID FORM OF QUESTION ANSWERING AND IT ESSENTIALLY CORRESPONDS TO A FORM OF PARSING WHERE YOU HAVE SOME TEXT AND YOU NEED TO TURN THAT INTO A LOGICAL FORM OR AN SQL STATEMENT. IF YOU PARSE IT CORRECTLY, YOU GET THE ANSWER.

INSTEAD, WE'RE GOING TO TALK ABOUT MORE OPEN-ENDED QUESTION ANSWERING THAT'S OFTEN CALLED "MACHINE READING" QUESTION ANSWERING

FOR REASONS THAT WILL BE CLEAR IN A SECOND. IN THESE QUESTION ANSWERING DATASETS, YOU NEED TO MAKE SENSE OF A LARGE BODY OF TEXT TO ANSWER QUESTIONS. MOREOVER, YOU USUALLY NEED TO BE ABLE TO ANSWER QUESTIONS ABOUT JUST ABOUT ANYTHING, NOT JUST A NARROW SET OF TOPICS.

WE'LL TALK ABOUT MODELS FOR ANSWERING THESE QUESTIONS LATER, BUT FOR THE MOMENT THINK ABOUT THE KINDS OF THINGS THAT YOUR MODEL WOULD NEED TO DO TO ANSWER THE QUESTION. IN MANY CASES, IT WILL BE EXACTLY THE SAME KIND OF THING YOU WOULD NEED TO DO TO ANSWER THE QUESTION: FIND A DOCUMENT THAT CONTAINS THE ANSWER, READ THROUGH THE DOCUMENT TO FIND THE ANSWER, AND THEN CONVEY THAT ANSWER.

OKAY, BACK TO THE DATASETS. RATHER THAN JUST LISTING THEM OFF, I'M GOING TO TRY TO ORGANIZE THE DATASETS BASED ON: HOW MUCH INFORMATION YOU HAVE TO ANSWER THE QUESTION AND HOW COMPLICATED THE ANSWER IS. LET'S SAY THAT WE CAN WRITE THIS AS A TWO DIMENSIONAL PLOT WITH PROVIDED INFORMATION AS THE X AXIS AND ANSWER COMPLEXITY AS THE Y AXIS. I'LL ROUGHLY GO CHRONOLOGICALLY THROUGH TIME.

IN ANOTHER VIDEO, THERE'S ANOTHER IMPORTANT DIMENSION: WHO'S WRITING THE QUESTION AND FOR WHAT PURPOSE, BUT I THINK THAT'S SO IMPORTANT THAT I WANT TO FOCUS ON IT BY ITSELF. SO IF YOU'RE INTERESTED IN THAT, CHECK OUT THE CRANFIELD VS. MANCHESTER DISCUSSION LINKED IN THE DESCRIPTION.

8.2 TREC QA

LET'S START WITH THE MOST SIMPLE EXAMPLE OF A QUESTION ANSWERING DATASET: YOU GET A QUESTION AS INPUT AND YOU NEED TO PRODUCE A SHORT ANSWER. THIS IS OFTEN CALLED FACTOID QA. THE PROTOTYPICAL EXAMPLE OF THIS IS TREC QA. THIS IS A SUPER OLD SCHOOL DATASET, BUT IMPORTANT TO REMEMBER, AS IT SET THE FORM OF EVERYTHING THAT CAME AFTER.

THE ACTUAL DATASET IS TINY: ONLY A COUPLE HUNDRED QUESTIONS. AND THE OFFICIAL EVALUATION PROCESS SEEMS IMPOSSIBLE BY TODAY'S STANDARDS: YOU HAD PROFESSIONAL ANNOTATORS JUDGE WHETHER A SYSTEM PROVIDED THE RIGHT ANSWER. WHILE EVERY OTHER DATASET WE'LL TALK ABOUT HAS MORE AUTOMATIC EVALUATIONS, THIS IS PROBABLY THE BEST WAY TO DO IT, AT LEAST FOR SOME TRICKIER ANSWERS.

THIS IS OUR FIRST QA DATASET, SO LET'S PUT THIS ONTO THE BOARD HERE. LET'S NOW GO INTO THE MODERN AGE!

8.3 CNN / Daily Mail

ONE OF THE FIRST MACHINE READING DATA SETS WAS THE CNN / DAILY MAIL DATA SET FROM KARL MORITZ HERMANN AND COMPANY. HERE YOU HAVE A BUNCH OF NEWS ARTICLES. THIS DATASET STARTED THE TREND OF CONDITIONING A QUESTION ON A DOCUMENT: IN THIS CASE A NEWS ARTICLE. THEY ASKED

YOU TO ANSWER QUESTIONS ABOUT THOSE NEWS ARTICLES. FOR EXAMPLE, “WHO REACHED THE CRICKET WORLD CUP”. SOMEWHERE IN THE DOCUMENT WAS THE ANSWER “NEW ZEALAND” BUT YOU NEED TO FIND THAT.

ONE THING THAT IS CONTROVERSIAL IS THAT YOU DIDN’T ACTUALLY FIND THE STRING “NEW ZEALAND” BUT RATHER AN ANONYMIZED SYMBOL. THIS MEANT THAT YOU COULDN’T JUST HAVE A LANGUAGE MODEL MEMORIZE ALL THE ANSWERS (WHICH IS EVEN MORE OF A PROBLEM NOW THAT IT WAS BACK THEN). HOWEVER, THIS HAD ANOTHER DRAWBACK! BECAUSE THERE WEREN’T THAT MANY SYMBOLS, THE TASK WAS ESSENTIALLY A MULTIPLE-CHOICE QUESTION: YOU IDENTIFIED THE ENTITIES IN THE PARAGRAPH AND THEN YOU PICKED WHICH ONE WAS CORRECT. DESPITE THIS PROBLEM AND THAT INTRODUCED A LITTLE BIT OF NOISE TO THE PROBLEM, YOU CAN SEE THE INFLUENCE OF THIS QUESTION ANSWERING DATASET ON THE REST OF THE FIELD, PARTICULARLY ON “THE NEXT BIG THING”: SQuAD.

8.4 SQuAD

THE STANFORD QUESTION ANSWERING DATASET HAS BECOME THE GOLD STANDARD. WHILE IT MIGHT NOT BE THE BEST DATASET, IT’S THE MOST POPULAR. IF YOU DON’T REPORT RESULTS ON SQuAD, YOU NEED TO EXPLAIN WHY. THIS 2016 DATASET SET HELPED POPULARIZED THE MODERN UPSURGE OF QUESTION ANSWERING. THESE QUESTIONS USE PARAGRAPHS FROM WIKIPEDIA ASK CROWD WORKERS TO WRITE QUESTIONS ABOUT THAT PARAGRAPH. THE QUESTIONS ARE INTENDED TO HAVE AN ANSWER THAT IS A SPAN WITHIN THAT WIKIPEDIA PARAGRAPH. ONE WAY OF THINKING ABOUT THIS IS THAT A SPAN IS A CONTIGUOUS SPAN OF CHARACTERS: YOU HIGHLIGHT THAT SPAN TO GIVE THE ANSWER.

THIS IS A VERY LARGE DATA SET SO IT DOESN’T HAVE MANY OF THE PROBLEMS OF THE SMALLER DATA SETS THAT WE’VE TALKED ABOUT BEFORE. THE DOWNSIDE OF THIS IS THAT THESE CROWD WORKERS CAN OFTEN BE REVERSE ENGINEERED THERE ARE TRICKS AND CHEATS THAT THEY USE TO CREATE QUESTIONS. FOR EXAMPLE, IF A QUESTION ASKS “WHEN”, THEN ALL A QA SYSTEM REALLY HAS TO DO IS FIND THE THING THAT IS BEING ASKED ABOUT AND THEN FIND THE CLOSEST THING THAT LOOKS LIKE A DATE NEARBY. IN OTHER WORDS, THE MACHINE LEARNING ALGORITHMS ARE ESSENTIALLY REVERSE ENGINEERING THOSE ANNOTATORS RATHER THAN SOLVING THE INTRINSIC PROBLEM.

THE REAL REASON, IN MY OPINION, THAT SQuAD HAS BECOME THE GOLD STANDARD IS THAT IT HAS A VERY WELL TRAFFICKED LEADERBOARD WHERE PEOPLE CAN SUBMIT THEIR SYSTEMS AND TO SEE HOW THEY STACK UP AGAINST THE REST OF THE WORLD. BECAUSE THE ANSWERS ARE SPANS WITHIN THE ORIGINAL DOCUMENT THIS REDUCES THE TASK TO MULTIPLE-CHOICE MUCH LIKE THE DAILY MAIL/CNN DATA SET THAT WE TALKED BEFORE. THUS, IT DOESN’T ALLOW ANSWERS THAT ARE LATENT IN THE TEXT THAT AREN’T A FIXED SPAN. FOR EXAMPLE, IF THE QUESTION IS “WHAT STATES SHARE THE DELMARVA PENINSULA WITH MARYLAND” AND THE WIKIPEDIA PAGE HAS “THE DELMARVA

PENINSULA IS MADE UP OF DELEWARE, MARYLAND, AND VIRGINIA”, IT CAN’T ANSWER “DELAWARE AND VIRIGINA” BECAUSE THERE’S A PESKY MARYLAND IN THE MIDDLE!

ONE THING THAT REALLY GRINDS MY GEARS ABOUT SQuAD IS THAT THE SQuAD DATASET PROVIDES A “HUMAN UPPER BOUND” (AND TO BE CLEAR THIS IS NOT THE FAULT OF THE PEOPLE WHO DESIGNED SQuAD BUT FROM PEOPLE WHO MISINTERPRET WHAT THEY DID) BUT THERE LESS WELL INFORMED NEWS ORGANIZATIONS OR PROMOTERS OF RESEARCH THAT MISCHARACTERIZE WHAT IT MEANS TO BE NEAR THE HUMAN BOUND. IT ALSO MISCHARACTERIZES WHAT IT MEANS TO “READ” A DOCUMENT. IF MACHINE LEARNING SYSTEMS ARE GETTING CLOSE TO THIS HUMAN BOUND IT DOESN’T MEAN THAT COMPUTERS ARE READING BETTER THAN HUMANS. RATHER, IT MEANS THAT COMPUTERS ARE DOING A BETTER JOB OF HIGHLIGHTING PARTICULAR PASSAGES BETTER THAN UNDERPAID CROWD WORKERS. I’LL TALK MORE ABOUT WHAT I THINK IS A FAIR HUMAN COMPARISON IN A MOMENT.

SO PUTTING THIS ON THE CHART, I’M GOING TO PUT SQuAD A LITTLE BIT ABOVE CNN/DAILYMAIL. THEY’RE BOTH IMPLICITLY MULTIPLE CHOICE, BUT THERE ARE MORE SPANS THAN ENTITIES, SO ANSWERING SQuAD QUESTIONS IS A LITTLE MORE DIFFICULT.

8.5 Search Engines: The Source of Natural Questions

ONE CRITICISM OF SQuAD QUESTIONS IS THAT THEIR QUESTIONS ARE UN-NATURAL. THEY’RE WRITTEN BY CROWDWORKERS PAID TO WRITE ARTIFICIAL QUESTIONS. TWO DATASETS THAT USE “REAL” QUESTIONS ARE MS MARCO AND NATURAL QUESTIONS. THEY’RE BOTH SOMEWHAT SIMILAR, SO I’LL MOSTLY TALK ABOUT NATURAL QUESTIONS, WHICH IS NEWER AS A STAND-IN FOR BOTH. THESE QUESTIONS ARE HARDER THAN SQuAD AND TREC QUESTIONS! THAT’S BECAUSE THE QUESTIONS COME FROM PEOPLE USING SEARCH ENGINES. THESE PEOPLE DON’T ALWAYS KNOW WHAT THEY’RE ASKING ABOUT (THAT’S WHY THEY’RE USING A SEARCH ENGINE AFTER ALL), SO SOMETIMES THEY HAVE FALSE ASSUMPTIONS, AMBIGUITIES, AND OTHER PROBLEMS.

ANOTHER FACET OF THESE DATASETS THAT MAKE THE QUESTIONS HARDER IS THAT SEARCH ENGINES HAVE TRAINED PEOPLE OVER YEARS TO ANSWER QUESTIONS THEMSELVES BY FINDING WEBPAGES. NOBODY GETS GOOD ANSWERS WHEN THEY TYPE IN A QUESTION LIKE “WHAT IS THE LAST CAPITAL IN WESTERN EUROPE TO BE INVaded BY A FOREIGN ARMY” SO IF YOU TYPE THAT INTO GOOGLE YOU LIKELY WILL NOT GET AN ANSWER TO THAT QUESTION DIRECTLY. INSTEAD, YOU’LL PIECE TOGETHER MULTIPLE QUERIES TO ANSWER YOUR QUESTION.

IN OTHER WORDS, THE FAILURE OF SEARCH ENGINES TO ANSWER QUESTIONS OVER THE LAST COUPLE OF DECADES LIMITS THE USEFULNESS OF THE KINDS OF QUESTIONS THAT YOU CAN GET FROM SEARCH ENGINES TODAY.

SO HOW DOES NATURAL QUESTIONS FIND THE ANSWERS TO THESE QUESTIONS? THEY FIND WIKIPEDIA GIVEN THE QUESTION AS A SEARCH STRING:

AN INFORMATION RETRIEVAL PROBLEM! THEY THEN HAVE ANNOTATORS VALIDATE THAT THIS IS A GOOD SEARCH RESULT AND THEN HAVE THE ANNOTATORS HIGHLIGHT THE ANSWER STRING. NOW THIS ISN'T PERFECT.

HERE ARE SOME EXAMPLES WHERE I DISAGREE WITH SOME OF THE IMPLICIT ASSUMPTIONS THAT GET MADE IN THE NQ DATASET. FOR EXAMPLE, IT ASSUMES THAT WHEN YOU'RE ASKING ABOUT "MICHIGAN", YOU'RE ASKING ABOUT THE "UNIVERSITY OF MICHIGAN FOOTBALL TEAM", WHEN YOU'RE ASKING ABOUT "HOCKEY", YOU'RE ASKING ABOUT "MEN'S ICE HOCKEY", AND WHEN YOU'RE ASKING ABOUT "SUPREME COURT" YOU'RE ASKING ABOUT THE "INDIAN SUPREME COURT".

AS WE PUT THESE DATASETS UP ON THE CHART, I WANT TO MAKE A DISTINCTION BETWEEN THE "TRADITIONAL" AND "OPEN" INTERPRETATIONS OF NATURAL QUESTIONS. WHILE INITIAL NATURAL QUESTIONS WAS TREATED AS A SQuAD-LIKE DATASET, RECENT INTERPRETATIONS HAVE NOT ASSUMED THAT YOU KNOW WHICH WIKIPEDIA PAGE HAD THE ANSWER: YOU ALSO NEED TO DO THE IR COMPONENT AS WELL. THIS IS WHAT WE DID IN THE EFFICIENTQA COMPETITION (LINK IN THE DESCRIPTION): YOU CAN SEE HOW I MADE THE HUMAN VS. COMPUTER COMPARISON A LITTLE MORE FAIR.

OBVIOUSLY, THE OPEN INTERPRETATION MAKES IT MUCH HARDER TO FIND THE ANSWER.

MOVING ON ... YOU KNOW WHO'S BETTER THAN UNDERPAID CROWD WORKERS OR RANDOM PEOPLE ON THE INTERNET? PROFESSIONAL TRIVIA WRITERS. AND A NUMBER OF DATASETS USE QUESTIONS HARVESTED FROM THEM.

LET'S TALK ABOUT THREE OF THEM: SEARCHQA, TRIVIAQA, AND QUIZ BOWL.

SEARCHQA AND TRIVIAQA ACTUALLY LOOK A LOT LIKE NATURAL QUESTIONS, BUT INSTEAD OF THE QUERIES COMING FROM RANDOM PEOPLE ON THE INTERNET, THEY COME FROM PROFESSIONAL WRITERS. RESEARCHERS GO OUT AND FIND PAGES GIVEN THE QUERY, HIGHLIGHT THE ANSWERS. I'M NOT GOING TO DWELL TOO MUCH ON THESE DATASETS BECAUSE THE COPYRIGHT IS A LITTLE MURKY.

INSTEAD, LET'S NOW TURN TO MY HOBBY HORSE: QUIZ BOWL, WHICH IS PROBABLY BEST KNOWN AS AN NBC GAME SHOW IN THE US CALLED "COLLEGE BOWL" AND "UNIVERSITY CHALLENGE" IN THE UK. AND PLAYED BY THOUSANDS OF TRIVIA NERDS EVERY WEEKEND, PARTICULARLY AT THE HIGH SCHOOL AND COLLEGE LEVEL. AND UNLIKE SEARCHQA AND TRIVIAQA, THERE'S AN EXPLICIT COMMUNITY CONSENSUS THAT OLD QUIZ BOWL QUESTIONS ARE PUBLIC DOMAIN. SO AS A RESULT WE GET A LOT OF FREE DATA FROM EXPERTS. THESE ARE NOT CROWD WORKERS: THESE ARE PEOPLE WHO ARE PASSIONATE ABOUT HUMAN QUESTION ANSWERING, AND WE'RE USING THEIR EXPERTISE AND THEIR EXPERIENCE TO IMPROVE COMPUTER QUESTION ANSWERING.

THIS IS THE GOLD STANDARD IN HUMAN QUESTION ANSWERING (I.E., TRIVIA TOURNAMENTS), AND MY GOAL IS TO MAKE IT MORE POPULAR THAN IT CURRENTLY IS FOR COMPUTER QUESTION ANSWERING. WHAT'S DIFFERENT ABOUT QUIZ BOWL IS THAT THERE ISN'T JUST ONE QUESTION THAT YOU NEED TO ANSWER. INSTEAD, THERE ARE MULTIPLE CLUES THAT YOU COULD USE TO ANSWER

THE QUESTION.

THIS IS HARD TO EXPLAIN, SO LET'S SEE AN EXAMPLE. AS I READ IT, THINK ABOUT WHEN YOU KNOW THE ANSWER TO THE QUESTION. WHAT CLUE GAVE IT AWAY?

WHILE PASSING THROUGH THIS STATE, THE 6TH MASSACHUSETTS SUFFERED THE FIRST UNION CASUALTIES OF THE CIVIL WAR WHILE SUPPRESSING RIOTS THAT BEGAN AT ITS PRESIDENT STREET STATION. THE LINCOLN ADMINISTRATION REJECTED A RULING BY ROGER TANEY THAT A CITIZEN OF THIS STATE COULD NOT BE HELD WITHOUT HABEAS CORPUS IN THE CASE *EX PARTE MERRYMAN*. AFTER CAPTURING A COPY OF ROBERT E. LEE'S SPECIAL ORDER 191, GEORGE McCLELLAN FOUGHT A BATTLE IN THIS STATE THAT ENDED THE FIRST CONFEDERATE INVASION OF THE NORTH AND LED TO THE ANNOUNCEMENT OF THE EMANCIPATION PROCLAMATION. FOR 10 POINTS, THE BATTLE OF ANTIETAM WAS FOUGHT IN WHAT BORDER STATE WHERE FEDERAL TROOPS BLOCKED SECESSIONIST EFFORTS IN BALTIMORE?

THE CLUES ARE PRESENTED ONE BY ONE UNTIL SOMEONE CAN ANSWER THE QUESTION WITH THE GREAT STATE OF MARYLAND. BECAUSE THIS IS A COMPETITION, WHOEVER CAN ANSWER THE QUESTION FASTER (ON HARDER CLUES) USUALLY KNOWS MORE ABOUT THE TOPIC. BECAUSE THIS IS A TRIVIA GAME, THE GOAL IS TO GIVE MORE POINTS TO THE SMARTER PLAYER. QUIZ BOWL DOES THIS REALLY WELL! BETTER THAN SOME OTHER FUN TO WATCH KNOWLEDGE COMPETITIONS.

THE DOWNSIDE IS THAT SOMETIMES WHAT'S DIFFICULT FOR A HUMAN ISN'T DIFFICULT FOR A COMPUTER AND—JUST LIKE SQUAD—THERE SOMETIMES EASILY SOLVED QUESTIONS THAT ARE INTENDED TO BE DIFFICULT (FOR HUMANS) BUT AREN'T ACTUALLY DIFFICULT (FOR COMPUTERS).

FOR EXAMPLE, IN ONE OF THE EARLY COMPETITIONS WE HAD WITH HUMANS, A COMPUTER COULD EASILY RECOGNIZE A PAINTING FROM A DESCRIPTION.

BUT COMPUTERS STRUGGLE WITH SIMPLE LANGUAGE,

THE COMPUTER INCORRECTLY THINKS THAT WHEN THE QUESTION SAYS “THIS AUTHOR OPENS CRIME AND PUNISHMENT” THAT THE ANSWER IS DOSTOYEVSKI.

AS YOU CAN TELL FROM THESE CLIPS, ONE OF THE REASONS THAT I'M A BIG FAN OF THIS COMPETITION IS THAT IT LENDS ITSELF TO COMPARISONS. LET'S SAY YOU HAVE TWO COMPETITORS A AND B. WHAT IF WE TRY TO COMPARE THEM ON NORMAL QUESTIONS LIKE FROM NATURAL QUESTIONS? IF BOTH GET THE QUESTION RIGHT OR WRONG, YOU DON'T LEARN ANYTHING. YOU ONLY LEARN SOMETHING IF ONE GETS IT RIGHT AND THE OTHER GETS IT WRONG. HUMAN QUESTION WRITERS LEARNED REALLY EARLY ON THAT IT'S REALLY HARD TO GET THE LEVEL OF DIFFICULTY JUST RIGHT: EVEN WITH FANCY STATISTICS, YOU NEED HUNDREDS OF QUESTIONS TO RELIABLY RANK PEOPLE IN COMPETITIVE EXAMS.

SO THIS IS WHY THE HUMAN TRIVIA COMMUNITY TURNED TO QUIZ BOWL. INSTEAD OF JUST GETTING ONE PIECE OF INFORMATION PER QUESTION, YOU NOW POTENTIALLY GET A SIGNAL FROM EVERY WORD IN THE QUESTION. BECAUSE THE QUESTIONS ARE PACKED TO THE GILLS WITH CLUES THAT GO FROM

SUPER DIFFICULT TO TO ALMOST PAINFULLY EASY, AFTER EACH CLUE YOU GET INFORMATION ABOUT HOW MUCH A AND B KNOWS. THUS, IT LENDS ITSELF TO HUMAN-COMPUTER COMPETITION OR COMPARING COMPUTER SYSTEMS.

ANOTHER REASON THAT I LIKE QUIZ BOWL IS THAT THERE ARE MULTIPLE WAYS TO APPROACH IT COMPUTATIONALLY. BECAUSE THE ANSWERS ARE USUALLY ENTITIES, YOU CAN TAKE A STRAIGHT INFORMATION RETRIEVAL APPROACH: WHAT'S BEING TALKED ABOUT IN THIS QUESTION? OR YOU CAN TAKE A MORE MACHINE READING APPROACH. BOTH WORK WELL!

OKAY, THAT'S ENOUGH ABOUT QUIZ BOWL. LET'S PUT IT ON THE BOARD: THE ANSWER STRUCTURE IS ABOUT THE SAME AS NQ, BUT THE QUESTION STRUCTURE IS VERY COMPLICATED.

8.6 Natural Questions Five Years Later

IT'S BEEN FIVE YEARS SINCE NATURAL QUESTIONS CAME OUT. IT'S HARD TO OVERSTATE ITS IMPACT ON QUESTION ANSWERING (QA). PERHAPS ONLY SECOND TO SQuAD. SQuAD INTRODUCED THE EXTRACTIVE FORMAT, BUT SQuAD WAS "FAKE" ... NATURAL QUESTIONS HAD TRUE COMPLEXITY AND VERISIMILITUDE: PEOPLE ASK CRAZY QUESTIONS, SOME OF THE QUESTIONS WERE UNANSWERABLE, AND THE ANSWERS WERE TIED TO WIKIPEDIA PAGES. AND IT HAS YET TO BE ECLIPSED ... AND PERHAPS IT NEVER WILL.

PART OF THE REASON THAT THIS WILL PROBABLY NEVER HAPPEN AGAIN IS THAT NATURAL QUESTIONS HAS TWO VERY UNIQUE CHARACTERISTICS: QUESTIONS FROM REAL USERS EXPENSIVELY VERIFIED BY EXPENSIVE ANNOTATORS. WHILE IT'S POSSIBLE THAT WE MIGHT GET SOMETHING LIKE THIS AGAIN IN PART—UW FOLKS CRAWLED GOOGLE N-GRAMS TO EXTRACT THE QUESTIONS, WE PAID QUIZBOWLERS TO FIND PASSAGES TO ANSWER QUESTIONS—GETTING THEM BOTH AT ONCE IS PROBABLY NOT GOING TO HAPPEN AGAIN. BUT I'D LOVE TO BE PROVED WRONG! I'LL END THE VIDEO WITH HOW I'D BUILD A QA DATASET WITH INFINITE RESOURCES.

HOWEVER, NATURAL QUESTIONS AREN'T PERFECT. SEWON MIN AND COMPANY CREATED AMBIGQA THAT SHOWED SOME OF THE QUESTIONS IN NATURAL QUESTIONS WERE AMBIGUOUS. HERE ARE SOME EXAMPLES THAT BEN AND I TALKED ABOUT IN OUR 2020 PAPER (AND I TALKED ABOUT AT MORE LENGTH IN THE QA DATASETS LECTURE, LINK IN THE DESCRIPTION). THERE ARE MULTIPLE INTERPRETATIONS TO THESE QUESTIONS, BUT NATURAL QUESTIONS ONLY SAYS THAT ONE OF THEM IS CORRECT.

HOW IS THE "CORRECT" INTERPRETATION CHOSEN IN THE ORIGINAL NATURAL QUESTIONS PAPER? THE ANNOTATORS SAW SEARCH RESULTS FROM AN INFORMATION RETRIEVAL SYSTEM, AND IF THEY FOUND AN ANSWER THERE, THEY CALLED IT THE CORRECT ANSWER. SO WHATEVER PROCESS BIASES SEARCH RESULTS TOWARD ENGLISH-SPEAKING JUDGES, MEN'S SPORTS, OR FLAGSHIP UNIVERSITIES GETS RECAPITULATED IN NATURAL QUESTIONS.

ANOTHER DEVELOPMENT IS THE REALIZATION THAT NOT ALL QUESTIONS ARE

VALID AND SHOULDN'T BE ANSWERED. THERE ARE MULTIPLE ASPECTS OF THIS: APPROPRIATELY CALIBRATING YOUR QA SYSTEM TO ABSTAIN WHEN IT DOESN'T KNOW THE ANSWER FOR EXAMPLE. NATURAL QUESTIONS ACTUALLY DOES A PRETTY GOOD JOB OF SUPPORTING THIS BY HAVING UNANSWERABLE QUESTIONS.

BUT INSTEAD, I WANT TO FOCUS ON LINGUISTIC PHENOMENA THAT *CAUSES* A QUESTION TO *BECOME* UNANSWERABLE. THIS LINE OF RESEARCH BEGAN WITH WHAT HAS COME TO BE KNOWN AS AS "THE LIGHTBULB PAPER" FROM THE TITLE "WHICH LINGUIST INVENTED THE LIGHTBULB", WHICH HIGHLIGHTS THE PROBLEM: THERE'S A WELL-DEFINED ANSWER TO THE QUESTION "WHO INVENTED THE FIRST COMMERCIALLY SUCCESSFUL LIGHT BULB", BUT IT SURE AS HECK WASN'T A LINGUIST ... ANSWERING THAT QUESTION WOULD PERHAPS NOT BE USEFUL TO A USER.

THERE ARE TWO TAKEAWAYS HERE: FIRST THAT MANY QA SYSTEMS TEND TO ROLL WITH QUESTIONS LIKE "WHEN DID PABLO PICASSO PAINT THE MONA LISA" AND ANSWER 1503, HUNDREDS OF YEARS BEFORE PICASSO WAS BORN. SECOND, THAT MANY OF THE UNANSWERABLE QUESTIONS IN NATURAL QUESTIONS ARE UNANSWERABLE BECAUSE OF THESE PHENOMENA. AND IT'S NOT JUST NATURAL QUESTIONS: YU ET AL SHOWED THAT THIS ALSO SHOWS UP ON REDDIT. IT'S AN IMPORTANT PROBLEM!

BUT WE DON'T ACTUALLY KNOW *WHY* A PARTICULAR NATURAL QUESTION IS UNANSWERABLE. IS IT BECAUSE WIKIPEDIA LACKS THE INFORMATION, BECAUSE THERE'S A FALSE ASSUMPTION, BECAUSE THERE'S ACTUALLY DISPUTE ABOUT THE CANONICAL ANSWER TO A QUESTION, OR BECAUSE IT'S A MATTER OF TASTE / OPINION? BECAUSE NATURAL QUESTIONS WAS FOCUSED ON FACTUAL QA (A REASONABLE POSITION FIVE YEARS AGO), WE DON'T HAVE A SINGLE COMPREHENSIVE DATASET TO ADDRESS ALL OF THESE ISSUES (THE NUMBERS HERE ARE JUST FOR A SMALL SAMPLE OF NQ).

AND PERHAPS WE DON'T NEED THAT ANY MORE. WITH RLHF, WE CAN JUST MIX A BUNCH OF DATASETS TOGETHER. PERHAPS WE DON'T NEED A MONOLITHIC QA TRAINING DATASET TO RULE THEM ALL. WE JUST NEED ROBUST EVALUATION PROTOCOLS (WHICH CAN BE MUCH SMALLER).

AND THIS IS, AGAIN, A WEAKNESS OF NATURAL QUESTIONS ... IT'S NOT A GREAT EVALUATION SET. LEWIS ET AL. DISCOVERED THAT MANY QA DATASETS, INCLUDING NATURAL QUESTIONS, HAVE A HUGE OVERLAP BETWEEN THE TRAINING SET AND THE TESTING SET. I'M A PROFESSOR, SO LET'S PUT THIS IN A CLASSROOM CONTEXT: THIS IS LIKE GIVING OUT TWO THIRDS OF THE ANSWERS AND ONE THIRD OF THE QUESTIONS AS A STUDY GUIDE BEFORE THE TEST. NOW, IT'S A GREAT SANITY CHECK THAT SYSTEMS CAN ANSWER THOSE QUESTIONS, BUT THAT MEANS THAT ONLY A THIRD OF THE QUESTIONS ARE PART OF A FUNCTIONAL TEST SET. ADD THAT TO THE PROBLEMS OF AMBIGUITY, AND FAR LESS OF THE TEST SET IS ACTUALLY USABLE, MAKING IT DIFFICULT TO KNOW IF WE'RE MAKING PROGRESS ON THE DATASET.

THIS IS SOMETHING THAT THE TRIVIA COMMUNITY HAS KNOWN FOR A WHILE. YOU CAN'T REUSE CLUES! FOR IMPORTANT TEST SETS (LIKE NATIONAL TRIVIA

TOURNAMENTS), YOU NEED TO MAKE SURE THAT YOU'RE FINDING NEW, INTERESTING CLUES THAT HAVEN'T BEEN USED BEFORE (OR AT LEAST FOR A WHILE). THIS IS WHY WE'VE ALWAYS USED THE FLAGSHIP NATIONAL TRIVIA TOURNAMENTS AS OUR TEST SETS FOR QUIZ BOWL.

SO THAT'S ALL I HAVE ABOUT THE PROBLEMS OF NATURAL QUESTIONS. BUT I WANT TO MAKE IT CLEAR THAT *NATURAL QUESTIONS REMAINS A GREAT RESOURCE*. IT HASN'T BEEN MATCHED. I STILL USE IT. I'M JEALOUS OF THE PEOPLE WHO GOT TO BUILD IT. BUT IT'S IMPORTANT TO KNOW THE LIMITATIONS. IS THERE ANYTHING I GOT WRONG, ANYTHING I MISSED? IF SO, LET ME KNOW IN THE COMMENTS.

NOW, FOR ALL OF THE REASONS I TALKED ABOUT BEFORE, IT PROBABLY WILL NEVER BE DONE AGAIN. NOBODY IS CRAZY ENOUGH TO DO IT. BUT HOW WOULD I DO IT IF SOMEONE WERE FOOLISH ENOUGH TO GIVE ME THE RESOURCES TO CREATE THE NEXT ITERATION OF SOMETHING LIKE IT?

BACK IN THE EARLY 2000S, THERE WAS A QA SERVICE CALLED AARDVARK. THE IDEA WAS THAT IT WOULD CONNECT UP REAL HUMANS WITH TOPICS THOSE SMART HUMANS COULD ANSWER QUESTIONS ABOUT. QUESTION COMES IN AND IT GETS ROUTED TO THAT PERSON. THAT WAS A GREAT IDEA BUT BEFORE IT'S TIME. WHAT YOU'D DO TODAY IS BUILD UP A LARGE COLLECTION OF THESE PEOPLE (AND PAY THEM WELL), A QUESTION FROM A MAJOR SEARCH ENGINE COMES IN ... WE TRY TO DETECT IF THE COMPUTER CAN ANSWER IT DIRECTLY. IF YES, WE ANSWER IT, IF NOT, WE FARM IT OUT TO THE EXPERTS. WE'VE TRAINED THE EXPERTS TO INTERROGATE THE QUESTION ASKERS TO REMOVE AMBIGUITIES, FALSE PRESUPPOSITIONS, ETC. TO CREATE A "CLEAN" VERSION OF THE QUESTION AND THEN A STRONG REFERENCE ANSWER SUPPORTED BY RELEVANT INFORMATION.

YOGESH RAUT'S RECREATIONAL THINKING PODCAST IS AN EXAMPLE OF THIS INTERROGATION PROCESS: THEY SPEND SEVERAL MINUTES INTERROGATING INDIVIDUAL QUESTIONS, FIGURING OUT POSSIBILITIES, AND THEN JUSTIFYING THE ANSWER.

SO FOR EXAMPLE, IF THE QUESTION COMES IN: "WHAT LINGUIST INVENTED THE LIGHTBULB", YOU MIGHT HAVE AN INTERACTION THAT GOES LIKE THIS:

"UM, ACTUALLY, MOST OF THE PEOPLE CREDITED WITH THE DEVELOPMENT OF THE LIGHTBULB WERE CHEMISTS, PHYSICISTS, OR ENGINEERS"

"OKAY, SMARTBUTT, WHO MADE THE FIRST LIGHTBULB"

"WELL, DO YOU MEAN CREATED THE FIRST ELECTRIC LIGHT (HUMPHRY DAVY), THE FIRST FILAMENT INSIDE A VACUUM TUBE (JOSPEH SWAN), THE FIRST METAL FILAMENT IN A VACUUM TUBE (WARREN DE LA RUE), OR THE FIRST COMMERCIAL SUCCESSFUL LIGHTBULB WITH THOSE PROPERTIES (THOMAS EDISON)?"

"HMM ... I GUESS I MEANT WHO GOT THE CREDIT FOR IT"

"THEN THAT WOULD BE THOMAS EDISON"

AND TO BE CLEAR, THAT'S NOT WHAT HAPPENS IF YOU ASK THIS TO CHATGPT ... IT DOESN'T REALLY PROBE THE USER'S INTENT BUT TRIES TO DO SOME CORRECTION (NOT INVENTED BY A LINGUIST) AND PROVIDE THE ANSWER TO

THE MOST LIKELY QUESTION. BUT STILL ERASES ALL OF THE PEOPLE WHOSE WORK EDISON BUILT ON.

BUT WE CAN BUILD SUCH A SYSTEM—AGAIN, MY DMs ARE OPEN—NATURAL QUESTIONS ARE STILL THE BEST WE HAVE.

8.7 Is IID the Right Choice of Data?

CREATING NEW MACHINE LEARNING SYSTEMS REQUIRES DATA. THESE DATA NEED TO COME FROM SOMEWHERE, AND NLP SYSTEMS ARE NO DIFFERENT! IF YOU WERE TO ASK A STATISTICIAN, THEY WOULD TELL YOU GET A SAMPLE FROM AN IID DISTRIBUTION: INDEPENDENT AND IDENTICALLY DISTRIBUTED. INDEPENDENT MEANS THAT EXAMPLE x_n DOES NOT DEPEND ON ANY OF THE PREVIOUS EXAMPLES 1 TO $n-1$, INDEED, ALL OF THE EXAMPLES COME FROM THE *SAME* UNDERLYING DISTRIBUTION.

BUT THAT'S THE PROBLEM. LANGUAGE ISN'T LIKE NEUTRINOS SPEWED FROM THE SUN OR WATER FROM A GEYSER. LANGUAGE IS INHERENTLY AN INTERACTIVE PROCESS: AND WHAT I'M GOING TO ARGUE IS THAT THIS HAS ALWAYS BEEN A PROBLEM FOR OUR NLP DATASETS, THAT IT'S GETTING WORSE WITH THE ADVENT OF MUPPET MODELS, AND THAT WE NEED TO ADDRESS THE PROBLEM HEAD ON GOING FORWARD.

The Ghost of NLP Data Past

ONE OF THE BIGGEST SOURCES OF NLP DATA ARE CROWDWORKERS. THESE ARE NOT IID! CROWDWORKERS ARE NOT REPRESENTATIVE OF THE GLOBAL POPULATION: BY DEFINITION THEY HAVE ACCESS TO A COMPUTER, THEY'RE TYPICALLY YOUNGER, AND THEY'RE FOCUSED IN PARTICULAR GEOGRAPHIC AREAS.

AND EVEN IF YOU RESTRICT TO US CROWDWORKERS (BECAUSE NOBODY HAS VIRTUAL PRIVATE NETWORKS OR VIRTUAL CREDIT CARDS), THIS DISTRIBUTION DOES NOT MATCH THE OVERALL US POPULATION.

AND BEYOND THE UNDERLYING DISTRIBUTION ISSUE, THERE'S A QUALITY ISSUE. CROWDSOURCED DATA OFTEN HAS “SHORTCUTS” THAT ALLOW SYSTEMS TO “CHEAT”. WE'VE SEEN THIS FOR: ENTAILMENT: SYSTEMS CAN DETECT WHETHER A CLAIM IS TRUE OR NOT JUST BY LOOKING AT THE PREMISE OR THE HYPOTHESIS (WHEN YOU SHOULD REALLY NEED BOTH). GRURANGAN ET AL SHOW THAT CROWDWORKERS USE THE SAME TRICKS TO CREATE EXAMPLES: “AT LEAST” IS OFTEN ENTAILED, IRRELEVANT INFORMATION IS OFTEN NEUTRAL, AND NEGATION IS OFTEN CONTRADICTED. POLIAK ET AL SHOWED THAT YOU CAN GET PRETTY HIGH ACCURACY BY IGNORING THE PREMISE AND CHARACTERIZED THE “TELLS” THAT OFTEN APPEARED IN THE DATA. AND WE HAD OUR OWN PAPER THAT EVEN WHEN THE PREMISE IS REQUIRED, IT DOESN'T REQUIRE DEEP REASONING. FOR FACT CHECKING: LIKE WITH ENTAILMENT, SOME LOGICAL OPERATORS LIKE NEGATIONS ARE MORE LIKELY ASSOCIATED WITH REFUTATION. WE TALKED ABOUT THIS IN OUR ADVERSARIAL FACT CHECKING PAPER, FOOL ME TWICE.

AND FOR QUESTION ANSWERING: FOR DATASETS LIKE SQuAD, THE ANSWER SELECTION PROCESS OFTEN DIDN'T NEED DEEP UNDERSTANDING OF THE SOURCE CONTEXT AS DESCRIBED BY WEISSENBORN ET AL. LOOK AT THE PASSAGE AND FIND THE FIRST THING THAT COULD POSSIBLY FIT THE ANSWER TYPE: WHEN FINDS THE YEAR, WHO FINDS A PERSON, WHERE FINDS A PLACE, ETC.

THIS IS BECAUSE CROWDWORKERS HAD TO LOOK AT A PASSAGE AND GENERATE A QUESTION THAT COULD BE ANSWERED BY THAT PASSAGE. THE QUESTIONS DIDN'T REQUIRE THAT MUCH REASONING BECAUSE THE CROWDWORKERS WANTED TO GET DONE AS QUICKLY AS POSSIBLE.

SPEAKING OF QUESTION ANSWERING, NATURAL QUESTIONS CLAIM TO BE "NATURAL" ... THESE ARE THE DATA FROM REAL USERS, THIS ACTUALLY IS IID RIGHT? AND I THINK THE ANSWER IS NO! THIS IS BECAUSE ALL LANGUAGE IS A NEGOTIATION—A GIVE AND TAKE BETWEEN TWO PARTIES—AND GOOGLE IS A PART OF THAT NEGOTIATION.

NOW, GOOGLE IS A FANTASTIC TOOL, BUT IT HAS ITS STRENGTHS AND WEAKNESSES. IT'S REALLY GOOD AT LOOKING UP WHO PLAYED ROLE X IN MOVIE Y. IMDB (AND OTHER STRUCTURED KNOWLEDGE BASES) MAKE ANSWERING THIS KIND OF QUESTION EASY ... SO PEOPLE KNOW THAT IF THEY ASK THIS KIND OF QUESTION FROM GOOGLE THEY'RE LIKELY TO GET A GOOD ANSWER. IN GENERAL, IF I CAN CLEARLY SPECIFY ENTITY X AND A RELATIONSHIP F(X) THAT'S GOING TO APPEAR IN SOME KB, THEN I'M GOING TO GET A GREAT ANSWER.

SO THIS MEANS THAT PEOPLE ARE MORE LIKELY TO ASK THIS KIND OF QUESTION. BUT LONGER QUESTIONS JUST HARDER TO ANSWER. PEOPLE WANT TO GET ANSWERS TO THEIR QUESTIONS, SO GOOGLE HAS TRAINED USERS OVER THE LAST 25 YEARS TO ASK THE KINDS OF THINGS IT CAN ANSWER (AND THE AVERAGE LENGTH OF QUERIES IS GOING UP!). BUT AT THE END OF THE DAY, THE KINDS OF QUESTIONS THAT GOOGLE ISN'T SO GOOD AT ARE NOT GOING TO BE ASKED SO THEY'RE NOT GOING TO BE A PART OF THE DATASET. AND THAT MEANS THAT IT WON'T BE PART OF THE ALL IMPORTANT NATURAL QUESTIONS LEADERBOARD. WHICH MEANS THAT PEOPLE WILL IGNORE THE QUESTIONS THAT GOOGLE ISN'T ALREADY GOOD AT.

NOW, JORDAN, AREN'T YOU JUST BEING A MEANIE BY PICKING ON NATURAL QUESTIONS? YOUR FAVORITE QA DATASET THAT YOU KEEP PUSHING DOWN EVERYBODY'S THROAT—QUIZ BOWL—AIN'T IID EITHER.

AND THAT'S ABSOLUTELY RIGHT! I LIKE THESE CRAZY QUIZ BOWL QUESTIONS WHERE THE DISTRIBUTION HAS BEEN CHOSEN BY HAND TO REFLECT THE UNDERGRADUATE CURRICULUM AT AMERICAN UNIVERSITIES. THIS ALSO ISN'T IID, BUT I THINK SPECIFYING THE DISTRIBUTION AND BEING ABLE TO DEFEND IT IS A GOOD THING!

AND IN ADDITION TO THE "KINDS" OF QUESTIONS IN NATURAL QUESTIONS, GOOGLE HAS TRAINED USERS TO ASK QUESTIONS IN A PARTICULAR WAY: INSTEAD OF ASKING "WHO IS AMAL ALAMUDDIN'S HUSBAND", PEOPLE TYPE "AMAL ALAMUDDIN HUSBAND" BECAUSE GOOGLE STARTED AS A SEARCH ENGINE THAT IGNORED FUNCTION WORDS. SO THIS ALSO AFFECTS WHAT GETS INTO NATU-

RAL QUESTIONS: GOOGLE HAS TRAINED PEOPLE TO WRITE TERSE QUERIES, NQ SUBSELECTS FROM THAT, AND THAT BECOMES THE LEADERBOARD.

FROM A SOCIO-TECHNICAL PERSPECTIVE, GOOGLE IS TO THE TURN OF THE CENTURY LIKE MUPPET MODELS ARE TO TODAY [LINK IN THE DESCRIPTION IF YOU DON'T KNOW WHY I CALL THINGS LIKE GPT MUPPET MODELS]. PEOPLE ARE LEARNING HOW TO SKILLFULLY PROMPT MUPPET MODELS, MUPPET MODELS ARE IMPERSONATING HUMANS ON CROWDWORKING PLATFORMS (ANOTHER REASON THAT DATASETS FROM CROWDWORKERS AREN'T IID), PEOPLE ARE USING THEM TO HELP WRITE THEIR ESSAYS (WHICH WILL IN TURN BE USED TO TRAIN THE NEXT GENERATION OF MUPPET MODELS WHEN THOSE ESSAYS GET POSTED ON THE WEB).

SO WHAT DOES THIS MEAN?

OKAY, SO I'M GOING TO GET ON MY SOAPBOX FOR A BIT ... WE NEED TO EMBRACE THE INTERACTION BETWEEN USERS AND SYSTEMS. AND THAT MEANS SOMETHING LIKE ADVERSARIAL EXAMPLES WHICH WE'LL TALK ABOUT LATER IN LATER LECTURES ... LINK TO A LECTURE IF YOU'RE IMPATIENT. BUT THE BASIC IDEA IS THAT YOU HAVE HUMANS AND MODELS INTERACT WITH EACH OTHER TO CREATE EXAMPLES THAT "PUSH THE FRONTIER" OF WHAT'S POSSIBLE.

SAM BOWMAN IN HIS PAPER ABOUT THE DANGERS OF UNDERCLAIMING ... IF YOU DO RESEARCH ON ADVERSARIAL DATA, THE ACCURACY NUMBERS ARE GOING TO BE LOWER THAN "NORMAL" DATASETS. BUT I THINK THAT THIS MAKES A CATEGORY ERROR! THERE'S NO SUCH THING AS A NORMAL, IID DATASET. ALL DATASETS HAVE THEIR OWN PATHOLOGIES, AND YOU CANNOT READILY COMPARE SQuAD NUMBERS TO NQ NUMBERS: THERE ARE FUNDAMENTALLY DIFFERENT UNDERLYING DISTRIBUTIONS AND DATA COLLECTION PRACTICES. THE SAME THING IS TRUE FOR MY BELOVED QUIZ BOWL FORMAT.

AND I THINK THIS IS A HEALTHY DEVELOPMENT. RATHER THAN OBSESSING ABOUT RANK ON A LEADERBOARD AND THE SINGLE FLOATING POINT NUMBER WITH YOUR ACCURACY THAT COMES OUT, WE SHOULD FOCUS ON WHAT ARE THE ABILITIES (AND WEAKNESSES) OF MODELS. AND ACCEPTING THAT THERE IS NO SUCH THING AS THE IID LANGUAGE BUNNY BRINGING YOU A BASKET OF NICELY DISTRIBUTED QUESTIONS WILL HELP US GET THERE FASTER.

MY GOAL WITH THIS VIDEO IS TO GIVE AN OVERVIEW OF THE KINDS OF QUESTION ANSWERING DATASETS THAT ARE OUT THERE, CONTRAST THEM, AND GIVE MY OPINION ABOUT WHICH IS THE BEST ONE.

FIRST, LET ME START OFF BY SAYING WHAT WE'RE NOT TALKING ABOUT! THERE ARE MANY QUESTION ANSWERING DATASETS OUT THERE THAT ARE ESSENTIALLY LOGIC PROBLEMS OR SQL QUERY LOOKUPS AND WE'RE NOT GOING TO TALK ABOUT THEM. THAT'S A VERY VALID FORM OF QUESTION ANSWERING AND IT ESSENTIALLY CORRESPONDS TO A FORM OF PARSING WHERE YOU HAVE SOME TEXT AND YOU NEED TO TURN THAT INTO A LOGICAL FORM OR AN SQL STATEMENT. IF YOU PARSE IT CORRECTLY, YOU GET THE ANSWER.

INSTEAD, WE'RE GOING TO TALK ABOUT MORE OPEN-ENDED QUESTION ANSWERING THAT'S OFTEN CALLED "MACHINE READING" QUESTION ANSWERING

FOR REASONS THAT WILL BE CLEAR IN A SECOND. IN THESE QUESTION ANSWERING DATASETS, YOU NEED TO MAKE SENSE OF A LARGE BODY OF TEXT TO ANSWER QUESTIONS. MOREOVER, YOU USUALLY NEED TO BE ABLE TO ANSWER QUESTIONS ABOUT JUST ABOUT ANYTHING, NOT JUST A NARROW SET OF TOPICS.

WE'LL TALK ABOUT MODELS FOR ANSWERING THESE QUESTIONS LATER, BUT FOR THE MOMENT THINK ABOUT THE KINDS OF THINGS THAT YOUR MODEL WOULD NEED TO DO TO ANSWER THE QUESTION. IN MANY CASES, IT WILL BE EXACTLY THE SAME KIND OF THING YOU WOULD NEED TO DO TO ANSWER THE QUESTION: FIND A DOCUMENT THAT CONTAINS THE ANSWER, READ THROUGH THE DOCUMENT TO FIND THE ANSWER, AND THEN CONVEY THAT ANSWER.

OKAY, BACK TO THE DATASETS. RATHER THAN JUST LISTING THEM OFF, I'M GOING TO TRY TO ORGANIZE THE DATASETS BASED ON: HOW MUCH INFORMATION YOU HAVE TO ANSWER THE QUESTION AND HOW COMPLICATED THE ANSWER IS. LET'S SAY THAT WE CAN WRITE THIS AS A TWO DIMENSIONAL PLOT WITH PROVIDED INFORMATION AS THE X AXIS AND ANSWER COMPLEXITY AS THE Y AXIS. I'LL ROUGHLY GO CHRONOLOGICALLY THROUGH TIME.

IN ANOTHER VIDEO, THERE'S ANOTHER IMPORTANT DIMENSION: WHO'S WRITING THE QUESTION AND FOR WHAT PURPOSE, BUT I THINK THAT'S SO IMPORTANT THAT I WANT TO FOCUS ON IT BY ITSELF. SO IF YOU'RE INTERESTED IN THAT, CHECK OUT THE CRANFIELD VS. MANCHESTER DISCUSSION LINKED IN THE DESCRIPTION.

LET'S START WITH THE MOST SIMPLE EXAMPLE OF A QUESTION ANSWERING DATASET: YOU GET A QUESTION AS INPUT AND YOU NEED TO PRODUCE A SHORT ANSWER. THIS IS OFTEN CALLED FACTOID QA. THE PROTOTYPICAL EXAMPLE OF THIS IS TREC QA. THIS IS A SUPER OLD SCHOOL DATASET, BUT IMPORTANT TO REMEMBER, AS IT SET THE FORM OF EVERYTHING THAT CAME AFTER.

THE ACTUAL DATASET IS TINY: ONLY A COUPLE HUNDRED QUESTIONS. AND THE OFFICIAL EVALUATION PROCESS SEEMS IMPOSSIBLE BY TODAY'S STANDARDS: YOU HAD PROFESSIONAL ANNOTATORS JUDGE WHETHER A SYSTEM PROVIDED THE RIGHT ANSWER. WHILE EVERY OTHER DATASET WE'LL TALK ABOUT HAS MORE AUTOMATIC EVALUATIONS, THIS IS PROBABLY THE BEST WAY TO DO IT, AT LEAST FOR SOME TRICKIER ANSWERS.

THIS IS OUR FIRST QA DATASET, SO LET'S PUT THIS ONTO THE BOARD HERE. LET'S NOW GO INTO THE MODERN AGE!

ONE OF THE FIRST MACHINE READING DATA SETS WAS THE CNN / DAILY MAIL DATA SET FROM KARL MORITZ HERMANN AND COMPANY. HERE YOU HAVE A BUNCH OF NEWS ARTICLES. THIS DATASET STARTED THE TREND OF CONDITIONING A QUESTION ON A DOCUMENT: IN THIS CASE A NEWS ARTICLE. THEY ASKED YOU TO ANSWER QUESTIONS ABOUT THOSE NEWS ARTICLES. FOR EXAMPLE, "WHO REACHED THE CRICKET WORLD CUP". SOMEWHERE IN THE DOCUMENT WAS THE ANSWER "NEW ZEALAND" BUT YOU NEED TO FIND THAT.

ONE THING THAT IS CONTROVERSIAL IS THAT YOU DIDN'T ACTUALLY FIND THE STRING "NEW ZEALAND" BUT RATHER AN ANONYMIZED SYMBOL. THIS MEANT THAT YOU COULDN'T JUST HAVE A LANGUAGE MODEL MEMORIZE ALL THE

ANSWERS (WHICH IS EVEN MORE OF A PROBLEM NOW THAT IT WAS BACK THEN). HOWEVER, THIS HAD ANOTHER DRAWBACK! BECAUSE THERE WEREN'T THAT MANY SYMBOLS, THE TASK WAS ESSENTIALLY A MULTIPLE-CHOICE QUESTION: YOU IDENTIFIED THE ENTITIES IN THE PARAGRAPH AND THEN YOU PICKED WHICH ONE WAS CORRECT. DESPITE THIS PROBLEM AND THAT INTRODUCED A LITTLE BIT OF NOISE TO THE PROBLEM, YOU CAN SEE THE INFLUENCE OF THIS QUESTION ANSWERING DATASET ON THE REST OF THE FIELD, PARTICULARLY ON “THE NEXT BIG THING”: SQuAD.

THE STANFORD QUESTION ANSWERING DATASET HAS BECOME THE GOLD STANDARD. WHILE IT MIGHT NOT BE THE BEST DATASET, IT'S THE MOST POPULAR. IF YOU DON'T REPORT RESULTS ON SQuAD, YOU NEED TO EXPLAIN WHY. THIS 2016 DATASET SET HELPED POPULARIZED THE MODERN UPSURGE OF QUESTION ANSWERING. THESE QUESTIONS USE PARAGRAPHS FROM WIKIPEDIA ASK CROWD WORKERS TO WRITE QUESTIONS ABOUT THAT PARAGRAPH. THE QUESTIONS ARE INTENDED TO HAVE AN ANSWER THAT IS A SPAN WITHIN THAT WIKIPEDIA PARAGRAPH. ONE WAY OF THINKING ABOUT THIS IS THAT A SPAN IS A CONTIGUOUS SPAN OF CHARACTERS: YOU HIGHLIGHT THAT SPAN TO GIVE THE ANSWER.

THIS IS A VERY LARGE DATA SET SO IT DOESN'T HAVE MANY OF THE PROBLEMS OF THE SMALLER DATA SETS THAT WE'VE TALKED ABOUT BEFORE. THE DOWNSIDE OF THIS IS THAT THESE CROWD WORKERS CAN OFTEN BE REVERSE ENGINEERED THERE ARE TRICKS AND CHEATS THAT THEY USE TO CREATE QUESTIONS. FOR EXAMPLE, IF A QUESTION ASKS “WHEN”, THEN ALL A QA SYSTEM REALLY HAS TO DO IS FIND THE THING THAT IS BEING ASKED ABOUT AND THEN FIND THE CLOSEST THING THAT LOOKS LIKE A DATE NEARBY. IN OTHER WORDS, THE MACHINE LEARNING ALGORITHMS ARE ESSENTIALLY REVERSE ENGINEERING THOSE ANNOTATORS RATHER THAN SOLVING THE INTRINSIC PROBLEM.

THE REAL REASON, IN MY OPINION, THAT SQuAD HAS BECOME THE GOLD STANDARD IS THAT IT HAS A VERY WELL TRAFFICKED LEADERBOARD WHERE PEOPLE CAN SUBMIT THEIR SYSTEMS AND TO SEE HOW THEY STACK UP AGAINST THE REST OF THE WORLD. BECAUSE THE ANSWERS ARE SPANS WITHIN THE ORIGINAL DOCUMENT THIS REDUCES THE TASK TO MULTIPLE-CHOICE MUCH LIKE THE DAILY MAIL/CNN DATA SET THAT WE TALKED BEFORE. THUS, IT DOESN'T ALLOW ANSWERS THAT ARE LATENT IN THE TEXT THAT AREN'T A FIXED SPAN. FOR EXAMPLE, IF THE QUESTION IS “WHAT STATES SHARE THE DELMARVA PENINSULA WITH MARYLAND” AND THE WIKIPEDIA PAGE HAS “THE DELMARVA PENINSULA IS MADE UP OF DELEWARE, MARYLAND, AND VIRGINIA”, IT CAN'T ANSWER “DELAWARE AND VIRGINIA” BECAUSE THERE'S A PESKY MARYLAND IN THE MIDDLE!

ONE THING THAT REALLY GRINDS MY GEARS ABOUT SQuAD IS THAT THE SQuAD DATASET PROVIDES A “HUMAN UPPER BOUND” (AND TO BE CLEAR THIS IS NOT THE FAULT OF THE PEOPLE WHO DESIGNED SQuAD BUT FROM PEOPLE WHO MISINTERPRET WHAT THEY DID) BUT THERE LESS WELL INFORMED NEWS ORGANIZATIONS OR PROMOTERS OF RESEARCH THAT MISCHARACTERIZE WHAT IT MEANS TO BE NEAR THE HUMAN BOUND. IT ALSO MISCHARACTERIZES WHAT IT

MEANS TO “READ” A DOCUMENT. IF MACHINE LEARNING SYSTEMS ARE GETTING CLOSE TO THIS HUMAN BOUND IT DOESN’T MEAN THAT COMPUTERS ARE READING BETTER THAN HUMANS. RATHER, IT MEANS THAT COMPUTERS ARE DOING A BETTER JOB OF HIGHLIGHTING PARTICULAR PASSAGES BETTER THAN UNDERPAID CROWD WORKERS. I’LL TALK MORE ABOUT WHAT I THINK IS A FAIR HUMAN COMPARISON IN A MOMENT.

SO PUTTING THIS ON THE CHART, I’M GOING TO PUT SQuAD A LITTLE BIT ABOVE CNN/DAILYMAIL. THEY’RE BOTH IMPLICITLY MULTIPLE CHOICE, BUT THERE ARE MORE SPANS THAN ENTITIES, SO ANSWERING SQuAD QUESTIONS IS A LITTLE MORE DIFFICULT.

ONE CRITICISM OF SQuAD QUESTIONS IS THAT THEIR QUESTIONS ARE UN-NATURAL. THEY’RE WRITTEN BY CROWDWORKERS PAID TO WRITE ARTIFICIAL QUESTIONS. TWO DATASETS THAT USE “REAL” QUESTIONS ARE MS MARCO AND NATURAL QUESTIONS. THEY’RE BOTH SOMEWHAT SIMILAR, SO I’LL MOSTLY TALK ABOUT NATURAL QUESTIONS, WHICH IS NEWER AS A STAND-IN FOR BOTH. THESE QUESTIONS ARE HARDER THAN SQuAD AND TREC QUESTIONS! THAT’S BECAUSE THE QUESTIONS COME FROM PEOPLE USING SEARCH ENGINES. THESE PEOPLE DON’T ALWAYS KNOW WHAT THEY’RE ASKING ABOUT (THAT’S WHY THEY’RE USING A SEARCH ENGINE AFTER ALL), SO SOMETIMES THEY HAVE FALSE ASSUMPTIONS, AMBIGUITIES, AND OTHER PROBLEMS.

ANOTHER FACET OF THESE DATASETS THAT MAKE THE QUESTIONS HARDER IS THAT SEARCH ENGINES HAVE TRAINED PEOPLE OVER YEARS TO ANSWER QUESTIONS THEMSELVES BY FINDING WEBPAGES. NOBODY GETS GOOD ANSWERS WHEN THEY TYPE IN A QUESTION LIKE “WHAT IS THE LAST CAPITAL IN WESTERN EUROPE TO BE INVaded BY A FOREIGN ARMY” SO IF YOU TYPE THAT INTO GOOGLE YOU LIKELY WILL NOT GET AN ANSWER TO THAT QUESTION DIRECTLY. INSTEAD, YOU’LL PIECE TOGETHER MULTIPLE QUERIES TO ANSWER YOUR QUESTION.

IN OTHER WORDS, THE FAILURE OF SEARCH ENGINES TO ANSWER QUESTIONS OVER THE LAST COUPLE OF DECADES LIMITS THE USEFULNESS OF THE KINDS OF QUESTIONS THAT YOU CAN GET FROM SEARCH ENGINES TODAY.

SO HOW DOES NATURAL QUESTIONS FIND THE ANSWERS TO THESE QUESTIONS? THEY FIND WIKIPEDIA GIVEN THE QUESTION AS A SEARCH STRING: AN INFORMATION RETRIEVAL PROBLEM! THEY THEN HAVE ANNOTATORS VALIDATE THAT THIS IS A GOOD SEARCH RESULT AND THEN HAVE THE ANNOTATORS HIGHLIGHT THE ANSWER STRING. NOW THIS ISN’T PERFECT.

HERE ARE SOME EXAMPLES WHERE I DISAGREE WITH SOME OF THE IMPLICIT ASSUMPTIONS THAT GET MADE IN THE NQ DATASET. FOR EXAMPLE, IT ASSUMES THAT WHEN YOU’RE ASKING ABOUT “MICHIGAN”, YOU’RE ASKING ABOUT THE “UNIVERSITY OF MICHIGAN FOOTBALL TEAM”, WHEN YOU’RE ASKING ABOUT “HOCKEY”, YOU’RE ASKING ABOUT “MEN’S ICE HOCKEY”, AND WHEN YOU’RE ASKING ABOUT “SUPREME COURT” YOU’RE ASKING ABOUT THE “INDIAN SUPREME COURT”.

AS WE PUT THESE DATASETS UP ON THE CHART, I WANT TO MAKE A DISTINCTION BETWEEN THE “TRADITIONAL” AND “OPEN” INTERPRETATIONS OF

NATURAL QUESTIONS. WHILE INITIAL NATURAL QUESTIONS WAS TREATED AS A SQuAD-LIKE DATASET, RECENT INTERPRETATIONS HAVE NOT ASSUMED THAT YOU KNOW WHICH WIKIPEDIA PAGE HAD THE ANSWER: YOU ALSO NEED TO DO THE IR COMPONENT AS WELL. THIS IS WHAT WE DID IN THE EFFICIENTQA COMPETITION (LINK IN THE DESCRIPTION): YOU CAN SEE HOW I MADE THE HUMAN VS. COMPUTER COMPARISON A LITTLE MORE FAIR.

OBVIOUSLY, THE OPEN INTERPRETATION MAKES IT MUCH HARDER TO FIND THE ANSWER.

MOVING ON ... YOU KNOW WHO'S BETTER THAN UNDERPAID CROWD WORKERS OR RANDOM PEOPLE ON THE INTERNET? PROFESSIONAL TRIVIA WRITERS. AND A NUMBER OF DATASETS USE QUESTIONS HARVESTED FROM THEM.

LET'S TALK ABOUT THREE OF THEM: SEARCHQA, TRIVIAQA, AND QUIZ BOWL.

SEARCHQA AND TRIVIAQA ACTUALLY LOOK A LOT LIKE NATURAL QUESTIONS, BUT INSTEAD OF THE QUERIES COMING FROM RANDOM PEOPLE ON THE INTERNET, THEY COME FROM PROFESSIONAL WRITERS. RESEARCHERS GO OUT AND FIND PAGES GIVEN THE QUERY, HIGHLIGHT THE ANSWERS. I'M NOT GOING TO DWELL TOO MUCH ON THESE DATASETS BECAUSE THE COPYRIGHT IS A LITTLE MURKY.

INSTEAD, LET'S NOW TURN TO MY HOBBY HORSE: QUIZ BOWL, WHICH IS PROBABLY BEST KNOWN AS AN NBC GAME SHOW IN THE US CALLED "COLLEGE BOWL" AND "UNIVERSITY CHALLENGE" IN THE UK. AND PLAYED BY THOUSANDS OF TRIVIA NERDS EVERY WEEKEND, PARTICULARLY AT THE HIGH SCHOOL AND COLLEGE LEVEL. AND UNLIKE SEARCHQA AND TRIVIAQA, THERE'S AN EXPLICIT COMMUNITY CONSENSUS THAT OLD QUIZ BOWL QUESTIONS ARE PUBLIC DOMAIN. SO AS A RESULT WE GET A LOT OF FREE DATA FROM EXPERTS. THESE ARE NOT CROWD WORKERS: THESE ARE PEOPLE WHO ARE PASSIONATE ABOUT HUMAN QUESTION ANSWERING, AND WE'RE USING THEIR EXPERTISE AND THEIR EXPERIENCE TO IMPROVE COMPUTER QUESTION ANSWERING.

THIS IS THE GOLD STANDARD IN HUMAN QUESTION ANSWERING (I.E., TRIVIA TOURNAMENTS), AND MY GOAL IS TO MAKE IT MORE POPULAR THAN IT CURRENTLY IS FOR COMPUTER QUESTION ANSWERING. WHAT'S DIFFERENT ABOUT QUIZ BOWL IS THAT THERE ISN'T JUST ONE QUESTION THAT YOU NEED TO ANSWER. INSTEAD, THERE ARE MULTIPLE CLUES THAT YOU COULD USE TO ANSWER THE QUESTION.

THIS IS HARD TO EXPLAIN, SO LET'S SEE AN EXAMPLE. AS I READ IT, THINK ABOUT WHEN YOU KNOW THE ANSWER TO THE QUESTION. WHAT CLUE GAVE IT AWAY?

WHILE PASSING THROUGH THIS STATE, THE 6TH MASSACHUSETTS SUFFERED THE FIRST UNION CASUALTIES OF THE CIVIL WAR WHILE SUPPRESSING RIOTS THAT BEGAN AT ITS PRESIDENT STREET STATION. THE LINCOLN ADMINISTRATION REJECTED A RULING BY ROGER TANEY THAT A CITIZEN OF THIS STATE COULD NOT BE HELD WITHOUT HABEAS CORPUS IN THE CASE EX PARTE MERRYMAN. AFTER CAPTURING A COPY OF ROBERT E. LEE'S SPECIAL ORDER 191, GEORGE McCLELLAN FOUGHT A BATTLE IN THIS STATE THAT ENDED THE FIRST CONFEDERATE INVASION OF THE NORTH AND LED TO THE ANNOUNCEMENT OF

THE EMANCIPATION PROCLAMATION. FOR 10 POINTS, THE BATTLE OF ANTI-ETAM WAS FOUGHT IN WHAT BORDER STATE WHERE FEDERAL TROOPS BLOCKED SECESSIONIST EFFORTS IN BALTIMORE?

THE CLUES ARE PRESENTED ONE BY ONE UNTIL SOMEONE CAN ANSWER THE QUESTION WITH THE GREAT STATE OF MARYLAND. BECAUSE THIS IS A COMPETITION, WHOEVER CAN ANSWER THE QUESTION FASTER (ON HARDER CLUES) USUALLY KNOWS MORE ABOUT THE TOPIC. BECAUSE THIS IS A TRIVIA GAME, THE GOAL IS TO GIVE MORE POINTS TO THE SMARTER PLAYER. QUIZ BOWL DOES THIS REALLY WELL! BETTER THAN SOME OTHER FUN TO WATCH KNOWLEDGE COMPETITIONS.

THE DOWNSIDE IS THAT SOMETIMES WHAT'S DIFFICULT FOR A HUMAN ISN'T DIFFICULT FOR A COMPUTER AND—JUST LIKE SQUAD—THERE SOMETIMES EASILY SOLVED QUESTIONS THAT ARE INTENDED TO BE DIFFICULT (FOR HUMANS) BUT AREN'T ACTUALLY DIFFICULT (FOR COMPUTERS).

FOR EXAMPLE, IN ONE OF THE EARLY COMPETITIONS WE HAD WITH HUMANS, A COMPUTER COULD EASILY RECOGNIZE A PAINTING FROM A DESCRIPTION.

BUT COMPUTERS STRUGGLE WITH SIMPLE LANGUAGE,

THE COMPUTER INCORRECTLY THINKS THAT WHEN THE QUESTION SAYS “THIS AUTHOR OPENS CRIME AND PUNISHMENT” THAT THE ANSWER IS DOSTOYEVSKI.

AS YOU CAN TELL FROM THESE CLIPS, ONE OF THE REASONS THAT I'M A BIG FAN OF THIS COMPETITION IS THAT IT LENDS ITSELF TO COMPARISONS. LET'S SAY YOU HAVE TWO COMPETITORS A AND B. WHAT IF WE TRY TO COMPARE THEM ON NORMAL QUESTIONS LIKE FROM NATURAL QUESTIONS? IF BOTH GET THE QUESTION RIGHT OR WRONG, YOU DON'T LEARN ANYTHING. YOU ONLY LEARN SOMETHING IF ONE GETS IT RIGHT AND THE OTHER GETS IT WRONG. HUMAN QUESTION WRITERS LEARNED REALLY EARLY ON THAT IT'S REALLY HARD TO GET THE LEVEL OF DIFFICULTY JUST RIGHT: EVEN WITH FANCY STATISTICS, YOU NEED HUNDREDS OF QUESTIONS TO RELIABLY RANK PEOPLE IN COMPETITIVE EXAMS.

SO THIS IS WHY THE HUMAN TRIVIA COMMUNITY TURNED TO QUIZ BOWL. INSTEAD OF JUST GETTING ONE PIECE OF INFORMATION PER QUESTION, YOU NOW POTENTIALLY GET A SIGNAL FROM EVERY WORD IN THE QUESTION. BECAUSE THE QUESTIONS ARE PACKED TO THE GILLS WITH CLUES THAT GO FROM SUPER DIFFICULT TO ALMOST PAINFULLY EASY, AFTER EACH CLUE YOU GET INFORMATION ABOUT HOW MUCH A AND B KNOWS. THUS, IT LENDS ITSELF TO HUMAN-COMPUTER COMPETITION OR COMPARING COMPUTER SYSTEMS.

ANOTHER REASON THAT I LIKE QUIZ BOWL IS THAT THERE ARE MULTIPLE WAYS TO APPROACH IT COMPUTATIONALLY. BECAUSE THE ANSWERS ARE USUALLY ENTITIES, YOU CAN TAKE A STRAIGHT INFORMATION RETRIEVAL APPROACH: WHAT'S BEING TALKED ABOUT IN THIS QUESTION? OR YOU CAN TAKE A MORE MACHINE READING APPROACH. BOTH WORK WELL!

OKAY, THAT'S ENOUGH ABOUT QUIZ BOWL. LET'S PUT IT ON THE BOARD: THE ANSWER STRUCTURE IS ABOUT THE SAME AS NQ, BUT THE QUESTION STRUCTURE IS VERY COMPLICATED.

IF YOU SUBSCRIBE TO THIS CHANNEL, YOU'LL HEAR ABOUT IT OFTEN ENOUGH

(AND SEE MATCHES BETWEEN HUMANS AND COMPUTERS).

BEFORE WE WRAP UP, I WANT TO TALK ABOUT A BUNCH OF OTHER DATASETS REALLY QUICKLY. I WON'T DISCUSS THEM IN DETAIL AND I WON'T TALK ABOUT ALL OF THE QA DATASETS OUT THERE ... THERE ARE JUST TOO MANY THESE DAYS! BUT IF ANY OF THESE SOUND INTERESTING TO YOU OR IF I'VE LEFT OUT A FAVORITE, PUT A LINK IN THE COMMENTS!

QUAC: THIS DATASET IS ABOUT CONVERSATIONAL QUESTION ANSWERING. AS A RESULT, MANY OF THE QUESTIONS DEPEND ON COREFERENCE OR OTHER AMBIGUITIES. WE CREATED A VERSION CALLED CANARD TO REMOVE THOSE AMBIGUITIES. AND OUR QBLINK DOES SOMETHING SIMILAR WITH QUIZ BOWL QUESTIONS.

NARRATIVEQA: THIS IS A GREAT DATASET, BUT IT'S NOT AS POPULAR AS IT SHOULD BE BECAUSE IT'S JUST TOO DARN HARD. YOU HAVE A SCRIPT OF AN INTERACTION AND YOU NEED TO ANSWER A QUESTION ABOUT WHAT HAPPENED IN THE STORY. IT DOESN'T LEND ITSELF TO THE MACHINE READING APPROACHES THAT WORK FOR OTHER DATASETS.

HOTPOTQA: REQUIRES YOU TO STRING TOGETHER MULTIPLE PIECES OF INFORMATION TO ANSWER THE QUESTION. THIS IS IMPLICITLY IN A LOT OF QUIZ BOWL QUESTIONS BUT NOT ANNOTATED EXCEPT IN A SMALL SUBSET OF QUESTIONS IN A GUHA ET AL 2015 DATASET. THE AUTHORS OF HOTPOTQA GOT AUTHORS TO EXPLICITLY MARK WHAT THE ANSWER IS TO THE INTERNAL QUESTIONS.

ALSO, THERE ARE WHOLE TYPES OF QUESTION ANSWERING THAT I HAVEN'T TALKED ABOUT AT ALL. WHAT ABOUT NON-ENGLISH DATASETS? THERE ARE A FEW (BUT NOT ENOUGH). WHAT ABOUT ANSWERING QUESTIONS ABOUT PICTURES? THERE ARE DATASETS FOR THOSE, BUT THIS VIDEO IS ALREADY LONG ENOUGH AS IT IS.

SO I SHOULD PROBABLY WRAP UP FOR REAL. QUESTION ANSWERING IS A REALLY EXCITING FIELD RIGHT NOW. THERE'S SO MUCH ENERGY BOTH FROM PEOPLE BUILDING DATASETS AND MODELS. AND DESPITE MY ENTHUSIASM FOR QUIZ BOWL, I ALSO RECOGNIZE THAT IT'S NOT THE PERFECT FORMAT. BUT BECAUSE ML DATASETS NEED TRAINING DATA TO IMPROVE, BEFORE WE CREATE MACHINES THAT CAN ANSWER ANY QUESTION, WE FIRST NEED TO CREATE DATASETS THAT WILL TEACH THE MACHINES.

8.8 *Surprise, this is a Trivia Tournament!*

"MY RESEARCH ISN'T A SILLY TRIVIA TOURNAMENT," YOU SAY. THAT MAY BE, BUT LET US FIRST TELL YOU A LITTLE ABOUT WHAT RUNNING A TOURNAMENT IS LIKE, AND PERHAPS YOU MIGHT SEE SIMILARITIES.

FIRST, THE QUESTIONS. EITHER YOU WRITE THEM YOURSELF OR YOU PAY SOMEONE TO WRITE QUESTIONS BY A PARTICULAR DATE (SOMETIMES PEOPLE ON THE INTERNET).

THEN, YOU ADVERTISE. YOU TALK ABOUT YOUR QUESTIONS: WHO IS WRITING THEM, WHAT SUBJECTS ARE COVERED, AND WHY PEOPLE SHOULD TRY TO ANSWER

THEM.

NEXT, YOU HAVE THE TOURNAMENT. YOU KEEP YOUR QUESTIONS SECURE UNTIL TEST TIME, COLLECT ANSWERS FROM ALL PARTICIPANTS, AND DECLARE A WINNER. AFTERWARD, PEOPLE USE THE QUESTIONS TO TRAIN FOR FUTURE TOURNAMENTS.

THESE HAVE NATURAL ANALOGS TO CROWD SOURCING QUESTIONS, WRITING THE PAPER, ADVERTISING, AND RUNNING A LEADERBOARD. TRIVIA NERDS CANNOT HELP YOU FORM HYPOTHESES OR WRITE YOUR PAPER, BUT THEY CAN TELL YOU HOW TO RUN A FUN, WELL-CALIBRATED, AND DISCRIMINATIVE TOURNAMENT.

SUCH TOURNAMENTS ARE DESIGNED TO EFFECTIVELY FIND A WINNER, WHICH MATCHES THE SCIENTIFIC GOAL OF KNOWING WHICH MODEL BEST ANSWERS QUESTIONS. OUR GOAL IS **not to encourage the QA community to adopt the quirks and gimmicks of trivia games**. INSTEAD, IT'S TO ENCOURAGE EXPERIMENTS AND DATASETS THAT **consistently and efficiently find the systems that best answer questions**.

Are we having fun?

MANY AUTHORS USE CROWDWORKERS TO ESTABLISH HUMAN ACCURACY (RAJPURKAR ET AL., 2016; CHOI ET AL., 2018). HOWEVER, THEY ARE NOT THE ONLY HUMANS WHO SHOULD ANSWER A DATASET'S QUESTIONS. SO SHOULD THE DATASET'S CREATORS.

IN THE TRIVIA WORLD, THIS IS CALLED A **play test**: GET IN THE SHOES OF SOMEONE ANSWERING THE QUESTIONS. IF YOU FIND THEM BORING, REPETITIVE, OR UNINTERESTING, SO WILL CROWDWORKERS. IF YOU CAN FIND SHORTCUTS TO ANSWER QUESTIONS (RONDEAU AND HAZEN, 2018; KAUSHIK AND LIPTON, 2018), SO WILL A COMPUTER.

CONCRETELY, WEISSENBORN ET AL. (2017) CATALOG ARTIFACTS IN SQUAD (RAJPURKAR ET AL., 2018), THE MOST POPULAR QA LEADERBOARD. IF YOU SEE A LIST LIKE “ALONG WITH CANADA AND THE UNITED KINGDOM, WHAT COUNTRY...”, YOU CAN IGNORE THE REST OF THE QUESTION AND JUST TYPE CTRL+F (YUAN ET AL., 2019; RUSSELL, 2020) TO FIND THE THIRD COUNTRY—AUSTRALIA IN THIS CASE—THAT APPEARS WITH “CANADA AND THE UK”. OTHER TIMES, A SQUAD PLAYTEST WOULD REVEAL FRUSTRATING QUESTIONS THAT ARE I) ANSWERABLE GIVEN THE INFORMATION BUT NOT WITH A DIRECT SPAN,¹ II) ANSWERABLE ONLY GIVEN FACTS BEYOND THE GIVEN PARAGRAPH,² III) UNINTENTIONALLY EMBEDDED IN A DISCOURSE, RESULTING IN ARBITRARY CORRECT ANSWERS,³ IV) OR NON-QUESTIONS.

SEARCHQA (DUNN ET AL., 2017), DERIVED FROM *JEOPARDY!*, ASKS “AN ARTICLE THAT HE WROTE ABOUT HIS RIVERBOAT DAYS WAS EVENTUALLY EXPANDED INTO *LIFE ON THE MISSISSIPPI*.” THE APPRENTICE AND NEWSPAPER WRITER WHO WROTE THE ARTICLE IS NAMED SAMUEL LANGHORNE CLEMENS; HOWEVER, THE REFERENCE ANSWER IS HIS LATER PEN NAME, MARK TWAIN. MOST QA

¹ A source paragraph says “In [Commonwealth countries]...the term is generally restricted to...Private education in North America covers the whole gamut...”; thus, “What is the term private school restricted to in the US?” has the information needed but not as a span.

² A source paragraph says “Sculptors [in the collection include] Nicholas Stone, Caius Gabriel Cibber, [...], Thomas Brock, Alfred Gilbert, [...] and Eric Gill”, i.e., a list of names; thus, the question “Which British sculptor whose work includes the Queen Victoria memorial in front of Buckingham Palace is included in the V&A collection?” should be unanswerable in SQUAD.

³ A question “Who else did Luther use violent rhetoric towards?” has the gold answer “writings condemning the Jews and in diatribes against Turks”.

EVALUATION METRICS WOULD COUNT SAMUEL CLEMENS AS INCORRECT. IN A REAL GAME OF *JEOPARDY!*, THIS WOULD NOT BE AN ISSUE (SECTION 5.2).

OF COURSE, FUN IS RELATIVE, AND ANY DATASET IS BOUND TO CONTAIN ERRORS. HOWEVER, PLAYTESTING IS AN EASY WAY TO FIND SYSTEMATIC PROBLEMS: UNFAIR, UNFUN PLAYTESTS MAKE FOR INEFFECTIVE LEADERBOARDS. EATING YOUR OWN DOG FOOD CAN HELP DIAGNOSE ARTIFACTS, SCORING ISSUES, OR OTHER SHORTCOMINGS EARLY IN THE PROCESS.

THE DEEPER ISSUES WHEN CREATING A QA TASK ARE: I) HAVE YOU DESIGNED A TASK THAT IS INTERNALLY CONSISTENT, II) SUPPORTED BY A SCORING METRIC THAT MATCHES YOUR GOALS, III) USING GOLD ANNOTATIONS THAT REWARD THOSE WHO DO THE TASK WELL? IMAGINE SOMEONE WHO LOVES ANSWERING THE QUESTIONS YOUR TASK POSES: WOULD THEY HAVE FUN ON YOUR TASK? THIS IS THE FOUNDATION OF GAMIFICATION ([VON AHN, 2006](#)), WHICH CAN CREATE QUALITY DATA FROM USERS MOTIVATED BY FUN RATHER THAN PAY. EVEN IF YOU PAY CROWDWORKERS, UNFUN QUESTIONS MAY UNDERMINE YOUR DATASET GOALS.

Am I measuring what I care about?

ANSWERING QUESTIONS REQUIRES MULTIPLE SKILLS: IDENTIFYING ANSWER MENTIONS ([HERMANN ET AL., 2015](#)), NAMING THE ANSWER ([YIH ET AL., 2015](#)), ABSTAINING WHEN NECESSARY ([RAJPURKAR ET AL., 2018](#)), AND JUSTIFYING AN ANSWER ([THORNE ET AL., 2018](#)). IN QA, THE EMPHASIS ON SOTA AND LEADERBOARDS HAS FOCUSED ATTENTION ON SINGLE AUTOMATICALLY COMPUTABLE METRICS—SYSTEMS TEND TO BE COMPARED BY THEIR ‘SQUAD SCORE’ OR THEIR ‘NQ SCORE’, AS IF THIS WERE ALL THERE IS TO SAY ABOUT THEIR RELATIVE CAPABILITIES. LIKE QA LEADERBOARDS, TRIVIA TOURNAMENTS NEED TO DECIDE ON A SINGLE WINNER, BUT THEY EXPLICITLY RECOGNIZE THAT THERE ARE MORE INTERESTING COMPARISONS.

A TOURNAMENT MAY RECOGNIZE DIFFERENT BACKGROUND/RESOURCES—HIGH SCHOOL, SMALL SCHOOL, UNDERGRADUATES ([HENTZEL, 2018](#)). SIMILARLY, MORE PRACTICAL LEADERBOARDS WOULD REFLECT TRAINING TIME OR RESOURCE REQUIREMENTS (SEE [DODGE ET AL., 2019](#)) INCLUDING ‘CONSTRAINED’ OR ‘UNCONSTRAINED’ TRAINING ([BOJAR ET AL., 2014](#)). TOURNAMENTS ALSO GIVE SPECIFIC AWARDS (E.G., HIGHEST SCORE WITHOUT INCORRECT ANSWERS). AGAIN, THERE ARE OBVIOUS LEADERBOARD ANALOGS THAT WOULD GO BEYOND A SINGLE NUMBER. IN SQUAD 2.0 ([RAJPURKAR ET AL., 2018](#)), ABSTAINING CONTRIBUTES THE SAME TO THE OVERALL F_1 AS A FULLY CORRECT ANSWER, OBSCURING WHETHER A SYSTEM IS MORE PRECISE OR AN EFFECTIVE ABSTAINER. IF THE TASK RECOGNIZES BOTH ABILITIES AS IMPORTANT, REPORTING A SINGLE SCORE RISKS IMPLICITLY PRIORITIZING ONE BALANCE OF THE TWO.

Do my questions separate the best?

ASSUME THAT YOU HAVE PICKED A METRIC (OR A SET OF METRICS) THAT CAPTURES WHAT YOU CARE ABOUT. A LEADERBOARD BASED ON THIS METRIC CAN RACK UP CITATIONS AS PEOPLE CHASE THE TOP SPOT. BUT YOUR LEADERBOARD IS ONLY USEFUL IF IT IS **discriminative**: THE BEST SYSTEM RELIABLY WINS.

THERE ARE MANY WAYS QUESTIONS MIGHT NOT BE DISCRIMINATIVE. IF EVERY SYSTEM GETS A QUESTION RIGHT (E.G., ABSTAIN ON NON-QUESTIONS LIKE “ASDF” OR CORRECTLY ANSWER “WHAT IS THE CAPITAL OF POLAND?”), THE DATASET DOES NOT SEPARATE PARTICIPANTS. SIMILARLY, IF EVERY SYSTEM FLUBS “WHAT IS THE OLDEST NORTH-FACING KOSHER RESTAURANT”, IT IS NOT DISCRIMINATIVE. SUGAWARA ET AL. (2018) CALL THESE QUESTIONS “EASY” AND “HARD”; WE INSTEAD ARGUE FOR A THREE-WAY DISTINCTION.

IN BETWEEN EASY QUESTIONS (SYSTEM ANSWERS CORRECTLY WITH PROBABILITY 1.0) AND HARD (PROBABILITY 0.0), QUESTIONS WITH PROBABILITIES NEARER TO 0.5 ARE MORE INTERESTING. TAKING A CUE FROM VYGOTSKY’S PROXIMAL DEVELOPMENT THEORY OF HUMAN LEARNING (CHAIKLIN, 2003), THESE DISCRIMINATIVE QUESTIONS—RATHER THAN THE EASY OR THE HARD ONES—SHOULD MOST IMPROVE QA SYSTEMS. THESE GOLDILOCKS⁴ QUESTIONS (NOT RANDOM NOISE) DECIDE WHO TOPS THE LEADERBOARD. UNFORTUNATELY, EXISTING DATASETS HAVE MANY EASY QUESTIONS. SUGAWARA ET AL. (2020) FIND THAT ABLATIONS LIKE SHUFFLING WORD ORDER (FENG ET AL., 2018), SHUFFLING SENTENCES, OR ONLY OFFERING THE MOST SIMILAR SENTENCE DO NOT IMPAIR SYSTEMS. NEWER DATASETS SUCH AS DROP (DUA ET AL., 2019) AND HELLA-SWAG (ZELLERS ET AL., 2019) ARE HARDER FOR TODAY’S SYSTEMS; BECAUSE GOLDILOCKS IS A MOVING TARGET, WE PROPOSE ANNUAL EVALUATIONS IN SECTION **?reference?** ??.

Why so few Goldilocks questions?

THIS IS A COMMON PROBLEM IN TRIVIA TOURNAMENTS, PARTICULARLY PUB QUIZZES (DIAMOND, 2009), WHERE CHALLENGING QUESTIONS CAN SCARE OFF PATRONS. MANY QUIZ MASTERS PREFER POPULARITY WITH PLAYERS AND THUS WRITE EASIER QUESTIONS.

SOMETIMES THERE ARE FEWER GOLDILOCKS QUESTIONS NOT BY CHOICE, BUT BY CHANCE: A DATASET BECOMES LESS DISCRIMINATIVE THROUGH ANNOTATION ERROR. ALL DATASETS HAVE SOME ANNOTATION ERROR; IF THIS ANNOTATION ERROR IS CONCENTRATED ON THE GOLDILOCKS QUESTIONS, THE DATASET WILL BE LESS USEFUL. AS WE WRITE THIS IN 2020, HUMANS AND COMPUTERS SOMETIMES STRUGGLE ON THE SAME QUESTIONS.

FIGURE 8.1 SHOWS TWO DATASETS OF THE SAME SIZE WITH THE SAME ANNOTATION ERROR. HOWEVER, THEY HAVE DIFFERENT DIFFICULTY DISTRIBUTIONS AND CORRELATION OF ANNOTATION ERROR AND DIFFICULTY. THE DATASET THAT HAS MORE DISCRIMINATIVE QUESTIONS AND CONSISTENT ANNOTATOR

⁴ In a British folktale first recorded by Robert Southey, the character Goldilocks finds three beds: one too hard, one not hard enough, and one “just right”.





Figure 8.1: Two datasets with 0.16 annotation error: the top, however, better discriminates QA ability. In the good dataset (top), most questions are challenging but not impossible. In the bad dataset (bottom), there are more trivial or impossible questions *and* annotation error is concentrated on the challenging, discriminative questions. Thus, a smaller fraction of questions decide who sits atop the leaderboard, requiring a larger test set.



ERROR HAS FEWER QUESTIONS THAT DO NOT DISCRIMINATE THE WINNER OF THE LEADERBOARD. WE CALL THIS THE EFFECTIVE DATASET PROPORTION ρ (HIGHER IS BETTER). FIGURE 8.2 SHOWS THE TEST SET SIZE REQUIRED TO RELIABLY DISCRIMINATE SYSTEMS FOR DIFFERENT ρ , BASED ON A SIMULATION (APPENDIX ?reference? ??).

AT THIS POINT, YOU MAY DESPAIR ABOUT HOW BIG A DATASET YOU NEED.⁵ THE SAME TERROR BESETS TRIVIA TOURNAMENT ORGANIZERS. INSTEAD OF WRITING MORE QUESTIONS, THEY USE PYRAMIDALITY (SECTION 5.3) TO MAKE EVERY QUESTION COUNT.

Figure 8.2: How much test data do you need to discriminate two systems with 95% confidence? This depends on both the difference in accuracy between the systems (x axis) and the average accuracy of the systems (closer to 50% is harder). Test set creators do not have much control over those. They do have control, however, over how many questions are discriminative. If all questions are discriminative (right), you only need 2500 questions, but if three-quarters of your questions are *too hard* (left), you need 15,000 (Tsochers, 2019).

THE TOOLS TO ANSWER QUESTIONS

HAVING DISCUSSED HOW WE CREATE DATASETS THAT CAN TEACH COMPUTERS HOW TO ANSWER QUESTIONS, WE NOW EXPLORE HOW THOSE DATASETS CAN TRAIN MODERN COMPUTER METHODS.

TOPICS IN CHAPTER:

1. **MODERN KNOWLEDGE.** THIS SECTION REVIEWS HOW WE CAN TURN NATURAL LANGUAGE QUERIES INTO ACTIONABLE QUERIES IN DATABASES.
2. **JOHNNY 5 CAN'T READ.** HOWEVER, PUTTING INFORMATION IN A DATABASE IS DIFFICULT AND TIME-CONSUMING... NOT ALL INFORMATION IS IN A DATABASE (INDEED, SOME INFORMATION DEFIES STRICT DATABASE SCHEMAS). Thus, we need to teach computers how to read. This chapter reviews the process of "MACHINE READING", where computers find information in a large text corpus and then extracts the answer from it.
3. **A VECTOR IN A HAYSTACK.** AS DEEP LEARNING BECAME PRACTICAL, THE FIELD HAS MOVED FROM REPRESENTING TEXT WITH DISCRETE WORDS TO CONTINUOUS VECTORS. THIS IS ALSO THE CASE FOR QUESTION ANSWERING. THIS CHAPTER REVIEWS HOW THESE REPRESENTATIONS CAN HELP US FIND RELEVANT ANSWERS TO QUESTIONS.
4. **WE COULDN'T FIND AN ANSWER, SO WE HAD TO GENERATE ONE.** THE DEEP LEARNING REVOLUTION NOT JUST HELPS US FIND ANSWERS BUT TO GENERATE THEM. THIS CHAPTER TALKS ABOUT THE PROMISE AND PERIL OF THESE APPROACHES: HOW WE CAN SYNTHESIZE MUCH RICHER INFORMATION, MAKE STUFF UP, ENCOURAGE THESE AGENTS TO ALIGN TO OUR WISHES, AND HOW IT'S HARD TO TELL IF AN ANSWER IS ANY GOOD.
5. **COOPERATIVE QA.** HUMANS AND COMPUTERS WORKING TOGETHER CAN ANSWER QUESTIONS BETTER

IN PREVIOUS VIDEOS, WE TALKED ABOUT THE GENERAL FRAMEWORK FOR HOW COMPUTERS ANSWER A QUESTION: YOU HAVE THE QUESTION, THAT GOES INTO AN INFORMATION RETRIEVAL SYSTEM, YOU GET SOME EVIDENCE, AND THEN YOU RUN A READER OVER THAT TO FIND THE ANSWER.

TODAY WE'RE GOING TO TALK ABOUT HOW WE CAN DO THAT A LITTLE BIT BETTER WITH THE FANCY TOOLS LIKE ENCODERS IN THE PREVIOUS LECTURES ... AND HOW THESE NEW TECHNIQUES FOR DOING THIS UNLOCKS EXCITING ADVANCES IN QUESTION ANSWERING. STUFF CALLED "DEEP RETRIEVAL": INSTEAD

OF USING A TRADITIONAL INFORMATION RETRIEVAL SYSTEM, WE'LL USE DEEP LEARNING TO TURN BOTH OUR QUERIES AND OUR EVIDENCE INTO VECTORS.

NOW, YOU REMEMBER HOW A TRADITIONAL IR SYSTEM WORKS: YOU GET A SPARSE VECTOR WITH WEIGHTS FOR THE INDIVIDUAL WORDS. HOW THOSE WEIGHTS ARE SET ISN'T IMPORTANT. IT COULD BE TF-IDF, BM25, OR JUST RAW WORD COUNTS. LET'S FOCUS ON THE DIMENSIONALITY AND WHAT KINDS OF THINGS CAN BE MATCHED.

FIRST, THE DIMENSIONALITY IS IN THE TENS OR EVEN HUNDREDS OF THOUSANDS. THERE ARE MANY WORDS IN ENGLISH, PARTICULARLY WHEN YOU WANT TO HAVE COVERAGE OF ALL OF THE RARE NAMED ENTITIES LIKE "THE TREATY OF WESTPHALIA" OR "ULURU" OR "MURASAKI". EVEN THOUGH THESE VECTORS ARE SPARSE, YOU PAY A PRICE FOR KEEPING TRACK OF ALL OF THESE THOUSANDS OF DIMENSIONS. IT'S OFTEN NOT WORTH IT TO DO EXACT SEARCH OVER THOUSANDS OF VOCABULARY WORDS.

SO THERE ARE GOOD ALGORITHMS FOR SEARCHING THROUGH THIS, MADE BY MARYLAND'S OWN DAVID MOUNT, BUT YOU CAN'T DO IT EXACTLY: ONCE THE DIMENSIONALITY GETS HIGH ENOUGH, YOU'RE DOING AN APPROXIMATE SEARCH.

IF YOU CAN'T DO AN EXACT SEARCH IN HIGH DIMENSIONS, THEN WHY KEEP THEM AROUND? A SEMINAL PAPER CALLED "MIN-WISE INDEPENDENT PERMUTATIONS" SAID THAT YOU SHOULDN'T DO THIS FOR COMPUTATIONS WITH THESE BIG SPARSE VECTORS: CREATE NEW VECTORS AND DO YOUR COMPUTATIONS ON THAT INSTEAD. IF YOU USE HASH FUNCTIONS BASED ON SOME DISTANCE FUNCTION (E.G., COSINE SIMILARITY) TO CREATE THESE REPRESENTATIONS, THIS IS CALLED "LOCALITY-SENSITIVE HASHING".

THIS IDEA WAS CREATED FOR THE SEARCH ENGINE ALTAVISTA. PLEASE LEAVE A COMMENT BELOW IF YOU REMEMBER ALTAVISTA AND TELL THESE DANG KIDS TO GET OFF OUR LAWN. I REMEMBER THAT LOCALITY-SENSITIVE HASHING WAS A REALLY COOL IDEA THAT WAS JUST BECOMING POPULAR WHEN I WAS A GRAD STUDENT (MOSES WAS A PROFESSOR AT PRINCETON WHEN I WAS THERE): YOU COULD DO FASTER COMPUTATIONS WITH LESS MEMORY WITH THESE REPRESENTATIONS.

9.1 *A Vector in a Haystack*

NOW WHY AM I TALKING ABOUT THESE ANCIENT ALGORITHMS FROM THE TURN OF THE CENTURY IN A LECTURE ABOUT BERT FOR DEEP RETRIEVAL? THE ALGORITHMS THAT WERE DEVELOPED TO DO COMPARISONS AND SEARCHES IN THIS LOWER DIMENSIONAL SPACE CAN BE REPURPOSED FOR THE REPRESENTATIONS CREATED BY DEEP LEARNING ALGORITHMS, AND THE IDEA IS THE SAME: CREATE REPRESENTATIONS SPECIFIC TO YOUR TASK.

BUT WHAT I REALLY WANT TO TALK ABOUT IS HOW THESE REPRESENTATIONS CAN BE NOT JUST FASTER BUT CAN ALSO BE BETTER! CAN WE CREATE REPRESENTATIONS TAILOR-MADE FOR A TASK LIKE QUESTION ANSWERING?

THE ANSWER IS YES, WE CAN DO THIS USING THE DEEP LEARNING APPROACHES WE'VE TALKED ABOUT BEFORE! BUT BEFORE WE TALK ABOUT HOW, LET'S TALK

ABOUT WHY YOU MIGHT WANT TO DO THIS. NOW, SOMETIMES TF-IDF IS GOING TO BE JUST FINE. IF YOU HAVE GOOD TF-IDF CLUES IN A QUESTION LIKE “WHAT DO YOU NEED THE HALL-HEROULT PROCESS FOR?”, THEN IT’S FAIRLY EASY TO MAP THAT UP TO THE PROCESS FOR SMELTING ALUMINUM. BUT SOMETIMES THERE ARE SEQUENCES OF WORDS THAT AREN’T HIGH TF-IDF LIKE “WHO SAID TO BE OR NOT TO BE”. NOW, YOU COULD ADD N-GRAM DIMENSIONS BUT THAT WOULD EXPLODE THE DIMENSIONALITY EVEN MORE.

EVEN WORSE, SOMETIMES PEOPLE ASKING QUESTIONS, PARTICULARLY IN THE MANCHESTER PARADIGM (THAT WE TALKED ABOUT IN A PREVIOUS VIDEO), CAN BE TRICKY! IF THE QUESTION IS ABOUT LUEBECK, A CITY THAT WAS PART OF THE HANSEATIC LEAGUE, THE QUESTION WRITER MIGHT BE TRICKY AND SAY SOMETHING LIKE “A CITY ON THE TRAVE THAT WAS A MEMBER OF A CONFEDERATION OF BALTIC CITIES” INSTEAD OF MENTIONING THE HANSEATIC LEAGUE BY NAME.

SO HOW CAN YOU LEARN HOW TO DO THIS? THIS IS WHERE FRAMEWORKS LIKE PYTORCH ARE REALLY USEFUL. YOU CAN CREATE A GIANT COMPUTATION GRAPH THAT NOT ONLY LEARNS A REPRESENTATION OF A QUESTION BUT ALSO OF THE EVIDENCE PASSAGES WHERE YOU CAN EASILY FIND THE ANSWER. SO NOW YOU LEARN TO PUSH “WHO SAID TO BE OR NOT TO BE” CLOSE TO EVIDENCE PASSAGES THAT CONTAIN “HAMLET” AND “CITY ON THE TRAVE THAT WAS A MEMBER OF A CONFEDERATION OF BALTIC CITIES” CLOSE TO EVIDENCE PASSAGES THAT CONTAIN LUEBECK. AND YOU CAN USE YOUR FAVORITE NEURAL NETWORK TO ENCODE BOTH SIDES. TODAY YOU’D USE BERT TO DO THIS.

BACK IN 2015, WE USED DEEP AVERAGING NETWORKS TO DO THIS ... HERE’S AN EXAMPLE OF WHAT THE SPACE LOOKS LIKE FOR THE REPRESENTATIONS OF QUESTIONS ABOUT US PRESIDENTS. WHEN YOU PROJECT DOWN TO TWO DIMENSIONS, YOU CAN SEE A CLOCKWISE PROGRESSION THROUGH TIME FROM THE FOUNDING FATHERS TO THE 19TH CENTURY TO THE 20TH CENTURY.

AT THE END OF THE 2010S, THERE WERE THREE PAPERS THAT CAME OUT THAT PROPOSED WAYS OF LEARNING ENCODINGS FOR EVIDENCE AND QUESTIONS TO FIND THE CORRECT VECTORS. THESE ARE ORQA (“OPEN-RETRIEVAL QUESTION-ANSWERING”), REALM (“RETRIEVAL-AUGMENTED LANGUAGE MODEL PRE-TRAINING”), AND DPR (“DENSE PASSAGE RETRIEVAL”).

SO WHAT DO THESE PAPERS SAY? ORQA SHOWED THAT IT WAS POSSIBLE TO DO NEURAL RETRIEVAL BETTER THAN AN OFF THE SHELF IR SYSTEM.

REALM IMPROVED THE PRETRAINING, DOING MASKING OF SPANS THAT ARE LIKELY TO BE THE ANSWER. INSTEAD OF BERT, WHICH MASKS OUT A RANDOM 15

THIS COULD BE A QUIZ BOWL QUESTION WITH THE ANSWER OR IT COULD COME FROM THAT WIKIPEDIA PAGE, AGAIN, WE’RE JUST VIEWING IT AS RAW TEXT FOR BERT PRETRAINING. IS THERE ANYTHING WE CAN DO TO IMPROVE BERT VIA PRETRAINING? REALM ARGUES THAT IF YOU MASK THESE TOKENS, RATHER THAN A RANDOM 15

LET’S TAKE A LOOK AT THIS FIGURE FROM THE REALM PAPER THAT DISCUSSES

HOW YOU SHOULD GET YOUR NEGATIVE SAMPLES.

THEIR SUGGESTION IS TO USE ALL OF THE NEGATIVE EXAMPLES FROM A SINGLE BATCH. SO LET'S SAY YOUR BATCH SIZE IS FOUR: THERE ARE FOUR QUESTIONS IN YOUR BATCH. OBVIOUSLY THE POSITIVE EXAMPLE FOR QUESTION 2 TRAINING IS THE EVIDENCE PASSAGE THAT CONTAINS THE ANSWER TO QUESTION 2. BUT WHAT SHOULD YOU USE AS NEGATIVE EVIDENCE? THE DPR PAPER ARGUES YOU SHOULD USE IN-BATCH NEGATIVES, AND TAKE ALL OF THE NEGATIVE EXAMPLES FOR EVERYTHING ELSE IN THE BATCH.

NOW HOW DO YOU GET THESE NEGATIVE EXAMPLES? AGAIN, THE DPR PAPER ARGUES THAT THIS SHOULD BE TRADITIONAL ANSWERS FROM AN IR SYSTEM THAT HAVE HIGH TF-IDF, BM25, ETC. SCORES BUT DON'T HAVE THE RIGHT ANSWER.

IT'S NOT A HUGE DIFFERENCE, BUT IT DOES SEEM TO HELP.

LET'S WRAP UP. THESE MODELS ALLOW US TO NOT JUST MATCH WORDS FROM QUESTIONS AND ANSWERS BUT TO LEARN ARBITRARY MAPPINGS BETWEEN QUESTIONS AND ANSWERS. I THINK THAT WE'VE ONLY BEGUN TO SCRATCH THE POSSIBILITIES FOR THIS KIND OF SYSTEM, AND RECENT RESEARCH IS LOOKING INTO THESE OPPORTUNITIES.

IN THE PREVIOUS VIDEO, WE TALKED ABOUT HOW BERT CAN FIND ANSWER SPANS IN QUESTIONS ... HOW DOES THIS COMBINE WITH DEEP RETRIEVAL? IN SOME WAYS, IT GOES AWAY! INSTEAD OF JUST INDEXING ENTIRE DOCUMENTS, YOU CAN JUST ENCODE THE THINGS THAT BERT CAN ENCODE WELL: WIKIPEDIA PARAGRAPHS FOR INSTANCE. THEN YOU RETRIEVE THE PARAGRAPHS (RATHER THAN THE PAGES) THAT ARE LIKELY TO HAVE THE ANSWER.

SO IF YOU'RE GOING TO USE BERT ANYWAY TO FIND THE SPANS (BECAUSE IT'S BETTER THAN THE SEQUENCE-BASED APPROACHES THAT CAME BEFORE IT), WHY NOT GO ALL THE WAY AND ALSO USE IT FOR THE RETRIEVAL? YOU'LL FIX TWO PROBLEMS AT ONCE!

THERE ARE OTHER WAYS THAT DEEP RETRIEVAL IS BETTER.

DENSE PASSAGE RETRIEVAL LETS YOU FIND THE ANSWER TO A QUESTION LIKE "HOW WAS THE MONKEY KING BORN". BUT SINCE WE'RE DOING THE LOOKUP IN AN ARBITRARY VECTOR SPACE, WE COULD TRAIN OUR QUESTION ENCODER TO TAKE INPUT IN ANY LANGUAGE: ALL THAT MATTERS IS THAT THE RESULTING VECTOR SPACE LINES UP. BUT THEN WHY DOES IT EVEN HAVE TO BE TEXT? WE CAN ENCODE ANY OBJECT AS A VECTOR, SO WHY NOT ANSWER QUESTIONS WITH AN IMAGE?

SO, TO RECAP, WE GOT TO DENSE PASSAGE RETRIEVAL BECAUSE WE DIDN'T WANT TO PUT UP WITH THOUSANDS OF TYPES DEFINING OUR VOCABULARY, BUT WE ENDED UP WITH SOMETHING THAT CAN DO IMAGE RETRIEVAL INSTEAD. SO PERHAPS A PICTURE ENCODING IS WORTH A THOUSAND WORDS.

10

SIRI TAKES THE SAT

HOW DO WE KNOW HOW SMART A MACHINE IS? LIKE A HUMAN, WE TYPICALLY GIVE IT A TEST; THE DIFFERENCE IS THAT TESTS FOR COMPUTERS ARE CALLED ‘LEADERBOARDS’. THIS CHAPTER TALKS ABOUT THE PROS AND CONS OF LEADERBOARDS AND HOW SOME OF THE METHODS USED TO ANALYZE HUMAN STANDARDIZED TESTS CAN HELP US UNDERSTAND THE STRENGTHS AND WEAKNESSES OF NOT JUST COMPUTER QUESTION ANSWERING SPECIFICALLY BUT AI GENERALLY. THIS ALSO MARKS THE REAPPEARANCE OF ITEM RESPONSE THEORY, WHICH CAN HELP MAKE DECIDING THE SMARTEST COMPUTER MORE EFFICIENT. PREVIOUSLY PUBLISHED AS [RODRIGUEZ ET AL.](#) (2021) AND ?.

TOPICS IN CHAPTER:

1. **WHEN 21 QUESTIONS ARE WORTHLESS.** WE COMPARE HOW A TRIVIA CHEATING SCANDAL EXTENDS TO AI LEADERBOARDS.
2. . IF EXISTING DATASETS AND GAME SHOW APPEARANCES AREN’T ENOUGH TO TELL WHETHER HUMANS OR COMPUTERS ARE BETTER AT ANSWERING QUESTIONS, WHAT CAN WE DO? WHILE THERE IS AN ARGUMENT FOR FOCUSING ON NATURAL DATA, MODERN LANGUAGE MODELS ARE CHANGING NOT JUST WHAT IS POSSIBLE COMPUTATIONALLY BUT CHANGING THE LANGUAGE ITSELF. THUS, WE NEED TO SELECT EXAMPLES SPECIFICALLY TO CHALLENGE COMPUTERS. THESE EXAMPLES ARE CALLED *ADVERSARIAL EXAMPLES*, AND THIS CHAPTER PRESENTS HOW TO GATHER THEM AND HOW THEY CAN REVEAL THE STRENGTHS AND WEAKNESSES OF COMPUTER QUESTION ANSWERING.

IN A PREVIOUS VIDEO, I TALKED ABOUT HOW QUESTIONS CREATED MODERN CIVILIZATION: THE USE OF EXAMS TO SELECT CIVIL SERVANTS. COMPARED TO THE ALTERNATIVE THAT PRECEDED THEM, I THINK THESE TESTS WERE AN UNMITIGATED GOOD.

TODAY, WE’RE TALKING ABOUT SOMETHING WHERE THE VERDICT OF HISTORY IS LESS CLEAR-CUT: STANDARDIZED TESTS AND THE PSYCHOMETRIC ANALYSIS THAT IS APPLIED ON TOP OF THEM. WHEN STANDARDIZED TESTS WERE INTRODUCED, THE IDEAL WAS THAT THESE TESTS COULD MAKE ADMISSIONS MORE SCIENTIFIC AND LESS RANDOM OR CONNECTED TO FAMILY TIES.

AT THE END OF THE 19TH CENTURY, A COUPLE OF UNIVERSITIES BANDED TOGETHER TO CREATE THE “COLLEGE BOARD” TO ADMINISTER TESTS TO SEE

WHO COULD GET INTO UNIVERSITIES LIKE HARVARD, YALE, AND PRINCETON. DO WELL ON THE TEST, AND YOU GET IN!

WE'LL TALK ABOUT THE MATH IN THE NEXT VIDEO, BUT FOR NOW I WANT TO FOCUS ON THE PROMISE AND THE BIG PICTURE OF STANDARDIZED TESTING AND HOW IT RELATES TO THE LEADERBOARDS THAT HAVE TAKEN OVER MACHINE LEARNING.

LET'S TALK ABOUT WHAT THE COLLEGE BOARD AND LATER EDUCATIONAL TESTING SERVICE PROMISES: ANYBODY CAN SIGN UP FOR A TEST WHENEVER THEY WANT, AND IT GIVES A NUMBER THAT POPS OUT THAT ACCURATELY DESCRIBES HOW SMART THEY ARE. LET'S CALL THAT NUMBER THETA TO STAND FOR THE SKILL OF A STUDENT.

SO CAN YOU JUST REPORT ACCURACY? WOULD THAT BE A GOOD WAY OF GETTING THETA? NO! THE TEST IN JANUARY MIGHT NOT BE THE SAME AS THE TEST IN JULY. SO THE MATHEMATICAL CHALLENGE IS HOW TO CREATE THAT ESTIMATE THETA SO THAT IT'S CONSISTENT AGAINST DIFFERENT TESTS AND TO SELECT QUESTIONS FROM THE POOL OF POSSIBLE QUESTIONS SO THAT YOU CAN RELIABLY GENERATE THAT NUMBER THETA. AND THAT'S WHERE THE "STANDARDIZED" IN STANDARDIZED TESTS COMES FROM.

THE PROCESS HAS CHANGED OVER THE YEARS...AT FIRST YOU HAD TO GIVE EVERYBODY THE SAME TEST AT EACH OFFERING OF THE TEST, BUT NOW THEY DO ADAPTIVE TESTING TO DO A BETTER JOB OF ESTIMATING THETA FOR A SPECIFIC TEST-TAKER. BUT THE TWO BIG PROBLEMS REMAIN THE SAME: MAKING THETA STANDARDIZED AND DESCRIBING IN GENERAL WHAT KIND OF QUESTIONS THEY SHOULD BE ASKING.

SO WHY AM I, A MACHINE LEARNING RESEARCHER, TALKING ABOUT THIS? IS IT BECAUSE I WANT TO TALK ABOUT HOW HARVARD ADDED ADMISSIONS ESSAYS TO STANDARDIZED TESTS BECAUSE THEY WERE CONCERNED THAT TOO MANY JEWS WERE GETTING IN? OR BECAUSE THE FEDERAL GOVERNMENT DEEMED UCLA'S ADMISSIONS POLICIES RACIST BECAUSE NOT ENOUGH QUALIFIED ASIANS WERE GETTING IN? OR HOW THERE'S NOW A BACKLASH AGAINST STANDARDIZED TESTS BECAUSE AFRICAN AMERICANS SYSTEMATICALLY HAVE LOWER SCORES THAN OTHER GROUPS? NOPE! AS ANIMATED AS I GET ABOUT THESE TOPICS, THESE ARE DECIDEDLY OUT OF SCOPE OF THIS VIDEO.

BUT I MENTION THEM BECAUSE THEY'RE DEFINITELY REAL PROBLEMS, AND CAUSES A LOT OF PEOPLE TO HATE STANDARDIZED TESTS. BUT IT'S NOT A PROBLEM WITH THE *MATH* BEHIND STANDARDIZED TESTS. IT'S A PROBLEM WITH THE PEOPLE WHO WRITE THE QUESTIONS AND WHAT HAPPENS WHEN THE RESULTS COME OUT. AND QA SYSTEMS AREN'T IMMUNE FROM THESE PROBLEMS EITHER, AS I TALK ABOUT ELSEWHERE. BUT THE MATH IS SOUND.

HOPEFULLY IT REMINDS YOU OF A LEADERBOARD: SYSTEMS COME IN, YOU NEED TO GIVE THEM A SCORE. BUT THE MACHINE LEARNING COMMUNITY HAS A DISTINCT ADVANTAGE HERE. YOU CAN WIPE THE MEMORY OF A COMPUTER SYSTEM, SO YOU CAN GIVE IT THE SAME TEST OVER AND OVER AGAIN. BUT AS I ARGUED BEFORE IN MY PLEA FOR AN EMBRACE OF THE MANCHESTER PARADIGM

FROM TRIVIA SHOWS, THIS ISN'T ALWAYS THE CASE: THERE'S STILL SOME INFORMATION THAT GETS LEAKED.

SO WE'LL GET INTO THE MATH SOON. BUT IT'S NOT JUST FOR CREATING LEADER BOARDS IN MACHINE LEARNING AND NATURAL LANGUAGE PROCESSING. VERY SIMILAR MODELS ARE USED FOR RATING PLAYERS IN CHESS AND OTHER GAMES AND FOR PREDICTING HOW POLITICIANS WILL VOTE ON PARTICULAR BILLS.

10.1 How Computers Cheat

PARTIAL INPUT, MEMORIZATION, OVERLAP

10.2 Leaderboards are Shiny

LEADERBOARDS ARE WIDELY USED IN NLP AND PUSH THE FIELD FORWARD. WHILE LEADERBOARDS ARE A STRAIGHTFORWARD RANKING OF NLP MODELS, THIS SIMPLICITY CAN MASK NUANCES IN EVALUATION ITEMS (EXAMPLES) AND SUBJECTS (NLP MODELS). RATHER THAN REPLACE LEADERBOARDS, WE ADVOCATE A RE-IMAGINING SO THAT THEY BETTER HIGHLIGHT IF AND WHERE PROGRESS IS MADE. BUILDING ON EDUCATIONAL TESTING, WE CREATE A BAYESIAN LEADERBOARD MODEL WHERE LATENT SUBJECT SKILL AND LATENT ITEM DIFFICULTY PREDICT CORRECT RESPONSES. USING THIS MODEL, WE ANALYZE THE RANKING RELIABILITY OF LEADERBOARDS. AFTERWARDS, WE SHOW THE MODEL CAN GUIDE WHAT TO ANNOTATE, IDENTIFY ANNOTATION ERRORS, DETECT OVERTFITTING, AND IDENTIFY INFORMATIVE EXAMPLES. WE CONCLUDE WITH RECOMMENDATIONS FOR FUTURE BENCHMARK TASKS.

LEADERBOARD EVALUATIONS—FOR BETTER OR WORSE—ARE THE *DE FACTO* STANDARD FOR MEASURING PROGRESS IN QUESTION ANSWERING (?) AND IN MANY NLP TASKS (?). AN UNFORTUNATE SIDE EFFECT OF LEADERBOARD POPULARITY IS SOTA-CHASING, OFTEN AT THE EXPENSE OF CAREFULLY INSPECTING DATA AND MODELS (?). FOR EXAMPLE, THE SAME “SUPER-HUMAN” MODELS THAT TOP QUESTION ANSWERING LEADERBOARDS ([NAJBERG, 2018](#)) OFTEN FAIL SPEC-TACULARLY (??) BY LEARNING NON-GENERALIZABLE STATISTICAL PATTERNS (??). FINALLY, FOCUSING SOLELY ON METRICS CONFLATES PROGRESS ON A SPECIFIC TASK WITH PROGRESS ON REAL-WORLD NLP PROBLEMS BEHIND THE TASK (?). PLAINLY, FOCUSING ON HEADLINE SOTA NUMBERS “PROVIDE(S) LIMITED VALUE FOR SCIENTIFIC PROGRESS ABSENT INSIGHT INTO WHAT DRIVES THEM” AND WHERE THEY FAIL (?).

IN THIS WORK WE TAKE LEADERBOARDS “AS THEY ARE,” AND IMAGINE HOW THEY MIGHT BETTER SUPPORT RESEARCH. LEADERBOARDS ESTABLISH DIFFERENCES BETWEEN MODELS ON A FIXED TASK. HENCE, LEADERBOARDS SHOULD ENABLE AND ENCOURAGE THE COMPARISON OF MODELS AND INSPECTION OF EXAMPLES. AND LEADERBOARDS SHOULD ALSO SIGNAL WHEN THEY HAVE OUTLIVED THEIR USEFULNESS (?).

Figure 10.1: Difficulty and Ability Discriminating (DAD) leaderboards infer the difficulty discriminativeness and feasibility of examples. Negative discriminability suggests an annotation error; for example, the question with most negative discriminability asks **DIFFICULTY AND ABILITY DISCRIMINATING (DAD)** LEADERBOARDS THAT EXPLICITLY MODEL BOTH TASK AND SUBMISSIONS JOINTLY, RATHER THAN EITHER IN ISOLATION. DAD'S UNDERLYING MODEL IS BASED ON ITEM RESPONSE THEORY (??, IRT, REVIEWED IN §10.3), A WIDELY USED (?) ALTERNATIVE IN EDUCATIONAL TESTING TO SIMPLE SUMMARY STATISTICS (?).

DAD CAN EXPLICITLY IDENTIFY THE DIFFICULTY AND DISCRIMINABILITY OF ITEMS (FIGURE 10.1),¹ WHICH IN TURN CAN LEAD TO A MORE NUANCED RANKING OF MODELS, IDENTIFYING POOR ITEMS, AND BETTER UNDERSTANDING OF A DATASET AND TASK. THROUGHOUT THE PAPER, WE USE THE QUESTION ANSWERING (QA) BENCHMARK SQUAD 2.0 (?). FOR EXAMPLE, DAD CAN IDENTIFY QUESTIONS THAT ARE CHALLENGING TO MODELS AND QUESTIONS THAT ARE WRONG (INCORRECTLY ANNOTATED). IN ADDITION TO BETTER UNDERSTANDING DATASETS, IT IS ALSO HELPFUL FOR EFFICIENTLY SELECTING EVALUATION ITEMS TO ANNOTATE. WE CONCLUDE WITH RECOMMENDATIONS FOR FUTURE LEADERBOARDS (§10.9) AND DISCUSS WHERE IRT IN NLP CAN GO NEXT (§10.10).

How the SAT can Help Leaderboards

IF YOU'VE EVER APPLIED TO A UNIVERSITY, YOU'VE PROBABLY HAD TO TAKE A STANDARDIZED TEST. AS WE TALKED ABOVE, THE "STANDARDIZED" IN STANDARDIZED TEST MEANS THAT THE SCORE YOU GET MAKES SENSE NO MATTER WHEN YOU TOOK THE TEST. SO HOW DOES THE MATH WORK THAT LETS THAT HAPPEN?

THE BIG IDEA IS THAT WE WANT TO MODEL THE PROBABILITY OF A PERSON j GETTING QUESTION i RIGHT. WE'RE GOING TO TREAT THIS AS A LOGISTIC REGRESSION. AS WITH EVERY LOGISTIC REGRESSION, YOU PASS SOME NUMBER THROUGH A SIGMOID FUNCTION TO GENERATE A PROBABILITY. THE NUMBER THAT YOU PASS IN THIS FUNCTION—LIKE A FEATURE IN A LOGISTIC REGRESSION—HELPS EXPLAIN YOUR DATA. AND IN THE CASE THE DATA ARE A BIG MATRIX OF WHO GOT WHAT QUESTIONS RIGHT AND WHO GOT WHAT QUESTIONS WRONG.

IN THE LINGO OF IRT, THE QUESTIONS ARE THE "ITEMS", THE AGENTS ANSWERING THE QUESTIONS ARE THE "SUBJECTS" AND THEN WHETHER THEY GOT A QUESTION RIGHT OR NOT IS THE "RESPONSE". SO NOW HOPEFULLY YOU SEE WHAT IT'S CALLED ITEM RESPONSE THEORY. THE GOAL IS TO PREDICT EACH OF THESE BINARY RESPONSES AS IF IT WERE A LOGISTIC REGRESSION.

SO LIKE ALL LOGISTIC REGRESSION, THE PREDICTION COMES OUT OF A SIGMOID FUNCTION THAT YOU HOPEFULLY REMEMBER. SO WHAT GOES *IN* TO THAT SIGMOID FUNCTION? WE'LL ADD MORE PARAMETERS SOON, BUT LET'S START WITH THE FIRST TWO: DIFFICULTY AND SKILL.

FOR CONCRETENESS, LET'S ASSUME THAT WE HAVE A QUESTION-ANSWERING

¹ Example and feasibility distribution in Appendix 10.11. Interactive visualization linked from <http://irt.pedro.ai>.

TASK AND TWO SUBJECTS: KEN, WHO IS GOOD AT TRIVIA, AND BURT, WHO IS NOT. IN THE SIMPLEST IRT MODELS, EACH SUBJECT j HAS A RANDOM VARIABLE θ_j CORRESPONDING TO THEIR SKILL: KEN'S IS BIG, BURT'S IS SMALL. THE HIGHER THE SKILL NUMBER (THETA), THE MORE LIKELY THEY ARE TO ANSWER QUESTIONS CORRECTLY.

BUT YOU CANNOT KNOW THAT UNTIL YOU START ASKING THEM QUESTIONS OF VARYING DIFFICULTY β_i . HARDER QUESTIONS HAVE A HIGHER DIFFICULTY ("WHAT IS THE AIRSPEED OF AN UNLADEN SWALLOW") THAN EASY ONES ("WHO IS BURIED IN GRANT'S TOMB"). THE BIGGER THE MARGIN BETWEEN A SUBJECT'S SKILL θ_j AND AN ITEM'S DIFFICULTY β_i , $\theta_j - \beta_i$, THE MORE LIKELY THAT SUBJECT j RESPONDS CORRECTLY $p_{i,j}(r_{i,j} = 1)$.

SO HOW DO THESE TWO NUMBERS FEED INTO A SIGMOID FUNCTION? YOU TAKE THE DIFFERENCE BETWEEN THE DIFFICULTY AND THE SKILL. WHEN THE SKILL MINUS THE DIFFICULTY IS HIGH, THE HIGHER THE PROBABILITY OF THE SUBJECT ANSWERING IT CORRECTLY. IF THE SKILL MATCHES THE DIFFICULTY, THEN THE SUBJECT HAS A 50-50 CHANCE OF ANSWERING CORRECTLY.

OKAY, SO LET'S MAKE THIS A LITTLE MORE COMPLICATED AND ADD TWO ADDITIONAL TERMS: DISCRIMINABILITY AND FEASIBILITY: DISCRIMINABILITY γ_i AND FEASIBILITY λ_i .

Discrimination DISCRIMINABILITY IS HOW WELL THE QUESTION DISTINGUISHES BETTER SUBJECTS FROM WORSE ONES. IT'S MULTIPLIED BY THE DIFFERENCE BETWEEN THE SKILL AND THE DIFFICULTY. MATHEMATICALLY, THIS CHANGES THE SLOPE OF THE RESPONSE CURVE. A DISCRIMINATIVE QUESTION IS CHALLENGING BUT CAN STILL BE ANSWERED CORRECTLY BY A STRONG SUBJECT. IF KEN'S ABILITY IS HIGHER THAN MOST ITEMS' DIFFICULTY ($\theta_j - \beta_i$ IS LARGE), ITEM DISCRIMINABILITY MULTIPLIES THIS GAP BY γ_i . QUESTIONS WITH LOW γ_i ARE LOW QUALITY: THEY HAVE ANNOTATION ERROR OR DO NOT MAKE SENSE.

HERE'S A PICTURE OF THAT FOR A QUESTION WITH ZERO DIFFICULTY. AS THE SKILL (X-AXIS GOES UP), THE MORE LIKELY THE SUBJECT IS TO ANSWER THE QUESTION. THE BIGGER THE DISCRIMINABILITY, THE STEEPER THE CURVE. HERE IT IS FOR 0.5 AND HERE IT IS FOR 2.0.

THIS WOULD NEVER HAPPEN IN PRACTICE, BUT LET'S IMAGINE YOU COULD CREATE A TEST SET WHERE EVERY QUESTION HAD PERFECT DISCRIMINATION. THEN THE RESPONSE CURVE WON'T BE A CURVE, IT'LL BE A STEP FUNCTION. SO LET'S SAY YOU HAVE A QUESTION WITH DIFFICULTY 0.2 AND PERFECT DISCRIMINATION. ANYBODY WITH SKILL GREATER THAN 0.2 WILL GET IT RIGHT 100% OF THE TIME.

BUT IF YOU CAN DO IT ONCE, YOU CAN DO IT AGAIN. LET'S CREATE SEVERAL PERFECT DISCRIMINATORS WITH A NICE GRID OF DIFFICULTY. SO NOW, WHEN A PERSON COMES IN, THEY'LL ANSWER EVERY QUESTION CORRECT UNTIL THEIR SKILL IS SURPASSED BY THE QUESTION. FIRST QUESTION THEY GET WRONG MEANS THEY CAN LEAVE THE TEST AND GO HOME EARLY.

AGAIN, THIS FANTASY IS IMPOSSIBLE, BUT YOU CAN IMAGINE WHY FOLKS

Figure 10.2: A DAD leaderboard uses IRT to jointly infer item difficulty β_i , discriminability or feasibility γ_i , and subject skill θ_j . These predict the likelihood $p_{ij}(r_{ij} = 1)$ of a correct response r_{ij} .

Feasibility ANOTHER WAY OF CAPTURING POOR QUALITY QUESTIONS IS THE FEASIBILITY λ_i . FOR EXAMPLE, IF THE QUESTION “WHO WAS THE FIRST PRESIDENT” HAS THE ANSWER RAJENDRA PRASAD, THE QUESTION HAS AN UNSTATED IMPLICIT ASSUMPTION THAT SUBJECTS MUST GUESS WHAT COUNTRY OR COMPANY THE QUESTION IS ABOUT. IN THE MODEL **TODO, FIX ME!!!!**, IF A LARGE FRACTION OF SUBJECTS ALL GET AN ITEM WRONG, EVERYONE’S PROBABILITY OF GETTING THE ITEM RIGHT IS CAPPED AT λ_i . IN NLP TERMS, $1 - \lambda_i$ CORRESPONDS TO THE PREVALENCE OF ANNOTATION ERRORS THAT LEAD TO UNSOLVABLE ITEMS.

LET’S NOW SEE WHAT FEASIBILITY DOES. FEASIBILITY IS RELEVANT TO NLP SINCE INFEASIBLE QUESTIONS SUGGEST AN ANNOTATION ERROR. LET’S GO BACK TO OUR FAVORITE “BAD” QUESTION: WHAT IS THE CAPITAL OF GEORGIA? IF THIS SORT OF QUESTION CAN HAVE THE ANSWER “ATLANTA” OR “TBLISI” WITH EQUAL PROBABILITY, THEN NOBODY, NO MATTER HOW SMART THEY ARE, CAN GET THIS RIGHT WITH MORE THAN 0.5 PROBABILITY.

MATHEMATICALLY, SEE HOW THE FEASIBILITY PARAMETER WORKS IS MORE STRAIGHTFORWARD. IT’S JUST A LINEAR SCALING, THAT BRINGS DOWN THE MAXIMUM VALUE THE PROBABILITY CAN TAKE ON. THE NORMAL LOGISTIC CURVE IS ON TOP, LAMBDA=1, BUT IT GETS SQUISHED DOWN AS LAMBDA GETS SMALLER AND SMALLER. IF LAMBDA= $\frac{1}{2}$, THEN NO MATTER HOW SMART YOU ARE, IT’S BASICALLY A COIN FLIP WHETHER YOU CAN ANSWER THE QUESTION.

10.3 A Generative Story for Leaderboards

LEADERBOARDS ARE A PRODUCT OF THE METRICS, EVALUATION DATA, AND SUBJECTS (MACHINE OR HUMAN) WHO ANSWER ITEMS (FIGURE 10.2).

GENERALLY, GIVEN n TEST ITEMS $\mathcal{X} = (X_1, \dots, X_n)$ AND m SUBJECTS $\mathcal{S} = (S_1, \dots, S_m)$, WHERE EACH SUBJECT ANSWERS EVERY ITEM, WE WANT TO ESTIMATE SUBJECT SKILLS AND ITEM DIFFICULTIES. TO DISCOVER THE RANDOM VARIABLES THAT BEST EXPLAIN THE DATA, WE TURN TO PROBABILISTIC INFERENCE [PEARL \(1988\)](#).

HAVING INTRODUCED ALL OF THE CONSTITUENT ELEMENTS OF THE MODEL, WE CAN NOW PRESENT THE FULL GENERATIVE MODEL:

1. FOR EACH SUBJECT j :
 - (A) DRAW SKILL $\theta_j \sim \mathcal{N}(\mu_\theta, \tau_\theta^{-1})$
2. FOR EACH ITEM i :
 - (A) DRAW DIFFICULTY $\beta_i \sim \mathcal{N}(\mu_\beta, \tau_\beta^{-1})$
 - (B) DRAW DISCRIMINABILITY $\gamma_i \sim \mathcal{N}(\mu_\gamma, \tau_\gamma^{-1})$
 - (C) DRAW FEASIBILITY $\lambda_i \sim U[0, 1]$

3. DRAW SUBJECT i RESPONSE ON ITEM j ,

$$r_{ij} \sim p_{ij}(r_{ij} | \theta_j, \beta_i, \lambda_i) =$$

$$p_{ij}(r_{ij} = 1 | \theta_j) = \frac{\lambda_i}{1 + e^{-\gamma_i(\theta_j - \beta_i)}}. \quad (10.1)$$

FOR TODO, FIX ME!!!, γ_i AND λ_i ARE FIXED TO 1.0, WHILE FOR TODO, FIX ME!!!, ONLY λ_i IS FIXED.²

MEANS $\mu_\theta, \mu_\beta, \mu_\gamma$ ARE DRAWN FROM $\mathcal{N}(0, 10^6)$ AND $\tau_\theta, \tau_\beta, \tau_\gamma$ FROM A $\Gamma(1, 1)$ PRIOR, AS IN ? AND RECOMMENDED BY ?.³

BECAUSE IT IS DIFFICULT TO COMPLETELY CODIFY SKILL AND DIFFICULTY INTO A SINGLE NUMBER, WE CAN REWRITE THE EXPONENT IN EQUATION 10.1 AS A SUM OVER DIMENSIONS $-\gamma_i(\sum_k \theta_{j,k} - \beta_{i,k})$, WHERE EACH DIMENSION CAPTURES THE INTERACTION BETWEEN AN ITEM'S DIFFICULTY AND A SUBJECT'S SKILL. FOR EXAMPLE, PERHAPS BURT COULD BETTER EXPLOIT ARTIFACTS IN ONE DIMENSION (THEIR SKILL FOR $\theta_{j,k=5}$ IS HIGH BUT EVERYWHERE ELSE IS LOW) WHILE KEN MIGHT NOT KNOW MUCH ABOUT A PARTICULAR TOPIC LIKE POTENT POTABLES ($\theta_{j,k=2}$ IS LOW BUT EVERYWHERE ELSE IS HIGH). WE CALL THIS MODEL **TODO, FIX ME!!!!**.⁴ MULTIDIMENSIONAL IRT MODELS (?) COULD—IN ADDITION TO BETTER MODELING DIFFICULTY—ALSO CLUSTER ITEMS FOR INTERPRETATION; WE BRIEFLY EXPERIMENT WITH THIS (APPENDIX 10.16), BUT LEAVE MORE TO FUTURE WORK (§10.10).

Examples are Not Equally Useful

IRT'S FUNDAMENTAL ASSUMPTION IS THAT NOT ALL ITEMS AND SUBJECTS ARE EQUAL. THIS EXPLAINS WHY LEADERBOARDS CAN FAIL WHILE HAVING "NORMAL LOOKING" ACCURACIES. AS A THOUGHT EXPERIMENT, CONSIDER A DATASET THAT IS ONE THIRD EASY ($\beta_i \in [0, 1]$), ONE THIRD MEDIUM DIFFICULTY ($\beta_i \in [2, 3]$), AND ONE THIRD HARD ($\beta_i \in [6, 7]$). SUPPOSE THAT KEN HAS SKILL $\theta_k = 4$ WHILE BURT HAS SKILL $\theta_b = 2$. A STANDARD LEADERBOARD WOULD SAY THAT KEN HAS HIGHER ACCURACY THAN BURT. BUT SUPPOSE THERE'S A NEW SUBJECT THAT WANTS TO CHALLENGE KEN; THEY ARE NOT GOING TO RELIABLY DETHRONE KEN UNTIL THEIR SKILL θ_c IS GREATER THAN SIX.

THIS IS A MORE MATHEMATICAL FORMULATION OF THE "EASY" AND "HARD" DATASET SPLITS IN QUESTION ANSWERING (??). IN TODO, FIX ME!!!, THIS RECAPITULATES THE OBSERVATION OF ? THAT ANNOTATION ERROR CAN HINDER EFFECTIVE LEADERBOARDS. DAD HELPS SYSTEMATIZE THESE OBSERVATIONS AND DIAGNOSE DATASET ISSUES.

Inference

TO ESTIMATE THE LATENT PARAMETERS OF OUR MODEL, WE USE MEAN-FIELD VARIATIONAL INFERENCE JORDAN ET AL. (1999). IN VARIATIONAL INFERENCE, WE PROPOSE A DISTRIBUTION OVER THE LATENT VARIABLES, $q_\phi(\cdot)$, THAT APPROXIMATES THE TRUE BUT INTRACTABLE POSTERIOR $p(\cdot)$. WE THEN MINIMIZE

² In psychometrics, TODO, FIX ME!!!! is called a Rasch (?) or 1 parameter logistic (1PL) model, TODO, FIX ME!!!! is a 2PL model, and TODO, FIX ME!!!! is a 4PL model with guessing set to zero.

³ We differ by allowing $\gamma < 0$ to identify bad items.

⁴ We do not incorporate feasibility into the TODO, FIX ME!!!! model since it already improves over 1D models without it.

THE KL-DIVERGENCE BETWEEN THESE DISTRIBUTIONS, EQUIVALENT TO MAXIMIZING THE EVIDENCE LOWER-BOUND (ELBO) WITH RESPECT TO THE VARIATIONAL PARAMETERS.

IN OUR CASE, $q_\phi(\cdot)$ IS A MEAN-FIELD DISTRIBUTION, WHICH MEANS IT FACTORIZES OVER EACH OF THE LATENT VARIABLES (THE PRODUCT IS OVER THE $n \times m$ SUBJECT-ITEM PAIRS)

$$q_\phi(\theta, \beta, \gamma, \mu, \tau) = q(\mu)q(\tau) \prod_{i,j} q(\theta_j)q(\beta_i)q(\gamma_i)$$

SPECIFICALLY, FOR OUR KEY LATENT VARIABLES $z \in \{\theta, \beta, \gamma\}$, THE ASSOCIATED VARIATIONAL DISTRIBUTIONS ARE OF THE FORM $q(z) = \mathcal{N}(u_z, t_z^{-1})$. RECALL THAT IN THE GENERATIVE DISTRIBUTION, EACH LATENT z IS DRAWN FROM A $\mathcal{N}(\mu_z, \tau_z^{-1})$ WHOSE PARAMETERS ARE ALSO LATENT VARIABLES; FOR THESE VARIABLES, WE USE THE VARIATIONAL DISTRIBUTIONS $q(\mu_z) = \mathcal{N}(u_{\mu_z}, t_{\mu_z}^{-1})$ AND $q(\tau_z) = \Gamma(a_{\tau_z}, b_{\tau_z})$. WE OPTIMIZE THE ELBO WITH RESPECT TO THE VARIATIONAL PARAMETERS

$$\phi = \{u_z, t_z, u_{\mu_z}, t_{\mu_z}, a_{\tau_z}, b_{\tau_z}, \lambda\}$$

FOR ALL z USING ADAM (?).

WITH DAD'S LEADERBOARD IRT MODEL INTRODUCED, WE NEXT DISCUSS HOW LEADERBOARD SUBJECTS ARE STATISTICALLY COMPARED AND ALTERNATIVE METHODS—SUCH AS USING IRT PARAMETERS—to EVALUATE WHETHER TWO MODELS ARE TRULY DIFFERENT.

10.4 Ranking and Comparing Subjects

FUNDAMENTALLY, THE OBJECTIVE OF COMPARATIVE EVALUATIONS LIKE LEADERBOARDS IS TO DECIDE WHETHER MODEL A IS BETTER THAN MODEL B . A THREAD OF NLP HAS RIGHTFULLY ADVOCATED FOR ADDING RIGOR TO THESE DECISIONS USING STATISTICS (?), CLASSICAL TESTING THEORY WHERE THE OBJECTIVE IS TO INFER A TRUE SCORE T FROM THE OBSERVED TEST SCORE $X = T + E$ GIVEN A MEASUREMENT ERROR E , UNIFORM ACROSS SUBJECTS. HOWEVER, IN EDUCATIONAL TESTING—A FIELD MEASURING SKILL AND KNOWLEDGE IN HUMANS—IRT IS A PRIMARY MEASUREMENT INSTRUMENT (?), p. 2). A MAJOR MOTIVATION FOR IRT IS THAT SUBJECTS OF DIFFERENT SKILL HAVE *DIFFERENT* ERRORS. IRT EXPLICITLY ACCOUNTS FOR THE BANDWIDTH-FIDELITY DILEMMA (?): ITEMS CAN EITHER ACCURATELY MEASURE A NARROW ABILITY RANGE (FIDELITY) OR INACCURATELY MEASURE LARGE ABILITY RANGES (BANDWIDTH).⁵ THIS SECTION AND THE NEXT CONTRAST METHODS FOR IDENTIFYING THE BEST MODEL AND ADVOCATE FOR IRT.

IMPLICIT IN NEARLY ALL LEADERBOARD EVALUATIONS IS RANKING MODELS BY A STATISTIC SUCH AS THE AVERAGE ACCURACY. AS WE SHOW IN §10.5, NAÏVE RANKINGS ARE NOISIER THAN IRT RANKINGS.

⁵ Estimation error of θ varies by position (Appendix 10.15).

10.5 IRT for Leaderboards

LEADERBOARDS SHOULD: (1) RELIABLY AND EFFICIENTLY RANK BETTER MODELS AHEAD OF WORSE MODELS (??) AND (2) GUIDE INSPECTION OF ITEMS AND SUBJECTS (§10.6). THE FIRST AMELIORATES THE UNAVOIDABLE RANDOMNESS OF FINITE EVALUATIONS WHILE THE SECOND ENABLES ERROR ANALYSIS (?) AND MODEL PROBING (??). FIRST WE VERIFY THAT IRT MODELS ACCURATELY PREDICT THE RESPONSES OF SUBJECTS (§10.5). NEXT, A RANKING STABILITY ANALYSIS SHOWS THAT IRT HAS MODESTLY BETTER RELIABILITY THAN CLASSICAL RANKINGS (§10.5). LASTLY, USING IRT TO ACTIVELY SAMPLE ITEMS FOR ANNOTATION YIELDS RANKINGS WITH BETTER CORRELATION TO COMPLETE TEST DATA (§10.5).

Why a Linear Model Baseline

AT FIRST BLUSH, THE DIFFERENCES BETWEEN IRT AND LOGISTIC REGRESSION ARE MINIMAL, BUT WE INCLUDE THE COMPARISON TO ADDRESS NATURAL QUESTIONS FROM THE NLP COMMUNITY: (1) DO THE IDIOSYNCRASIES OF THE IRT FORMULATION HURT ACCURACY? (2) SHOULD WE ADD FEATURES TO BETTER UNDERSTAND PHENOMENA IN THE QUESTIONS? (3) WHY NOT USE DEEP MODELS?

THE NEXT SECTION ARGUES THAT BOTH IRT AND LOGISTIC REGRESSION ARE ACCURATE EVEN WITHOUT LABORIOUSLY ENGINEERED TASK-SPECIFIC FEATURES. ADDING OBVIOUS FEATURES SUCH AS ITEM WORDS (E.G., QUESTIONS) ONLY MINIMALLY IMPROVES THE ACCURACY. WE EXPLICITLY OMIT LESS INTERPRETABLE DEEP MODELS SINCE OUR GOAL IS TO MAKE LEADERBOARDS MORE INTERPRETABLE.

Response Prediction is Accurate

JUST AS EDUCATIONAL TESTING RESEARCHERS VALIDATE IRT MODELS BY SEEING IF THEY PREDICT SUBJECT RESPONSES CORRECTLY (?), WE VALIDATE HOW WELL DAD PREDICTS WHETHER SQUAD MODELS GET QUESTIONS RIGHT.

WE COMPARE AGAINST A LOGISTIC REGRESSION LINEAR MODEL (LM) IMPLEMENTED WITH VOWPAL WABBIT (?). SINCE INTEGRATING HAND-CRAFTED FEATURES IS EASY, WE INCORPORATE FEATURES DERIVED FROM SUBJECT IDS; ITEM IDS; FUNCTIONS OF THE SQUAD QUESTION, ANSWER, AND TITLE; AND IRT PARAMETERS (DETAILS IN APPENDIX 10.12). AS IN IRT, LOGISTIC REGRESSION PREDICTS WHETHER A SUBJECT CORRECTLY RESPONDS TO AN ITEM. LATER, WE DISCUSS WAYS TO INTEGRATE MORE FEATURES INTO IRT (§10.10).

SQuAD Leaderboard Data

EXPERIMENTS ARE ON THE SQUAD 2.0 LEADERBOARD. DEVELOPMENT DATA ARE PUBLICLY AVAILABLE, AND ORGANIZERS PROVIDE TEST SET RESPONSES. THERE ARE 161 DEVELOPMENT SUBJECTS, 115 TEST SUBJECTS, AND 11,873 ITEMS (1.9 MILLION TOTAL PAIRS). EXPERIMENTS THAT DO NOT NEED TEST RESPONSES USE ALL DEVELOPMENT SUBJECTS; THOSE THAT DO USE THE SMALLER TEST SUBSET.

Evaluation Scheme

FOLLOWING PRIOR WORK (?), WE EVALUATE IRT AND LINEAR MODELS BY HOLDING OUT 10% OF RESPONSES AND COMPUTING CLASSIFICATION METRICS.⁶ IN SQUAD, PREDICTING WHETHER A RESPONSE IS CORRECT IS AN IMBALANCED CLASSIFICATION PROBLEM (80.4% OF RESPONSES IN THE DEVELOPMENT SET ARE CORRECT). THUS, WE USE ROC AUC, MACRO F1, AND ACCURACY.

IRT Response Prediction is Accurate

IRT MODELS THAT INCORPORATE MORE PRIORS INTO THE GENERATIVE STORY SHOULD BE BETTER, BUT ARE THEY? WE COMPARE FOUR IRT MODELS: TODO, FIX ME!!!, TODO, FIX ME!!!, TODO, FIX ME!!!, AND TODO, FIX ME!!! (§10.3). THE MORE SOPHISTICATED MODELS ARE BETTER AND ALL IMPROVE OVER THE LM (FIGURE 10.3) AND CORRELATE WELL WITH EACH OTHER (APPENDIX 10.13). TO BE CLEAR, WHILE HIGHER ACCURACY THAN LM IS GOOD, OUR GOAL IS TO VALIDATE THAT IRT MODELS ARE ACCURATE; LATER, WE INSPECT MODEL ERRORS AND IDENTIFY ANNOTATION ERRORS (§10.6).

What Model Features are Predictive?

INTEGRATING ADDITIONAL FEATURES INTO BAYESIAN MODELS IS NOT TRIVIAL, SO WE INSTEAD USE THE FLEXIBILITY OF LINEAR MODELS TO IDENTIFY USEFUL FEATURES. OUR LEAVE-ONE-IN ABLATION COMPARES FEATURES (FIGURE 10.3): THE TOP ABLATIONS BOTH USE IRT FEATURES, FURTHER VALIDATING IRT PARAMETERS. THE SUBJECT AND ITEM IDENTIFIER FEATURES ARE ALSO STRONGLY PREDICTIVE, BUT ITEM IS THE STRONGER OF THE TWO. TEXT-BASED FEATURES ARE WEAKER, BUT THIS SUGGESTS FUTURE WORK TO BETTER INTEGRATE THEM INTO IRT MODELS (§10.10).

Ranking with IRT

LEADERBOARDS SHOULD PRODUCE RELIABLE SUBJECT RANKINGS: CAN DAD RANK SYSTEMS EVEN WITH A TINY TEST SET? THUS, WE COMPARE THE CORRELATION BOTH OF TRADITIONAL AVERAGE ACCURACY (§10.4) AND IRT RANKINGS ON THE WHOLE TEST SET COMPARED TO THE RANKINGS OF THE SAME METRIC ON A SMALLER TEST SET. OUR FIRST EXPERIMENT (§10.5) EXAMINES THE STABILITY OF EXISTING ITEMS AND SUBJECTS WHILE THE SECOND (§10.5) INVESTIGATES STABILITY OF “NEW” EVALUATION DATA USING SAMPLING STRATEGIES.

IRT Rankings Have Better Reliability

RANKINGS SHOULD BE RELIABLE WITHIN THE SAME DATASET (E.G., ON DEV SET) AND GENERALIZE TO SIMILAR DATASETS (E.G., WITH A TEST DATASET). TO TEST

⁶ Everywhere else in the paper, we train on all responses.

Figure 10.3: We compare each IRT and linear model (LM) by how well they predict subject responses. We focus on ROC AUC since predicting responses is an imbalanced classification problem (most subjects are correct). Under that metric, all IRT models improve over the best LM, and the strongest LM ablation only uses IRT features. That textual features are predictive in the LM suggests they could improve future models.

THE FIRST, WE MEASURE THE RANKING STABILITY OF MUTUALLY EXCLUSIVE SAMPLES OF THE DEVELOPMENT DATA $?$. TO TEST THE SECOND, WE MEASURE THE CORRELATION BETWEEN DEVELOPMENT SET SAMPLE RANKINGS TO TEST SET RANKINGS $(?)$.

SPECIFICALLY, FOR A RANGE OF SAMPLE SIZES⁷ WE (1) SAMPLE TWO PARTITIONS OF THE DATA, (2) COMPUTE THE CLASSICAL RANKING⁸ AND THE IRT RANKING FROM A REFIT TODO, FIX ME!!! MODEL, THEN (3) COMPUTE KENDALL'S CORRELATION $(?)$ BETWEEN THE SAMPLES FOR EACH RANKING (DETAILS IN APPENDIX 10.14). IN BOTH CASES IRT RANKINGS HAVE HIGHER CORRELATION THAN CLASSICAL RANKINGS (FIGURE 10.4, LEFT). SINCE THE BENEFIT IS STRONGEST AT LOW SAMPLE SIZES, IRT CAN IMPROVE THE RELIABILITY OF SMALL-SCALE EVALUATIONS.

THE SECOND EXPERIMENT EXAMINES RANKING GENERALIZATION: IRT YIELDS MORE RELIABLE MEASURES OF SUBJECT SKILL, IMPLYING A GREATER CONSISTENCY IN SUBJECT RANKINGS ACROSS EVALUATION SETTINGS. FIGURE 10.4 COMPARES THE DEVELOPMENT SET SAMPLE RANKINGS COMPUTED ABOVE TO RANKINGS OBTAINED USING SUBJECTS' TEST SET RESPONSES (WITH THE SAME IRT MODEL).

ACROSS ALL SAMPLE SIZES, SUBJECTS' IRT ABILITY ESTIMATED ON THE DEVELOPMENT SET CORRELATES WELL TEST SET ABILITY. CRUCIALLY, THIS IS BETTER THAN THE CORRESPONDING CLASSICAL METRICS LIKE ACCURACY (APPENDIX 10.14 QUANTIFIES THE STATISTICAL SIGNIFICANCE OF THE DIFFERENCE), SUPPORTING OUR ORIGINAL MOTIVATION FOR USING IRT.⁹

IRT Improves Cold Start Reliability

IRT CAN ALSO GUIDE THE CONSTRUCTION OF TESTS. JUST AS IRT PRACTITIONERS PREPARE TESTS FOR HUMANS, WE TOO CONSTRUCT TESTS FOR MACHINES. IN EDUCATIONAL TESTING, COLLECTING RESPONSES FROM HUMANS IS EXPENSIVE; LIKEWISE, ALTHOUGH QUESTIONS ARE CHEAP IN SEARCH-BASED QA TASKS $(??)$, ANNOTATING ANSWERS IS EXPENSIVE. LIKEWISE, "GRADING" MACHINE DIALOG RESPONSES IS EXPENSIVE AND IRT HELPS $(?)$. TO EMULATE THIS SETTING, WE USE COMPUTERIZED ADAPTIVE TESTING $(?)$ TO ITERATIVELY SELECT SQUAD ITEMS TO "ANNOTATE."

AS IN HUMAN TEST PREPARATION, WE USE EXISTING ANNOTATIONS TO INFER ITEM PARAMETERS AND ITERATIVELY INFER THE ABILITY OF NEW SUBJECTS. THIS EXPERIMENT SPLITS m SUBJECTS INTO A TRAINING GROUP (80%) AND A TESTING GROUP (20%). THE TRAINING GROUP REPRESENTS SUBJECTS FOR WHICH WE HAVE FULL ITEM PREDICTIONS AND ANNOTATIONS; THE TESTING GROUP REPRESENTS A NEW GROUP OF SUBJECTS THAT WE NEED TO RANK. TO EFFICIENTLY RANK, WE SHOULD ITERATIVELY CHOOSE ITEMS TO ANNOTATE THAT YIELD THE MOST INFORMATION ABOUT THE RANKING IF ALL THE DATA WERE ANNOTATED.

THIS EXPERIMENT COMPARES HOW WELL SEVERAL ITEM SELECTION STRATE-

Figure 10.4: Compared to the final ranking over a large test set, how well does a small test set correlate? The left shows correlation between mutually exclusive development set samples and the right between development samples and the full test set. In both experiments (panes), ranking systems by IRT ability is more stable—across all sample sizes—
 $?$ The sample size is much less than half the size of the development data so that we can obtain two samples.
 shows 95% confidence intervals of rank correlations across ten trials per sample size.
 exact match score.

⁷ Since the maximum trial size was limited, we train one final model with the full data, see Table 10.3 in the Appendix 10.14.

Figure 10.5: Suppose we need to cold start and collect annotations for a new subject: what order would most rapidly increase correlation to the full test data? As we expect, the correlations eventually converge, but what order does each item selection method converge in?

WHICH ITEM SELECTION STRATEGIES SHOULD WE COMPARE? AS A BASELINE, WE USE NAIVE RANDOM SAMPLING. LIKE PRIOR WORK, WE COMPARE SELECTING ITEMS WITH THE HIGHEST DIFFICULTY AND THE HIGHEST DISCRIMINABILITY (?) AS WELL AS THE SUM OF THE TWO.¹¹ WE PROPOSE THAT ITEMS SHOULD BE SELECTED ACCORDING TO THEIR FISHER INFORMATION CONTENT (?)

$$I_i(\theta_j) = \frac{(p'_{ij})^2}{p_{ij}(1-p_{ij})} = \gamma_i^2 p_{ij}(1-p_{ij}) \quad (10.2)$$

AS DERIVED BY ?, P. 70.

INTUITIVELY, IF WE DO NOT YET KNOW THE TRUE SKILL θ_j , WE SHOULD PICK ITEMS WHOSE EXPECTED RESPONSE WE ARE MOST UNCERTAIN ABOUT. OUR UNCERTAINTY (ENTROPY) IS MAXIMIZED WHEN THE LIKELIHOOD OF A CORRECT RESPONSE p_{ij} IS THE SAME AS THE LIKELIHOOD OF AN INCORRECT RESPONSE $1 - p_{ij}$, WHICH CORRESPONDS TO THE MAXIMAL VALUE OF $I_i(\theta_j)$; IT IS ALSO SENSIBLE THIS VALUE INCREASES AS DISCRIMINABILITY γ_i INCREASES.

TO INFER THE MAXIMALLY INFORMATIVE ITEMS, WE ESTIMATE THE ABILITY θ_j OF EACH SUBJECT USING THE CURRENTLY SELECTED ITEMS, USE THE ABILITY TO COMPUTE THE INFORMATION OF EACH YET-TO-BE-ANNOTATED ITEM FOR EACH SUBJECT, AND THEN AGGREGATE THE INFORMATIVENESS

$$\text{INFO}(i) = \sum_j I_i(\theta_j) \quad (10.3)$$

BY ITEM i SUMMED OVER SUBJECTS j . THIS APPROACH IS SIMILAR TO UNCERTAINTY SAMPLING AND REDUCES TO IT FOR THE TODO, FIX ME!!! MODEL (?). WE INITIALLY SEED WITH THE TWENTY-FIVE MOST DISCRIMINATIVE ITEMS (DETAILS IN APPENDIX 10.14).

LIKE COMPUTERIZED ADAPTIVE TESTING (?), FIGURE 10.5 SHOWS THAT AT LOWER SAMPLE SIZES THREE OF THE IRT SAMPLING METHODS ARE BETTER THAN RANDOM SAMPLING—DIFFICULTY DOES WORSE. THE OTHER IRT METHODS HAVE COMPARABLE CORRELATION. THUS, BY USING IRT, DAD CAN BOTH IMPROVE RANKINGS AND GUIDE ANNOTATION.

10.6 Qualitative Insights on Leaderboards

DAD ALSO HELPS QUALITATIVE ANALYSIS OF ITEMS AND SUBJECTS. FIRST, IRT IDENTIFIES OVERTFITTNG AND GENERALIZES PARTITIONING DATASETS BY DIFFICULTY. THEN WE SHOW THAT—LIKE IN EDUCATIONAL TESTING—IRT IDENTIFIES GOOD AND BAD ITEMS.

¹⁰ We compute correlations with the complete development set on ten trials to build 95% confidence intervals.

¹¹ We train an TODO, FIX ME!!!! model to simplify sampling (e.g., avoiding a tradeoff between feasibility and discriminability).

Guiding Analysis with IRT

SEVERAL WORKS CURATE EASY AND HARD QA SUBSETS BASED ON HOW MANY MODELS ANSWER CORRECTLY (?) OR HEURISTICS (?). IRT CAN CREATE SIMILAR SUBSETS USING TODO, FIX ME!!!, THE BEST 1D MODEL. DIFFICULTY FINDS WHERE SUBJECTS IMPROVE WHILE DISCRIMINABILITY AND FEASIBILITY CAN SURFACE ITEMS THAT MAY BE INVALID. FOR EXAMPLE, ONE LOW FEASIBILITY QUESTION (FIGURE 10.10) ASKS “WHAT ARE TWO EXAMPLES OF TYPES OF TURING MACHINES?” WHICH HAS TWO PROBLEMS: (1) THE ANSWER OMITS FIVE TYPES AND (2) SPAN-BASED EVALUATION PRECLUDES SELECTING NON-CONTIGUOUS TYPES.

AFTER EXCLUDING ITEMS WITH NEGATIVE DISCRIMINABILITY—THEY ARE LIKELY ERRONEOUS—WE SORT ITEMS INTO BINS. WE BREAK BOTH DIFFICULTY AND DISCRIMINABILITY INTO FOUR BINS—TAKING THE 25TH, 50TH, AND 75TH PERCENTILES—CREATING EIGHT TOTAL BINS. THEN WE SELECT REPRESENTATIVE SQUAD SUBJECTS WITH THEIR EXACT MATCH SCORES (FIGURE 10.6). LET’S EXAMINE A FEASIBLE ITEM WITH POSITIVE DIFFICULTY AND DISCRIMINABILITY LIKE “WHAT REFORM WAS ATTEMPTED FOLLOWING THE NICE TREATY?”¹² IN THIS CASE, THE ANNOTATOR’S SPAN IS TOO LONG—RESULTING IN ALMOST NO CORRECT ANSWERS AND A LOW FUZZY MATCH (TOKEN F1). IN CONTRAST, ONE HIGHLY DISCRIMINATIVE QUESTION SUCCEEDS BECAUSE THERE ARE MULTIPLE PLAUSIBLE GUESSES TO “WHO DID THE NORMANS TEAM UP WITH IN ANATOLIA?”¹³ WHILE BOTH THE ARMENIAN STATE AND TURKISH FORCES ARE SUPERFICIALLY PLAUSIBLE ANSWERS, ONLY TURKISH FORCES IS CORRECT; NONETHELESS, SOME MODELS ARE FOOLED. USING IRT TO GUIDE SUBJECT ANALYSIS IS HELPFUL; NEXT, WE TEST HOW EFFICIENT IT IS IN IDENTIFYING ANNOTATION ERROR.

Identifying Annotation Error

TO TEST IF IRT CAN IDENTIFY ANNOTATION ERROR, WE INSPECT SIXTY SQUAD DEVELOPMENT SET ITEMS. WE SELECT TEN ITEMS FROM EACH OF THESE GROUPS: THE MOST NEGATIVE DISCRIMINABILITY, DISCRIMINABILITY NEAREST TO ZERO, THE HIGHEST DISCRIMINABILITY, THE LEAST DIFFICULT, MOST DIFFICULT, AND IRT MODEL ERRORS. FOR EACH, WE ANNOTATE WHETHER THE ITEM WAS CORRECT, WAS “CORRECT” YET FLAWED IN SOME WAY, OR SIMPLY WRONG (FIGURE 10.7).¹⁴ INTER-ANNOTATOR AGREEMENT BETWEEN THREE AUTHORS ON THIS THREE-WAY ANNOTATION WITH KRIPPENDORFF’S α (?) IS 0.344. DESPITE ONLY MODEST AGREEMENT, JUST AS IN THE DEVELOPMENT OF EDUCATION TESTS, NEGATIVE DISCRIMINABILITY IS PREDICTIVE OF BAD ITEMS. WHEN DISCRIMINABILITY IS NEGATIVE, THEN THE PROBABILITY OF GETTING THE ANSWER RIGHT IS HIGHER WHEN ABILITY IS LOWER, WHICH IS UNDESIRABLE: KEN CONSISTENTLY LOSES TO BURT ON THOSE ITEMS. THIS COULD IDENTIFY BAD ITEMS IN EVALUATION SETS

Figure 10.6: We partition evaluation data by IRT difficulty and discriminability with accuracy in each quartile. Most improvements in high-accuracy systems come from getting high-difficulty questions right. Items with low discriminability (and thus prone to annotation errors) are difficult for all subjects except the overfit ARGS-BERT model. We include top-performing SQUAD SUBJECTS, SEVERAL NOTABLE SUBJECTS (SYSTEMS), AND A PAIR FROM THE BOTTOM OF THE LEADERBOARD.

¹² A: “there was an attempt to reform the constitutional law of the EU and make it more transparent.” (Appendix Figure 10.11)

¹³ Example with statistics in Appendix Figure 10.12.

¹⁴ Annotation guidelines provided in supplementary materials; Figure 10.7 uses the first set of annotations which were later augmented by two additional sets of annotations.

FOR REMOVAL.

10.7 A Re-Imagined Leaderboard Dashboard

BASING A NEW LEADERBOARD ON IRT MODELS HAS ANOTHER SIGNIFICANT BENEFIT: AS THEY ARE PURPOSEFULLY INTERPRETABLE, THEY ARE VERY USEFUL FOR VISUALIZING LEADERBOARD DATA. THE FINAL COMPONENT OF THIS WORK IS A DEMONSTRATION—BASED ON SQuAD—OF COMPONENTS THAT A RE-IMAGINED LEADERBOARD MAY HAVE.¹⁵ OUR PROPOSED VISUALIZATION HAS THREE MAIN COMPONENTS: (1) AN IMPROVED MODEL LISTING, (2) A MODEL-CENTRIC VIEW, (3) AN EXAMPLE-CENTRIC VIEW.

THE MAJORITY OF LEADERBOARDS LIST ONLY ONE RANKING METRIC OR LIST MULTIPLE WHILE ONLY ORDERING MODELS BY A PRIMARY METRIC. HOWEVER, MULTIPLE FACTORS—INCLUDING MULTIPLE METRICS—SHOULD BE INTEGRATED SUCH AS MODEL SIZE AND EFFICIENCY (?). THIS PRESENTS A CHALLENGE: HOW SHOULD ALL THESE BE VISUALIZED? IN OUR PROPOSED MODEL LISTING, WHILE THERE WOULD BE A DEFAULT RANKING, OUR VISUALIZATION WOULD SUPPORT RANKING BY ANY METRIC OR COMBINATIONS OF RANKINGS WHILE SHOWING THE INFLUENCE OF EACH AS IN THE PARALLEL AXIS PLOTS IN LINEUP (?).

Figure 10.8: A proof-of-concept visualization

that compares two models—a BERT and IR
SIMILARLY TO MANIFOLD (?). THE GOAL OF THE MODEL-CENTRIC VIEW IS TO
INTERVIEW system by visualizing which exam-
UNDERSTAND THE DIFFERENCES WEARDON PREDICTIONS—BETWEEN MODELS.
extending this to incorporate more than two
AN INITIAL APPROACH (FIGURE 10.8) COMPARES HOW PAIRS OF MODELS DIFFER,
models.
BUT COULD BE EXTENDED TO SUPPORT MULTIPLE MODELS BY ENHANCING THE
UNDERLYING SOFTWARE. ANOTHER APPROACH TO A MODEL-CENTRIC VIEWS
INSTEAD ATTEMPTS AT CHARACTERIZING THE PERFORMANCE ALONG PARTITIONS
OF THE DATA (?). FOR EXAMPLE, ? PROPOSE FACETED HISTOGRAMS AS A WAY TO
“UN-AGGREGATE” METRICS TO INSPECT QA ACCURACY BY FACTORS LIKE QUESTION
TYPE AND CONFIDENCE. IN ADDITION TO VISUALIZING WHICH INSTANCES PAIRS
OF MODELS AGREE AS IN MANIFOLD, WE PROPOSE INTEGRATING IRT PARAMETERS
TO BETTER IDENTIFY INSTANCES THAT HAVE A LARGER INFLUENCE OVER THE
OUTPUT RANKING.

A DOMINANT APPROACH TO INSPECTING MODELS EXAMINES THEIR PERFORMANCE ON INDIVIDUAL EXAMPLES OR PARTITIONS OF EXAMPLES SHARING A TRAIT (?). IN ADDITION TO DOMAIN-RELEVANT TRAITS, WE CAN GUIDE EXAMPLE SAMPLE SELECTION WITH IRT WHICH MAKES IT MORE LIKELY TO IDENTIFY IMPORTANT ISSUES OVER RANDOM SAMPLING (?) WHILE STILL COMBINING CAREFUL SPECIFICATION OF ERROR TYPES (?). SIMILARLY, VISUALIZATIONS CAN USE THE MULTIDIMENSIONAL IRT PARAMETERS FOR CLUSTERING WHICH CAN ALSO BE COMBINED WITH DOMAIN-SPECIFIC ATTRIBUTES LIKE QUESTION TYPE OR TOPIC.

Figure 10.7: We annotate SQuAD items by discriminability, difficulty, and IRT prediction errors. For example, one question with negative discriminability was classified as “WRONG” with the explanation that the annotated answer indicates it is *not* answerable, but the question actually *is* answerable. Items with negative discriminability or where IRT’s prediction is wrong have a much higher rate of annotation error (“Flawed” or “Wrong”). Using similar methodology, errors in datasets could be more rapidly identified.

¹⁵ The visualization doubles as my course project for CMSC734: Information Visualization.

A MAJOR ADVANTAGE OF OUR APPROACH—EMBEDDING THESE COMPARISONS IN LEADERBOARDS—is that model developers can compare the outputs of many models on individual or small sets of examples. While there are numerous tools for comparing models, error analysis, and related tasks, they face the headwinds of convincing practitioners that they are worthwhile time investments. Instead, we aim to provide real benefits of leaderboard-wide analysis without incurring up-front time commitment from practitioners. Although this does require commitment from task organizers, we believe that the trade-off of inherently improving the evaluation while providing practitioners with “shiny tools” is well worth it.

10.8 Related Work

DAD DRAWS TOGETHER TWO PRIMARY THREADS: WE USE IRT TO UNDERSTAND DATASETS, WHICH HAS BEEN APPLIED TO OTHER NLP TASKS, AND APPLY IT TO IMPROVING LEADERBOARDS. FINALLY, WE EXPLORE HOW THE INSIGHTS OF IRT CAN IMPROVE NOT JUST THE ANALYSIS OF TEST SETS BUT TO IMPROVE THE CONSTRUCTION OF TEST SETS.

IRT in NLP IRT IS GAINING TRACTION IN MACHINE LEARNING RESEARCH (??) WHERE AUTOMATED METRICS CAN BE MISLEADING (?): MACHINE TRANSLATION (?) AND CHATBOT EVALUATION (?). CONCURRENT WITH OUR WORK, ? COMPARE NLP TEST SETS WITH IRT. CLOSEST TO OUR WORK IN NLP IS ?, WHO RANK MACHINE TRANSLATION SUBJECTS AND COMPUTE CORRELATIONS WITH GOLD SCORES. SIMILARLY, ? USE IRT ON NON-LANGUAGE AI VIDEO GAME BENCHMARKS. JUST AS WE USE IRT TO IDENTIFY DIFFICULT OR EASY ITEMS, ? CREATE CHALLENGE SETS FOR TEXTUAL ENTAILMENT. WE TEST IRT AS A WAY TO GUIDE ANNOTATION, BUT IT CAN ALSO TRAIN NLP MODELS; FOR EXAMPLE, DEEP MODELS LEARN “EASY” EXAMPLES FASTER (?) AND MAINTAIN TEST ACCURACY WHEN TRAINING DATA ARE DOWN-SAMPLED (?).

Improving Leaderboards THE RISE NLP LEADERBOARDS HAS ENCOURAGED CRITICAL THOUGHT INTO IMPROVING THEM (?), IMPROVING EVALUATION MORE BROADLY (?), AND THOUGHTFUL CONSIDERATION OF THEIR INFLUENCE ON THE DIRECTION OF RESEARCH (??). DAD AIMS MAKE LEADERBOARD YARDSTICKS (?) MORE RELIABLE, INTERPRETABLE, AND PART OF CURATING THE BENCHMARK ITSELF. IN LINE WITH OUR RELIABILITY GOAL, JUST AS STATISTICAL TESTS SHOULD APPEAR IN PUBLICATIONS (??), THEY SHOULD BE “FREEBIES” FOR LEADERBOARD PARTICIPANTS (?). ALTERNATIVELY, ? POSIT THAT LEADERBOARDS COULD BE AUTOMATICALLY EXTRACTED FROM PUBLICATIONS. HOW TO AGGREGATE MULTI-TASK BENCHMARKS (???) AND MULTI-METRIC BENCHMARKS (?) IS AN OPEN QUESTION WHICH—ALTHOUGH WE DO NOT ADDRESS—is one use for IRT.

THIS WORK IMPLICITLY ARGUES THAT LEADERBOARDS SHOULD BE CONTIN-

UALLY UPDATED. AS A (STATIC) LEADERBOARD AGES, THE TASK(S) OVERFIT (?) WHICH—ALTHOUGH MITIGABLE (??)—IS BEST SOLVED BY CONTINUALLY COLLECTING NEW DATA (?). IDEALLY, NEW DATA SHOULD CHALLENGE MODELS THROUGH ADVERSARIAL COLLECTION (??) AND RELATED METHODS (?). HOWEVER, IF MAKING AN EASY LEADERBOARD MORE DIFFICULT IS POSSIBLE, THE LEADERBOARD HAS OUTLIVED ITS HELPFULNESS AND SHOULD BE RETIRED (?).

PART OF OUR WORK CENTERS ON ALTERNATE TASK EFFICACY RANKINGS, BUT THIS NAÏVELY ASSUMES THAT TASK EFFICACY IS THE SOLE USE CASE OF LEADERBOARDS. INDEED, FOCUSING SOLELY THESE FACTORS CAN MISLEAD THE PUBLIC (?) AND MAY NOT REFLECT HUMAN LANGUAGE CAPABILITIES (?). LEADERBOARDS ARE ALSO WELL POSITIONED TO PROVIDE INCENTIVE STRUCTURES FOR PARTICIPANTS TO PRIORITIZE FAIRNESS (?) AND EFFICIENCY (???) OR INCORPORATE TESTING OF SPECIFIC CAPABILITIES ([?DUNIETZ ET AL., 2020](#)). TO ENABLE THESE MORE NUANCED ANALYSES, LEADERBOARDS SHOULD ACCEPT RUNNABLE MODELS RATHER THAN STATIC PREDICTIONS (?).

Active Learning BEYOND IRT, THE ANALYSIS OF TRAINING DYNAMICS AND ACTIVE LEARNING (?) IS HELPFUL FOR ACTIVELY SAMPLING SPECIFIC ITEMS OR IDENTIFYING LOW-QUALITY ITEMS (?). FOR EXAMPLE, ? AND ? PROPOSE ALTERNATIVE TRAINING DYNAMICS-BASED METHODS FOR IDENTIFYING DIFFICULT ITEMS AS WELL ANNOTATION ERRORS. EVEN CLOSER TO GOALS, ? USE ACTIVE LEARNING TO BUILD A TEST COLLECTION. EXPLICITLY MEASURING HOW EFFICIENTLY EXAMPLES SEPARATE THE BEST SUBJECT FROM THE REST ALLOWS TEST SET CURATORS TO “FOCUS ON THE BUBBLE” (?), PRIORITIZING EXAMPLES MOST LIKELY TO REVEAL INTERESTING DISTINCTIONS BETWEEN SUBMITTED SYSTEMS.

Alternate Formulations IRT IS AN EXAMPLE OF CONVERGENT EVOLUTION OF MODELS THAT PREDICT SUBJECT ACTION GIVEN AN ITEM. IDEAL POINT MODELS ? CONSIDER HOW A LEGISLATOR (SUBJECT) WILL VOTE ON A BILL (ITEM) AND USE A SIMILAR MATHEMATICAL FORMULATION. THE VENERABLE ELO MODEL [GLICKMAN AND JONES \(1999\)](#) AND MODERN EXTENSIONS ? PREDICT WHETHER A PLAYER (SUBJECT) WILL DEFEAT AN OPPONENT (ITEM) WITH, AGAIN, A SIMILAR MATHEMATICAL MODEL. CERTAIN IRT MODELS CAN ALSO BE FORMULATED AS NONLINEAR MIXED MODELS ?, WHERE THE ITEM PARAMETERS ARE FIXED EFFECTS AND THE LATENT SUBJECT PARAMETERS ARE RANDOM EFFECTS. THIS ALLOWS FOR COMPARISONS BETWEEN IRT MODELS AND OTHER MIXED EFFECTS MODELS UNDER A CONSISTENT FRAMEWORK. **TODO, FIX ME!!!!** AND **TODO, FIX ME!!!!** CAN BE FORMULATED AS NONLINEAR MIXED MODELS, AND **TODO, FIX ME!!!!** CAN BE FORMULATED AS A DISCRETE MIXTURE MODEL OVER ITEMS. AS WE DISCUSS FURTHER IN THE NEXT SECTION, DAD’S APPLICATION OF IRT CAN FURTHER BE IMPROVED BY ADOPTING INTERPRETABLE EXTENSIONS OF THESE MODELS.

10.9 Conclusion

THIS PAPER ADVOCATES INCORPORATING DECADES OF RESEARCH IN CRAFTING EDUCATION TESTS TO IMPROVE HOW WE EVALUATE THE CAPABILITIES OF NLP MODELS. WE PROPOSE AND VALIDATE AN ALTERNATE IRT RANKING METHOD FOR LEADERBOARD EVALUATIONS, SHOW IT CAN GUIDE ANNOTATION, DETECT ANNOTATION ERROR, AND NATURALLY PARTITION EVALUATION DATA. JUST AS EDUCATORS MOVED FROM CLASSICAL TESTING TO IRT, THE NLP COMMUNITY SHOULD CONSIDER FUTURE EVALUATIONS WITH IRT.

Limitations

ALTHOUGH THERE IS MUCH TO GAIN THROUGH IRT EVALUATION, THERE ARE LIMITATIONS WHICH MAKE IT HARD TO IMPLEMENT. FIRST, IT REQUIRES ACCESS TO ITEM-LEVEL RESPONSES FOR ALL EXAMPLES FOR ALL SUBJECTS WHICH ARE OFTEN ONLY AVAILABLE TO ORGANIZERS. SECOND, ? NOTES THAT SAMPLING MUTUALLY EXCLUSIVE SUBSETS HAS DRAWBACKS—SAMPLES ARE NOT ENTIRELY INDEPENDENT. LASTLY, OUR WORK IS A PROOF OF CONCEPT USING SQUAD 2.0 AS A TEST BED AND OUR RESULTS MAY NOT GENERALIZE.

10.10 Future Work

WE SEE A FEW DIRECTIONS FOR FUTURE WORK. FIRST, THIS PAPER IS INTENDED TO VALIDATE IRT AND ITS USEFULNESS AS AN ACTIVE PART OF THE LEADERBOARD LIFECYCLE; THE NATURAL NEXT STEP IS TO IMPLEMENT IT IN A LEADERBOARD. SECOND, OUR IRT MODELS DO NOT INCORPORATE THE ITEM CONTENT (E.G., EXAMPLE TEXT) TO PREDICT RESPONSES, BUT IN PRINCIPLE COULD; BAYESIAN MODELS WITH METADATA (?) AND IDEAL POINT MODELS FROM POLITICAL SCIENCE (?) THAT INCORPORATE BILLS AND SPEECHES DO EXACTLY THIS (??). ANALOGOUSLY, IRT FOR LEADERBOARDS CAN AND SHOULD ALSO INCORPORATE TEXT FROM PASSAGES, QUESTIONS, AND ANSWERS TO BETTER MODEL WHAT MAKES QUESTIONS DIFFICULT. SUCH A MODEL CAN ALSO PREDICT WHICH CHARACTERISTICS WOULD CREATE DISCRIMINATING OR DIFFICULT ITEMS. LASTLY, MULTIDIMENSIONAL IRT MODELS TO EVALUATE MULTIPLE SKILLS COULD AID MULTI-TASK OR MULTI-METRIC LEADERBOARDS LIKE MRQA (?) AND DYNABOARD (?).

Acknowledgements

FOR THEIR WORK ON EARLY ITERATIONS OF LEADERBOARD VISUALIZATIONS, WE THANK JACOB BREMERMAN AND Wei Wei CHI. FOR INSIGHTFUL DISCUSSIONS AND IDEAS WE THANK SHI FENG, DOUG OARD, JOÃO SEDOC, MIKE WU, AND PATRICK LEWIS. WE THANK PETER RANKEL FOR RECOMMENDATIONS ON STATISTICAL TESTING METHODS. FOR DISCUSSION AND FEEDBACK ON VISUALIZATIONS, WE THANK LEO ZHICHENG LIU, CALVIN BAO, AND CLASSMATES IN UMD'S FALL 2020 "INFORMATION VISUALIZATION" COURSE. FOR SUGGESTIONS ON TOPIC

Discriminability: -9.63 **Difficulty:** -0.479 **Feasibility:** 0.614 **Mean Exact Match:** 0.472

Wikipedia Page: [Economic inequality](#) **Question ID:** 572a1c943f37b319004786e3

Question: Why did the demand for rentals decrease?

Official Answer: [demand for higher quality housing](#)

Context: A number of researchers (David Rodda, Jacob Vigdor, and Janna Matlack), argue that a shortage of affordable housing – at least in the US – is caused in part by income inequality. David Rodda noted that from 1984 and 1991, the number of quality rental units decreased as the [demand for higher quality housing increased](#) (Rhoda 1994:148). Through gentrification of older neighbourhoods, for example, in East New York, rental prices increased rapidly as landlords found new residents willing to pay higher market rate for housing and left lower income families without rental units. The ad valorem property tax policy combined with rising prices made it difficult or impossible for low income residents to keep pace.

Figure 10.9: The example from SQuAD WITH THE LOWEST DISCRIMINABILITY. SURPRISINGLY, IT HAD A NEGATIVE DISCRIMINABILITY,

[042301C](#) THAT THE LESS SKILLED A SUBJECT IS, THE MORE LIKELY THEIR RESPONSE IS TO BE CORRECT.

Discriminability: 3.24 **Difficulty:** 3.86 **Feasibility:** 0 **Mean Exact Match:** 0

Wikipedia Page: [Computational Complexity Theory](#) **Question ID:** 56e1b00ce3433e140042301C

Question: In the determination of complexity classes, what are two examples of types of [Turing machines](#)?

Official Answer: [probabilistic Turing machines, non-deterministic Turing machines](#)

Context: Many types of Turing machines are used to define complexity classes, such as deterministic Turing machines, probabilistic Turing machines, non-deterministic Turing machines, quantum Turing machines, symmetric Turing machines and alternating Turing machines. They are all equally powerful in principle, but when resources (such as time or space) are bounded, some of these may be more powerful than others.

MODELING, WE THANK PHILIP RESNIK AND MARIA ANTONIAK. FOR FEEDBACK ON PRIOR VERSIONS OF THIS PAPER, WE THANK OUR ANONYMOUS ACL REVIEWERS AND MEMBERS OF THE UMD CLIP LAB. BOYD-GABER AND RODRIGUEZ'S WORK AT UMD WERE SUPPORTED BY NSF GRANT IIS-1822494. THE VIEWS AND CONCLUSIONS CONTAINED HEREIN ARE THOSE OF THE AUTHORS AND SHOULD NOT BE INTERPRETED AS NECESSARILY REPRESENTING THE OFFICIAL POLICIES, EITHER EXPRESSED OR IMPLIED, OF THE SPONSOR. THE U.S. GOVERNMENT IS AUTHORIZED TO REPRODUCE AND DISTRIBUTE REPRINTS FOR GOVERNMENTAL PURPOSES NOTWITHSTANDING ANY COPYRIGHT ANNOTATION THEREIN.

Figure 10.10: This question is regarded as infeasible by the IRT model. Upon further inspection, the answer omits five acceptable answers, but more importantly does not permit all combinations of Turing machines.

10.11 SQuAD Item Examples

FIGURES 10.9, 10.10, 10.11, AND 10.12 SHOW PREVIOUSLY DISCUSSED SQuAD EXAMPLES (§10.6) IN FULL. ON THE SAME PAGE, WE PROVIDE A WEB INTERFACE FOR INSPECTING THE PARAMETERS OF THE IRT MODELS. FIGURE 10.13 SHOWS THE FEASIBILITY DISTRIBUTION CORRESPONDING TO FIGURE 10.1.

10.12 Logistic Regression Features

THE LINEAR MODEL (§10.5) INCLUDES FEATURES BASED ON ITEM IDS, SUBJECT IDS, TEXTUAL FEATURES OF THE QUESTION, CONTEXT, AND ANSWER, AND TOPIC MODEL FEATURES. TABLE 10.1 LISTS THE FEATURE NAMES FROM FIGURE 10.3 WITH DESCRIPTIONS OF EACH. WHEN IRT FEATURES OR THE STATISTICS FEATURES ARE USED, THEY INCLUDE INTERACTION TERMS WITH THEMSELVES.

Discriminability: 2.1 **Difficulty:** 2.38 **Feasibility:** 0.995 **Mean Exact Match:** 0.00621 **Mean F1:** 0.546

Wikipedia Page: European Union Law **Question ID:** 57268f2bf1498d1400e8e3c4

Question: What reform was attempted following the Nice Treaty?

Official Answer: an attempt to reform the constitutional law of the European Union and make it more transparent

Context: Following the Nice Treaty, there was an attempt to reform the constitutional law of the European Union and make it more transparent; this would have also produced a single constitutional document. However, as a result of the referendum in France and the referendum in the Netherlands, the 2004 Treaty establishing a Constitution for Europe never came into force. Instead, the Lisbon Treaty was enacted. Its substance was very similar to the proposed constitutional treaty, but it was formally an amending treaty, and – though it significantly altered the existing treaties – it did not completely replace them.

Figure 10.11: This example shows that the answer span is likely too large, causing models to fail in both SQuAD's EXACT MATCH AND F1 METRICS.

Discriminability: 8.01 **Difficulty:** -1.41 **Feasibility:** 0.939 **Mean Exact Match:** 0.64 **Mean F1:** 0.667

Wikipedia Page: Normas **Question ID:** 56de10b44396321400ee2595

Question: Who did the Normans team up with in Anatolia?

Official Answer: Turkish forces

Context: Some Normans joined Turkish forces to aid in the destruction of the Armenians vassal-states of Sassoun and Taron in far eastern Anatolia. Later, many took up service with the Armenian state further south in Cilicia and the Taurus Mountains. A Norman named Oursel led a force of "Franks" into the upper Euphrates valley in northern Syria....

Figure 10.12: This highly discriminative question succeeds because there are many plausible answers. For example, although only "Turkish forces" is correct, some models answer "the Armenian state."

10.13 IRT Model Type Correlation

ALTHOUGH EACH IRT MODEL DIFFERS IN EXPRESSIVENESS, THEY SHOULD—IN GENERAL—PRODUCE SIMILAR RESULTS. THIS IS CONFIRMED BY COMPUTING THE KENDALL'S RANK CORRELATION BETWEEN THE SUBJECT ABILITIES AND ITEM DIFFICULTIES (TABLE 10.2).

10.14 Ranking Stability Experiments

HERE WE PROVIDE FURTHER DETAILS FOR THE RANKING STABILITY EXPERIMENTS (§10.5). FIRST, WE FILTER FROM THE 161 SUBJECTS THAT HAVE DEVELOPMENT SET SCORES TO THE 115 THAT ALSO HAVE TEST SET SCORES.¹⁶ IN OUR SIMULATION, WE RUN 10 TRIALS FOR EVERY SAMPLE SIZE; SAMPLE SIZE BEGINS AT 100 AND WITH STEPS OF 100. IN ADDITION TO THESE, WE ALSO RUN TRIALS FOR SAMPLE SIZES 25, 50, AND 75. SINCE EACH SAMPLE CAN BE NO LARGER THAN HALF THE DATASET, WE STOP AT HALF THE DATASET.

¹⁶ The SQuAD ORGANIZERS CURATE THE TEST SET SUBJECTS TO AVOID OVERFIT, GARBAGE, OR DUPLICATE SUBMISSIONS.

Development and Test Set Correlations

TABLE 10.3 USES A TODO, FIX ME!!! MODEL SINCE WE NOTICED THAT IN COMPARISON TODO, FIX ME!!! OVERFIT THE DATA, YIELDING WORSE RESULTS. THE CORRELATIONS WITH THE FULL DATA ARE ALL STRONG, BUT NOT THE SAME. WE CONCLUDE THAT—at least on SQuAD—IRT RANKINGS ARE MODESTLY MORE RELIABLE THAN CLASSICAL RANKINGS.

Figure 10.13: The feasibility parameter λ of our IRT model represents the probability that	
Example is unsolvable. For example, annotation error could lead to an example always	
All being scored incorrectly. All the features show	
good the model is. In square values for difficulty, discriminability, feasibility, THE 5% PERCENTILE, $\lambda \leq .698$ FOR THE 7.5%,	
AND $\lambda < .931$ IN THE 10% PERCENTILE.	
Item ID	The item's ID
Subject ID	The subject's ID
Question	Question words
Context	Context words
Stats	Question & context lengths; answerability, answer position & length; difficulty from ?
Subject & Item ID	Item and Subject ID
Topics 1K	Topic weights of question words
Title	Wikipedia page title words
Baseline	No features, majority class baseline

Table 10.1: The linear model integrates a variety of features to determine which are most predictive of a subject responding correctly to an item.

Statistical Significance of Difference in Kendall Tau Coefficients

WHILE FIGURE 10.4 SHOWS A CONSISTENT DIFFERENCE IN CORRELATION BETWEEN RANKING METHODS, IT IS UNCLEAR WHETHER THIS DIFFERENCE IS STATISTICALLY SIGNIFICANT. WE ESTIMATE THE STATISTICAL SIGNIFICANCE OF THE DIFFERENCE THROUGH BOOTSTRAP SAMPLING (?).

SINCE THE NULL CASE IS NO DIFFERENCE IN CORRELATION COEFFICIENTS, WE SEEK A SYMMETRIC SAMPLING DISTRIBUTION CENTERED AT ZERO THAT REPRESENTS A REALISTIC DENSITY FUNCTION. EACH RANKING STABILITY EXPERIMENT¹⁷ TRIAL RESULTS IN TWO LISTS OF NUMBER PAIRS. THE LISTS CORRESPOND TO SUBJECT SCORES ON TWO DATASETS;¹⁸ EACH NUMBER PAIR IS THE SUBJECT'S ACCURACY AND IRT SCORE. TO CREATE THE BOOTSTRAP DISTRIBUTION, WE (1) SAMPLE WITH REPLACEMENT PAIRS FROM ONE LIST, (2) COMPUTE THE CORRELATION BETWEEN THE RESAMPLED RANKING AND UNUSED RANKING WHEN USING ACCURACY VERSUS IRT SCORE, AND (3) COMPUTE AND STORE THE IRT CORRELATION SCORE MINUS THE ACCURACY CORRELATION SCORE. WE REPEAT THIS PROCESS 1000 TIMES FOR EACH OF THE 10 TRIALS IN THE ORIGINAL EXPERIMENT AND AGGREGATE ALL THE DIFFERENCES TO BUILD THE BOOTSTRAP DISTRIBUTION. FOR EACH SAMPLE SIZE WE COMPUTE THE EMPIRICAL P-VALUE ON EACH TRIAL WHICH WE SHOW IN BOX AND WHISKER PLOTS (FIGURE 10.14).

10.15 The IRT Statistical Test

THE IRT TEST DIFFERS IN TWO SUBSTANTIAL WAYS FROM OTHER TESTS: (1) IT DOES NOT ASSUME THAT ITEMS ARE EQUALLY INFORMATIVE AND (2) IT DOES ASSUME THAT THE INFORMATIVENESS OF ITEMS IS A FUNCTION OF THE SUBJECT'S SKILL θ_j . IN THE LITERATURE, THIS IS CLOSELY CONNECTED TO RELIABILITY (?) AND EACH ITEM PROVIDES INFORMATION ABOUT THE LOCATION OF θ_j ; AS WE ACCUMULATE MORE EVIDENCE FOR THE LOCATION OF θ_j , THE ERROR OF ESTIMA-

¹⁷ One experiment for development sample to development sample and one for development sample to test set.

¹⁸ In the first experiment, development set samples; in the second, a development set sample and the full test set.

Ability	TODO, FIX ME!!!!	TODO, FIX ME!!!!	TODO, FIX ME!!!!
TODO, FIX ME!!!!	1.00	0.947	0.895
TODO, FIX ME!!!!	0.947	1.00	0.907
TODO, FIX ME!!!!	0.895	0.907	1.00

	EM _{dev}	EM _{test}	Ability _{dev}	Ability _{test}
EM _{dev}	1.00	0.953	0.954	0.931
EM _{test}	0.953	1.00	0.944	0.947
Ability _{dev}	0.954	0.944	1.00	0.950
Ability _{test}	0.931	0.947	0.950	1.00

TION DECREASES. IT IS A WELL KNOWN RESULT IN IRT THAT STANDARD ERROR OF ESTIMATE (SEE) $\sigma(\hat{\theta}|\theta)$ VARIES WITH RESPECT TO THE AGENT LOCATION PARAMETER θ (? , p. 30) AND IS CONNECTED TO THE FISHER INFORMATION

$$I_i(\theta) = \frac{(p'_i)^2}{p_i(1-p_i)} \quad (10.4)$$

OF EACH ITEM. FOR A 2PL MODEL, INFORMATION

$$I_i(\theta) = \gamma^2 p_i(1-p_i) \quad (10.5)$$

IS MAXIMIZED WHEN $p_i = (1 - p_i)$. SINCE FISHER INFORMATION IS ADDITIVE, THE INFORMATION OF THE EVALUATION SET IS MAXIMAL WHEN ITEMS HAVE A 50% CHANCE OF BEING RESPONDED TO CORRECTLY. AS DERIVED BY ?, p. 102, THE STANDARD ERROR OF ESTIMATION

$$\text{SEE}(\theta) = \sqrt{\frac{1}{\sum_i I_i(\theta)}}. \quad (10.6)$$

IS COMPUTED BY ACCUMULATING THE INFORMATION GAINED FROM EACH ITEM. GIVEN TWO SUBJECTS X AND Y , ONE CAN USE THE PROBABILITY DISTRIBUTION OF SCORE DIFFERENCES

$$N(\theta_Y - \theta_X, \text{SEE}(\theta_X)^2 + \text{SEE}(\theta_Y)^2) \quad (10.7)$$

TO COMPUTE THE PROBABILITY THAT THE DIFFERENCE IN SKILL IS GREATER THAN TWO STANDARD ERRORS WHICH CORRESPONDS TO AN $\alpha \leq .05$ SIGNIFICANCE LEVEL.

10.16 Multidimensional IRT Clustering

WHILE WE ACHIEVE STRONG HELD-OUT ACCURACY WITH 10 DIMENSIONAL IRT (TODO, FIX ME!!!!), WE HAD LIMITED SUCCESS IN INTERPRETING PARAMETERS. WE USE TSNE¹⁹ PLOTS OVERLAYED WITH FEATURES LIKE ITEM ACCURACY, THE QUESTION'S WIKIPEDIA PAGE, IF THE QUESTION WAS ANSWERABLE, LENGTH OF QUESTIONS, AND TOPIC MODEL WEIGHTS. OF THESE, ITEM ACCURACY AND ANSWERABILITY SHOWED THE MOST OBVIOUS PATTERNS (FIGURE 10.15).

Table 10.2: Table entries are Kendall's τ rank correlation of IRT subject ability between rows and columns. Generally, the models agree on the ranking with the TODO, FIX ME!!!! and TODO, FIX ME!!!! having the strongest correlation.

Table 10.3: Entries are Kendall's rank correlation between rows and columns. Scores are SQUAD EXACT MATCH (EM) AND TODO, FIX ME!!!! ABILITY.

¹⁹We use openTSNE (?) with default parameters.

Figure 10.15: In SQuAD, tSNE shows a relationship between mean exact match (item accuracy) and answerability with respect to multidimensional difficulty (shared task MRQA (?)). However, instead of using 10 dimensions we use 6 to match the number of development set tasks in MRQA. Although questions in NarrativeQA standout (Figure 10.16), there is not a discernible pattern amongst the other tasks. We leave more sophisticated methods for making multidimensional IRT models interpretable to future work.

10.17 Reproducibility Checklist

HERE WE PROVIDE REPRODUCIBILITY DETAILS TO COMPLEMENT OUR SOURCE CODE ([HTTPS://IRT.PEDRO.AI](https://irt.pedro.ai)).

Software and Parameters

ALL IRT MODELS ARE IMPLEMENTED IN PYTORCH (?) AND PYRO (?). LINEAR MODELS ARE TRAINED WITH VOWPAL WABBIT (?). THE TOPIC MODEL THAT GENERATES FEATURES FOR THE LINEAR MODEL USES MALLET (?).

THE NUMBER OF IRT MODEL PARAMETERS IS PROPORTIONAL TO THE NUMBER OF SUBJECTS m AND THE NUMBER OF ITEMS n . THE TODO, FIX ME!!! HAS ONE PARAMETER PER SUBJECT AND ONE PER ITEM. THE TODO, FIX ME!!! HAS ONE PARAMETER PER SUBJECT AND TWO PER ITEM. THE TODO, FIX ME!!! HAS ONE PARAMETER PER SUBJECT AND THREE PER ITEM. THE TODO, FIX ME!!! HAS TEN PARAMETERS PER SUBJECT AND THIRTY PER ITEM.

Hyperparameters

WE DID NOT INVEST SIGNIFICANT EFFORT IN HYPER-PARAMETER TUNING THE IRT MODELS AND INSTEAD USED THE DEFAULTS IN THE PY-IRT SOFTWARE²⁰ PROVIDED BY ?. THE TODO, FIX ME!!!, TODO, FIX ME!!!, AND TODO, FIX ME!!! MODELS WERE TRAINED FOR 1000 EPOCHS WITH NO EARLY STOPPING CONDITIONS AND A LEARNING RATE OF 0.1 WITH ADAM (?). THE TODO, FIX ME!!! MODEL WAS TRAINED FOR 2500 EPOCHS AND USED 10 DIMENSIONS.

IN THE LINEAR MODEL, WE USED A HYPEROPT-BASED (?) TOOL PROVIDED BY VOWPAL WABBIT²¹ FOR HYPER PARAMETER SEARCH. FOR EACH LM, THE TOOL SPENT 20 ITERATIONS OPTIMIZING THE LEARNING RATE, L2 REGULARIZATION, AND NUMBER OF BITS AGAINST THE LOGUNIFORM SAMPLING, L2 REGULARIZATION FROM $1e - 8$ TO 1, AND BITS FROM 20 TO 23 AS CATEGORICAL VARIABLES.

THE TOPIC MODEL THAT GENERATED FEATURES FOR THE LINEAR MODEL USED

Figure 10.14: P-values of the rank correlation difference for each sample size and trial in Figure 10.4. The inherent noise in dev set sampling makes inferring significance difficult (left); test set driven results (right) are more significant.

²⁰ github.com/jplalor/py-irt

²¹ github.com/VowpalWabbit/vowpal_wabbit

Figure 10.16: In mRQA, tSNE shows a relationship between whether the task is NarrativeQA or MRCQA with respect to difficulty and discriminability. SET HYPER PARAMETERS.²² SPECIFICALLY, WE USED AN OPTIMIZATION INTERVAL OF 10, REMOVED STOP WORDS, TRAINED FOR 1000 ITERATIONS, AND USED A DOCUMENT-TOPIC THRESHOLD OF 0.05. EACH DOCUMENT WAS COMPRISED OF THE WIKIPEDIA PAGE TITLE AND THE QUESTION TEXT.

²² mallet.cs.umass.edu/topics.php

Computational Resources

THE MAJORITY OF EXPERIMENTS WERE CONDUCTED ON A SINGLE WORKSTATION WITH AN INTEL i7-7700K CPU, 47GB OF RAM, AND AN NVIDIA 1080Ti. THE AVERAGE RUNTIME FOR THE TODO, FIX ME!!! MODEL ON CPU IS 113 SECONDS WITH A STANDARD DEVIATION OF 2.31 OVER 5 TRIALS. THE AVERAGE RUNTIME OF THE TODO, FIX ME!!! MODEL ON GPU IS 110 SECONDS WITH A STANDARD DEVIATION OF 0.5 OVER 5 TRIALS.

SINCE EACH RANKING STABILITY EXPERIMENT REQUIRED (§10.5) RE-TRAINING AN TODO, FIX ME!!! MODEL ON EACH SUBSET, WE PARALLELIZED THIS EXPERIMENT ON A CPU CLUSTER WHERE EACH TRIAL RECEIVED TWO CPU CORES AND 16GB OF RAM. IN TOTAL, THIS INCLUDED 520 TRIALS WHICH CORRESPONDS TO TWICE THAT MANY TRAINED IRT MODELS SINCE ONE MODEL IS TRAINED ON EACH SUBSET OF THE DATA.

10.18 Adversarial Questions

WE FIRST LEARNED ABOUT TERRY WINOGRAD IN CHAPTER

10.19 Human–Computer Cooperation

11

FINDING AND DEALING WITH BAD QUESTIONS

THIS CHAPTER DISCUSSES HOW QUESTIONS CAN BE IMPERFECT AND HOW BOTH HUMANS AND COMPUTERS DEAL WITH THOSE IMPERFECT QUESTIONS. NAVIGATING AN IMPERFECT QUESTION REQUIRES MORE SKILL THAN ANSWERING A PERFECTLY POSED ONE. NOT ONLY MUST THE ANSWERER HAVE THE REQUISITE INFORMATION ABOUT THE CORRECT ANSWER, THEY MUST EXCLUDE INCORRECT INTERPRETATIONS AND NAVIGATE HOW TO DECIDE WHAT THE ASKER ACTUALLY NEEDS.

THIS CHAPTER DISCUSSES HOW QUESTIONS CAN BE IMPERFECT AND HOW BOTH HUMANS AND COMPUTERS DEAL WITH THOSE IMPERFECT QUESTIONS.

Part III

Question Answering Future

12

FULFILING TURING'S GAMESHOW

THE YEAR IS 2025, AND THERE'S A NEW GAMESHOW THAT NOT ONLY SHOWCASES THE MOST ADVANCED AI AVAILABLE BUT ALSO KEEPS THE PUBLIC INFORMED ABOUT THE LIMITATIONS AND STRUGGLES OF CURRENT TECHNOLOGY. THIS CHAPTER OUTLINES THE TEN SEASONS OF THE SHOW AND HOW IT TRACKS THE DEVELOPMENT OF MACHINE INTELLIGENCE, LEADING TO PASSING TURING TEST.

13

WHAT AI DYSTOPIAS YOU SHOULD BE AFRAID OF

WHILE THE BOOK BEGAN WITH HOW QUESTION ANSWERING IN MYTH HELPED THE ANCIENTS GRAPPLE WITH A CHANGING AND UNCERTAIN WORLD, TODAY'S FICTIONS REPRESENT OUR CONTEMPORARY ATTEMPTS TO UNDERSTAND THE ADVENT OF INTELLIGENT COMPUTERS. THIS CHAPTER REVIEWS HUMAN-COMPUTER QUESTION ANSWERING IN 2001, STAR TREK, THE TERMINATOR, BLADE RUNNER, FUTURAMA, AND WHAT THESE DEPICTIONS REVEAL ABOUT BOTH HUMAN CONCEPTIONS OF ARTIFICIAL INTELLIGENCE AND HOW IT MIGHT SHAPE FUTURE DEPLOYMENTS OF QUESTION ANSWERING.

TOPICS IN CHAPTER:

1. **COMPUTER, SPECULATE.** A REVIEW OF THE ROLE OF QA IN SCIENCE FICTION: VOIGT-KAMPF, STAR TREK, ETC.
2. **P(DOOM).** BEYOND THE IMAGINING OF SCIENCE FICTION, WHAT ARE THE ACTUAL DOWNSIDES TO THE WIDESPREAD AVAILABILITY AND DEPLOYMENT OF 'INTELLIGENT' AI? THIS CHAPTER TALKS ABOUT AI'S ABILITY TO AMPLIFY EXISTING NEGATIVE HUMAN TENDENCIES IN RESPONDING TO QUESTIONS AND HOW—WHILE DANGEROUS AND WORTHY OF MITIGATION—DON'T DEMAND US TO HALT THE DEVELOPMENT OF MODERN AI.

THE PREVIOUS STORY SAW OUR PROTAGONIST START HER CAREER, AND OUR ESSAYS WILL ALSO BEGIN WITH FOUNDATIONAL MATERIAL: WHAT A LAY USER NEEDS TO UNDERSTAND WHAT'S HAPPENING IN THE STORY AND HOW REALISTIC IT IS. WE'LL SAVE THE MORE FANCIFUL AND COMPLICATED STUFF FOR LATER CHAPTERS.

FOR THE MOMENT, WE FOCUS ON A FACET OF ARTIFICIAL INTELLIGENCE CALLED MACHINE LEARNING. MACHINE LEARNING IS A SMALL SUBSET OF ARTIFICIAL INTELLIGENCE,¹ BUT UNLIKE MUCH OF THE FANCIFUL CLAIMS OF ARTIFICIAL INTELLIGENCE PROPHETS, IT ACTUALLY WORKS NOW.

WE FIRST INTRODUCE A KEY COMPONENT OF MACHINE LEARNING—OBJECTIVE FUNCTIONS—that define why systems act the way they do. WE THEN SHOW HOW BAD HUMAN BEHAVIOR CAN BE MIMICKED BY THESE ALGORITHMS (IT'S ALREADY HAPPENING!). FINALLY, WE CLOSE THE CHAPTER WITH A CALL TO ACTION: KEEP YOUR COMPUTERS SAFE!

¹ Research is organized around conferences; these conferences often form a community that is protective of its turf. While there is substantial overlap between the communities (e.g., the Neural Information Processing Systems conference), many machine learning researchers (e.g., the International Conference on Machine Learning) prefer to remain distinct from "pure" artificial intelligence.

13.1 Learning by Doing: Objective Functions

LET US BEGIN WITH AN ESSENTIAL AND UNDER-APPRECIATED EXAMPLE OF MACHINE LEARNING THAT SIMPLY WORKS AND MAKES MODERN LIFE LIVABLE: SPAM FILTERS. AN E-MAIL IS EITHER SPAM OR NOT; COMPUTERS THINK IN ONES AND ZEROS, SO LET'S CALL SPAM A ONE AND NOT SPAM A ZERO.

AN ALGORITHM DOES NOT JUST SAY “YES” OR “NO” TO WHETHER AN E-MAIL IS SPAM. IT TYPICALLY GIVES A SCORE FOR EACH E-MAIL: THE HIGHER THE SCORE, THE MORE CONFIDENT IT IS THE E-MAIL IS SPAM. FOR THE MOMENT, LET'S ASSUME THAT ALL OF THE NUMBERS ARE BETWEEN ZERO AND ONE. THIS IS CONVENIENT BECAUSE WE CAN THINK OF THIS AS A PROBABILITY: A ZERO MEANS THAT THERE IS NO CHANCE THE E-MAIL IS SPAM AND A ONE IS COMPLETE CONFIDENCE THAT IT IS SPAM.

A PERFECT SCORE WOULD BE IF ALL SPAM E-MAILS GOT A 1.0 AND ALL OF THE GOOD E-MAILS GOT 0.0. BOTH PERFECTION AND ABSOLUTE CERTAINTY ARE UNATTAINABLE GOALS, SO WE NEED TO DEAL WITH BOTH ERRORS AND UNCERTAINTY. SO WE SUM ALL OF THE SPAM DOCUMENTS—HOW FAR THE SCORES ARE FROM 1.0—and we sum all of the good e-mails—how far the scores are from 0.0. This sum is our *objective function*: a perfect score is 0.0—no mistakes—and the worst score is the number of e-mails.

MACHINE LEARNING ALGORITHMS ARE DEFINED BY PARAMETERS. FOR EXAMPLE, A PARAMETER MIGHT DEFINE WHAT THE ALGORITHM DOES WHEN IT SEES THE WORD “VIAGRA” OR THE PHRASE “I AM A NIGERIAN PRINCE”. THESE PARAMETERS EFFECTIVELY DEFINE THE ALGORITHM; MACHINE LEARNING ALGORITHMS LEARN HOW TO SET THESE PARAMETERS FROM DATA, A PROCESS THAT CAN TAKE MULTIPLE NAMES,² BUT WE'LL CALL IT “OPTIMIZATION”.

Optimization

OPTIMIZATION IS THE PROCESS WHERE MACHINE LEARNING USES DATA TO FIND THE PARAMETERS THAT GIVE THE BEST SCORE ON THE OBJECTIVE FUNCTION. TYPICALLY, THE PROCESS STARTS WITH A RANDOM SETTING OF THE PARAMETERS. THIS WILL DO HORRIBLY, BUT IT PROVIDES A PLACE TO START: THE EARLY IMPROVEMENTS WILL BE EASY.

THE LEARNING PROCESS WILL SLOWLY CHANGE THE PARAMETERS TO EVER SO SLIGHTLY IMPROVE THE OBJECTIVE FUNCTION.³ THE LEARNING PROCESS GOES DOCUMENT BY DOCUMENT TO SLOWLY IMPROVE THE OBJECTIVE FUNCTION. EVENTUALLY, MODERN OPTIMIZATION TECHNIQUES FIND A SETTING OF THE PARAMETERS THAT WORK WELL FOR THIS PROBLEM.⁴

LET'S TAKE A CONCRETE EXAMPLE. LET'S SAY THAT THIS IS THE FIRST E-MAIL THE ALGORITHM HAS SEEN WITH THE PHRASE “SPECIAL OFFER” AND THE ALGORITHM SAYS THAT IT IS SPAM WITH SCORE 0.7. THIS IS A SPAM E-MAIL, SO THE ALGORITHM COULD DO BETTER IN TERMS OF ITS OBJECTIVE FUNCTION IF THE PARAMETER ASSOCIATED WITH THE PHRASE “SPECIAL OFFER” WERE MORE STRONGLY ASSOCIATED WITH SPAM. SO OUR OPTIMIZATION WILL PUSH THE

² Some of the alternate names are field specific; for example, in statistics this is called inference and the objective function is typically called the likelihood. However, the mechanics are similar. For some applications, it's called a “loss function”, but we'll stick with the more general term.

³ Mathematically, this is usually by looking at the derivative (if you remember that from your calculus class) of the objective function with respect to the parameters.

⁴ The form of the objective function determines whether this answer will be the best possible or merely a “good” solution.

PARAMETER A LITTLE MORE IN THAT DIRECTION SO THAT THE SCORE IS 0.75 THE NEXT TIME IT SEES A SIMILAR DOCUMENT. THIS CONTINUES UNTIL THE PARAMETERS CANNOT GET ANY BETTER.

ALL OF THIS SEEKS FAIRLY STRAIGHTFORWARD, BUT THERE IS SOME ART TO CREATING THESE SYSTEMS. DEFINING THE PARAMETERS REQUIRES A BIT OF EXPERTISE: EITHER HAND-CRAFTING PARAMETERS THAT FIT THE PROBLEM WELL OR USING AUTOMATIC APPROACHES THAT LEARN EFFECTIVE REPRESENTATIONS. DOING THIS WELL REQUIRES QUITE A BIT OF WORK, AND OFTEN YOUR FIRST (AND SECOND, ETC.) ATTEMPT WILL USUALLY FAIL.

Is this Artificial Intelligence?

A REASONABLE QUESTION IS WHETHER THIS IS *INTELLIGENCE*. BEFORE WE GET INTO THE NUANCE OF THE *QUESTION*, THE ANSWER—GIVEN THE SYSTEMS WE HAVE DISCUSSED—is definitely **NO**. BUT BECAUSE WE’LL DISCUSS THINGS THAT ARE ARTIFICIAL INTELLIGENCE LATER IN THE BOOK, IT’S USEFUL TO DISCUSS WHAT IT MEANS FOR A SYSTEM TO BE INTELLIGENT. THIS WAY, WE CAN AT LEAST SEE WHY THESE SYSTEMS ARE NOT INTELLIGENT.

IN SOME WAYS, IMPRESSIVE ACHIEVEMENTS IN TECHNOLOGY ARE A MOVING TARGET. AFTER A FEW DECADES OF LIVING WITH MACHINES THAT DO A FORMERLY IMPRESSIVE TASK, WE CEASE TO BE IMPRESSED BY IT AND ITS COMPARISON TO HUMANS.⁵ THE LUSTER OF ADDING MACHINES HAS FADED INTO BEIGE QUOTIDIAN BORINGNESS THROUGH THE TWENTIETH CENTURY.

EVEN IN MY OWN LIFE, MY EXCITEMENT ABOUT HALFWAY DECENT SPEECH RECOGNITION—TAKING A STREAM OF SPEECH AND TRANSCRIBING IT INTO WORDS ON A SCREEN—HAS GIVEN WAY TO ANNOYANCE AT MY PHONES’ (INCREASINGLY RARE) MISTAKES. BOTH OF THESE WERE CONSIDERED TASKS THAT ONLY HUMANS COULD DO⁶ BUT NOW ARE SO UBIQUITOUS THAT HALF OF HUMANITY IS JOINED AT THE HIP TO DEVICES THAT DO BOTH THESE THINGS.

A SKEPTIC WOULD SAY THAT MACHINE LEARNING IS A MERE MECHANICAL MATHEMATICAL EXERCISE. THE ALGORITHM THAT DECIDES WHETHER AN E-MAIL IS SPAM OR NOT DOESN’T ACTUALLY UNDERSTAND WHAT IT’S READING, AND THE PARAMETERS WERE DEFINED BY A HUMAN. WHILE THE VALUES OF THE PARAMETERS COME FROM DATA, THERE’S NO REAL LEARNING GOING ON.

A PROONENT WOULD COUNTER THAT INTELLIGENCE IS ACTUALLY A COMBINATION OF MANY SMALLER PROCESSES. TO UNDERSTAND AN E-MAIL, WE NEED TO UNDERSTAND PARTS OF SPEECH, SYNTAX, DISCOURSE, AND PRAGMATICS. EACH OF THESE CAN BE THOUGHT OF AS INDIVIDUAL PROBLEMS LIKE SPAM CLASSIFICATION. EACH ONE IS MECHANISTIC, BUT TOGETHER THEY FORM A PROCESS THAT CAN BE VIEWED AS INTELLIGENT.

FOR THE MOMENT, WE WILL LEAVE THE DEBATE THERE. CHAPTER ?? TAKES UP THE *QUESTION* ABOUT WHAT IT MEANS TO BE INTELLIGENT AND HOW TO MEASURE THE INTELLIGENCE OF ALGORITHMS.

NEVERTHELESS, MACHINE LEARNING HAS EARNED THE REPUTATION OF BEING

⁵ This is not just for intelligence; *The Ballad of John Henry* celebrates a competition between a steam drill and “a man who ain’t nothing but a man...with a hammer in [his] hand.” Today, such comparisons and competitions seem silly.

⁶ The term “computer” originally referred to specially trained humans proficient at high-stakes, accurate calculations. ? documents how labor shortages during World War II increasingly a female profession. However, like the protagonists of *Hidden Figures* by Margot Lee Shetterly, the women remain in the background.

?reference? ??

“ARTIFICIAL INTELLIGENCE THAT WORKS”. ITS MECHANISTIC DEFINITION IS A VIRTUE: IT IS EASY TO SEE WHAT WORKS AND WHAT DOES NOT. JUST AS THE OBJECTIVE FUNCTION ALLOWS THE ALGORITHM TO TUNE INDIVIDUAL PARAMETERS TO IMPROVE A SINGLE TASK, THE CLEAR DEFINITION ALLOWS RESEARCHERS TO TWEAK PROBLEMS, MODELS, AND DATA TO GET BETTER AT IMPORTANT TASKS.

MACHINE LEARNING HAS DONE WELL AT MANY TASKS: IDENTIFYING OBJECTS IN A PICTURE (?), PLAYING GAMES (?), ANSWERING QUESTIONS (FERRUCCI ET AL., 2010), AND DRIVING CARS (?). THESE TASKS ARE IMPRESSIVE, BUT EACH TASK IS A COMPOSITION OF SMALLER OBJECTIVE FUNCTIONS. AS WE MOVE INTO “REAL” ARTIFICIAL INTELLIGENCE, THESE OBJECTIVE FUNCTIONS WILL PLAY AN OUTSIZED ROLE. WHILE THE UNDERLYING MECHANISM HERE IS RELATIVELY DUMB, IT IS WORTH GRAPPLING WITH THE IMPORTANT QUESTIONS OF DEFINING A GOOD OBJECTIVE FUNCTION.

IN OUR EXAMPLE IN BOULDER, THE OBJECTIVE FUNCTION IS TYPICALLY ALIGNED WITH BUSINESS GOALS. FOR EXCEL, THIS COULD BE THE NUMBER OF NEW SUBSCRIBERS, REVENUE, OR PROFITS. OTHER COMPANIES HAVE DIFFERENT OBJECTIVES: MANY WEB SITES OPTIMIZE ENGAGEMENT (?), WHILE LOGISTICS COMPANIES WANT TO MINIMIZE TRAVEL TIME TO FULFILL PROMISES MADE TO CUSTOMERS (?).

Training Data

THIS MACHINE LEARNING SETUP,⁷ HOWEVER, ONLY WORKS WITH GENEROUS TRAINING DATA. MODERN MACHINE LEARNING ALGORITHMS ARE ALSO NOTORIously DATA HUNGRy. THEY CAN ALWAYS DO BETTER, BUT THEY CAN ONLY DO BETTER AT THE COST OF ADDITIONAL TRAINING DATA: IF YOU SEE INPUT x , WHAT SHOULD YOUR OUTPUT y BE?

HOWEVER, MANY TASKS CAN BE FRAMED AS MECHANICAL CALCULATIONS; THEY DO NOT REQUIRE THE ABSTRACT REASONING OR GENERALIZATION THAT PEOPLE TODAY CONSIDER “INTELLIGENCE”. IN CONTRAST, THEY JUST REQUIRE HUNDREDS OR THOUSANDS OF EXAMPLES TO KNOW WHAT TO DO. EVIL DUDE IS A PROLIFIC DEMONSTRATOR OF WHAT (NOT) TO DO. A HUMAN COULD LEARN EVIL DUDE’S TRICKS FROM A HIGH-LEVEL DESCRIPTION. AN ALGORITHM LEARNS TO COPY.

THIS SIMPLE MIMICRY CAN SEEM SOPHISTICATED. IT’S COUCED IN TERMS THAT SOUND SOPHISTICATED BUT ONLY FACILITATE SUPERFICIAL COPYING (FIGURE 13.1). EXCEL HAS A RELATIVELY SIMPLE OBJECTIVE FUNCTION TO OPTIMIZE: ARE PEOPLE BUYING ITS PRODUCTS OR NOT? IF SALES GO UP, IT’S GOOD. IF SALES GO DOWN, IT’S BAD. OPTIMIZING THIS FUNCTION WITHOUT ANY HELP IS A DIFFICULT PROBLEM (THIS IS WHY CEOS ARE PAID SO MUCH), BUT EXCEL DOESN’T HAVE TO DO IT ALONE: IT LEARNED IT FROM EVIL DUDE.

⁷ Formally, this setup is called *supervised machine learning*, where there are specific input-output examples the algorithm must replicate. Unsupervised machine learning is another active area of research, but it is much harder to see when things are working well. We talk about applications and evaluations of unsupervised machine learning in Section

?reference? ??



Figure 13.1: Much of “Machine Learning” is copying in the form of complicated equations defined by mathematical formulas using fancy words like “linear algebra”, “gradients”, and “non-convex optimization”.

13.2 Humans are Jerks

EVERYBODY'S A JERK. YOU, ME, THAT JERK OVER THERE... THAT'S MY PHILOSOPHY.

BENDER BENDING RODRIGUEZ, *FUTURAMA: I, ROOMMATE*

PRESUMABLY I DON'T NEED TO CONVINCE YOU THAT HUMANS CAN BE JERKS. INSTEAD, I WANT TO CONVINCE YOU THAT HUMANS CAN CREATE SYSTEMATIC, PREDICTABLE JERKITUDE THAT CAN BE CARRIED ON, EMULATED, AND PERFECTED BY MACHINE LEARNING.

LET'S TAKE A SPECIFIC EXAMPLE OF INSTITUTIONALIZED JERKDOM: REDLINING. REDLINING IS THE PRACTICE THAT PREVENTED MORTGAGES FROM BEING ISSUED TO SPECIFIC NEIGHBORHOODS: ITS NAME CAME FROM THE MAPS CREATED FOR CHICAGO's AUSTIN NEIGHBORHOOD (?). THE NEIGHBORHOODS EXCLUDED FROM LENDING WERE TYPICALLY MINORITY-MAJORITY. THIS WAS A DESPICABLE PRACTICE AND EVEN THOUGH IT HAS OSTENSIBLY ENDED, IT HAS PERSISTENT EFFECTS ON FAMILIAL WEALTH, EDUCATION, VOTING REDISTRICTING, AND SEGREGATION.

OF COURSE, LAWS PROHIBITED OUTRIGHT DISCRIMINATION IN THE PUBLIC SECTOR. HOWEVER, AN IMPORTANT PART OF THE HOUSING MARKET IS CONTROLLED BY PRIVATE BANKS:⁸ WHETHER OR NOT A FAMILY CAN GET A MORTGAGE FOR A HOUSE. AN OVERLY SIMPLIFIED VERSION OF THE REDLINING STORY IS THAT IF A FAMILY COMES IN AND ASKS FOR A LOAN TO BUY A HOUSE IN THE "WRONG" NEIGHBORHOOD, THE BANK WILL DENY THEM THE LOAN.

MANY AMERICANS—PARTICULARLY THOSE WHOSE FAMILIES WERE SPARED BY THE PRACTICE—AREN'T AWARE OF THE HISTORY OF THE PRACTICE. UNFORTUNATELY, THESE ARE OFTEN THE PEOPLE IN POSITIONS OF POWER WHEN, FOR INSTANCE, A BANK DECIDES TO CREATE A MACHINE LEARNING ALGORITHM TO HELP DECIDE WHETHER OR NOT TO GIVE A FAMILY A LOAN FOR A HOUSE.

I Learned it from You!

MACHINE LEARNING ALGORITHMS LEARN FROM EXAMPLES. EVIL DUDE'S MONKEYING WITH THE POWER GRID, BANKERS DENYING EVERY MINORITY BORROWER, OR OLD BOYS' NETWORKS WHO ONLY HIRE FROM IVY LEAGUE SCHOOLS. JUST AS HUMANS CAN JUSTIFY THEIR ACTIONS THROUGH EXCEL'S VIRTUOUS MISSION, PRESERVING NEIGHBORHOODS, OR PRESERVING A COMPANY'S CULTURE, MACHINE LEARNING ALGORITHMS CAN "JUSTIFY" THEIR ACTIONS WITH A VENEER OF RESPECTABILITY.

AN ALGORITHM HAS NO PREJUDICE OR ILL INTENT—IT AFTER ALL LACKS INTELLIGENCE. BUT THAT DOESN'T MEAN THAT IT CANNOT RECAPITULATE THE MALICE THAT IT "LEARNS" BY RECREATING MAN'S INHUMANITY TO MAN.⁹

ONE OF MACHINE LEARNING'S STRENGTHS IS THAT IT CAN DISCOVER SURPRISING CORRELATIONS: E-MAILS SENT AT THIS HOUR FROM TOWN ARE OFTEN SPAM, PEOPLE STOP TALKING ABOUT THE FUTURE WHEN THEY'RE ENDING A RELATIONSHIP. THIS STRENGTH CAN CONCEAL WHEN A MACHINE LEARNING

⁸ After the 2008 housing market implosion and the federal government increasingly becoming the lender of last resort, this fiction is a flimsy veneer. Although private (sometimes local) banks originate loans, the government is usually the ultimate lender.

⁹ I was going to document that this is a reference to *The Dirge* by Robert Burns, but in documenting that, I learned that this phrase might have itself been a reference to 1673's *The Whole Duty of Man According to the Law of Nature* by von Pufendorf.

ALGORITHM IS DOING SOMETHING IT “SHOULDN’T”. AGAIN, THESE ALGORITHMS ARE TOO DUMB TO HAVE A MORALITY; THIS JUDGMENT IS AN EXTERNAL ONE BASED ON SOCIETY’S MORES.

TAKE THE CASE OF REDLINING AGAIN. THE ALGORITHM MAY NOT KNOW THE RACE OF LOAN APPLICANTS. BUT IT MIGHT KNOW THEIR CURRENT ADDRESS, AGE, OCCUPATION, AND NAME. THAT IS OFTEN ENOUGH TO INFER THE RACE OF AN APPLICANT AND FROM THAT CAN REPLICATE PAST HUMANS’ PREJUDICE WHILE GIVING CURRENT HUMANS THE EXCUSE OF “I’M JUST DOING WHAT THE MACHINE SAYS”.

THIS IS NOT JUST AN ACADEMIC DISCUSSION. IT IS THE SUBJECT OF A SUBFIELD OF MACHINE LEARNING FOCUSING ON FAIRNESS, ACCOUNTABILITY, TRANSPARENCY, AND ETHICS (FATE). THIS SUBFIELD FOCUSES ON MAKING MACHINE LEARNING MORE UNDERSTANDABLE (?) AND QUANTIFYING UNFAIRNESS (?).

ONE OF THE LESSONS LEARNED FROM THE REDLINING DEBACLE IS HOW TO CHARACTERIZE THESE POOR OUTCOMES. IT’S OFTEN WORTHWHILE TO DISTINGUISH DISPARATE IMPACTS VS. DISPARATE TREATMENT (?).

Objective Function Mismatch

ANOTHER LESSON FROM EXCEL’S CHOICES IS THEIR OBJECTIVE FUNCTION: OPTIMIZING SALES. A DIE-HARD SKEPTIC OF CAPITALISM MIGHT SAY THAT THIS IS AN INHERENT FAILING OF A FREE MARKET SYSTEM, BUT THESE PROBLEMS ARE ALSO PRESENT IN PLANNED ECONOMIES (?): BUT THE OBJECTIVE FUNCTION IS BEING APPLIED TO A SPRAWLING COUNTRY INSTEAD OF AN INDIVIDUAL.¹⁰

OBJECTIVE FUNCTIONS CAN SEEM REASONABLE AT FIRST BLUSH BUT FAIL IN REALITY. IN NEW YORK CITY, THE RAND CORPORATION HELPED OPTIMIZE FIRE STATION PLACEMENT AND STAFFING BASED ON MINIMIZING RESPONSE TIME. THIS WAS ALREADY RECORDED AND SEEMS REASONABLE: YOU CANNOT PUT OUT A FIRE UNTIL THE FIRST FIRE TRUCK GETS THERE.

HOWEVER, ? IN *THE FIRES: HOW A COMPUTER FORMULA, BIG IDEAS, AND THE BEST OF INTENTIONS BURNED DOWN NEW YORK CITY—AND DETERMINED THE FUTURE OF CITIES* ARGUES THAT THIS LED TO POOR OUTCOMES FOR NEW YORK: NOT EVERYONE CALLED IN FIRES AT THE SAME RATE (OR AS QUICKLY AFTER A FIRE STARTS). NOT ALL CALLS ARE EQUAL: IT’S OKAY IF YOU WAIT TO ADDRESS A FALSE ALARM IN FAVOR OF A BLAZING INFERNO. MOREOVER, THE FIRST RESPONSE ISN’T THE WHOLE STORY: WHAT ACTUALLY MATTERS IS PROTECTING LIFE AND PROPERTY, BUT THIS IS HARDER TO QUANTIFY.

Prevention is Tough

PREVENTING THESE ISSUES IS DIFFICULT. IT REQUIRES DETECTING MALICE IN HUMANS, UNDERSTANDING WHAT FACTORS LEAD TO THE MALICE, AND THEN ESPECIALLY ENGINEERING YOUR MACHINE LEARNING ALGORITHM TO AVOID THAT MALICE. ORGANIZATIONS CAN DELIBERATELY OBFUSCATE THEIR MALICE OR INADVERTENTLY OVERLOOK IT (E.G., THEY DON’T KNOW THE HISTORY OF

¹⁰ *Red Plenty* (?) is an academic/fiction mashup much like this book that realizes the dream of a planned economy with high-tech information technology. In real life, even Trotsky recognized that central planning “is checked and, to a considerable degree, realized through the market” (?).

THEIR DATA).

LEGAL REGIMES MAY BE A MECHANISM FOR PREVENTING ALGORITHMIC DISCRIMINATION. FOR EXAMPLE, THE PRICE OF INSURANCE CAN ONLY BE BASED ON AGE. LIMITING THE DATA THAT AN ALGORITHM CAN USE PREVENTS IT FROM FINDING OBSCURE CORRELATIONS BUT ALSO PREVENTS THE ALGORITHM FROM DOING ANYTHING USEFUL THAT COULD, SAY, LOWER OVERALL HEALTHCARE COSTS.¹¹ THESE TRADEOFFS MIGHT BE WORTH MAKING, HOWEVER.

RICH CARUNA ARGUES THAT YOU SHOULD NOT HIDE SENSITIVE FEATURES TO ELIMINATE BIAS. INSTEAD, HE ARGUES THAT YOU SHOULD EXPLICITLY INCLUDE THAT DATA AND PREVENT IT FROM BEING PART OF IMPORTANT DECISIONS. OTHER, MORE NUANCED MODELS, ARGUE FOR ADAPTING THE “DISPARATE IMPACT” FROM THE LEGAL WORLD INTO MACHINE LEARNING (?): SCRUB A DATASET SO THAT INFORMATION ABOUT PROTECTED CLASSES (RACE, GENDER, AGE, ETC.) DO NOT LEAK INTO PARAMETERS THAT CAN BE USED BY A MACHINE LEARNING ALGORITHM.

¹¹ This is also connected to the concept of *overfitting* in machine learning: fewer features might lead to a worse objective function score but might be better in the “real world”.

Today’s Cybersecurity Crisis: Tomorrow’s Machine Learning Crisis

LET’S FOCUS, HOWEVER, ON EVIL DUDE’S PARTICULAR BAG OF TRICKS. EXCEL’S ALGORITHM WAS ONLY ABLE TO RECREATE THE JERK BY HAVING ACCESS TO MACHINES AND DEVICES THAT IT SHOULDN’T.

STUXNET IS AN EXAMPLE OF HOW FAR THIS CAN GO: A COMPUTER VIRUS THAT SPECIFICALLY TARGETED CENTRIFUGES IN IRAN’S NUCLEAR PROGRAM. STUXNET WAS A “PRECISION WEAPON—A TECHNIQUE FOR ANONYMOUSLY TARGETING A SPECIFIC SYSTEM OR ORGANIZATION WITHOUT MASSIVE COLLATERAL DAMAGE” (?). IT WAS ABLE TO EXPLOIT SECURITY FLAWS IN OPERATING SYSTEMS, AND—LIKE COFFEE MACHINES IN BOULDER—HOW SPECIFIC EVERYDAY DEVICES INTERACT AND INTERFERE WITH EACH OTHER.

HOWEVER, STUXNET WAS CRAFTED BY EXPERTS OVER THE COURSE OF YEARS (ONE ASSUMES, GIVEN THE COMPLEXITY OF THE SOURCE CODE... ITS ORIGINS ARE SHROUDED IN SECRECY). IN CONTRAST, MACHINE LEARNING IS FAST, EFFICIENT, AND SCALABLE. THE SAME ARTISINAL MISCHIEF THAT HUMAN JERKS CAN DO LOCALLY, MACHINE LEARNING CAN INSTANTLY REPRODUCE GLOBALLY.

THIS IS NOT TO ARGUE FOR ABANDONING ALL TECHNOLOGY. INSTEAD, THE WORLD OF RONALD D. MOORE’S *BATTLESTAR GALACTICA* REBOOT PROVIDES A USEFUL TEMPLATE. THEY NEEDED COMPUTERS FOR JUMPS AND HAD PLENTY OF ELECTRONICS. THEY RELIGIOUSLY MADE SURE THAT COMPUTERS WERE NEVER EVER LINKED TOGETHER.¹²

HOWEVER, IN THE EARLY TWENTY-FIRST CENTURY, EVERY DEVICE IS LINKED TO AN **UNSECURED** INTERNET. THE “INTERNET OF THINGS” MOVEMENT HAS PROLIFERATED THE DEVICES THAT ARE CONNECTED TO THE INTERNET: REFRIGERATORS (?), CARS, AND MEDICAL EQUIPMENT (?). THESE GADGETS ARE CATNIP TO YUPPIES WHO LIVE IN PLACES LIKE BOULDER. AND LET’S FACE IT, THEY CAN DO SOME COOL STUFF.

¹² The downfall of the colonies was Gaius Baltar circumventing these preventative measures.

BUT WE CREATED THE DEVICES BEFORE WE MADE OUR HOMES AND INTERNET READY FOR THEM. IDEALLY, EACH HOME SHOULD BE A NICE, INSULATED ENVIRONMENT WITH CONTROLLED ENTRY AND EXIT. JUST LIKE YOU WOULDN'T LEAVE YOUR FRONT DOOR OPEN ALL DAY AND ALL NIGHT, YOU WOULD NOT LET ANYONE IN THE WORLD POTENTIALLY FIDDLE WITH THE GADGETS IN YOUR HOME.

THE INTERNET WAS DESIGNED AS A PLACE WHERE EVERYONE COULD BE TRUSTED. THE INTERNET BEGAN AS A RESOURCE FOR RESEARCHERS FUNDED BY AMERICA'S DEFENSE ADVANCED RESEARCH PROJECTS AGENCY (DARPA)¹³ TO SHARE DATA AND COMMUNICATE QUICKLY (?). WHEN THE INTERNET HAD DOZENS OF NODES WHERE EVERYONE KNEW AND TRUSTED EACH OTHER, YOU DIDN'T HAVE TO BE PARANOID.

THIS MIGHT BE OKAY IF THE UNDERLYING DEVICES WERE SECURE OR UNIMPORTANT. WE'RE ATTACHING EVERYTHING TO THE INTERNET AND NOT DOING A GOOD JOB OF IT. BAD GUYS CAN SPY ON YOU THROUGH YOUR "COMPUTER, HOME SECURITY SYSTEM, OR DEVICES SUCH AS BABY MONITORS" (?), REMOTELY HIJACK YOUR CAR (?), AND STEAL YOUR MONEY.

UNFORTUNATELY, NEITHER CONSUMERS NOR COMPANIES REALLY SEE THE PROBLEM. HIGHLY SECURE DEVICES ARE ANNOYING TO THE CONSUMER.¹⁴ YOU NEED TO SET UP PASSWORDS, FIREWALLS, AND UPDATE THE FIRMWARE CONSISTENTLY. THE IDEAL DEVICE IS ONE THAT THE USER PLUGS IN, IT WORKS, AND THEN THE USER NEVER HAS TO THINK ABOUT IT AGAIN. MANY COMPANIES CAN ONLY SURVIVE IF THEIR NEXT PRODUCT SELLS, AND YOU'RE NOT GOING TO SELL MANY UNITS IF YOUR PRODUCT IS SECURE (AND THUS ANNOYING). HUMANS ARE BAD ABOUT ESTIMATING LOW-PROBABILITY EVENTS AND ASSUME THAT THEY'LL NEVER HAVE A SECURITY ISSUE.

EVENTUALLY, BIG FIRMS MIGHT BE COERCED INTO DOING SOMETHING BY THE THREAT OF LITIGATION. COMPANIES THAT HAVE DEEP POCKETS AND STRONG REPUTATIONS MIGHT EVENTUALLY BE SCARED OF BAD PRESS AND BIG LAWSUITS INTO CARING ABOUT SECURITY. AND CONSUMERS TOO MIGHT BE WILLING TO PUT UP WITH THE ANNOYANCE OF DOING CYBERSECURITY RIGHT IF THAT'S SIMPLY "WHAT YOU HAVE TO DO".

BUT AS LONG AS THE INTERNET REMAINS AS IT IS, WE'LL BE AS VULNERABLE AS THE WEAKEST LINK. WHILE THE FREE MARKET MIGHT GET SOME FIRMS TO MAKE SECURE DEVICES, SOME CONSUMERS, SOME COMPANIES, OR SOME BAD INTERACTION BETWEEN THE TWO MIGHT LEAD TO UNSECURE ENVIRONMENTS. AND THIS IS THE REAL RISK: A HOME OR A BUSINESS IS ONLY AS STRONG AS ITS WEAKEST LINK. AND THE STATUS QUO HASN'T FORGED MANY STRONG LINKS; INDEED, IT'S ENCOURAGED AND REWARDED THE PROLIFERATION OF WEAK LINKS.

Securing the Beachheads

BECAUSE THIS IS SOCIETY'S PROBLEM, I SUSPECT THAT THE ONLY LONG-TERM SOLUTION IS FROM SOCIETY: LEGISLATION. INDIVIDUAL INCENTIVES AREN'T ENOUGH TO MAKE SURE THAT EVERYONE IS SAFE, SO TO PROTECT SOCIETY WE

¹³ In future chapters, much of the research into artificial intelligence will be directed by fictional governments for military ends. This isn't just to make the story more exciting; it also reflects reality. Much of the breakthrough research in AI has happened in America, and much that—like last century's breakthroughs in communications—has been funded by military organizations such as the Defense Advanced Research Projects Agency.

¹⁴ This applies to commercial and residential consumers. In both cases, devices will be set up by people, often non-experts, who will not want to put up with the hassle of making sure devices are secure. This often leads to shortcuts ()�.

NEED TO MAKE SURE PEOPLE ARE PROTECTING THEMSELVES.

DOING IT RIGHT IS DIFFICULT FOR MANY REASONS, BUT HERE I'LL TALK ABOUT THE MOST SALIENT. WHILE INNOVATION MOVES QUICKLY, LEGISLATION DECIDEDLY DOES NOT; MOREOVER, WHEN LEGISLATION DOES HAPPEN, IT OFTEN HAS UNINTENDED CONSEQUENCES. REFINING LEGISLATION—FIXING THE UNINTENDED CONSEQUENCES—REQUIRES INPUT FROM INDUSTRY (THE EXPERTS), BUT THIS BRINGS ITS OWN ISSUES.

PEOPLE WORKING IN TECHNOLOGY BROADLY LIKE TO CLAIM THAT THINGS MOVE FAST—IN OTHER WORDS, APPLYING THE EARLY FACEBOOK ETHOS TO “MOVE FAST AND BREAK THINGS” (?). WHILE SOME OF THIS IS HYPE (IT STILL TAKES TIME TO GET THINGS RIGHT), THERE IS SOME TRUTH TO THIS CLAIM. NEW TECHNOLOGIES REALLY DO ALLOW NEW BUSINESS MODELS TO EMERGE, AND THESE OFTEN ARE INCOMPATIBLE WITH EXISTING LEGAL STRUCTURES (?).

COPYRIGHT LAW HAS BEEN VEXED BY THE REMIX CULTURE MADE POSSIBLE BY PLATFORMS LIKE YOUTUBE (?). SPAM PHONE AND E-MAIL FILTERS HAVE NOT KEPT UP WITH THE REALITY OF ROBODIALERS AND TEXT-TO-SPEECH. THE INTRINSICALLY INTERNATIONAL CONCEPTION OF THE INTERNET GENERALLY IS AT ODDS WITH REGIONAL-BASED LAWS AND REGULATION.

AROUND THE WORLD, GOVERNMENT HAS NOT DONE A STELLAR JOB OF RESPONDING TO THIS CHALLENGE. IN AMERICA, GRIDLOCK HAS REIGNED. COMPANIES PROPOSE THEIR OWN FRAMEWORKS AND GUIDELINES, BUT THEY LACK LEGAL FORCE AND ARE ULTIMATELY VOLUNTARY. LIKE THE OVER-HYPED MYTH OF POLISH CALVARY CHARGING GERMAN TANKS WITH LANCES,¹⁵ EUROPE'S RESPONSE HAS BEEN MOSTLY TO FIGHT THE LAST WAR: FOCUSING ON PRIVACY AND TRACKING (REAL PROBLEMS IN MANY EUROPEAN REGIMES) RATHER THAN THAT OF PROTECTING SYSTEMS.

EUROPE'S WELL-INTENTIONED GDPR RESPONSE ALSO HIGHLIGHTS ANOTHER CRITICAL ISSUE: UNINTENDED CONSEQUENCES (?). REGULATORY BURDENS CAN FAVOR ENTRENCHED FIRMS OVER UPSTARTS WHO CAN'T AFFORD LAWYERS TO OVERSEE COMPLIANCE AND CAN PREVENT YOUR CITIZENS FROM HAVING ACCESS TO THE REST OF THE WORLD: COMPANIES MIGHT PREFER TO NOT OFFER SERVICE IN YOUR COUNTRY IF THEY DON'T LIKE YOUR LAWS.

IF THE UNINTENDED CONSEQUENCES FAVOR PARTICULAR SECTORS, THOSE SECTORS MIGHT THEN TRY TO KEEP THOSE REGULATIONS ENTRENCHED AND STATIC (?). THEN THE REGULATIONS HELP THAT SECTOR RATHER THAN THE PUBLIC. WHATEVER FORM LEGISLATION TAKES TO SOLVE THESE UNDERLYING ISSUES, IT MUST BE WELL-CRAFTED TO ADDRESS NOT JUST THE EMERGENT ISSUES OF TODAY BUT TO PROVIDE A FRAMEWORK TO SOLVE FUTURE ISSUES. IN AMERICA, THE REPUBLICAN¹⁶ SYSTEM AND THE IMPLICIT VETO POWER OF MINORITY POLITICAL PARTIES IN THE COOLING SAUCER OF THE SENATE; IN EUROPE THE FEDERAL SYSTEM OF EUROPEAN REGULATION WITH EXPLICIT OR IMPLICIT VETO FROM MEMBER STATES; IN CHINA, THE OPAQUE CONSENSUS BUILDING BETWEEN PARTY, PROVINCE, AND INDUSTRY... NO MODERN SOCIETY MOVES QUICKLY. THE DISCREPANCY IN TEMPO IS MORE ACUTE WHEN IT COMES TO TECHNOLOGY. THE

¹⁵ The charge at Krojanty did involve infantry, tanks, and cavalry, but the Polish calvary was used effectively given the circumstances.

¹⁶ The “r” here is lower-case: in the American republic, voters select representatives from oddly drawn districts.

SILVER HAIR MANDARINS IN WHITEHALL, DC, BRUSSELS, OR BEIJING ARE OFTEN NOT EQUIPPED TO UNDERSTAND THE INTERACTION BETWEEN SOCIETY AND NEW TECHNOLOGY.

THE TEMPTATION THEN IS TO ABDICATE THE ISSUE TO MORE NIMBLE AGENTS: THE MILITARY OR INDUSTRY. THESE ARE IMPORTANT PARTS OF HOW ARTIFICIAL INTELLIGENCE WILL ENTER SOCIETY, BUT THEY HAVE THEIR OWN INTEREST SOMEONE AT ODDS WITH THE REST OF SOCIETY. THE NEXT CHAPTERS WILL CONSIDER HOW MARTIAL OR COMMERCIAL AI COULD GO WRONG.

But what about Artificial Intelligence?

THE ASTUTE READER HAS PROBABLY NOTICED THAT VERY LITTLE IN THIS SECTION HAS TO DO WITH AI *PER SE*. IT'S MORE ABOUT PREVENTING THE EVERYDAY JERKS FROM SCREWING WITH OUR LIVES.

THE SAME GOOD DIGITAL HYGIENE THAT PROTECTS US FROM THE COMMON JERK ALSO PREVENTS THE SCALED-UP, FAST-PACED EPIDEMIC OF AI-DRIVEN TERROR.

THIS IS WHY WHEN A STUDENT COMES TO ME WHO IS INTERESTED IN COMPUTERS AND HAS NO IDEA WHAT TO DO, I DON'T STEER THEM TO THE BRIGHT SHINY FIELD OF AI (WHERE I WORK) BUT RATHER TO CYBERSECURITY. IT IS PERHAPS NOT AS SEXY, BUT IT IS WHAT WILL AVERT THE FIRST ROBOT APOCALYPSE.

BUT FRANKLY, THESE COMPUTERS ARE PRETTY DUMB. WHAT HAPPENS WHEN THEY GET A LITTLE SMARTER... SOMETHING THAT COULD ACTUALLY BE CONSIDERED **intelligent?**

BIBLIOGRAPHY

THE PAPER CHASE, 1973. BASED ON THE 1971 NOVEL BY JOHN JAY OSBORN JR.

YING BAI AND RUIXUE JIA. ELITE RECRUITMENT AND POLITICAL STABILITY: THE IMPACT OF THE ABOLITION OF CHINA'S CIVIL SERVICE EXAM. *ECONOMETRICA*, 84(2):677–733, 2016. ISSN 00129682, 14680262. URL <http://www.jstor.org/stable/43866446>.

ONDŘEJ BOJAR, CHRISTIAN BUCK, CHRISTIAN FEDERMANN, BARRY HADDOW, PHILIPP KOEHN, JOHANNES LEVELING, CHRISTOF MONZ, PAVEL PECINA, MATT POST, HERVE SAINT-AMAND, RADU SORICUT, LUCIA SPECIA, AND ALEŠ TAMCHYNA. FINDINGS OF THE 2014 WORKSHOP ON STATISTICAL MACHINE TRANSLATION. IN *PROCEEDINGS OF THE NINTH WORKSHOP ON STATISTICAL MACHINE TRANSLATION*, JUNE 2014.

LÉON BOTTOU. STOCHASTIC LEARNING. IN OLIVIER BOUSQUET AND ULRIKE VON LUXBURG, EDITORS, *ADVANCED LECTURES ON MACHINE LEARNING*, LECTURE NOTES IN ARTIFICIAL INTELLIGENCE, LNAI 3176, PAGES 146–168. SPRINGER VERLAG, BERLIN, 2004. URL <http://leon.bottou.org/papers/bottou-mlss-2004>.

SAMUEL R. BOWMAN, GABOR ANGELI, CHRISTOPHER POTTS, AND CHRISTOPHER D. MANNING. A LARGE ANNOTATED CORPUS FOR LEARNING NATURAL LANGUAGE INFERENCE. IN LLUÍS MÀRQUEZ, CHRIS CALLISON-BURCH, AND JIAN SU, EDITORS, *PROCEEDINGS OF THE 2015 CONFERENCE ON EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, PAGES 632–642, LISBON, PORTUGAL, SEPTEMBER 2015. ASSOCIATION FOR COMPUTATIONAL LINGUISTICS. DOI: 10.18653/v1/D15-1075. URL <https://aclanthology.org/D15-1075/>.

JORDAN BOYD-GRABER AND BENJAMIN BÖRSCHINGER. WHAT QUESTION ANSWERING CAN LEARN FROM TRIVIA NERDS. IN *ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2020. URL http://umiacs.umd.edu/~jbg/docs/2020_acl_trivia.pdf.

JORDAN BOYD-GRABER, BRIANNA SATINOFF, HE HE, AND HAL DAUME III. BESTING THE QUIZ MASTER: CROWDSOURCING INCREMENTAL CLASSIFICATION

GAMES. IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, 2012.

SERGEY BRIN AND LAWRENCE PAGE. THE ANATOMY OF A LARGE-SCALE HYPERTEXTUAL WEB SEARCH ENGINE. *COMPUTER NETWORKS AND ISDN SYSTEMS*, 30(1):107–117, 1998. ISSN 0169-7552. doi: [HTTPS://DOI.ORG/10.1016/S0169-7552\(98\)00110-X](https://doi.org/10.1016/S0169-7552(98)00110-X). URL <https://www.sciencedirect.com/science/article/pii/S016975529800110X>. PROCEEDINGS OF THE SEVENTH INTERNATIONAL WORLD WIDE WEB CONFERENCE.

DANA L. BURGESS. THE GIFT OF THE SPHINX.

SETH CHAIKLIN. *THE ZONE OF PROXIMAL DEVELOPMENT IN VYGOTSKY'S ANALYSIS OF LEARNING AND INSTRUCTION*, PAGE 39–64. LEARNING IN DOING: SOCIAL, COGNITIVE AND COMPUTATIONAL PERSPECTIVES. CAMBRIDGE UNIVERSITY PRESS, 2003. doi: 10.1017/CBO9780511840975.004.

EUNSOL CHOI, HE HE, MOHIT IYYER, MARK YATSKAR, WEN-TAU YIH, YEJIN CHOI, PERCY LIANG, AND LUKE ZETTLEMoyer. QUAC: QUESTION ANSWERING IN CONTEXT. IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, 2018. URL <https://www.aclweb.org/anthology/D18-1241>.

SHIUON CHU. FAILURE STORIES: INTERPRETATIONS OF REJECTED PAPERS IN THE LATE IMPERIAL CIVIL SERVICE EXAMINATIONS. *T'OUNG PAO*, 101(1/3):168–207, 2015. ISSN 00825433. URL <http://www.jstor.org/stable/24755031>.

CYRIL CLEVERDON. THE CRANFIELD TESTS ON INDEX LANGUAGE DEVICES. *ASLIB PROCEEDINGS*, 19(6):173–194, 06 1967. ISSN 0001-253X. doi: 10.1108/eb050097. URL <https://doi.org/10.1108/eb050097>.

CARLOS A. CUADRA. RELEVANCE ASSESSMENT PROJECT. IN *APA ANNUAL MEETING*, 1968.

ANTHONY CUTHBERTSON. ROBOTS CAN NOW READ BETTER THAN HUMANS, PUTTING MILLIONS OF JOBS AT RISK. *NEWSWEEK*, 2018. URL <https://www.newsweek.com/robots-can-now-read-better-humans-putting-millions-jobs-risk-781393>.

PAUL DIAMOND. *HOW TO MAKE 100 POUNDS A NIGHT (OR MORE) AS A PUB QUIZMASTER*. DP QUIZ, 2009.

JESSE DODGE, SUCHIN GURURANGAN, DALLAS CARD, ROY SCHWARTZ, AND NOAH A. SMITH. SHOW YOUR WORK: IMPROVED REPORTING OF EXPERIMENTAL RESULTS. IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*. ASSOCIATION FOR COMPUTATIONAL LINGUISTICS,

2019. doi: 10.18653/v1/D19-1224. URL <https://www.aclweb.org/anthology/D19-1224>.

DHEERU DUA, YIZHONG WANG, PRADEEP DASIGI, GABRIEL STANOVSKY, SAMEER SINGH, AND MATT GARDNER. DROP: A READING COMPREHENSION BENCHMARK REQUIRING DISCRETE REASONING OVER PARAGRAPHS. IN *CONFERENCE OF THE NORTH AMERICAN CHAPTER OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2019. URL <https://www.aclweb.org/anthology/N19-1246>.

JESSE DUNIETZ, GREG BURNHAM, AKASH BHARADWAJ, OWEN RAMBOW, JENNIFER CHU-CARROLL, AND DAVE FERRUCCI. TO TEST MACHINE COMPREHENSION, START BY DEFINING COMPREHENSION. IN DAN JURAFSKY, JOYCE CHAI, NATALIE SCHLUTER, AND JOEL TETREAULT, EDITORS, *PROCEEDINGS OF THE 58TH ANNUAL MEETING OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, PAGES 7839–7859, ONLINE, JULY 2020. ASSOCIATION FOR COMPUTATIONAL LINGUISTICS. DOI: 10.18653/v1/2020.acl-main.701. URL <https://aclanthology.org/2020.acl-main.701/>.

MATTHEW DUNN, LEVENT SAGUN, MIKE HIGGINS, V. UGUR GÜNEY, VOLKAN CIRIK, AND KYUNGHYUN CHO. SEARCHQA: A NEW Q&A DATASET AUGMENTED WITH CONTEXT FROM A SEARCH ENGINE. *CoRR*, abs/1704.05179, 2017. URL <http://arxiv.org/abs/1704.05179>.

R. DWAN. *As Long as They're Laughing: Groucho Marx and You Bet Your Life*. MIDNIGHT MARQUEE PRESS, 2000. ISBN 9781887664363. URL <https://books.google.ch/books?id=h00KAQAAQAAJ>.

STEPHEN ELTINGE. QUIZBOWL LEXICON, 2013. URL <http://www.pace-nsc.org/pace-quizzbowl-lexicon/>.

MANAAL FARUQUI AND DIPANJAN DAS. IDENTIFYING WELL-FORMED NATURAL LANGUAGE QUESTIONS. IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, 2018.

SHI FENG, ERIC WALLACE, ALVIN GRISOM II, PEDRO RODRIGUEZ, MOHIT IYYER, AND JORDAN BOYD-GRABER. PATHOLOGIES OF NEURAL MODELS MAKE INTERPRETATION DIFFICULT. IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, 2018.

DAVID FERRUCCI, ERIC BROWN, JENNIFER CHU-CARROLL, JAMES FAN, DAVID GONDEK, ADITYA A. KALYANPUR, ADAM LALLY, J. WILLIAM MURDOCK, ERIC NYBERG, JOHN PRAGER, NICO SCHLAEFER, AND CHRIS WELTY. BUILDING WATSON: AN OVERVIEW OF THE DEEPQA PROJECT. *AI MAGAZINE*, 31(3), 2010.

FLIGHT STANDARDS SERVICE. *AVIATION INSTRUCTOR'S HANDBOOK*, VOLUME FAA-H-8083-9A. FEDERAL AVIATION ADMINISTRATION, DEPARTMENT OF TRANSPORTATION, 2008.

MORRIS FREEDMAN. THE FALL OF CHARLIE VAN DOREN. *THE VIRGINIA QUARTERLY REVIEW*, 73(1):157–165, 1997. ISSN 0042675X. URL <http://www.jstor.org/stable/26439004>.

MOR GEVA, YOAV GOLDBERG, AND JONATHAN BERANT. ARE WE MODELING THE TASK OR THE ANNOTATOR? AN INVESTIGATION OF ANNOTATOR BIAS IN NATURAL LANGUAGE UNDERSTANDING DATASETS. IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, 2019.

MARK E GLICKMAN AND ALBYN C JONES. RATING THE CHESS RATING SYSTEM. *CHANCE*, 12, 1999. URL http://scholar.google.de/scholar.bib?q=info:r4leZxFt1lwJ:scholar.google.com/&output=citation&scisig=AAGBfm0AAAAAUm2XckUm12ofzJwFf4A_kI-bLVVMwfDI&scisf=4&hl=en&scfhb=1.

BERT F. GREEN, ALICE K. WOLF, CAROL CHOMSKY, AND KENNETH LAUGHERY. BASEBALL: AN AUTOMATIC QUESTION-ANSWERER. IN *PAPERS PRESENTED AT THE MAY 9-11, 1961, WESTERN JOINT IRE-AIEE-ACM COMPUTER CONFERENCE*, IRE-AIEE-ACM '61 (WESTERN), PAGE 219–224, NEW YORK, NY, USA, 1961. ASSOCIATION FOR COMPUTING MACHINERY. ISBN 9781450378727. doi: 10.1145/1460690.1460714. URL <https://doi.org/10.1145/1460690.1460714>.

MERV GRIFFIN. *MERV : MAKING THE GOOD LIFE LAST*. THORNDIKE PRESS, WATERVILLE, ME, 2003. ISBN 0786253533.

ALVIN GRISSOM II, HE HE, JORDAN BOYD-GABER, AND JOHN MORGAN. DON'T UNTIL THE FINAL VERB WAIT: REINFORCEMENT LEARNING FOR SIMULTANEOUS MACHINE TRANSLATION. IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, 2014.

MANSI GUPTA, NITISH KULKARNI, RAGHUVeer CHANDA, ANIRUDHA RAYASAM, AND ZACHARY C. LIPTON. AMAZONQA: A REVIEW-BASED QUESTION ANSWERING TASK. IN *INTERNATIONAL JOINT CONFERENCE ON ARTIFICIAL INTELLIGENCE*, 2019.

THOMAS M. HALADYNA. *DEVELOPING AND VALIDATING MULTIPLE-CHOICE TEST ITEMS*. LAWRENCE ERLBAUM ASSOCIATES, 2004.

BOB HARRIS. *PRISONER OF TREBEKISTAN : A DECADE IN JEOPARDY*. CROWN PUBLISHERS, NEW YORK, 2006. ISBN 0307339564.

ROBERT X. D. HAWKINS, ANDREAS STUHLMÜLLER, JUDITH DEGEN, AND NOAH D. GOODMAN. WHY DO YOU ASK? GOOD QUESTIONS PROVOKE INFORMATIVE ANSWERS. IN *PROCEEDINGS OF THE 37TH ANNUAL MEETING OF THE COGNITIVE SCIENCE SOCIETY*, 2015.

HE HE, JORDAN BOYD-GRABER, KEVIN KWOK, AND HAL DAUMÉ III. OPPONENT MODELING IN DEEP REINFORCEMENT LEARNING. IN *PROCEEDINGS OF THE INTERNATIONAL CONFERENCE OF MACHINE LEARNING*, 2016.

R. ROBERT HENTZEL. NAQT ELIGIBILITY OVERVIEW, 2018. URL <https://www.naqt.com/rules/eligibility.html>.

KARL MORITZ HERMANN, TOMAS KOCISKY, EDWARD GREFENSTETTE, LASSE ESPEHOLT, WILL KAY, MUSTAFA SULEYMAN, AND PHIL BLUNSMON. TEACHING MACHINES TO READ AND COMPREHEND. IN *PROCEEDINGS OF ADVANCES IN NEURAL INFORMATION PROCESSING SYSTEMS*, 2015. URL <http://papers.nips.cc/paper/5945-teaching-machines-to-read-and-comprehend.pdf>.

DOUGLAS R. HOFSTADTER. *LE TON BEAU DE MAROT: IN PRAISE OF THE MUSIC OF LANGUAGE*. BASIC BOOKS, NEW YORK, 1997. ISBN 978-0465086450. ALSO PUBLISHED BY BLOOMSBURY (UK) AND BASIC BOOKS (1998 PAPERBACK).

FENG-HSIUNG HSU. *BEHIND DEEP BLUE: BUILDING THE COMPUTER THAT DEFEATED THE WORLD CHESS CHAMPION*. PRINCETON UNIVERSITY PRESS, USA, 2002. ISBN 0691090653.

SHANKAR IYER, NIKHIL DANDEKAR, AND KORNÉL CSERNAI. QUORA QUESTION PAIRS, 2017. URL <https://www.quora.com/q/quoradata/First-Quora-Dataset-Release-Question-Pairs>.

THOMAS JACKSON, EDITOR. *ON DRESS*. NORTHWEST NAZARENE UNIVERSITY, 1872.

KEN JENNINGS. *BRAINIAC: ADVENTURES IN THE CURIOUS, COMPETITIVE, COMPUTIVE WORLD OF TRIVIA BUFFS*. VILLARD, 2006. ISBN 9781400064458.

MARK JOHNSON, THOMAS L. GRIFFITHS, AND SHARON GOLDWATER. BAYESIAN INFERENCE FOR PCFGs VIA MARKOV CHAIN MONTE CARLO. IN *CONFERENCE OF THE NORTH AMERICAN CHAPTER OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2007.

MICHAEL I. JORDAN, ZOUBIN GHAHRAMANI, TOMMI S. JAAKKOLA, AND LAWRENCE K. SAUL. AN INTRODUCTION TO VARIATIONAL METHODS FOR GRAPHICAL MODELS. *MACHINE LEARNING*, 37(2):183–233, 1999. ISSN 0885-6125. doi: 10.1023/A:1007665907178.

MANDAR JOSHI, EUNSOL CHOI, DANIEL WELD, AND LUKE ZETTLEMOYER. TRIVIAQA: A LARGE SCALE DISTANTLY SUPERVISED CHALLENGE DATASET FOR READING COMPREHENSION. IN *PROCEEDINGS OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2017. URL <https://www.aclweb.org/anthology/P17-1147>.

BORIS KATZ, SUE FELSHIN, DENIZ YURET, ALI IBRAHIM, JIMMY J. LIN, GREGORY MARTON, ALTON JEROME MCFARLAND, AND BARIS TEMELKURAN. OMNIBASE: UNIFORM ACCESS TO HETEROGENEOUS DATA FOR QUESTION ANSWERING. IN *PROCEEDINGS OF THE 6TH INTERNATIONAL CONFERENCE ON APPLICATIONS OF NATURAL LANGUAGE TO INFORMATION SYSTEMS-REVISED PAPERS*, NLDB '02, PAGE 230–234, BERLIN, HEIDELBERG, 2002. SPRINGER-VERLAG. ISBN 354000307X.

DIVYANSH KAUSHIK AND ZACHARY C. LIPTON. HOW MUCH READING DOES READING COMPREHENSION REQUIRE? A CRITICAL INVESTIGATION OF POPULAR BENCHMARKS. IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, 2018.

MICHAEL KEEN. TECHNICAL REPORT, COLLEGE OF AERONAUTICS, UNIVERSITY OF CRANFIELD, 1966.

CHRISTOF KOCH. HOW THE COMPUTER BEAT THE GO PLAYER. *SCIENTIFIC AMERICAN MIND*, 27(4):20–23, 2016. ISSN 15552284, 2331379X. URL <https://www.jstor.org/stable/24945452>.

TOM KWIATKOWSKI, JENNIMARIA PALOMAKI, OLIVIA REDFIELD, MICHAEL COLLINS, ANKUR PARikh, CHRIS ALBERTI, DANIELLE EPSTEIN, ILLIA POLOSUKHIN, MATTHEW KELCEY, JACOB DEVLIN, KENTON LEE, KRISTINA N. TOUTANOVA, LLION JONES, MING-WEI CHANG, ANDREW DAI, JAKOB USZKOREIT, QUOC LE, AND SLAV PETROV. NATURAL QUESTIONS: A BENCHMARK FOR QUESTION ANSWERING RESEARCH. *TRANSACTIONS OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2019. URL <https://tomkwiat.users.x20web.corp.google.com/papers/natural-questions/main-1455-kwiatkowski.pdf>.

YANN LEcUN, LéON BOTTOU, YOSHUA Bengio, AND PATRICK HAFFNER. GRADIENT-BASED LEARNING APPLIED TO DOCUMENT RECOGNITION. IN *PROCEEDINGS OF THE IEEE*, VOLUME 86, PAGES 2278–2324, 1998. URL <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.42.7665>.

KENTON LEE, MING-WEI CHANG, AND KRISTINA TOUTANOVA. LATENT RETRIEVAL FOR WEAKLY SUPERVISED OPEN DOMAIN QUESTION ANSWERING. IN *PROCEEDINGS OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2019.

ROGER LEVY. INTEGRATING SURPRISAL AND UNCERTAIN-INPUT MODELS IN ONLINE SENTENCE COMPREHENSION: FORMAL TECHNIQUES AND EMPIRICAL RESULTS. IN *PROCEEDINGS OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2011.

ROGER P. LEVY, FLORENCIA REALI, AND THOMAS L. GRIFFITHS. MODELING THE EFFECTS OF MEMORY ON HUMAN ONLINE SENTENCE PROCESSING WITH

PARTICLE FILTERS. IN *PROCEEDINGS OF ADVANCES IN NEURAL INFORMATION PROCESSING SYSTEMS*, 2008.

PATRICK LEWIS, PONTUS STENETORP, AND SEBASTIAN RIEDEL. QUESTION AND ANSWER TEST-TRAIN OVERLAP IN OPEN-DOMAIN QUESTION ANSWERING DATASETS. IN *PROCEEDINGS OF THE 16TH CONFERENCE OF THE EUROPEAN CHAPTER OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS: MAIN VOLUME*, PAGES 1000–1008, ONLINE, APRIL 2021. ASSOCIATION FOR COMPUTATIONAL LINGUISTICS. URL <https://aclanthology.org/2021.eacl-main.86>.

J.F. LIGHT. *THE CULTURAL ENCYCLOPEDIA OF BASEBALL*, 2D ED. EBL-SCHWEITZER. MCFARLAND, INCORPORATED, PUBLISHERS, 2016. ISBN 9781476617442. URL <https://books.google.de/books?id=ii0-CgAAQBAJ>.

ZACHARY C. LIPTON AND JACOB STEINHARDT. TROUBLING TRENDS IN MACHINE LEARNING SCHOLARSHIP. *QUEUE*, 17(1), FEBRUARY 2019. ISSN 1542-7730. doi: 10.1145/3317287.3328534. URL <https://doi.org/10.1145/3317287.3328534>.

ELAINE LOW. JEOPARDY! GOAT KEN JENNINGS ON TREBEK, TRASH TALK AND WHY HE'S HANGING UP HIS BUZZER. *VARIETY*, JANUARY 2020.

MINGBO MA, LIANG HUANG, HAO XIONG, RENJIE ZHENG, KAIBO LIU, BAIGONG ZHENG, CHUANQIANG ZHANG, ZHONGJUN HE, HAIRONG LIU, XING LI, HUA WU, AND HAIFENG WANG. STACL: SIMULTANEOUS TRANSLATION WITH IMPLICIT ANTICIPATION AND CONTROLLABLE LATENCY USING PREFIX-TO-PREFIX FRAMEWORK. IN *PROCEEDINGS OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2019.

BILL MACCARTNEY AND CHRISTOPHER D. MANNING. MODELING SEMANTIC CONTAINMENT AND EXCLUSION IN NATURAL LANGUAGE INFERENCE. IN DONIA SCOTT AND HANS USZKOREIT, EDITORS, *PROCEEDINGS OF INTERNATIONAL CONFERENCE ON COMPUTATIONAL LINGUISTICS*, PAGES 521–528, 2008. URL <https://aclanthology.org/C08-1066/>.

KENNY MALONE. HOW UNCLE JAMIE BROKE JEOPARDY. *PLANET MONEY*, (912), MAY 2019.

DAVID MARCHESE. IN CONVERSATION: ALEX TREBEK THE JEOPARDY! ICON ON RETIREMENT, HIS LEGACY, AND WHY KNOWLEDGE MATTERS. *VULTURE*, NOVEMBER 2018.

STEPHEN MITCHELL. *OLD NORSE RIDDLES AND OTHER VERBAL CONTESTS IN PERFORMANCE*, PAGES 123–136. ARC HUMANITIES PRESS, 2020. ISBN 9781641893381. URL <http://www.jstor.org/stable/j.ctv22d4zhq.17>.

ADAM NAJBERG. ALIBABA AI MODEL TOPS HUMANS IN READING COMPREHENSION, 2018. URL <https://www.alizila.com/alibaba-ai-model-tops-humans-in-reading-comprehension/>.

BO PANG AND LILLIAN LEE. *OPINION MINING AND SENTIMENT ANALYSIS*. Now PUBLISHERS INC, 2008. ISBN 1601981503.

DAVID A. PATTERSON. ROBOTS IN THE DESERT: A RESEARCH PARABLE FOR OUR TIMES. *COMMUN. ACM*, 48(12):31–33, DEC 2005. ISSN 0001-0782. doi: 10.1145/1101779.1101800. URL <https://doi.org/10.1145/1101779.1101800>.

JUDEA PEARL. *PROBABILISTIC REASONING IN INTELLIGENT SYSTEMS: NETWORKS OF PLAUSIBLE INFERENCE*. MORGAN KAUFMANN PUBLISHERS INC., SAN FRANCISCO, CA, USA, 1988. ISBN 1558604790.

ETHAN PEREZ, PATRICK LEWIS, WEN-TAU YIH, KYUNGHYUN CHO, AND DOUWE KIELA. UNSUPERVISED QUESTION DECOMPOSITION FOR QUESTION ANSWERING. IN BONNIE WEBBER, TREVOR COHN, YULAN HE, AND YANG LIU, EDITORS, *PROCEEDINGS OF THE 2020 CONFERENCE ON EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING (EMNLP)*, PAGES 8864–8880, ONLINE, NOVEMBER 2020. ASSOCIATION FOR COMPUTATIONAL LINGUISTICS. doi: 10.18653/v1/2020.emnlp-main.713. URL [https://aclanthology.org/2020.emnlp-main.713/](https://aclanthology.org/2020.emnlp-main.713).

PLATO. *THE COLLECTED DIALOGUES OF PLATO, INCLUDING THE LETTERS*. 1949.

M. QIAN AND G.O. TOTTEN. *TRADITIONAL GOVERNMENT IN IMPERIAL CHINA: A CRITICAL ANALYSIS*. TRADITIONAL GOVERNMENT IN IMPERIAL CHINA: A CRITICAL ANALYSIS. CHINESE UNIVERSITY PRESS, 1982. ISBN 9789622012547. URL <https://books.google.com.tw/books?id=A64gui61Zg0C>.

DRAGOMIR R. RADEV, HONG QI, ZHIPING ZHENG, SASHA BLAIR-GOLDENSOHN, ZHU ZHANG, WEIGUO FAN, AND JOHN PRAGER. MINING THE WEB FOR ANSWERS TO NATURAL LANGUAGE QUESTIONS. IN *PROCEEDINGS OF THE TENTH INTERNATIONAL CONFERENCE ON INFORMATION AND KNOWLEDGE MANAGEMENT*, CIKM ’01, PAGE 143–150, NEW YORK, NY, USA, 2001. ASSOCIATION FOR COMPUTING MACHINERY. ISBN 1581134363. doi: 10.1145/502585.502610. URL <https://doi.org/10.1145/502585.502610>.

PRANAV RAJPURKAR, JIAN ZHANG, KONSTANTIN LOPYREV, AND PERCY LIANG. SQuAD: 100,000+ QUESTIONS FOR MACHINE COMPREHENSION OF TEXT. IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, 2016.

PRANAV RAJPURKAR, ROBIN JIA, AND PERCY LIANG. KNOW WHAT YOU DON’T KNOW: UNANSWERABLE QUESTIONS FOR SQuAD. IN *PROCEEDINGS OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, JULY 2018.

DOI: 10.18653/v1/P18-2124. URL <https://www.aclweb.org/anthology/P18-2124>.

ALMUT-BARBARA RENGER. *OEDIPUS AND THE SPHINX: THE THRESHOLD MYTH FROM SOPHOCLES THROUGH FREUD TO COCTEAU*. UNIVERSITY OF CHICAGO PRESS, 2013. ISBN 9780226048116. URL <https://books.google.com/books?id=IRE4AAAAQBAJ>.

D.M. RITCHIE AND B.W. KERNIGHAN. *THE C PROGRAMMING LANGUAGE*. BELL LABORATORIES, 1988.

STEPHEN ROBERTSON AND HUGO ZARAGOZA. THE PROBABILISTIC RELEVANCE FRAMEWORK: BM25 AND BEYOND. *FOUND. TRENDS INF. RETR.*, 3(4):333–389, APRIL 2009. ISSN 1554-0669. DOI: 10.1561/1500000019. URL <https://doi.org/10.1561/1500000019>.

PEDRO RODRIGUEZ, SHI FENG, MOHIT IYYER, HE HE, AND JORDAN BOYD-GABER. QUIZBOWL: THE CASE FOR INCREMENTAL QUESTION ANSWERING. *CoRR*, abs/1904.04792, 2019. URL <http://arxiv.org/abs/1904.04792>.

PEDRO RODRIGUEZ, JOE BARROW, ALEXANDER HOYLE, JOHN P. LALOR, ROBIN JIA, AND JORDAN BOYD-GABER. EVALUATION EXAMPLES ARE NOT EQUALLY INFORMATIVE: HOW SHOULD THAT CHANGE NLP LEADERBOARDS? IN *ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2021. URL http://umiacs.umd.edu/~jbg//docs/2021_acl_leaderboard.pdf.

LISA ROGAK. *WHO IS ALEX TREBEK?: A BIOGRAPHY*. THOMAS DUNNE BOOKS, AN IMPRINT OF ST. MARTIN’S PUBLISHING GROUP, NEW YORK, NY, 2020. ISBN 1250773660.

MARC-ANTOINE RONDEAU AND T.J. HAZEN. SYSTEMATIC ERROR ANALYSIS OF THE STANFORD QUESTION ANSWERING DATASET. IN *PROCEEDINGS OF THE WORKSHOP ON MACHINE READING FOR QUESTION ANSWERING*, JULY 2018. URL <https://www.aclweb.org/anthology/W18-2602>.

DAN RUSSELL. WHY CONTROL-F IS THE SINGLE MOST IMPORTANT THING YOU CAN TEACH SOMEONE ABOUT SEARCH. *SEARCHRESEARCH*, 2020. URL <http://searchresearch1.blogspot.com/2010/02/why-control-f-is-single-most-important.html>.

PETER SCHICKELE. *P.D.Q. BACH ; OEDIPUS TEX AND OTHER CHORAL CALAMITIES*. TELARC, 1990. COMPACT DISC.

DAVID SEDMAN. *ABBOTT AND COSTELLO: WHO'S ON FIRST?*, PAGES 12–32. RUTGERS UNIVERSITY PRESS, 2011. ISBN 9780813549637. URL <http://www.jstor.org/stable/j.ctt5hj328.5>.

SOPHOCLES. *OEDIPUS THE KING*. WASHINGTON SQUARE PRESS, NEW YORK, 1959.

SAKU SUGAWARA, KENTARO INUI, SATOSHI SEKINE, AND AKIKO AIZAWA. WHAT MAKES READING COMPREHENSION QUESTIONS EASIER? IN *PROCEEDINGS OF EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING*, 2018. doi: 10.18653/v1/D18-1453. URL <https://www.aclweb.org/anthology/D18-1453>.

SAKU SUGAWARA, PONTUS STENETORP, KENTARO INUI, AND AKIKO AIZAWA. ASSESSING THE BENCHMARKING CAPACITY OF MACHINE READING COMPREHENSION DATASETS. IN *ASSOCIATION FOR THE ADVANCEMENT OF ARTIFICIAL INTELLIGENCE*, 2020.

IDAN SZPEKTOR AND GIDEON DROR. FROM QUERY TO QUESTION IN ONE CLICK: SUGGESTING SYNTHETIC QUESTIONS TO SEARCHERS. IN *PROCEEDINGS OF THE WORLD WIDE WEB CONFERENCE*, 2013.

GERALD TESAURO, DAVID C. GONDEK, JONATHAN LENCHER, JAMES FAN, AND JOHN M. PRAGER. *JOURNAL OF ARTIFICIAL INTELLIGENCE RESEARCH*, 21, 2013.

JAMES THORNE, ANDREAS VLACHOS, OANA COCARASCU, CHRISTOS CHRISTODOULOPOULOS, AND ARPIT MITTAL, EDITORS. *PROCEEDINGS OF THE FIRST WORKSHOP ON FACT EXTRACTION AND VERIFICATION (FEVER)*, 2018. ASSOCIATION FOR COMPUTATIONAL LINGUISTICS.

CHRISTOPHER TOLKIEN. *THE SAGA OF KING HEIDREK THE WISE*. NELSON'S ICELANDIC TEXTS. THOMAS NELSON AND SONS, LONDON, 1960. TRANSLATED FROM THE ICELANDIC WITH INTRODUCTION, NOTES AND APPENDICES.

LEE TREPANIER. *THE SOCRATIC METHOD TODAY: STUDENT-CENTERED AND TRANSFORMATIVE TEACHING IN POLITICAL SCIENCE*. TAYLOR & FRANCIS, 2017. ISBN 9781351245807. URL <https://books.google.com/books?id=9TFDDwAAQBAJ>.

LUIS VON AHN. GAMES WITH A PURPOSE. *COMPUTER*, 39:92–94, 07 2006. doi: 10.1109/MC.2006.196.

ELLEN VOORHEES. *THE EVOLUTION OF CRANFIELD*. CLEF@20 - INFORMATION RETRIEVAL EVALUATION IN A CHANGING WORLD: LESSONS LEARNED FROM 20 YEARS OF CLEF, SPRINGER-VERLAG, BERLIN, 2019-08-14 04:08:00 2019. doi: https://doi.org/10.1007/978-3-030-22948-1_2. URL https://tsapps.nist.gov/publication/get_pdf.cfm?pub_id=926510.

ELLEN M. VOORHEES. EVALUATING THE EVALUATION: A CASE STUDY USING THE TREC 2002 QUESTION ANSWERING TRACK. IN *CONFERENCE OF THE NORTH AMERICAN CHAPTER OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2003. URL <https://www.aclweb.org/anthology/N03-1034>.

ELLEN M. VOORHEES. *EVALUATING QUESTION ANSWERING SYSTEM PERFORMANCE*, PAGES 409–430. SPRINGER NETHERLANDS, DORDRECHT, 2008. ISBN 978-1-4020-4746-6. DOI: 10.1007/978-1-4020-4746-6_13. URL https://doi.org/10.1007/978-1-4020-4746-6_13.

ELLEN M VOORHEES AND DAWN M TICE. BUILDING A QUESTION ANSWERING TEST COLLECTION. IN *PROCEEDINGS OF THE ACM SIGIR CONFERENCE ON RESEARCH AND DEVELOPMENT IN INFORMATION RETRIEVAL*, 2000.

ERIC WALLACE, PEDRO RODRIGUEZ, SHI FENG, IKUYA YAMADA, AND JORDAN BOYD-GRABER. TRICK ME IF YOU CAN: HUMAN-IN-THE-LOOP GENERATION OF ADVERSARIAL QUESTION ANSWERING EXAMPLES. *TRANSACTIONS OF THE ASSOCIATION OF COMPUTATIONAL LINGUISTICS*, 10, 2019A.

ERIC WALLACE, YIZHONG WANG, SUJIAN LI, SAMEER SINGH, AND MATT GARDNER. DO NLP MODELS KNOW NUMBERS? PROBING NUMERACY IN EMBEDDINGS. IN *PROCEEDINGS OF THE 2019 CONFERENCE ON EMPIRICAL METHODS IN NATURAL LANGUAGE PROCESSING AND THE 9TH INTERNATIONAL JOINT CONFERENCE ON NATURAL LANGUAGE PROCESSING (EMNLP-IJCNLP)*, PAGES 5307–5315, HONG KONG, CHINA, NOVEMBER 2019B. ASSOCIATION FOR COMPUTATIONAL LINGUISTICS. DOI: 10.18653/v1/D19-1534. URL <https://aclanthology.org/D19-1534>.

DIRK WEISSENBORN, GEORG WIESE, AND LAURA SEIFFE. MAKING NEURAL QA AS SIMPLE AS POSSIBLE BUT NOT SIMPLER. IN *CONFERENCE ON COMPUTATIONAL NATURAL LANGUAGE LEARNING*, 2017. DOI: 10.18653/v1/K17-1028. URL <https://www.aclweb.org/anthology/K17-1028>.

R WILKINSON, JON GEIST, STANLEY JANET, PATRICK GROTHÉR, CHRISTOPHER BURGES, ROBERT CREECY, BOB HAMMOND, JONATHAN HULL, NORMAN LARSEN, THOMAS VOGL, AND CHARLES WILSON. THE FIRST CENSUS OPTICAL CHARACTER RECOGNITION SYSTEMS CONFERENCE, 1992-01-06 00:01:00 1992. URL https://tsapps.nist.gov/publication/get_pdf.cfm?pub_id=905664.

W.A. WOODS, R.M. KAPLAN, AND B. NASH-WEBBER. *THE LUNAR SCIENCES NATURAL LANGUAGE INFORMATION SYSTEM*. BBN REPORT 2378. PREPARED FOR THE LANGUAGE RESEARCH FOUNDATION BY BOLT, BERANEK AND NEWMAN, 1971. URL <https://books.google.com.tw/books?id=tUcCHQACAAJ>.

ZHILIN YANG, ZIHANG DAI, YIMING YANG, JAIME G. CARBONELL, RUSLAN SALAKHUTDINOV, AND QUOC V. LE. XLNET: GENERALIZED AUTOREGRESSIVE PRETRAINING FOR LANGUAGE UNDERSTANDING. 2019.

WEN-TAU YIH, MING-WEI CHANG, XIAODONG HE, AND JIANFENG GAO. SEMANTIC PARSING VIA STAGED QUERY GRAPH GENERATION: QUESTION ANSWERING

WITH KNOWLEDGE BASE. IN *PROCEEDINGS OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2015.

XINGDI YUAN, JIE FU, MARC-ALEXANDRE COTE, YI TAY, CHRISTOPHER PAL, AND ADAM TRISCHLER. INTERACTIVE MACHINE COMPREHENSION WITH INFORMATION SEEKING AGENTS, 2019.

ROWAN ZELLERS, ARI HOLTZMAN, YONATAN BISK, ALI FARHADI, AND YEJIN CHOI. HELLA SWAG: CAN A MACHINE REALLY FINISH YOUR SENTENCE? IN ANNA KORHONEN, DAVID R. TRAUM, AND LLUÍS MÀRQUEZ, EDITORS, *PROCEEDINGS OF THE ASSOCIATION FOR COMPUTATIONAL LINGUISTICS*, 2019.